

Scalable Robot Learning

by

Ashvin Nair

A dissertation submitted in partial satisfaction

of the requirements for the degree of

Doctor of Philosophy

in

Engineering - Electrical Engineering & Computer Science

in the Graduate Division of the

University of California, Berkeley

Summer 2022

Committee:

Professor Sergey Levine, Chair
Professor Pieter Abbeel
Professor Alison Gopnik

ABSTRACT

For robots to perform tasks in the unstructured environments of the real world, they must be able to be tasked with a desired objective in a general way, and learn to perform the desired task quickly if the robot does not already know how to accomplish it. In this thesis, we explore deep reinforcement learning as a solution to enable this vision for scalable learning-based real-world robotics. We cover two main themes: accelerating reinforcement learning from prior data and self-supervised RL. Accelerating RL from prior data or prior knowledge is important for making reinforcement learning algorithms sufficiently sample-efficient to run directly in the real world. We discuss utilizing human demonstrations to accelerate reinforcement learning, using human-designed residual controllers in combination with reinforcement learning for industrial insertion tasks, and algorithms for offline reinforcement learning that can also benefit from a small amount of online fine-tuning. Concurrently, while sample-efficiency of reinforcement learning algorithms is a well-appreciated problem, additional problems arise around agents that can learn from rich observations such as images: in particular, reward supervision and collecting data autonomously. We will discuss our work on self-supervised RL through goal reaching with a generative model, allowing agents to evaluate their own success at reaching goals and autonomously propose and practice skills. In the final section, we combine offline policy learning with self-supervised practice, allowing robots to practice and perfect skills in novel environments. These directions enable robots to supervise their own data collection, learning complex and general manipulation skills from interaction.

ACKNOWLEDGMENTS

TODO

CONTENTS

1	INTRODUCTION	1
I	REINFORCEMENT LEARNING WITH IMAGINED GOALS	4
2	VISUAL REINFORCEMENT LEARNING WITH IMAGINED GOALS	5
2.1	Introduction	5
2.2	Related Work	7
2.3	Background	8
2.4	Goal-Conditioned Policies with Unsupervised Representation Learning	10
2.4.1	Sample-Efficient RL with Learned Representations	10
2.4.2	Reward Specification	11
2.4.3	Improving Sample Efficiency with Latent Goal Relabeling	12
2.4.4	Automated Goal-Generation for Exploration	12
2.4.5	Algorithm Summary	13
2.5	Experiments	13
2.5.1	Visual RL with Physical Robots	18
2.6	Discussion and Future Work	19
3	SKEWFIT: STATE-COVERING SELF-SUPERVISED REINFORCEMENT LEARNING	21
3.1	Introduction	21
3.2	Problem Formulation	23
3.2.1	Minimizing $\mathcal{H}(\mathbf{G} \mathbf{S})$: Goal-Conditioned Reinforcement Learning	23
3.2.2	Maximizing $\mathcal{H}(\mathbf{G})$: Setting Diverse Goals	24
3.3	Advantage Weighted Actor Critic: Learning a Maximum Entropy Goal Distribution	25
3.3.1	Advantage Weighted Actor Critic Algorithm	25
3.3.2	Advantage Weighted Actor Critic Analysis	28
3.4	Training Goal-Conditioned Policies with Advantage Weighted Actor Critic	29
3.5	Related Work	30
3.6	Experiments	31
3.7	Conclusion	39
4	CONTEXTUAL IMAGINED GOALS FOR SELF-SUPERVISED ROBOTIC LEARNING	41
4.1	Introduction	41

4.2	Related Work	43
4.3	Background	45
	4.3.1 Goal-Conditioned Reinforcement Learning	45
	4.3.2 Variational Auto-Encoders	45
	4.3.3 Conditional Variational Auto-Encoders	46
	4.3.4 Reinforcement Learning with Imagined Goals	46
4.4	Self-Supervised Learning with Context-Conditioned Representations	47
	4.4.1 Context-Conditioned VAEs	47
	4.4.2 Context-Conditioned Reinforcement Learning with Imagined Goals	48
4.5	Experiments	49
	4.5.1 Self-Supervised Learning in Simulation	49
	4.5.2 Generalizing to Varying Appearance and Dynamics with Self-Supervised Learning	50
	4.5.3 Context-Conditioned VAE Goal Sampling	51
	4.5.4 Real-World Robotic Evaluation	52
4.6	Conclusion	54
II	UTILIZING PRIOR DATA FOR CONTROL	55
5	OVERCOMING EXPLORATION IN REINFORCEMENT LEARNING WITH DEMONSTRATIONS	56
5.1	Introduction	56
5.2	Related Work	58
5.3	Background	60
	5.3.1 Reinforcement Learning	60
	5.3.2 DDPG	60
	5.3.3 Multi-Goal RL	61
	5.3.4 Hindsight Experience Replay (HER)	62
5.4	Method	62
	5.4.1 Demonstration Buffer	62
	5.4.2 Behavior Cloning Loss	62
	5.4.3 Q-Filter	63
	5.4.4 Resets to demonstration states	63
5.5	Experimental Setup	64
	5.5.1 Environments	64
	5.5.2 Training Details	65
	5.5.3 Overview of Experiments	65
5.6	Comparison With Prior Work	65

5.6.1	Tasks	65
5.6.2	Results	66
5.7	Multi-Step Experiments	66
5.7.1	Block Stacking Task	67
5.7.2	Rewards	67
5.7.3	Network architectures	68
5.7.4	Baselines	68
5.7.5	Results	69
5.8	Ablation Experiments	69
5.8.1	Behavior Cloning Loss	70
5.8.2	Q-Filter	71
5.8.3	Resets From Demonstrations	72
5.9	Discussion and Future Work	72
6	RESIDUAL REINFORCEMENT LEARNING FOR ROBOT CONTROL	74
6.1	Introduction	74
6.2	Preliminaries	76
6.2.1	Problem Statement - System Theoretic Interpretation	76
6.2.2	Interpretation as a Reinforcement Learning Problem	77
6.3	Method	78
6.3.1	Residual Reinforcement Learning	78
6.3.2	Method Summary	79
6.4	Experimental Setup	80
6.4.1	Simulated Environment	80
6.4.2	Real-World Environment	80
6.4.3	Overview of Experiments	81
6.5	Experiments	81
6.5.1	Sample Efficiency of Residual RL	81
6.5.2	Effect of Environment Variation	81
6.5.3	Recovering from Control Noise	82
6.5.4	Sim-to-Real with Residual RL	83
6.6	Results	83
6.6.1	Sample Efficiency of Residual RL	83
6.6.2	Effect of Environment Variation	83
6.6.3	Recovering from Control Noise	85
6.6.4	Sim-to-Real with Residual RL	85
6.7	Related Work	85

6.8	Conclusion	87
7	DEEP REINFORCEMENT LEARNING FOR INDUSTRIAL INSERTION TASKS WITH VISUAL INPUTS AND NATURAL REWARDS	88
7.1	Introduction	88
7.2	Related Work	90
7.3	Electric Connector Plug Insertion Tasks	91
7.3.1	Adapters	92
7.3.2	Experimental Settings	92
7.4	Methods	94
7.4.1	Preliminaries	94
7.4.2	Efficient Off-Policy Reinforcement Learning	94
7.4.3	Residual Reinforcement Learning	95
7.4.4	Learning from Demonstrations	96
7.5	Experiments	96
7.6	Results	97
7.6.1	Vision-based Learning	98
7.6.2	Learning From Sparse Rewards	99
7.6.3	Perfect State Information	100
7.6.4	Robustness	100
7.6.5	Exploration Comparison	101
7.7	Conclusion	101
8	AWAC: ACCELERATING ONLINE REINFORCEMENT LEARNING WITH OFFLINE DATASETS	102
8.1	Introduction	102
8.2	Preliminaries	104
8.3	Challenges in Offline RL with Online Fine-tuning	106
8.3.1	Problem Definition	106
8.3.2	Data Efficiency	106
8.3.3	Bootstrap Error in Offline Learning with Actor-Critic Methods	107
8.3.4	Excessively Conservative Online Learning	108
8.4	Advantage Weighted Actor Critic: A Simple Algorithm for Fine-tuning from Offline Datasets	109
8.5	Related Work	113
8.6	Experimental Evaluation	115
8.6.1	Simulated Experiments	115
8.6.2	Real-World Robot Learning with Prior Data	118

8.7	Discussion and Future Work	120
9	OFFLINE REINFORCEMENT LEARNING WITH IMPLICIT Q-LEARNING	122
9.1	Related work	123
9.2	Preliminaries	125
9.3	Implicit Q-Learning	125
9.3.1	Expectile Regression	127
9.3.2	Learning the Value Function with Expectile Regression	128
9.3.3	Policy Extraction and Algorithm Summary	129
9.3.4	Analysis	130
9.4	Experimental Evaluation	131
9.4.1	The Difference Between One-Step Policy Improvement and IQN	132
9.4.2	Comparisons on Offline RL Benchmarks	132
9.4.3	Online Fine-tuning after Offline RL	134
9.5	Conclusion	135
III AFFORDANCE LEARNING IN UNSEEN ENVIRONMENTS		138
10	LEARNING NEW SKILLS BY IMAGINING VISUAL AFFORDANCES	139
10.1	Introduction	139
10.2	Related Work	141
10.3	Preliminaries	143
10.4	Problem Setting	144
10.5	Visuomotor Affordance Learning	145
10.5.1	Affordance Learning	145
10.5.2	Offline Behavior Learning	147
10.5.3	Online Behavior Learning	147
10.6	Real-World Experimental Evaluation	148
10.6.1	Generative Models for Affordance Learning	149
10.6.2	Real-World Visuomotor Affordance Learning	150
10.7	Experimental Evaluation in Simulation	151
10.7.1	Self-Supervised Online Fine-Tuning from Prior Data	152
10.7.2	Scalable Robot Learning with VAL	153
10.8	Conclusion	154
11	PLANNING TO PRACTICE: EFFICIENT ONLINE FINE-TUNING BY COMPOSING GOALS IN LATENT SPACE	155
11.1	Introduction	155
11.2	Related Work	158

11.3	Problem Statement	160
11.4	Preliminaries	160
11.5	Planning to Practice	161
11.5.1	Conditional Subgoal Generation	162
11.5.2	Efficient Planning in the Latent Space	163
11.5.3	Cost Function For Feasible Subgoals	165
11.6	Experiments	166
11.6.1	Experimental Setup	167
11.6.2	Implementation Details	168
11.6.3	Quantitative Comparisons	168
11.6.4	Generated Subgoals	170
11.7	Conclusion and Discussion	171

1

INTRODUCTION

Robots are becoming ubiquitous in manufacturing and other industries, for a variety of tasks such as bin picking, welding, painting, assembly, and so on. Yet, the autonomous capability of present-day robotics systems are still quite limited. Settings where robots operate are carefully controlled; they often require very specific end-effector tooling combined with high precision motions to do a particular task. In effect, robots rely on human ingenuity and engineering in order to do their job. But such systems are brittle, and the hardware and software must often be redesigned for slight variations of a task. If a manufacturing task actually requires adaptability or robustness to varying environment conditions, designing a working system becomes much more difficult. And beyond manufacturing, we will expect the robots of tomorrow to do a lot more: cook meals, assist the elderly in homes and other human-centric environments, navigate unmapped terrain, operate machinery and appliances, manipulate objects, and interact safely in presence of humans. This kind of open-world capability requires adaptability, generalization, and is beyond the reach of most robots today.

In contrast, humans perform highly skillful dexterous manipulation so easily that it is sometimes hard to conceive the difficulty of replicating this capability in a robot. Most humans within the first five years of their life have developed complex fine motor skills, successfully performing bimanual dexterous manipulation of all kinds of unfamiliar and dynamic objects, and using tools with a tight sensorimotor loop that entails perception, functional grasping, and control ([Adolph and Franchak, 2017](#)). It remains a challenge to develop equivalently robust feedback controllers for robots that can adapt to a wide variety of situations to accomplish goals. If robots were equally skillful, it would be incredibly economically valuable - they could be used to automate many of the tasks that humans have to do today. How can we develop methods to allow robots to become similarly skillful?

The past decade of deep learning suggests that learning from large datasets is the

key to such open-world generalization for robots. Expressive models that are trained on broad datasets have driven recent progress in artificial intelligence research. These models are trained on a broad enough dataset to capture corner cases within its training distribution. Beginning with the breakthrough of AlexNet in image classification, the recipe of large datasets combined with large models has resulted in human-level performance in vision (Krizhevsky et al., 2012). Similarly, in NLP, pretraining on large datasets dominates (Devlin et al., 2019). These models have contributed to significant progress in protein folding. If we could achieve such generality for control - the problem of selecting actions over time in order to maximize a reward function - it could enable truly general robots in the wild.

But control introduces two new challenges not found in the supervised learning setting. The first challenge is credit assignment: actions taken in the past affect the future. The second is exploration: actions taken change the distribution of data visited. To address these challenges, a promising line of work is deep reinforcement learning. Deep reinforcement learning has been applied successfully for super-human performance on games such as Atari (Mnih et al., 2015), Go (D. Silver et al., 2016a), Dota, and Starcraft. It has been used for robotics, stratospheric balloons, and even control of plasma in a nuclear fusion reactor. Yet, while algorithms for RL have been steadily advancing, becoming more sample efficient and stable, there are still significant obstacles towards truly solving robotics with RL. What challenges remain toward endowing robots with human-level skills?

Challenges in robot learning. The central issue is that the world is so varied that a robotic agent acquiring skills via learning based methods needs to experience both the diversity of perceptual inputs, combined with the complexity of control including exploration and credit assignment. collect its own data - but there's a chicken and the egg problem between the policy and data. You need a policy that generalizes to collect useful data in varied environments, but you need good varied data to obtain such a policy. To do so, we face several challenges. When a robot is dropped in a new environment, it must be able to use its prior knowledge to think of potentially useful behaviors that the environment affords. Then, the robot has to be able to actually practice these behaviors informatively. To now improve itself in the new environment, the robot must then be able to evaluate its own success somehow without an externally provided reward.

First, speeding up policy optimization from prior knowledge: usually prior data like offline experience or demonstrations, but also things like environment models and human-engineered controllers. The key idea is that instead of learning tasks from scratch, we ought to be able to utilize existing knowledge to make RL algorithms sample efficient

and practical for running in the real world.

Second, where does the reward/supervision come from? which leads to the framework of self-supervised RL. So here, we deal with being able to evaluate a reward function for many goals g that the agent may be tasked with, how to infer the distribution of potential goals to train on so you can generalize to test tasks, and how this interacts with policy learning.

If we can overcome these challenges reliably, we can bootstrap a cycle in which our agents use prior experience to learn a policy to collect high quality interaction data, improve the policy from that data, and so on.

Outline. This thesis is organized as follows. In Part I (Chapter 2-4), we explore utilizing goal-conditioned reinforcement learning for self-supervised exploration from raw observations. In chapter 2, we describe the framework of reinforcement learning with imagined goals (RIG), which enables self-supervised practice. In chapter 3, we discuss an extension of RIG to improve exploration in self-supervised goal-conditioned RL. In chapter 4, we discuss extending RIG to learn in novel situations, using a context-conditioned variation autoencoder (CCVAE). In Part 2 (Chapter 5-8), we discuss utilizing prior data and prior knowledge in order to initialize and accelerate reinforcement learning. In chapter 5, we discuss utilizing demonstrations to solve long-horizon tasks with reinforcement learning. In chapter 6, we discuss how to incorporate prior knowledge such as an expert controller using residual reinforcement learning. In chapter 7, we discuss an algorithm, advantage weighted actor critic (AWAC) to utilize arbitrary offline data, and finetune policies and Q functions online. In chapter 8, we discuss a significant improvement to AWAC, implicit Q-learning (IQL) that is state of the art in both offline RL and online finetuning. Finally, in Part 3 (Chapter 9-10), we show how these two directions come together to enable robots in the real world to explore novel situations. In chapter 9, we cover visuomotor affordance learning (VAL), a method to allow self-supervised learning from prior data. In chapter 10, we extend VAL with planning to enable finetuning of more complex skills. In chapter 11, we discuss some future directions.

Part I

REINFORCEMENT LEARNING WITH IMAGINED GOALS

2

VISUAL REINFORCEMENT LEARNING WITH IMAGINED GOALS

2.1 INTRODUCTION

Reinforcement learning (RL) algorithms hold the promise of allowing autonomous agents, such as robots, to learn to accomplish arbitrary tasks. However, the standard RL framework involves learning policies that are specific to individual tasks, which are defined by hand-specified reward functions. Agents that exist persistently in the world can prepare to solve diverse tasks by setting their own goals, practicing complex behaviors, and learning about the world around them. In fact, humans are very proficient at setting abstract goals for themselves, and evidence shows that this behavior is already present from early infancy (Smith and Gasser, 2005), albeit with simple goals such as reaching. The behavior and representation of goals grows more complex over time as they learn how to manipulate objects and locomote. How can we begin to devise a reinforcement learning system that sets its own goals and learns from experience with minimal outside intervention and manual engineering?

In this paper, we take a step toward this goal by designing an RL framework that jointly learns representations of raw sensory inputs and policies that achieve arbitrary goals under this representation by practicing to reach self-specified random goals during training. To provide for automated and flexible goal-setting, we must first choose how a general goal can be specified for an agent interacting with a complex and highly variable environment. Even providing the state of such an environment to a policy is a challenge. For instance, a task that requires a robot to manipulate various objects would require a combinatorial representation, reflecting variability in the number and type of objects in the current scene. Directly using raw sensory signals, such as images, avoids this challenge, but learning from raw images is substantially harder. In particular, pixel-wise Euclidean distance is not an effective reward function for visual tasks since distances between images do not correspond to meaningful distances between states (Ponomarenko

et al., 2015; Richard Zhang et al., 2018). Furthermore, although end-to-end model-free reinforcement learning can handle image observations, this comes at a high cost in sample complexity, making it difficult to use in the real world.

We propose to address both challenges by incorporating unsupervised representation learning into goal-conditioned policies. In our method, which is illustrated in Figure 26, a representation of raw sensory inputs is learned by means of a latent variable model, which in our case is based on the variational autoencoder (VAE) (D. P. Kingma and Welling, 2014). This model serves three complementary purposes. First, it provides a more structured representation of sensory inputs for RL, making it feasible to learn from images even in the real world. Second, it allows for sampling of new states, which can be used to set synthetic goals during training to allow the goal-conditioned policy to practice diverse behaviors. We can also more efficiently utilize samples from the environment by relabeling synthetic goals in an off-policy RL algorithm, which makes our algorithm substantially more efficient. Third, the learned representation provides a space where distances are more meaningful than the original space of observations, and can therefore provide well-shaped reward functions for RL. By learning to reach random goals sampled from the latent variable model, the goal-conditioned policy learns about the world and can be used to achieve new, user-specified goals at test-time.

The main contribution of our work is a framework for learning general-purpose goal-conditioned policies that can achieve goals specified with target observations. We call our method reinforcement learning with imagined goals (RIG). RIG combines sample-efficient off-policy goal-conditioned reinforcement learning with unsupervised representation learning. We use representation learning to acquire a latent distribution that can be used to sample goals for unsupervised practice and data augmentation, to provide a well-shaped distance function for reinforcement learning, and to provide a more structured representation for the value function and policy. While several prior methods, discussed in the following section, have sought to learn goal-conditioned policies, we can do so with image goals and observations without a manually specified reward signal. Our experimental evaluation illustrates that our method substantially improves the performance of image-based reinforcement learning, can effectively learn policies for complex image-based tasks, and can be used to learn real-world robotic manipulation skills with raw image inputs. Videos of our method in simulated and real-world environments can be found at <https://sites.google.com/site/visualrlwithimaginedgoals/>.

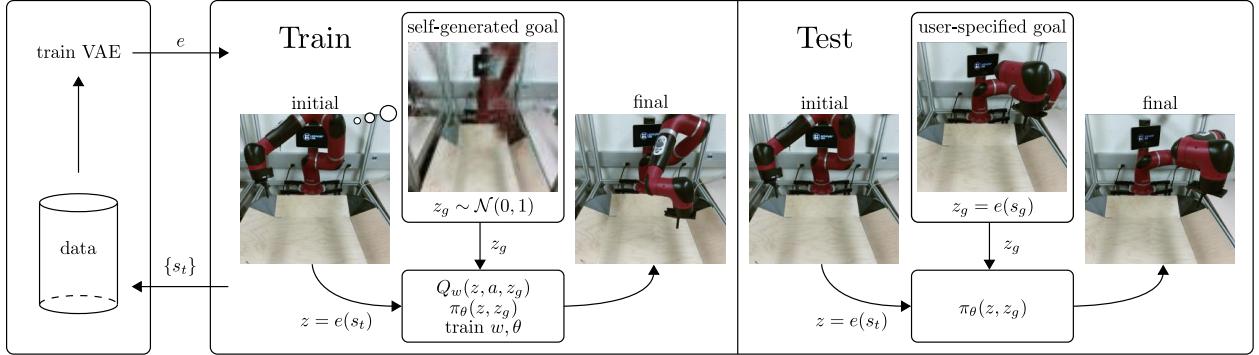


Figure 1: We train a VAE using data generated by our exploration policy (left). We use the VAE for multiple purposes during training time (middle): to sample goals to train the policy, to embed the observations into a latent space, and to compute distances in the latent space. During test time (right), we embed a specified goal observation o_g into a goal latent z_g as input to the policy.

2.2 RELATED WORK

While prior works on vision-based deep reinforcement learning for robotics can efficiently learn a variety of behaviors such as grasping (Pinto et al., 2018; Pinto and Abhinav Gupta, 2016; Levine et al., 2017), pushing (Agrawal et al., 2016; Ebert et al., 2017; Finn and Levine, 2016), navigation (Pathak et al., 2018; Lange et al., 2012c), and other manipulation tasks (Lillicrap et al., 2016; Levine et al., 2016a; Pathak et al., 2018), they each make assumptions that limit their applicability to training general-purpose robots. Levine et al. (2016a) uses time-varying models, which requires an episodic setup that makes them difficult to extend to non-episodic and continual learning scenarios. Pinto et al. (2018) proposed a similar approach that uses goal images, but requires instrumented training in simulation. Lillicrap et al. (2016) uses fully model-free training, but does not learn goal-conditioned skills. As we show in our experiments, this approach is very difficult to extend to the goal-conditioned setting with image inputs. Model-based methods that predict images (Watter et al., 2015; Finn and Levine, 2016; Ebert et al., 2017; Oh et al., 2015) or learn inverse models (Agrawal et al., 2016) can also accommodate various goals, but tend to limit the horizon length due to model drift. To our knowledge, no prior method uses model-free RL to learn policies conditioned on a single goal image with sufficient efficiency to train directly on real-world robotic systems, without access to ground-truth state or reward information during training.

Our method uses a goal-conditioned value function (Schaul et al., 2015b) in order to solve more general tasks (Sutton et al., 2011; L P Kaelbling, 1993). To improve the sample-

efficiency of our method during off-policy training, we retroactively relabel samples in the replay buffer with goals sampled from the latent representation. Goal relabeling has been explored in prior work (L P Kaelbling, 1993; Andrychowicz et al., 2017b; Rauber et al., 2017; Levy et al., 2017; V. Pong et al., 2018). Andrychowicz et al. (2017b) and Levy et al. (2017) use goal relabeling for sparse rewards problems with known goal spaces, restricting the resampled goals to states encountered along that trajectory, since almost any other goal will have no reward signal. We sample random goals from our learned latent space to use as replay goals for off-policy Q-learning rather than restricting ourselves to states seen along the sampled trajectory, enabling substantially more efficient learning. We use the same goal sampling mechanism for exploration in RL. Goal setting for policy learning has previously been discussed (Baranes and Oudeyer, 2012) and recently Pérez et al. (2018a) have also proposed using unsupervised learning for setting goals for exploration. However, we use a model-free Q-learning method that operates on raw state observations and actions, allowing us to solve visually and dynamically complex tasks.

A number of prior works have used unsupervised learning to acquire better representations for RL. These methods use the learned representation as a substitute for the state for the policy, but require additional information, such as access to the ground truth reward function based on the true state during training time (Higgins et al., 2017b; Ha and Schmidhuber, 2018; Watter et al., 2015; Finn et al., 2016b; Lange et al., 2012c; Jonschkowski et al., 2017), expert trajectories (Srinivas et al., 2018), human demonstrations (Sermanet et al., 2017), or pre-trained object-detection features (A. Lee et al., 2017). In contrast, we learn to generate goals and use the learned representation to obtain a reward function for those goals without any of these extra sources of supervision. Finn et al. (2016b) combine unsupervised representation learning with reinforcement learning, but in a framework that trains a policy to reach a single goal. Many prior works have also focused on learning controllable and disentangled representations (Schmidhuber, 1992; X. Chen et al., 2016; Cheung et al., 2014; Reed et al., 2014; Desjardins et al., 2012; V. Thomas et al., 2017). We use a method based on variational autoencoders, but these prior techniques are complementary to ours and could be incorporated into our method.

2.3 BACKGROUND

Our method combines reinforcement learning with goal-conditioned value functions and unsupervised representation learning. Here, we briefly review the techniques that we build on in our method.

GOAL-CONDITIONED REINFORCEMENT LEARNING. In reinforcement learning, the goal is to learn a policy $\pi(s_t) = a_t$ that maximizes expected return, which we denote as $R_t = \mathbb{E}[\sum_{i=t}^T \gamma^{(i-t)} r_i]$, where $r_i = r(s_i, a_i, s_{i+1})$ and the expectation is under the current policy and environment dynamics. Here, $s \in \mathcal{S}$ is a state observation, $a \in \mathcal{A}$ is an action, and γ is a discount factor. Standard model-free RL learns policies that achieve a single task. If our aim is instead to obtain a policy that can accomplish a variety of tasks, we can construct a goal-conditioned policy and reward, and optimize the expected return with respect to a goal distribution: $\mathbb{E}_{g \sim G}[\mathbb{E}_{r_i, s_i \sim E, a_i \sim \pi}[R_0]]$, where G is the set of goals and the reward is also a function of g . A variety of algorithms can learn goal-conditioned policies, but to enable sample-efficient learning, we focus on algorithms that acquire goal-conditioned Q-functions, which can be trained off-policy. A goal-conditioned Q-function $Q(s, a, g)$ learns the expected return for the goal g starting from state s and taking action a . Given a state s , action a , next state s' , goal g , and corresponding reward r , one can train an approximate Q-function parameterized by w by minimizing the following Bellman error

$$\mathcal{E}(w) = \frac{1}{2} \|Q_w(s, a, g) - (r + \gamma \max_{a'} Q_{\bar{w}}(s', a', g))\|^2 \quad (1)$$

where \bar{w} indicates that \bar{w} is treated as a constant. Crucially, one can optimize this loss using off-policy data (s, a, s', g, r) with a standard actor-critic algorithm (Lillicrap et al., 2016; Fujimoto et al., 2018b; Mnih et al., 2016).

VARIATIONAL AUTOENCODERS. Variational autoencoders (VAEs) have been demonstrated to learn structured latent representations of high dimensional data (D. P. Kingma and Welling, 2014). The VAE consists of an encoder p_ϕ , which maps states to latent distributions, and a decoder p_ψ , which maps latents to distributions over states. The encoder and decoder parameters, ψ and ϕ respectively, are jointly trained to maximize

$$\mathcal{L}(\psi, \phi; s^{(i)}) = -\beta D_{KL}(q_\phi(z|s^{(i)})||p(z)) + \mathbb{E}_{q_\phi(z|s^{(i)})}[\log p_\psi(s^{(i)} | z)], \quad (2)$$

where $p(z)$ is some prior, which we take to be the unit Gaussian, D_{KL} is the Kullback-Leibler divergence, and β is a hyperparameter that balances the two terms. The use of β values other than one is sometimes referred to as a β -VAE (Higgins et al., 2017a). The encoder q_ϕ parameterizes the mean and log-variance diagonal of a Gaussian distribution, $q_\phi(s) = \mathcal{N}(\mu_\phi(s), \sigma_\phi^2(s))$. The decoder p_ψ parameterizes a Bernoulli distribution for each pixel value. This parameterization corresponds to training the decoder with cross-entropy loss on normalized pixel values. Full details of the hyperparameters are in the Supplementary Material.

2.4 GOAL-CONDITIONED POLICIES WITH UNSUPERVISED REPRESENTATION LEARNING

To devise a practical algorithm based on goal-conditioned value functions, we must choose a suitable goal representation. In the absence of domain knowledge and instrumentation, a general-purpose choice is to set the goal space \mathcal{G} to be the same as the state observations space \mathcal{S} . This choice is fully general as it can be applied to any task, and still permits considerable user control since the user can choose a “goal state” to set a desired goal for a trained goal-conditioned policy. But when the state space \mathcal{S} corresponds to high-dimensional sensory inputs such as images,¹ learning a goal-conditioned Q-function and policy becomes exceedingly difficult as we illustrate empirically in Section 8.6.

Our method jointly addresses a number of problems that arise when working with high-dimensional inputs such as images: sample efficient learning, reward specification, and automated goal-setting. We address these problems by learning a latent embedding using a β -VAE. We use this latent space to represent the goal and state and retroactively relabel data with latent goals sampled from the VAE prior to improve sample efficiency. We also show that distances in the latent space give us a well-shaped reward function for images. Lastly, we sample from the prior to allow an agent to set and “practice” reaching its own goal, removing the need for humans to specify new goals during training time. We next describe the specific components of our method, and summarize our complete algorithm in Section 2.4.5.

2.4.1 Sample-Efficient RL with Learned Representations

One challenging problem with end-to-end approaches for visual RL tasks is that the resulting policy needs to learn both perception and control. Rather than operating directly on observations, we embed the state s_t and goals g into a latent space \mathcal{Z} using an encoder e to obtain a latent state $z_t = e(s_t)$ and latent goal $z_g = e(g)$. To learn a representation of the state and goal space, we train a β -VAE by executing a random policy and collecting state observations, $\{s^{(i)}\}$, and optimize Equation equation 2. We then use the mean of the encoder as the state encoding, i.e. $z = e(s) \triangleq \mu_\phi(s)$.

After training the VAE, we train a goal-conditioned Q-function $Q(z, a, z_g)$ and corresponding policy $\pi_\theta(z, z_g)$ in this latent space. The policy is trained to reach a goal

¹ We make the simplifying assumption that the system is Markovian with respect to the sensory input, and one could incorporate memory into the state for partially observed tasks.

z_g using the reward function discussed in Section 2.4.2. For the underlying RL algorithm, we use twin delayed deep deterministic policy gradients (TD3) (Fujimoto et al., 2018b), though any value-based RL algorithm could be used. Note that the policy (and Q-function) operates completely in the latent space. During test time, to reach a specific goal state g , we encode the goal $z_g = e(g)$ and input this latent goal to the policy.

As the policy improves, it may visit parts of the state space that the VAE was never trained on, resulting in arbitrary encodings that may not make learning easier. Therefore, in addition to procedure described above, we fine-tune the VAE using both the randomly generated state observations $\{s^{(i)}\}$ and the state observations collected during exploration. We show in Section ?? that this additional training helps the performance of the algorithm.

2.4.2 Reward Specification

Training the goal-conditioned value function requires defining a goal-conditioned reward $r(s, g)$. Using Euclidean distances in the space of image pixels provides a poor metric, since similar configurations in the world can be extremely different in image space. In addition to compactly representing high-dimensional observations, we can utilize our representation to obtain a reward function based on a metric that better reflects the similarity between the state and the goal. One choice for such a reward is to use the negative Mahalanobis distance in the latent space:

$$r(s, g) = -\|e(s) - e(g)\|_A = -\|z - z_g\|_A,$$

where the matrix A weights different dimensions in the latent space. This approach has an appealing interpretation when we set A to be the precision matrix of the VAE encoder, q_ϕ . Since we use a Gaussian encoder, we have that

$$r(s, g) = -\|z - z_g\|_A \propto \sqrt{\log q_\phi(z_g | s)} \quad (3)$$

In other words, minimizing this squared distance in the latent space is equivalent to rewarding reaching states that maximize the probability of the latent goal z_g . In practice, we found that setting $A = I$, corresponding to Euclidean distance, performed better than Mahalanobis distance, though its effect is the same — to bring z close to z_g and maximize the probability of the latent goal z_g given the observation. This interpretation would not be possible when using normal autoencoders since distances are not trained to have any probabilistic meaning. Indeed, we show in Section 8.6 that using distances in a normal

autoencoder representation often does not result in meaningful behavior.

2.4.3 Improving Sample Efficiency with Latent Goal Relabeling

To further enable sample-efficient learning in the real world, we use the VAE to relabel goals. Note that we can optimize Equation equation 35 using any valid (s, a, s', g, r) tuple. If we could artificially generate these tuples, then we could train our entire RL algorithm without collecting any data. However, we do not know the system dynamics, and therefore have to sample transitions (s, a, s') by interacting with the world. However, we have the freedom to relabel the goal and reward synthetically. In particular, if we have a mechanism for generating goals and computing rewards, then given (s, a, s') , we can generate a new goal g and new reward $r(s, a, s', g)$ to produce a new tuple (s, a, s', g, r) . By artificially generating and recomputing rewards, we can convert a single (s, a, s') transition into potentially infinitely many valid training datums.

For image-based tasks, this procedure would require generating goal images, an onerous task on its own. However, our reinforcement learning algorithm operates directly in the latent space for goals and rewards. So rather than generating goals g , we generate latent goals z_g by sampling from the VAE prior $p(z)$. We then recompute rewards using Equation equation 3. By retroactively relabeling the goals and rewards, we obtain much more data to train our value function. This sampling procedure is made possible by our use of a latent variable model, which is explicitly trained so that sampling from the latent distribution is straightforward.

Retroactively generating goals is also explored in tabular domains by L P Kaelbling (1993) and in continuous domains by Andrychowicz et al. (2017b) using hindsight experience replay (HER). However, HER is limited to sampling goals seen along a trajectory, which greatly limits the number and diversity of goals with which one can relabel a given transition. Our final method uses a mixture of the two strategies: half of the goals are generated from the prior and half from goals seen along the trajectory. We show in Section 8.6 that relabeling the goal with samples from the VAE prior results in significantly better sample-efficiency.

2.4.4 Automated Goal-Generation for Exploration

If we do not know which particular goals will be provided at test time, we would like our RL agent to carry out a self-supervised “practice” phase during training, where the algorithm proposes its own goals, and then practices how to reach them. Since the VAE

Algorithm 1 RIG: Reinforcement learning with imagined goals

Require: VAE encoder q_ϕ , VAE decoder p_ψ , policy π_θ , goal-conditioned value function Q_w .

1: Collect $\mathcal{D} = \{s^{(i)}\}$ using exploration policy.
2: Train β -VAE on \mathcal{D} by optimizing equation 2.
3: **for** $n = 0, \dots, N - 1$ episodes **do**
4: Sample latent goal from prior $z_g \sim p(z)$.
5: Sample initial state $s_0 \sim E$.
6: **for** $t = 0, \dots, H - 1$ steps **do**
7: Get action $a_t = \pi_\theta(e(s_t), z_g) + \text{noise}$.
8: Get next state $s_{t+1} \sim p(\cdot | s_t, a_t)$.
9: Store (s_t, a_t, s_{t+1}, z_g) into replay buffer \mathcal{R} .
10: Sample transition $(s, a, s', z_g) \sim \mathcal{R}$.
11: Encode $z' = e(s')$.
12: (Probability 0.5) replace z_g with $z'_g \sim p(z)$.
13: Compute new reward $r = -\|z' - z_g\|$.
14: Minimize equation 35 using (z, a, z', z_g, r) .
15: **end for**
16: Fine-tune β -VAE every K episodes on mixture of \mathcal{D} and \mathcal{R} .
17: **end for=**

prior represents a distribution over latent goals and state observations, we again sample from this distribution to obtain plausible goals. After sampling a goal latent from the prior $z_g \sim p(z)$, we give this to our policy $\pi(z, z_g)$ to collect data.

2.4.5 Algorithm Summary

We call the complete algorithm reinforcement learning with imagined goals (RIG) and summarize it in Algorithm 3. We first collect data with a simple exploration policy, though any exploration strategy could be used for this stage, including off-the-shelf exploration bonuses (Pathak et al., 2017; M. Bellemare et al., 2016) or unsupervised reinforcement learning methods (Eysenbach et al., 2018; Florensa et al., 2017). Then, we train a VAE latent variable model on state observations and finetune it over the course of training. We use this latent variable model for multiple purposes: We sample a latent goal z_g from the model and condition the policy on this goal. We embed all states and goals using the model’s encoder. When we train our goal-conditioned value function, we resample goals from the prior and compute rewards in the latent space using Equation equation 3. Any RL algorithm that trains Q-functions could be used, and we use TD3 (Fujimoto et al., 2018b) in our implementation.

2.5 EXPERIMENTS

Our experiments address the following questions:

1. How does our method compare to prior model-free RL algorithms in terms of sample efficiency and performance, when learning continuous control tasks from

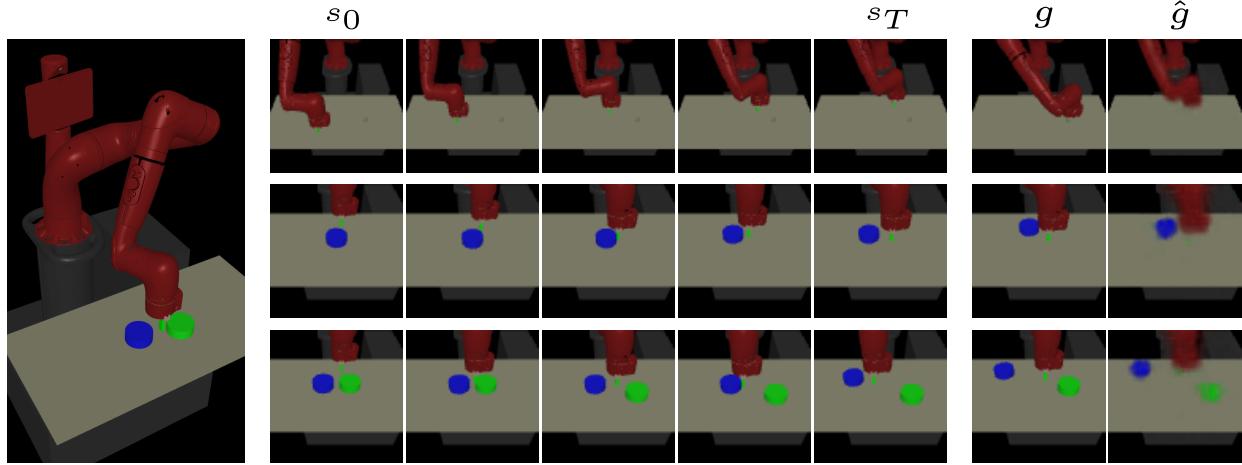


Figure 2: (Left) The simulated environment. (Right) Test rollouts from our learned policy on the three simulated environments. Each row is one rollout. The middle shows the goal image g and its VAE reconstruction \hat{g} . The right columns shows frames from the trajectory to reach the given goal.

images?

2. How critical is each component of our algorithm for efficient learning?
3. Does our method work on tasks where the state space cannot be easily specified ahead of time, such as tasks that require interaction with variable numbers of objects?
4. Can our method scale to real world vision-based robotic control tasks?

For the first two questions, we evaluate our method against a number of prior algorithms and ablated versions of our approach on a suite of simulated and real-world tasks: *Visual Reacher*: a MuJoCo (Todorov et al., 2012) environment with a 7-dof Sawyer arm reaching goal positions. The arm is shown the left of Figure 2. The end-effector (EE) is constrained to a 2-dimensional rectangle parallel to a table. The action controls EE velocity within a maximum velocity. *Visual Pusher*: a MuJoCo environment with a 7-dof Sawyer arm and a small puck on a table that the arm must push to a target push. *Visual Multi-Object Pusher*: a copy of the Visual Pusher environment with two pucks. Detailed descriptions of the environments are provided in the Supplementary Material.

Solving these tasks directly from images poses a challenge since the controller must learn both perception and control. The evaluation metric is the distance of objects (including the arm) to their respective goals. To evaluate our policy, we set the environment to a sampled goal position, capture an image, and encode the image to use as the goal. Although we use the ground-truth positions for evaluation, **we do not use the ground-**

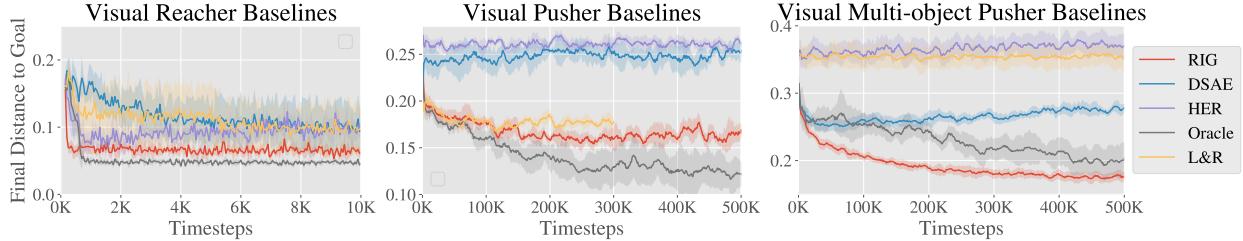


Figure 3: Simulation results, final distance to goal vs simulation steps². RIG (red) consistently outperforms the baselines, except for the oracle which uses ground truth object state for observations and rewards. On the hardest task, only our method and the oracle discover viable solutions.

truth positions for training the policies. The only inputs from the environment that our algorithm receives are the image observations. For Visual Reacher, we pretrained the VAE with 100 images. For other tasks, we used 10,000 images.

We compare our method with the following prior works. *L&R*: Lange and Riedmiller ([Lange and M. A. Riedmiller, 2010](#)) trains an autoencoder to handle images. *DSAE*: Deep spatial autoencoders ([Finn et al., 2016b](#)) learns a spatial autoencoder and uses guided policy search ([Levine et al., 2016a](#)) to achieve a single goal image. *HER*: Hind-sight experience replay ([Andrychowicz et al., 2017b](#)) utilizes a sparse reward signal and relabeling trajectories with achieved goals. *Oracle*: RL with direct access to state information for observations and rewards.

To our knowledge, no prior work demonstrates policies that can reach a variety of goal images without access to a true-state reward function, and so we needed to make modifications to make the comparisons feasible. L&R assumes a reward function from the environment. Since we have no state-based reward function, we specify the reward function as distance in the autoencoder latent space. HER does not embed inputs into a latent space but instead operates directly on the input, so we use pixel-wise mean squared error (MSE) as the metric. DSAE is trained only for a single goal, so we allow the method to generalize to a variety of test goal images by using a goal-conditioned Q-function. To make the implementations comparable, we use the same off-policy algorithm to train L&R, HER, and our method (TD3 ([Fujimoto et al., 2018b](#))). Unlike our method, prior methods do not specify how to select goals during training, so we favorably give them real images as goals for rollouts, sampled from the same distribution that we use to test.

We see in Figure 3 that our method can efficiently learn policies from visual inputs to perform simulated reaching and pushing, without access to the object state. Our ap-

² In all our simulation results, each plot shows a 95% confidence interval of the mean across 5 seeds.

proach substantially outperforms the prior methods, for which the use of image goals and observations poses a major challenge. HER struggles because pixel-wise MSE is hard to optimize. Our latent-space rewards are much better shaped and allow us to learn more complex tasks. Finally, our method is close to the state-based “oracle” method in terms of sample efficiency and performance, without having any access to object state. Notably, in the multi-object environment, our method actually outperforms the oracle, likely because the state-based reward contains local minima. Overall, these result show that our method is capable of handling raw image observations much more effectively than previously proposed goal-conditioned RL methods. Next, we perform ablations to evaluate our contributions in isolation. Results on Visual Pusher are shown but see the Supplementary Material (section ??) for experiments on all three simulated environments.

REWARD SPECIFICATION COMPARISON We evaluate how effective distances in the VAE latent space are for the Visual Pusher task. We keep our method the same, and only change the reward function that we use to train the goal-conditioned valued function. We include the following methods for comparison: *Latent Distance*, which is the reward used in RIG, i.e. $A = I$ in Equation equation 3; *Log Probability*, which uses the Mahalanobis distance in Equation equation 3, where A is the precision matrix of the encoder; and *Pixel MSE*, which computes mean-squared error (MSE) between state and goal in pixel space.³ In Figure 4, we see that latent distance significantly outperforms the log probability. We suspect that small variances of the VAE encoder results in drastically large rewards, making the learning more difficult. We also see that latent distances results in faster learning when compared to pixel MSE.

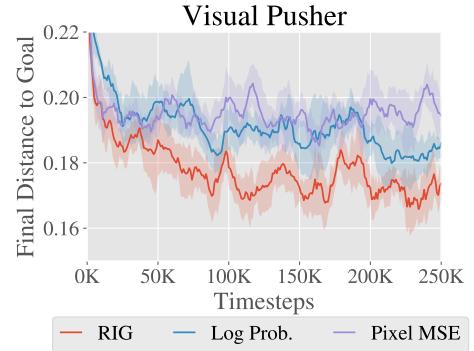


Figure 4: Reward type ablation results. RIG (red), which uses latent Euclidean distance, outperforms the other methods.

³ To compute the pixel MSE for a sampled latent goal, we decode the goal latent using the VAE decoder, p_ψ , to generate the corresponding goal image.

RELABELING STRATEGY COMPARISON As described in section 2.4.3, our method uses a novel goal relabeling method based on sampling from the generative model. To isolate how much our new goal relabeling method contributes to our algorithm, we vary the resampling strategy while fixing other components of our algorithm. The resampling strategies that we consider are: *Future*, relabeling the goal for a transition by sampling uniformly from future states in the trajectory as done in Andrychowicz et al. (2017b); *VAE*, sampling goals from the VAE only; *RIG*, relabeling goals with probability 0.5 from the VAE and probability 0.5 using the future strategy; and *None*, no relabeling. In Figure 5, we see that sampling from the VAE and Future is significantly better than not relabeling at all. In RIG, we use an equal mixture of the VAE and Future sampling strategies, which performs best by a large margin. Appendix section ?? contains results on all simulated environments, and section ?? considers relabeling strategies with a known goal distribution.

LEARNING WITH VARIABLE NUMBERS OF OBJECTS
A major advantage of working directly from pixels is that the policy input can easily represent combinatorial structure in the environment, which would be difficult to encode into a fixed-length state vector even if a perfect perception system were available. For example, if a robot has to interact with different combinations and numbers of objects, picking a single MDP state representation would be challenging, even with access to object poses. By directly processing images for both the state and the goal, no modification is needed to handle the combinatorial structure: the number of pixels always remains the same, regardless of how many objects are in the scene.

We demonstrate that our method can handle this difficult scenario by evaluating on a task where the environment, based on the Visual Multi-Object Pusher, randomly contains zero, one, or two objects in each episode during testing. During training, each episode still always starts with both objects in the scene, so the experiments tests whether a trained policy can handle variable numbers of objects at test time. Figure 6 shows that our method can learn to solve this task successfully, without decrease in performance

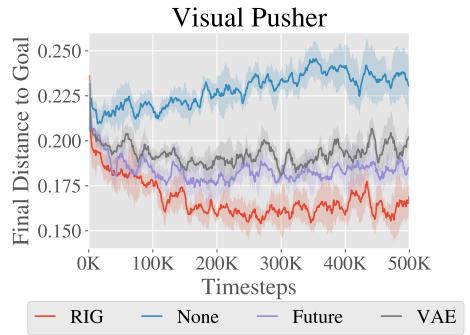


Figure 5: Relabeling ablation results. RIG (red), which uses a mixture of VAE and HER, outperforms the other methods.

Figure 5: Relabeling ablation results. RIG (red), which uses a mixture of VAE and HER, outperforms the other methods.

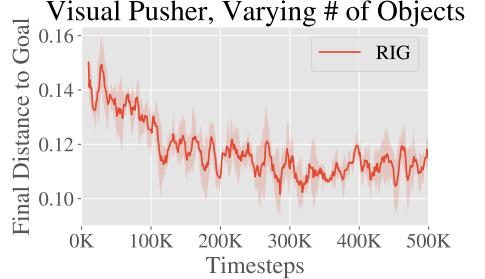


Figure 6: Training curve for learning with varying number of objects.

from the base setting where both objects are present (in Figure 3). Developing and demonstrating algorithms that solve tasks with varied underlying structure is an important step toward creating autonomous agents that can handle the diversity of tasks present “in the wild.”

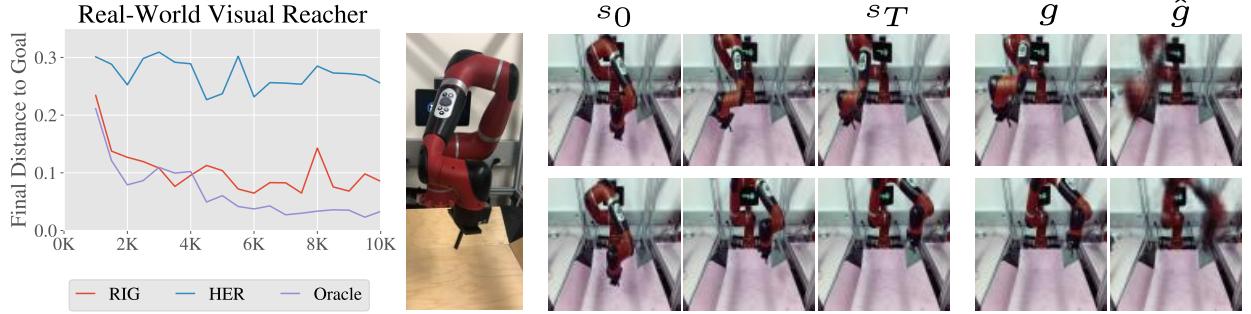


Figure 7: (Left) Our method compared to the HER baseline and oracle on a real-world visual reaching task. (Middle) Our robot setup is pictured. (Right) Test rollouts of our learned policy.

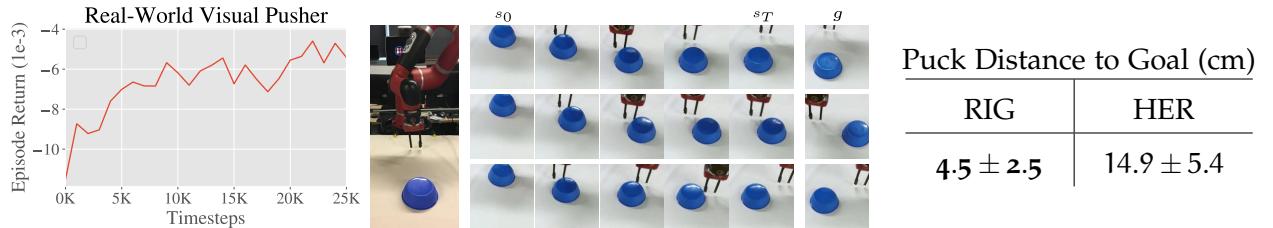


Figure 8: (Left) The learning curve for real-world pushing. (Middle) Our robot pushing setup is pictured, with frames from test rollouts of our learned policy. (Right) Our method compared to the HER baseline on the real-world visual pushing task. We evaluated the performance of each method by manually measuring the distance between the goal position of the puck and final position of the puck for 15 test rollouts, reporting mean and standard deviation.

2.5.1 Visual RL with Physical Robots

RIG is a practical and straightforward algorithm to apply to real physical systems: the efficiency of off-policy learning with goal relabeling makes training times manageable, while the use of image-based rewards through the learned representation frees us from the burden of manually design reward functions, which itself can require hand-engineered perception systems (Rusu et al., 2017). We trained policies for visual reaching

and pushing on a real-world Sawyer robotic arm, shown in Figure 7. The control setup matches Visual Reacher and Visual Pusher respectively, meaning that **the only input from the environment consists of camera images**.

We see in Figure 7 that our method is applicable to real-world robotic tasks, almost matching the state-based oracle method and far exceeding the baseline method on the reaching task. Our method needs just 10,000 samples or about an hour of real-world interaction time to solve visual reaching.

Real-world pushing results are shown in Figure 20. To solve visual pusher, which is more visually complicated and requires reasoning about the contact between the arm and object, our method requires about 25,000 samples, which is still a reasonable amount of real-world training time. Note that unlike previous results, we do not have access to the true puck position during training so for the learning curve we report test episode returns on the VAE latent distance reward. We see RIG making steady progress at optimizing the latent distance as learning proceeds.

2.6 DISCUSSION AND FUTURE WORK

In this paper, we present a new RL algorithm that can efficiently solve goal-conditioned, vision-based tasks without access to any ground truth state or reward functions. Our method trains a generative model that is used for multiple purposes: we embed the state and goals using the encoder; we sample from the prior to generate goals for exploration; we also sample latents to retroactively relabel goals and rewards; and we use distances in the latent space for rewards to train a goal-conditioned value function. We show that these components culminate in a sample efficient algorithm that works directly from vision. As a result, we are able to apply our method to a variety of simulated visual tasks, including a variable-object task that cannot be easily represented with a fixed length vector, as well as real world robotic tasks. Algorithms that can learn in the real world and directly use raw images can allow a single policy to solve a large and diverse set of tasks, even when these tasks require distinct internal representations.

The method we presented can be extended in a number of ways. First, an exciting line of future work would be to combine our method with existing work on exploration and intrinsic motivation. In particular, our method already provides a natural mechanism for autonomously generating goals by sampling from the prior. Modifying this procedure to not only be goal-oriented but also, e.g., be information seeking or uncertainty aware could provide better and safer exploration. Second, since our method operates directly from images, a single policy could potentially solve a large diverse set of visual

tasks, even if those tasks have different underlying state representations. Combining these ideas with methods from multitask learning and meta-learning is a promising path to creating general-purpose agents that can continuously and efficiently acquire skills. Lastly, while RIG uses goal images, extending the method to allow goals specified by demonstrations or more abstract representations such as language would enable our system to be much more flexible in interfacing with humans and therefore more practical.

3

SKEWFIT: STATE-COVERING SELF-SUPERVISED REINFORCEMENT LEARNING

3.1 INTRODUCTION

Reinforcement learning (RL) provides an appealing formalism for automated learning of behavioral skills, but separately learning every potentially useful skill becomes prohibitively time consuming, both in terms of the experience required for the agent and the effort required for the user to design reward functions for each behavior. What if we could instead design an unsupervised RL algorithm that automatically explores the environment and iteratively distills this experience into general-purpose policies that can accomplish new user-specified tasks at test time?

In the absence of any prior knowledge, an effective exploration scheme is one that visits as many states as possible, allowing a policy to autonomously prepare for user-specified tasks that it might see at test time. We can formalize the problem of visiting as many states as possible as one of maximizing the *state entropy* $\mathcal{H}(\mathbf{S})$ under the current policy.² Unfortunately, optimizing this objective alone does not result in a policy that can solve new tasks: it only knows how to maximize state entropy. In other words, to develop principled unsupervised RL algorithms that result in useful policies, maximizing $\mathcal{H}(\mathbf{S})$ is not enough. We need a mechanism that allows us to reuse the resulting policy to achieve new tasks at test-time.

We argue that this can be accomplished by performing *goal-directed exploration*: a policy should autonomously visit as many states as possible, but after autonomous exploration, a user should be able to reuse this policy by giving it a goal \mathbf{G} that corresponds to a state that it must reach. While not all test-time tasks can be expressed as reaching a goal state, a wide range of tasks can be represented in this way. Mathematically, the

² We consider the distribution over terminal states in a finite horizon task and believe this work can be extended to infinite horizon stationary distributions.

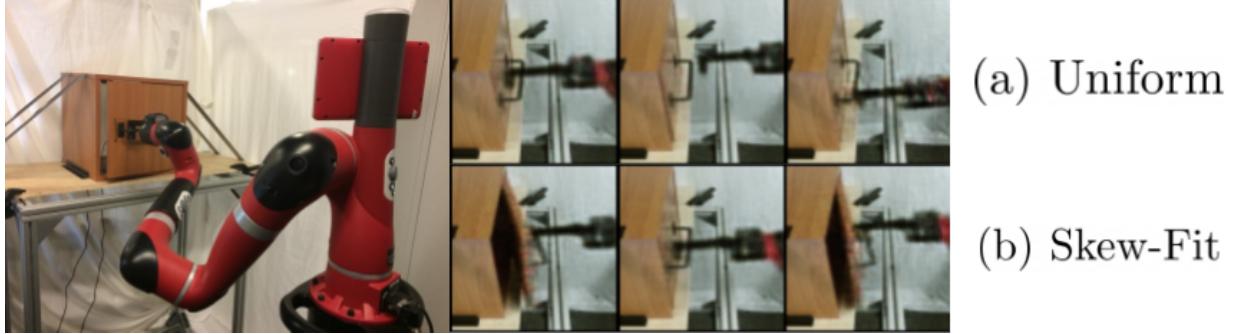


Figure 9: Left: Robot learning to open a door with Advantage Weighted Actor Critic, without any task reward. Right: Samples from a goal distribution when using (a) uniform and (b) Advantage Weighted Actor Critic sampling. When used as goals, the diverse samples from Advantage Weighted Actor Critic encourage the robot to practice opening the door more frequently.

goal-conditioned policy should minimize the conditional entropy over the states given a goal, $\mathcal{H}(\mathbf{S} | \mathbf{G})$, so that there is little uncertainty over its state given a commanded goal. This objective provides us with a principled way to train a policy to explore all states (maximize $\mathcal{H}(\mathbf{S})$) such that the state that is reached can be determined by commanding goals (minimize $\mathcal{H}(\mathbf{S} | \mathbf{G})$).

Directly optimizing this objective is in general intractable, since it requires optimizing the entropy of the marginal state distribution, $\mathcal{H}(\mathbf{S})$. However, we can sidestep this issue by noting that the objective is the mutual information between the state and the goal, $I(\mathbf{S}; \mathbf{G})$, which can be written as:

$$\mathcal{H}(\mathbf{S}) - \mathcal{H}(\mathbf{S}|\mathbf{G}) = I(\mathbf{S}; \mathbf{G}) = \mathcal{H}(\mathbf{G}) - \mathcal{H}(\mathbf{G}|\mathbf{S}). \quad (4)$$

Equation 4 thus gives an equivalent objective for an unsupervised RL algorithm: the agent should set diverse goals, maximizing $\mathcal{H}(\mathbf{G})$, and learn how to reach them, minimizing $\mathcal{H}(\mathbf{G} | \mathbf{S})$.

While learning to reach goals is the typical objective studied in goal-conditioned RL ([L P Kaelbling, 1993; Andrychowicz et al., 2017b](#)), setting goals that have maximum diversity is crucial for effectively learning to reach all possible states. Acquiring such a maximum-entropy goal distribution is challenging in environments with complex, high-dimensional state spaces, where even knowing which states are valid presents a major challenge. For example, in image-based domains, a uniform goal distribution requires sampling uniformly from the set of realistic images, which in general is unknown a

priori.

Our paper makes the following contributions. First, we propose a principled objective for unsupervised RL, based on [Equation 4](#). While a number of prior works ignore the $\mathcal{H}(\mathbf{G})$ term, we argue that jointly optimizing the entire quantity is needed to develop effective exploration. Second, we present a general algorithm called Advantage Weighted Actor Critic and prove that under regularity conditions Advantage Weighted Actor Critic learns a sequence of generative models that converges to a uniform distribution over the goal space, even when the set of valid states is unknown (e.g., as in the case of images). Third, we describe a concrete implementation of Advantage Weighted Actor Critic and empirically demonstrate that this method achieves state of the art results compared to a large number of prior methods for goal reaching with visually indicated goals, including a real-world manipulation task, which requires a robot to learn to open a door from scratch in about five hours, directly from images, and without any manually-designed reward function.

3.2 PROBLEM FORMULATION

To ensure that an unsupervised reinforcement learning agent learns to reach all possible states in a controllable way, we maximize the mutual information between the state \mathbf{S} and the goal \mathbf{G} , $I(\mathbf{S}; \mathbf{G})$, as stated in [Equation 4](#). This section discusses how to optimize [Equation 4](#) by splitting the optimization into two parts: minimizing $\mathcal{H}(\mathbf{G} | \mathbf{S})$ and maximizing $\mathcal{H}(\mathbf{G})$.

3.2.1 Minimizing $\mathcal{H}(\mathbf{G} | \mathbf{S})$: Goal-Conditioned Reinforcement Learning

Standard RL considers a Markov decision process (MDP), which has a state space \mathcal{S} , action space \mathcal{A} , and unknown dynamics $\rho(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t) : \mathcal{S} \times \mathcal{S} \times \mathcal{A} \mapsto [0, +\infty)$. Goal-conditioned RL also includes a goal space \mathcal{G} . For simplicity, we will assume in our derivation that the goal space matches the state space, such that $\mathcal{G} = \mathcal{S}$, though the approach extends trivially to the case where \mathcal{G} is a hand-specified subset of \mathcal{S} , such as the global XY position of a robot. A goal-conditioned policy $\pi(\mathbf{a} | \mathbf{s}, \mathbf{g})$ maps a state $\mathbf{s} \in \mathcal{S}$ and goal $\mathbf{g} \in \mathcal{S}$ to a distribution over actions $\mathbf{a} \in \mathcal{A}$, and its objective is to reach the goal, i.e., to make the current state equal to the goal.

Goal-reaching can be formulated as minimizing $\mathcal{H}(\mathbf{G} | \mathbf{S})$, and many practical goal-reaching algorithms ([L P Kaelbling, 1993](#); [Lillicrap et al., 2016](#); [Schaul et al., 2015b](#); [Andrychowicz et al., 2017b](#); [A. Nair et al., 2018b](#); [V. Pong et al., 2018](#); [Florensa et al.,](#)

2018a) can be viewed as approximations to this objective by observing that the optimal goal-conditioned policy will deterministically reach the goal, resulting in a conditional entropy of zero: $\mathcal{H}(\mathbf{G} | \mathbf{S}) = 0$. See ?? for more details. Our method may thus be used in conjunction with any of these prior goal-conditioned RL methods in order to jointly minimize $\mathcal{H}(\mathbf{G} | \mathbf{S})$ and maximize $\mathcal{H}(\mathbf{G})$.

3.2.2 Maximizing $\mathcal{H}(\mathbf{G})$: Setting Diverse Goals

We now turn to the problem of setting diverse goals or, mathematically, maximizing the entropy of the goal distribution $\mathcal{H}(\mathbf{G})$. Let U_S be the uniform distribution over S , where we assume S has finite volume so that the uniform distribution is well-defined. Let q_ϕ^G be the goal distribution from which goals \mathbf{G} are sampled, parameterized by ϕ . Our goal is to maximize the entropy of q_ϕ^G , which we write as $\mathcal{H}(\mathbf{G})$. Since the maximum entropy distribution over S is the uniform distribution U_S , maximizing $\mathcal{H}(\mathbf{G})$ may seem as simple as choosing the uniform distribution to be our goal distribution: $q_\phi^G = U_S$. However, this requires knowing the uniform distribution over valid states, which may be difficult to obtain when S is a subset of \mathbb{R}^n , for some n . For example, if the states correspond to images viewed through a robot's camera, S corresponds to the (unknown) set of valid images of the robot's environment, while \mathbb{R}^n corresponds to all possible arrays of pixel values of a particular size. In such environments, sampling from the uniform distribution \mathbb{R}^n is unlikely to correspond to a valid image of the real world. Sampling uniformly from S would require knowing the set of all possible valid images, which we assume the agent does not know when starting to explore the environment.

While we cannot sample arbitrary states from S , we can sample states by performing goal-directed exploration. To derive and analyze our method, we introduce a simple model of this process: a goal $\mathbf{G} \sim q_\phi^G$ is sampled from the goal distribution q_ϕ^G , and then the goal-conditioned policy π attempts to achieve this goal, which results in a distribution of terminal states $\mathbf{S} \in S$. We abstract this entire process by writing the resulting marginal distribution over S as $p_\phi^S(S) \triangleq \int_S q_\phi^G(G)p(S | G)dG$, where the subscript ϕ indicates that the marginal p_ϕ^S depends indirectly on q_ϕ^G via the goal-conditioned policy π . We assume that p_ϕ^S has full support, which can be accomplished with an epsilon-greedy goal reaching policy in a communicating MDP. We also assume that the entropy of the resulting state distribution $\mathcal{H}(p_\phi^S)$ is no less than the entropy of the goal distribution $\mathcal{H}(q_\phi^G)$. Without this assumption, a policy could ignore the goal and stay in a single state,

no matter how diverse and realistic the goals are.³ This simplified model allows us to analyze the behavior of our goal-setting scheme separately from any specific goal-reaching algorithm. We will however show in Section 8.6 that we can instantiate this approach into a practical algorithm that jointly learns the goal-reaching policy. In summary, our goal is to acquire a maximum-entropy goal distribution q_ϕ^G over valid states \mathcal{S} , while only having access to state samples from p_ϕ^S .

3.3 ADVANTAGE WEIGHTED ACTOR CRITIC: LEARNING A MAXIMUM ENTROPY GOAL DISTRIBUTION

Our method, Advantage Weighted Actor Critic, learns a maximum entropy goal distribution q_ϕ^G using samples collected from a goal-conditioned policy. We analyze the algorithm and show that Advantage Weighted Actor Critic maximizes the goal distribution entropy, and present a practical instantiation for unsupervised deep RL.

3.3.1 *Advantage Weighted Actor Critic Algorithm*

To learn a uniform distribution over *valid* goal states, we present a method that iteratively increases the entropy of a generative model q_ϕ^G . In particular, given a generative model $q_{\phi_t}^G$ at iteration t , we want to train a new generative model, $q_{\phi_{t+1}}^G$ that has higher entropy.

While we do not know the set of valid states \mathcal{S} , we could sample states $s_n \sim p_{\phi_t}^S$ using the goal-conditioned policy, and use the samples to train $q_{\phi_{t+1}}^G$. However, there is no guarantee that this would increase the entropy of $q_{\phi_{t+1}}^G$.

The intuition behind our method is simple: rather than fitting a generative model to these samples s_n , we *skew* the samples so that rarely visited states are given more weight. See Figure 10 for a visualization of this process. How should we skew the samples if we want to maximize the entropy of $q_{\phi_{t+1}}^G$? If we had access to the density of each state, $p_{\phi_t}^S(S)$, then we could simply weight each state by $1/p_{\phi_t}^S(S)$. We could then perform maximum likelihood estimation (MLE) for the uniform distribution by using the follow-

³ Note that this assumption does **not** require that the entropy of p_ϕ^S is strictly larger than the entropy of the goal distribution, q_ϕ^G .

ing importance sampling (IS) loss to train ϕ_{t+1} :

$$\begin{aligned}\mathcal{L}(\phi) &= \mathbb{E}_{\mathbf{S} \sim U_S} [\log q_\phi^G(\mathbf{S})] \\ &= \mathbb{E}_{\mathbf{S} \sim p_{\phi_t}^S} \left[\frac{U_S(\mathbf{S})}{p_{\phi_t}^S(\mathbf{S})} \log q_\phi^G(\mathbf{S}) \right] \\ &\propto \mathbb{E}_{\mathbf{S} \sim p_{\phi_t}^S} \left[\frac{1}{p_{\phi_t}^S(\mathbf{S})} \log q_\phi^G(\mathbf{S}) \right]\end{aligned}$$

where we use the fact that the uniform distribution $U_S(\mathbf{S})$ has constant density for all states in S . However, computing this density $p_{\phi_t}^S(\mathbf{S})$ requires marginalizing out the MDP dynamics, which requires an accurate model of both the dynamics and the goal-conditioned policy.

We avoid needing to model the entire MDP process by approximating $p_{\phi_t}^S(\mathbf{S})$ with our previous learned generative model: $p_{\phi_t}^S(\mathbf{S}) \approx q_{\phi_t}^G(\mathbf{S})$. We therefore weight each state by the following weight function

$$w_{t,\alpha}(\mathbf{S}) \triangleq q_{\phi_t}^G(\mathbf{S})^\alpha, \quad \alpha < 0. \quad (5)$$

where α is a hyperparameter that controls how heavily we weight each state. If our approximation $q_{\phi_t}^G$ is exact, we can choose $\alpha = -1$ and recover the exact IS procedure described above. If $\alpha = 0$, then this skew step has no effect. By choosing intermediate values of α , we trade off the reliability of our estimate $q_{\phi_t}^G(\mathbf{S})$ with the speed at which we want to increase the goal distribution entropy.

VARIANCE REDUCTION As described, this procedure relies on IS, which can have high variance, particularly if $q_{\phi_t}^G(\mathbf{S}) \approx 0$. We therefore choose a class of generative models where the probabilities are prevented from collapsing to zero, as we will describe in [Section 3.4](#) where we provide generative model details. To further reduce the variance, we train $q_{\phi_{t+1}}^G$ with sampling importance resampling (SIR) ([Rubin, 1988](#)) rather than IS. Rather than sampling from $p_{\phi_t}^S$ and weighting the update from each sample by $w_{t,\alpha}$, SIR

explicitly defines a skewed empirical distribution as

$$p_{\text{skewed}_t}(s) \triangleq \frac{1}{Z_\alpha} w_{t,\alpha}(s) \delta(s \in \{s_n\}_{n=1}^N) \quad (6)$$

$$Z_\alpha = \sum_{n=1}^N w_{t,\alpha}(s_n), \quad s_n \stackrel{\text{iid}}{\sim} p_{\phi_t}^S,$$

where δ is the indicator function and Z_α is the normalizing coefficient. We note that computing Z_α adds little computational overhead, since all of the weights already need to be computed. We then fit the generative model at the next iteration $q_{\phi_{t+1}}^G$ to p_{skewed_t} using standard MLE. We found that using SIR resulted in significantly lower variance than IS. See ?? for this comparison.

GOAL SAMPLING ALTERNATIVE Because $q_{\phi_{t+1}}^G \approx p_{\text{skewed}_t}$, at iteration $t + 1$, one can sample goals from either $q_{\phi_{t+1}}^G$ or p_{skewed_t} . Sampling goals from p_{skewed_t} may be preferred if sampling from the learned generative model $q_{\phi_{t+1}}^G$ is computationally or otherwise challenging. In either case, one still needs to train the generative model $q_{\phi_t}^G$ to create p_{skewed_t} . In our experiments, we found that both methods perform well.

SUMMARY Overall, Advantage Weighted Actor Critic collects states from the environment and resamples each state in proportion to Equation 5 so that low-density states are resampled more often. Advantage Weighted Actor Critic is shown in Figure 10 and summarized in Algorithm 6. We now discuss conditions under which Advantage Weighted Actor Critic converges to the uniform distribution.

Algorithm 2 Advantage Weighted Actor Critic

- 1: **for** Iteration $t = 1, 2, \dots$ **do**
 - 2: Collect N states $\{s_n\}_{n=1}^N$ by sampling goals from $q_{\phi_t}^G$ (or $p_{\text{skewed}_{t-1}}$) and running goal-conditioned policy.
 - 3: Construct skewed distribution p_{skewed_t} (Equation 5 and Equation 6).
 - 4: Fit $q_{\phi_{t+1}}^G$ to skewed distribution p_{skewed_t} using MLE.
 - 5: **end for**=o
-

3.3.2 Advantage Weighted Actor Critic Analysis

This section provides conditions under which $q_{\phi_t}^G$ converges in the limit to the uniform distribution over the state space \mathcal{S} . We consider the case where $N \rightarrow \infty$, which allows us to study the limit behavior of the goal distribution p_{skewed_t} . Our most general result is stated as follows:

Lemma 1. *Let \mathcal{S} be a compact set. Define the set of distributions $\Omega = \{p : \text{support}(p) \subseteq \mathcal{S}\}$. Let $\mathcal{F} : \Omega \mapsto \Omega$ be continuous with respect to the pseudometric $d_{\mathcal{H}}(p, q) \triangleq |\mathcal{H}(p) - \mathcal{H}(q)|$ and $\mathcal{H}(\mathcal{F}(p)) \geq \mathcal{H}(p)$ with equality if and only if p is the uniform probability distribution on \mathcal{S} , denoted as $U_{\mathcal{S}}$. Define the sequence of distributions $P = (p_1, p_2, \dots)$ by starting with any $p_1 \in \Omega$ and recursively defining $p_{t+1} = \mathcal{F}(p_t)$. The sequence P converges to $U_{\mathcal{S}}$ with respect to $d_{\mathcal{H}}$. In other words, $\lim_{t \rightarrow \infty} |\mathcal{H}(p_t) - \mathcal{H}(U_{\mathcal{S}})| \rightarrow 0$.*

Proof. See Appendix Section ??.

We will apply Lemma 1 to be the map from p_{skewed_t} to $p_{\text{skewed}_{t+1}}$ to show that p_{skewed_t} converges to $U_{\mathcal{S}}$. If we assume that the goal-conditioned policy and generative model learning procedure are well behaved (i.e., the maps from $q_{\phi_t}^G$ to $p_{\phi_t}^S$ and from p_{skewed_t} to $q_{\phi_{t+1}}^G$ are continuous), then to apply Lemma 1, we only need to show that $\mathcal{H}(p_{\text{skewed}_t}) \geq \mathcal{H}(p_{\phi_t}^S)$ with equality if and only if $p_{\phi_t}^S = U_{\mathcal{S}}$. For the simple case when $q_{\phi_t}^G = p_{\phi_t}^S$ identically at each iteration, we prove the convergence of Advantage Weighted Actor Critic true for any value of $\alpha \in [-1, 0)$ in ???. However, in practice, $q_{\phi_t}^G$ only approximates $p_{\phi_t}^S$. To address this more realistic situation, we prove the following result:

Lemma 2. *Given two distribution $p_{\phi_t}^S$ and $q_{\phi_t}^G$ where $p_{\phi_t}^S \ll q_{\phi_t}^G$ and*

$$\text{Cov}_{\mathbf{S} \sim p_{\phi_t}^S} [\log p_{\phi_t}^S(\mathbf{S}), \log q_{\phi_t}^G(\mathbf{S})] > 0, \quad (7)$$

define the p_{skewed_t} as in Equation 6 and take $N \rightarrow \infty$. Let $\mathcal{H}_\alpha(\alpha)$ be the entropy of p_{skewed_t} for a fixed α . Then there exists a constant $a < 0$ such that for all $\alpha \in [a, 0)$,

$$\mathcal{H}(p_{\text{skewed}_t}) = \mathcal{H}_\alpha(\alpha) > \mathcal{H}(p_{\phi_t}^S).$$

a $p \ll q$ means that p is absolutely continuous with respect to q , i.e. $p(s) = 0 \implies q(s) = 0$.

Proof. See Appendix Section ??.

This lemma tells us that our generative model $q_{\phi_t}^G$ does not need to exactly fit the sampled states. Rather, we merely need the log densities of $q_{\phi_t}^G$ and $p_{\phi_t}^S$ to be correlated, which we expect to happen frequently with an accurate goal-conditioned policy, since $p_{\phi_t}^S$ is the set of states seen when trying to reach goals from $q_{\phi_t}^G$. In this case, if we choose negative values of α that are small enough, then the entropy of p_{skewed_t} will be higher than that of $p_{\phi_t}^S$. Empirically, we found that α values as low as $\alpha = -1$ performed well.

In summary, p_{skewed_t} converges to U_S under certain assumptions. Since we train each generative model $q_{\phi_{t+1}}^G$ by fitting it to p_{skewed_t} with MLE, $q_{\phi_t}^G$ will also converge to U_S .

3.4 TRAINING GOAL-CONDITIONED POLICIES WITH ADVANTAGE WEIGHTED ACTOR CRITIC

Thus far, we have presented Advantage Weighted Actor Critic assuming that we have access to a goal-reaching policy, allowing us to separately analyze how we can maximize $\mathcal{H}(G)$. However, in practice we do not have access to such a policy, and this section discusses how we concurrently train a goal-reaching policy.

Maximizing $I(S; G)$ can be done by simultaneously performing Advantage Weighted Actor Critic and training a goal-conditioned policy to minimize $\mathcal{H}(G | S)$, or, equivalently, maximize $-\mathcal{H}(G | S)$. Maximizing $-\mathcal{H}(G | S)$ requires computing the density $\log p(G | S)$, which may be difficult to compute without strong modeling assumptions. However, for any distribution q , the following lower bound on $-\mathcal{H}(G | S)$:

$$\begin{aligned} -\mathcal{H}(G | S) &= \mathbb{E}_{(G,S) \sim q} [\log q(G | S)] + p q \\ &\geq \mathbb{E}_{(G,S) \sim q} [\log q(G | S)], \end{aligned}$$

where D_{KL} denotes Kullback–Leibler divergence as discussed by Barber and Agakov (2004). Thus, to minimize $\mathcal{H}(G | S)$, we train a policy to maximize the reward

$$r(S, G) = \log q(G | S).$$

The RL algorithm we use is reinforcement learning with imagined goals (RIG) (A. Nair et al., 2018b), though in principle any goal-conditioned method could be used. RIG

is an efficient off-policy goal-conditioned method that solves vision-based RL problems in a learned latent space. In particular, RIG fits a β -VAE (Higgins et al., 2017a) and uses it to encode observations and goals into a latent space, which it uses as the state representation. RIG also uses the β -VAE to compute rewards, $\log q(\mathbf{G} \mid \mathbf{S})$. Unlike RIG, we use the goal distribution from Advantage Weighted Actor Critic to sample goals for exploration and for relabeling goals during training (Andrychowicz et al., 2017b). Since RIG already trains a generative model over states, we reuse this β -VAE for the generative model q_{ϕ}^G of Advantage Weighted Actor Critic. To make the most use of the data, we train q_{ϕ}^G on all visited state rather than only the terminal states, which we found to work well in practice. To prevent the estimated state likelihoods from collapsing to zero, we model the posterior of the β -VAE as a multivariate Gaussian distribution with a fixed variance and only learn the mean. We summarize RIG and provide details for how we combine Advantage Weighted Actor Critic and RIG in ?? and describe how we estimate the likelihoods given the β -VAE in ??.

3.5 RELATED WORK

Many prior methods in the goal-conditioned reinforcement learning literature focus on training goal-conditioned policies and assume that a goal distribution is available to sample from during exploration (L P Kaelbling, 1993; Schaul et al., 2015b; Andrychowicz et al., 2017b; V. Pong et al., 2018), or use a heuristic to design a non-parametric (Colas et al., 2018b; Warde-Farley et al., 2018; Florensa et al., 2018a) or parametric (Péré et al., 2018b; A. Nair et al., 2018b) goal distribution based on previously visited states. These methods are largely complementary to our work: rather than proposing a better method for training goal-reaching policies, we propose a principled method for maximizing the entropy of a goal sampling distribution, $\mathcal{H}(\mathbf{G})$, such that these policies cover a wide range of states.

Our method learns without any task rewards, directly acquiring a policy that can be reused to reach user-specified goals. This stands in contrast to exploration methods that modify the reward based on state visitation frequency (M. Bellemare et al., 2016; Ostrovski et al., 2017; Tang et al., 2017; Chentanez et al., 2005; Lopes et al., 2012; Stadie et al., 2016; Pathak et al., 2017; Burda et al., 2018; Burda et al., 2019; Mohamed and Rezende, 2015; Tang et al., 2017; Fu et al., 2017). While these methods can also be used without a task reward, they provide no mechanism for distilling the knowledge gained from visiting diverse states into flexible policies that can be applied to accomplish new goals at test-time: their policies visit novel states, and they quickly forget about them as

other states become more novel. Similarly, methods that provably maximize state entropy without using goal-directed exploration (Hazan et al., 2019) or methods that define new rewards to capture measures of intrinsic motivation (Mohamed and Rezende, 2015) and reachability (Savinov et al., 2018) do not produce reusable policies.

Other prior methods extract reusable skills in the form of latent-variable-conditioned policies, where latent variables are interpreted as options (Sutton et al., 1999) or abstract skills (Hausman et al., 2018; Abhishek Gupta et al., 2018b; Eysenbach et al., 2018; Abhishek Gupta et al., 2018a; Florensa et al., 2017). The resulting skills are diverse, but have no grounded interpretation, while Advantage Weighted Actor Critic policies can be used immediately after unsupervised training to reach diverse user-specified goals.

Some prior methods propose to choose goals based on heuristics such as learning progress (Adrien Baranes and Pierre-Yves Oudeyer, 2012a; Veeriah et al., 2018; Colas et al., 2018a), how off-policy the goal is (Nachum et al., 2018), level of difficulty (Held et al., 2018), or likelihood ranking (Zhao and Tresp, 2019). In contrast, our approach provides a principled framework for optimizing a concrete and well-motivated exploration objective, can provably maximize this objective under regularity assumptions, and empirically outperforms many of these prior work (see Section 8.6).

3.6 EXPERIMENTS

Our experiments study the following questions: (1) Does Advantage Weighted Actor Critic empirically result in a goal distribution with increasing entropy? (2) Does Advantage Weighted Actor Critic improve exploration for goal-conditioned RL? (3) How does Advantage Weighted Actor Critic compare to prior work on choosing goals for vision-based, goal-conditioned RL? (4) Can Advantage Weighted Actor Critic be applied to a real-world, vision-based robot task?

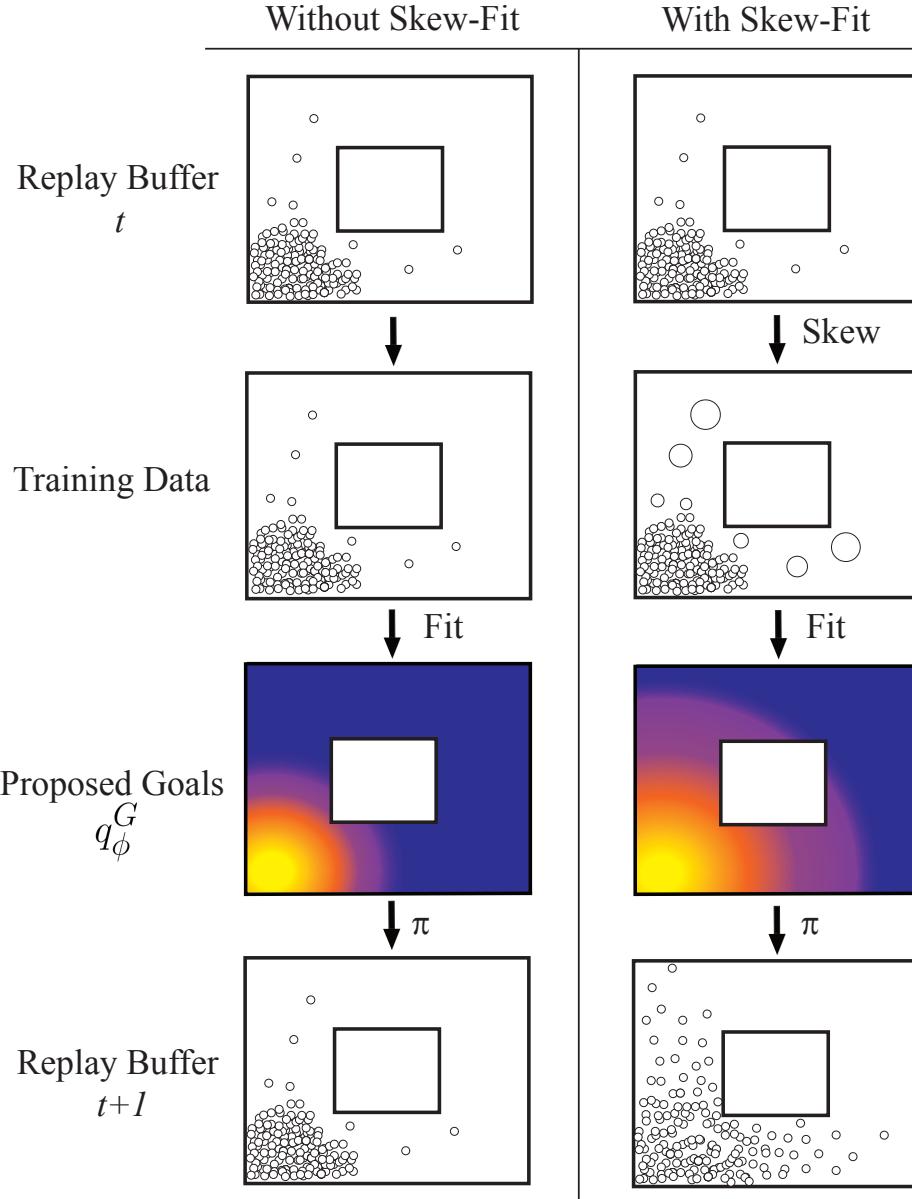


Figure 10: Our method, Advantage Weighted Actor Critic, samples goals for goal-conditioned RL. We sample states from our replay buffer, and give more weight to rare states. We then train a generative model $q_{\phi t+1}^G$ with the weighted samples. By sampling new states with goals proposed from this new generative model, we obtain a higher entropy state distribution in the next iteration.

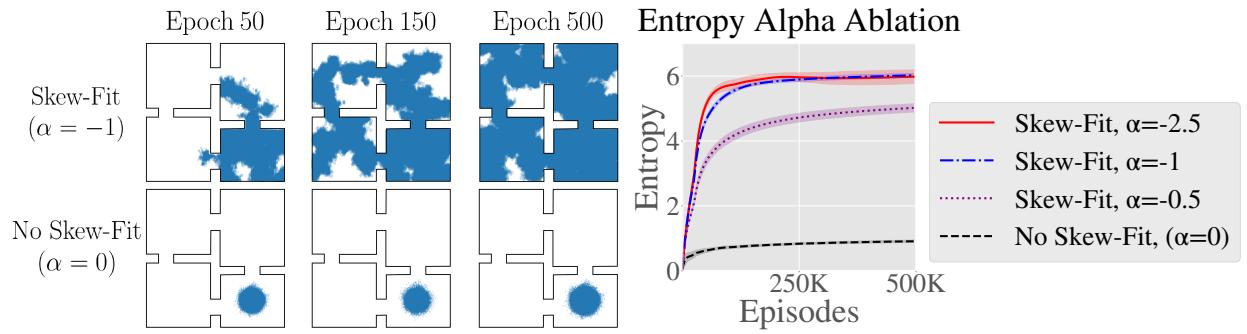


Figure 11: Illustrative example of Advantage Weighted Actor Critic on a 2D navigation task. (Left) Visited state plot for Advantage Weighted Actor Critic with $\alpha = -1$ and uniform sampling, which corresponds to $\alpha = 0$. (Right) The entropy of the goal distribution per iteration, mean and standard deviation for 9 seeds. Entropy is calculated via discretization onto an 11×11 grid. Advantage Weighted Actor Critic steadily increases the state entropy, reaching full coverage over the state space.

DOES ADVANTAGE WEIGHTED ACTOR CRITIC MAXIMIZE ENTROPY? To see the effects of Advantage Weighted Actor Critic on goal distribution entropy in isolation of learning a goal-reaching policy, we study an idealized example where the policy is a near-perfect goal-reaching policy. The environment consists of four rooms (Sutton et al., 1999). At the beginning of an episode, the agent begins in the bottom-right room and samples a goal from the goal distribution $q_{\phi_t}^G$. To simulate stochasticity of the policy and environment, we add a Gaussian noise with standard deviation of 0.06 units to this goal, where the entire environment is 11×11 units. The policy reaches the state that is closest to this noisy goal and inside the rooms, giving us a state sample s_n for training $q_{\phi_t}^G$. Due to the relatively small noise, the agent cannot rely on this stochasticity to explore the different rooms and must instead learn to set goals that are progressively farther and farther from the initial state. We compare multiple values of α , where $\alpha = 0$ corresponds to not using Advantage Weighted Actor Critic. The β -VAE hyperparameters used to train $q_{\phi_t}^G$ are given in ???. As seen in Figure 11, sampling uniformly from previous experience ($\alpha = 0$) to set goals results in a policy that primarily sets goal near the initial state distribution. In contrast, Advantage Weighted Actor Critic results in quickly learning a high entropy, near-uniform distribution over the state space.

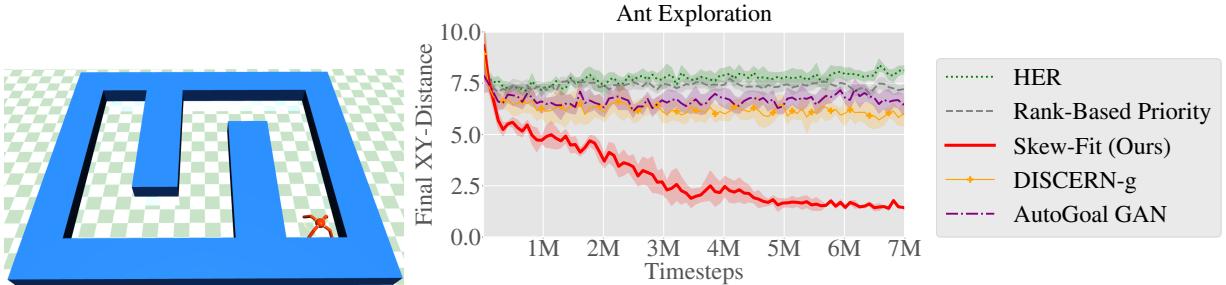


Figure 12: (Left) Ant navigation environment. (Right) Evaluation on reaching target XY position. We show the mean and standard deviation of 6 seeds. Advantage Weighted Actor Critic significantly outperforms prior methods on this exploration task.

EXPLORATION WITH ADVANTAGE WEIGHTED ACTOR CRITIC We next evaluate Advantage Weighted Actor Critic while concurrently learning a goal-conditioned policy on a task with state inputs, which enables us study exploration performance independently of the challenges with image observations. We evaluate on a task that requires training a simulated quadruped “ant” robot to navigate to different XY positions in a labyrinth, as shown in Figure 46. The reward is the negative distance to the goal XY-position, and additional environment details are provided in ???. This task presents a challenge for

goal-directed exploration: the set of valid goals is unknown due to the walls, and random actions do not result in exploring locations far from the start. Thus, Advantage Weighted Actor Critic must set goals that meaningfully explore the space while simultaneously learning to reach those goals.

We use this domain to compare Advantage Weighted Actor Critic to a number of existing goal-sampling methods. We compare to the relabeling scheme described in the hindsight experience replay (labeled **HER**). We compare to curiosity-driven prioritization (**Ranked-Based Priority**) (Zhao et al., 2019), a variant of HER that samples goals for relabeling based on their ranked likelihoods. Held et al. (2018) samples goals from a GAN based on the difficulty of reaching the goal. We compare against this method by replacing q_ϕ^G with the GAN and label it **AutoGoal GAN**. We also compare to the non-parametric goal proposal mechanism proposed by Warde-Farley et al., 2018, which we label **DISCERN-g**. Lastly, to demonstrate the difficulty of the exploration challenge in these domains, we compare to **#-Exploration** (Tang et al., 2017), an exploration method that assigns bonus rewards based on the novelty of new states. We train the goal-conditioned policy for each method using soft actor critic (SAC) (Haarnoja et al., 2018b). Implementation details of SAC and the prior works are given in ??.

We see in Figure 46 that Advantage Weighted Actor Critic is the only method that makes significant progress on this challenging labyrinth locomotion task. The prior methods on goal-sampling primarily set goals close to the start location, while the extrinsic exploration reward in #Exploration dominated the goal-reaching reward. These results demonstrate that Advantage Weighted Actor Critic accelerates exploration by setting diverse goals in tasks with unknown goal spaces.

VISION-BASED CONTINUOUS CONTROL TASKS

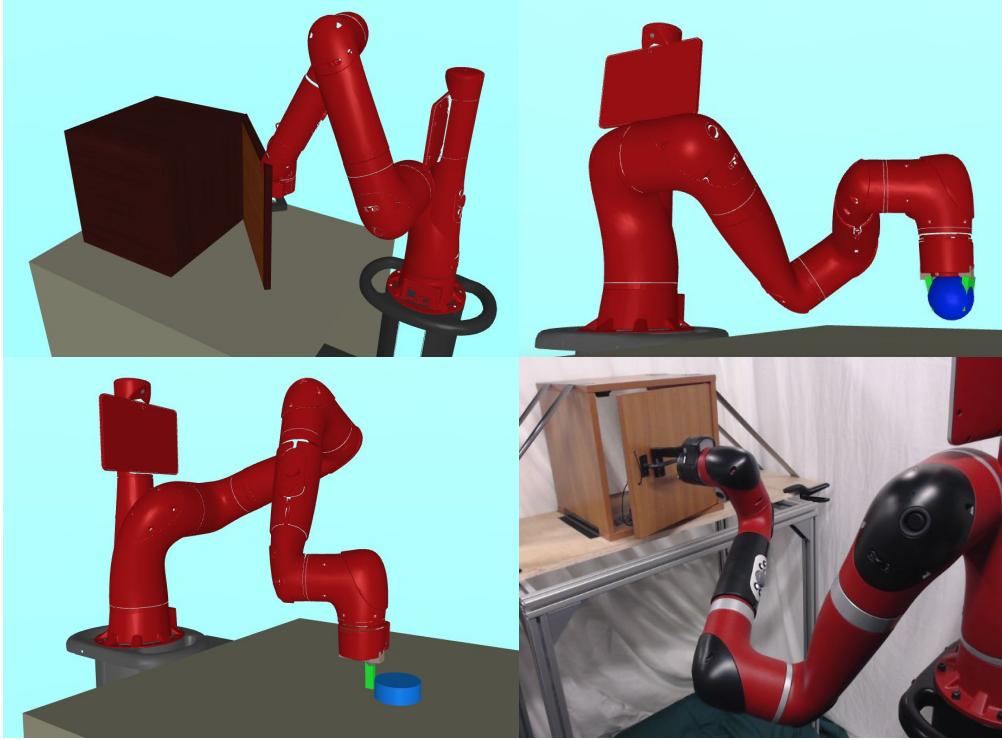


Figure 13: We evaluate on these continuous control tasks, from left to right: *Visual Door*, a door opening task; *Visual Pickup*, a picking task; *Visual Pusher*, a pushing task; and *Real World Visual Door*, a real world door opening task. All tasks are solved from images and without any task-specific reward. See Appendix ?? for details.

We now evaluate Advantage Weighted Actor Critic on a variety of image-based continuous control tasks, where the policy must control a robot arm using only image observations, there is no state-based or task-specific reward, and Advantage Weighted Actor Critic must directly set image goals. We test our method on three different image-based simulated continuous control tasks released by the authors of RIG ([A. Nair et al., 2018b](#)): *Visual Door*, *Visual Pusher*, and *Visual Pickup*. These environments contain a robot that can open a door, push a puck, and lift up a ball to different configurations, respectively. To our knowledge, these are the only goal-conditioned, vision-based continuous control environments that are publicly available and experimentally evaluated in prior work, making them a good point of comparison. See [Figure 13](#) for visuals and ?? for environment details. The policies are trained in a completely unsupervised manner, without access to any prior information about the image-space or any pre-defined goal-sampling distribution. To evaluate their performance, we sample goal images from a uniform dis-

tribution over valid states and report the agent’s final distance to the corresponding simulator states (e.g., distance of the object to the target object location), but the agent never has access to this true uniform distribution nor the ground-truth state information during training. While this evaluation method is only practical in simulation, it provides us with a quantitative measure of a policy’s ability to reach a broad coverage of goals in a vision-based setting.

We compare Advantage Weighted Actor Critic to a number of existing methods on this domain. First, we compare to the methods described in the previous experiment (HER, Rank-Based Priority, #Exploration, Autogoal GAN, and DISCERN-g). These methods that we compare to were developed in non-vision, state-based environments. To ensure a fair comparison across methods, we combine these prior methods with a policy trained using RIG. We additionally compare to Hazan et al. (2019), an exploration method that assigns bonus rewards based on the likelihood of a state (labeled Hazan et al.). Next, we compare to RIG without Advantage Weighted Actor Critic. Lastly, we compare to DISCERN (Warde-Farley et al., 2018), a vision-based method which uses a non-parametric clustering approach to sample goals and an image discriminator to compute rewards.

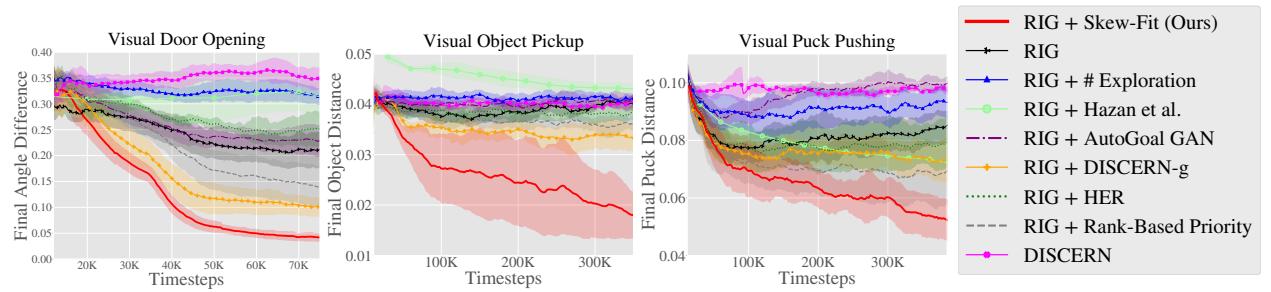


Figure 14: Learning curves for simulated continuous control tasks. Lower is better. We show the mean and standard deviation of 6 seeds and smooth temporally across 50 epochs within each seed. Advantage Weighted Actor Critic consistently outperforms RIG and various prior methods. See text for description of each method.

We see in Figure 14 that Skew-Fit significantly outperforms prior methods both in terms of task performance and sample complexity. The most common failure mode for prior methods is that the goal distributions collapse, resulting in the agent learning to reach only a fraction of the state space, as shown in Figure 9. For comparison, additional samples of q_ϕ^G when trained with and without Advantage Weighted Actor Critic are shown in ???. Those images show that without Advantage Weighted Actor Critic, q_ϕ^G produces a small, non-diverse distribution for each environment: the object is in the same

place for pickup, the puck is often in the starting position for pushing, and the door is always closed. In contrast, Advantage Weighted Actor Critic proposes goals where the object is in the air and on the ground, where the puck positions are varied, and the door angle changes.

We can see the effect of these goal choices by visualizing more example rollouts for RIG and Advantage Weighted Actor Critic. These visuals, shown in Figure ?? in ??, show that RIG only learns to reach states close to the initial position, while Advantage Weighted Actor Critic learns to reach the entire state space. For a quantitative comparison, Figure 15 shows the cumulative total exploration pickups for each method. From the graph, we see that many methods have a near-constant rate of object lifts throughout all of training. Advantage Weighted Actor Critic is the only method that significantly increases the rate at which the policy picks up the object during exploration, suggesting that only Advantage Weighted Actor Critic sets goals that encourage the policy to interact with the object.

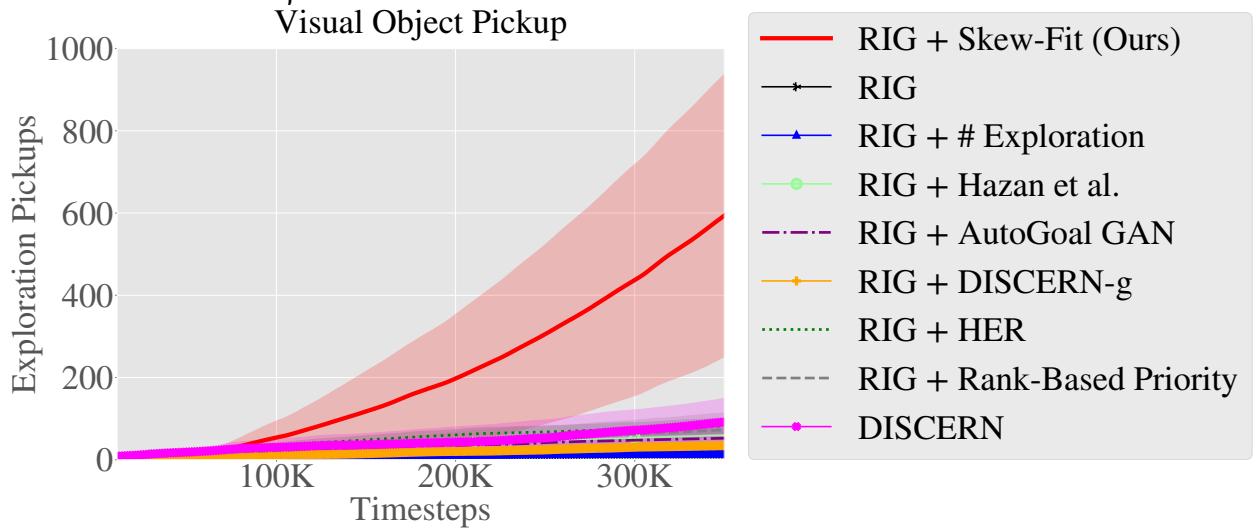


Figure 15: Cumulative total pickups during exploration for each method. Prior methods fail to pay attention to the object: the rate of pickups hardly increases past the first 100 thousand timesteps. In contrast, after seeing the object picked up a few times, Advantage Weighted Actor Critic practices picking up the object more often by sampling the appropriate exploration goals.

REAL-WORLD VISION-BASED ROBOTIC MANIPULATION We also demonstrate that Advantage Weighted Actor Critic scales well to the real world with a door opening task, *Real World Visual Door*, as shown in Figure 13. While a number of prior works have

studied RL-based learning of door opening (Kalakrishnan et al., 2011b; Chebotar et al., 2017b), we demonstrate the first method for autonomous learning of door opening without a user-provided, task-specific reward function. As in simulation, we do not provide any goals to the agent and simply let it interact with the door, without any human guidance or reward signal. We train two agents using RIG and RIG with Advantage Weighted Actor Critic. Every seven and a half minutes of interaction time, we evaluate on 5 goals and plot the cumulative successes for each method. Unlike in simulation, we cannot easily measure the difference between the policy’s achieved and desired door angle. Instead, we visually denote a binary success/failure for each goal based on whether the last state in the trajectory achieves the target angle. As Figure 16 shows, standard RIG only starts to open the door after five hours of training. In contrast, Advantage Weighted Actor Critic learns to occasionally open the door after three hours of training and achieves a near-perfect success rate after five and a half hours of interaction. Figure 16 also shows examples of successful trajectories from the Advantage Weighted Actor Critic policy, where we see that the policy can reach a variety of user-specified goals. These results demonstrate that Advantage Weighted Actor Critic is a promising technique for solving real world tasks without any human-provided reward function. Videos of Advantage Weighted Actor Critic solving this task and the simulated tasks can be viewed on our website.⁴

ADDITIONAL EXPERIMENTS To study the sensitivity of Advantage Weighted Actor Critic to the hyperparameter α , we sweep α across the values $[-1, -0.75, -0.5, -0.25, 0]$ on the simulated image-based tasks. The results are in ?? and demonstrate that Advantage Weighted Actor Critic works across a large range of values for α , and $\alpha = -1$ consistently outperform $\alpha = 0$ (i.e. outperforms no Advantage Weighted Actor Critic). Additionally, ?? provides a complete description our method hyperparameters, including network architecture and RL algorithm hyperparameters.

3.7 CONCLUSION

We presented a formal objective for self-supervised goal-directed exploration, allowing researchers to quantify and compare progress when designing algorithms that enable agents to autonomously learn. We also presented Advantage Weighted Actor Critic, an algorithm for training a generative model to approximate a uniform distribution over an initially unknown set of valid states, using data obtained via goal-conditioned reinforce-

⁴ <https://sites.google.com/view/skew-fit>

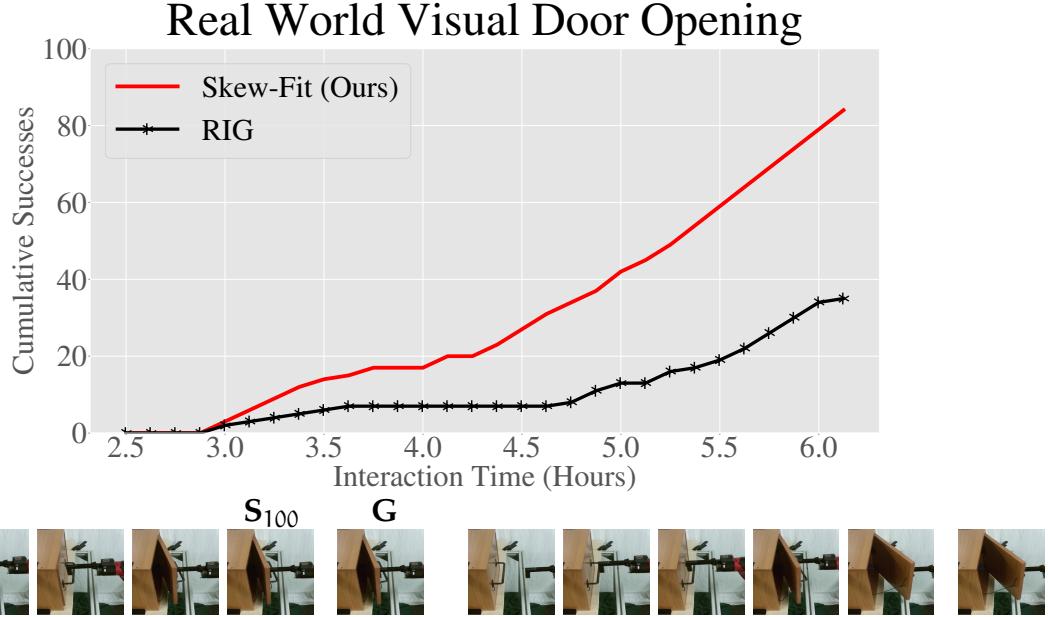


Figure 16: (Top) Learning curve for Real World Visual Door. Advantage Weighted Actor Critic results in considerable sample efficiency gains over RIG on this real-world task. (Bottom) Each row shows the Advantage Weighted Actor Critic policy starting from state S_1 and reaching state S_{100} while pursuing goal G . Despite being trained from only images without any user-provided goals during training, the Advantage Weighted Actor Critic policy achieves the goal image provided at test-time, successfully opening the door.

ment learning, and our theoretical analysis gives conditions under which Advantage Weighted Actor Critic converges to the uniform distribution. When such a model is used to choose goals for exploration and to relabeling goals for training, the resulting method results in much better coverage of the state space, enabling our method to explore effectively. Our experiments show that when we concurrently train a goal-reaching policy using self-generated goals, Advantage Weighted Actor Critic produces quantifiable improvements on simulated robotic manipulation tasks, and can be used to learn a door opening skill to reach a 95% success rate directly on a real-world robot, without any human-provided reward supervision.

4

CONTEXTUAL IMAGINED GOALS FOR SELF-SUPERVISED ROBOTIC LEARNING

4.1 INTRODUCTION

In order for robots to truly become *generalists*, they must be readily taskable by humans, handle raw sensory inputs without instrumentation, and be equipped with a range of skills that generalize effectively to new situations. Reinforcement learning autonomously learns policies that maximize a reward function and is a promising approach towards such generalist robots. However, in a general setting involving diverse objects and tasks, prior information about what tasks to learn is hard to come by without manually designing object detectors and reward functions. How can a robot explore the world in order to learn and fine-tune useful skills on diverse objects, only from acting and observing how its actions affect its sensory stream?

Prior methods have proposed to let an agent learn from its sensor stream by automatically generating plausible goals during an unsupervised training phase, and then learning policies that reach those goals (A. Nair et al., 2018b; Nachum et al., 2018; Warde-Farley et al., 2019; Vitchyr H. Pong et al., 2019). Such goals can be defined in a variety of ways, but a simple choice is to use goal observations, such that each proposed task requires reaching a different observation. When the robot observes the world via raw camera images, this corresponds to using images as goals. At test time, a user then provides the robot with a new goal image.

While such methods have been demonstrated in both simulated and real-world settings, they are typically used to learn behaviors in domains with relatively little visual diversity. In the real world, a robot might interact with highly diverse scenes and objects, and the tasks that it can perform from each of the many possible initial states will be different. If the robot is presented with an object, it can learn to pick up or grasp it, and when it is presented with a door, it can learn to open it. However, it must generate and

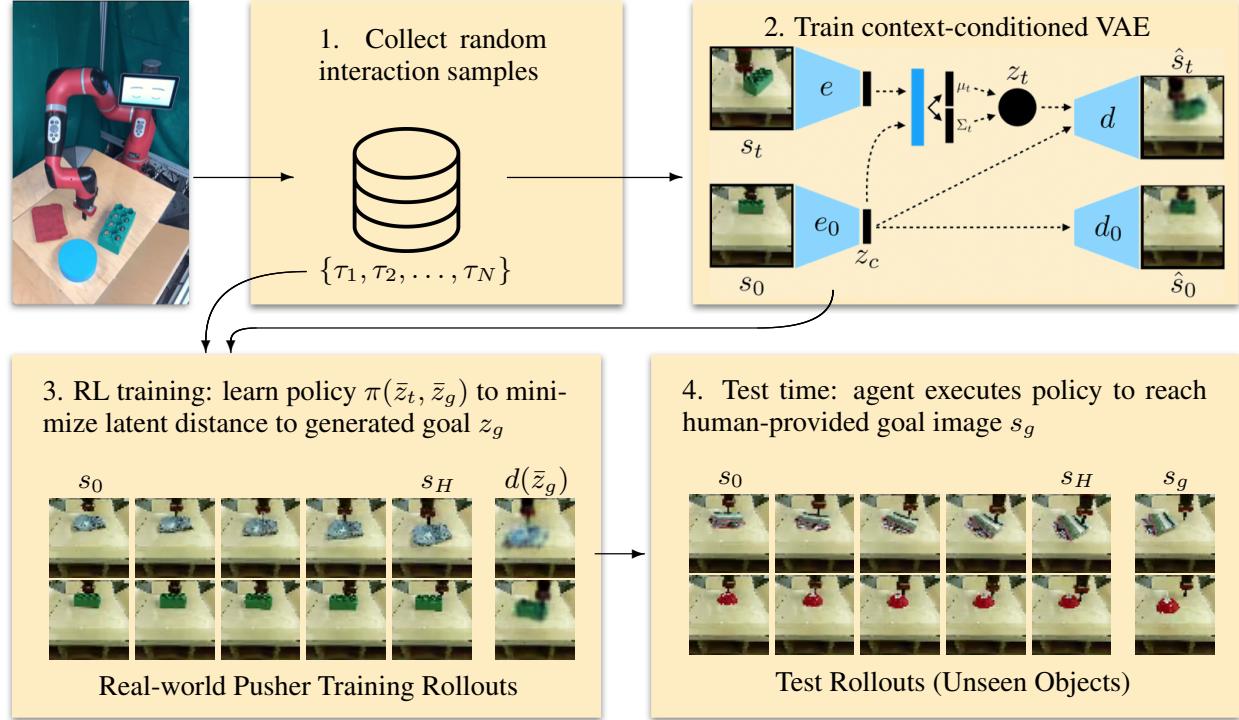


Figure 17: System overview of our self-supervised learning algorithm. (1) The agent collects random interaction data, to be used for both representation learning and as additional off-policy data for RL. (2) We propose a context-conditioned generative model (CC-VAE) to learn generalizable skills. In order to improve the generation of plausible goal states that result from a starting state s_0 , our model allows information to flow freely through the context z_c while an information bottleneck on z_t gives us the ability to generate samples by re-sampling only z_t . This architecture provides a compact representation of the scene disentangling information that changes within a rollout (z_t) and information that changes between rollouts (z_c). We then use $\bar{z}_t = (z_t, z_c)$ as the representation for RL. (3) Our proposed CC-RIG algorithm samples latent goals, using the above representation, and learns a policy to minimize the latent distance to the goal with off-policy RL. Rollouts are shown on the real-world Sawyer robot pusher environment with visual variation. We include the initial image s_0 , selected frames from the rollout, final image s_H , and the decoded goal latent $d(\bar{z}_g)$. (4) At test time, the agent is given a goal image s_g and executes the policy to reach it. Our method successfully handles pushing novel objects that were unseen at training time. Example rollouts can be found at <https://ccrig.github.io/>

practice goals that are suitable for each scene. In this paper, we propose and evaluate a self-supervised policy learning method that learns to propose goals that are suitable to the current scene via a conditional goal generation model, allowing it to learn in visually varied settings that prove challenging for prior algorithms.

The key idea in our work is that representing every element in a visually complex scene is often not necessary for control. A scene is a visual form of context that can be factored out, while only the controllable entities in the environment need to be captured for goal setting and representing the state. To this end, we propose learning a context-conditioned generative model that learns a smooth, compressed latent variable with an information bottleneck, while allowing the context, in the form of the initial state image, to be used freely to reconstruct other images during the task. This context-conditioned generative model architecture is shown in Figure 26.

The main contribution in this paper builds on this context-conditioned generative model to devise a complete self-supervised goal-conditioned reinforcement learning algorithm, which can handle visual variability in the scene via context-conditioned goal setting. Our method can learn policies that reach visually indicated goals without any additional supervision during training, using the context-conditioned generative model to set goals that are appropriate to the current scene. We show that our approach learns coherent representations of visually varied environments, capturing controllable dimensions of variation while ignoring dimensions that vary but cannot be influenced by the agent, such as lighting and object appearance. We further show that our approach can learn policies to solve tasks in visually varied environments, including in a real-world robotic pushing task with a wide variety of distinct objects in Figure 20.

4.2 RELATED WORK

While many practical robots today perform tasks by executing hand-engineered sequences of motor commands, machine learning is opening up a new avenue to train a wide variety of robotic tasks from interaction. This body of work includes grasping (Ekvall and Kragic, 2004; Kroemer et al., 2010; Bohg and Kragic, 2010), and general tasks (Peters and Schaal, 2008b; Kober et al., 2013), multi-task learning (M. P. Deisenroth et al., 2014), baseball (Peters and Schaal, 2008c), ping-pong (Peters et al., 2010), and various other tasks (Marc Peter Deisenroth and Rasmussen, 2011). More recently, using expressive function approximators such as neural networks has reduced manual feature engineering and has increased task complexity and diversity, finding use in decision-making domains, such as solving Atari games (Mnih et al., 2013) and Go (D. Silver et al.,

2016a). Deep learning for robotics has proved to be difficult due to a host of challenges including noisy state estimation, specifying reward functions, and handling continuous action spaces, but has been used to investigate grasping (Pinto and Abhinav Gupta, 2016), pushing (Agrawal et al., 2016), manipulation of 3D object models (Krainin et al., 2011), active learning (Martinez et al., 2014) and pouring liquids (Schenck and Fox, 2017). Deep reinforcement learning, which autonomously maximizes a given reward function, has been used to solve precise manipulation tasks (Levine et al., 2016a), grasping (Pinto et al., 2017; Levine et al., 2017), door opening (S. Gu et al., 2016), and navigation (Kahn et al., 2018). These methods have succeeded on specific tasks, often with hard-coded reward functions. However, to scale task generalization robots may need to learn methods that can handle significant environment variation and require relatively little external supervision.

Several works have investigated self-supervised robotic interaction with varied objects in the deep learning setting with the goal of generalizing between objects. For example, in the domain of robotic grasping, several works have studied autonomous data collection to learn to grasp from a hand-specified grasping reward (Pinto and Abhinav Gupta, 2016; Levine et al., 2017). However, hand-specifying such rewards in general settings and for arbitrary manipulation skills is very cumbersome. Other work has focused on self-supervised learning with visual forward models, either by enforcing a simplified dynamical structure (Watter et al., 2015; M. Zhang et al., 2019) or with pixel transformer architectures (Finn and Levine, 2016; Ebert et al., 2017; Ebert et al., 2018; A. X. Lee et al., n.d.; Ebert et al., n.d.). However, these methods rely on accurate visual forward modelling, which is itself a very challenging problem. Instead, we build on self-supervised model-free approaches, which allow the agent to efficiently reach visual goals without planning with a visual forward model.

Prior work has also sought to perform self-supervised learning with model-free approaches. Using visual inverse models (Agrawal et al., 2016) is one such approach, but may not work well for complex interaction dynamics or longer horizon planning. Most closely related to our approach are prior methods on goal-conditioned reinforcement learning (L P Kaelbling, 1993; Schaul et al., 2015b; Andrychowicz et al., 2017b). The methods have been extended to frame self-supervised RL as learning goal reaching with automatically proposed goals, including visually-specified goals (A. Nair et al., 2018b; Vitchyr H. Pong et al., 2019; Warde-Farley et al., 2019; Florensa et al., 2019; Lin et al., 2019). However, they generally focus on learning in narrow environments with little between-trial variability. In this setting, any previously visited state represents a valid goal. However, in the general case, this is no longer true: when the robot is presented

with a different scene or different objects on each trial, it must only set those goals that can be accomplished in the current scene. In contrast, we focus on enabling self-supervised learning from off-policy data in heterogeneous environments with increased factors of variability.

4.3 BACKGROUND

In this section, we provide an overview of relevant prior work necessary to understand our method, including goal-conditioned reinforcement learning and representation learning with variational auto-encoders.

4.3.1 Goal-Conditioned Reinforcement Learning

In an MDP consisting of states $s_t \in \mathcal{S}$, actions $a_t \in \mathcal{A}$, dynamics $p(s_{t+1}|s_t, a_t)$, rewards r_t , horizon H , and discount factor γ , reinforcement learning addresses optimizing a policy $a_t = \pi_\theta(s_t)$ to maximize expected return $\mathbb{E}[\sum_{t=0}^H \gamma^t r_t]$. To learn a variety of skills, it is instead convenient to optimize over a family of reward functions parametrized by goals, as in the framework of goal-conditioned RL ([L P Kaelbling, 1993](#)). A variety of RL algorithms exist, but as we are primarily interested in sample-efficient off-policy learning, we consider goal-conditioned Q-learning algorithms ([Schaul et al., 2015b](#)). These algorithms learn a parametrized Q-function $Q_w(s, a, g)$ that estimates the expected return of taking action a from state s with goal g . Q-learning methods rely on minimizing the Bellman error:

$$\mathcal{L} = |Q_w(s, a, g) - (r + \max_{a'} Q_w(s', a', g))| \quad (8)$$

This objective can be optimized using standard actor-critic algorithms using a set of transitions (s, a, s', g, r) which can be collected off-policy ([Lillicrap et al., 2016](#)). In practice, a target network is often used for the second Q-function.

4.3.2 Variational Auto-Encoders

Above, the goal description can take many forms. To handle high-dimensional goals, in this work, we learn a latent representation of the state using a variational auto-encoder (VAE). A VAE is a probabilistic generative model that has been shown to learn structured representations of high-dimensional data ([D. P. Kingma and Welling, 2014](#)) successfully.

It is trained by reconstructing states s with a parametrized encoder that converts the state into a normal random distribution $q_\phi(z|s)$ with a parametrized decoder that converts the latent variable z into a predicted state distribution $p_\psi(s|z)$, while keeping the latent z close (in KL divergence) to its prior $p(z)$, a standard normal distribution. To train the encoder and decoder parameters ϕ and ψ , we jointly optimize both objectives when minimizing the negative evidence lower bound:

$$\mathcal{L}_{\text{VAE}} = -\mathbb{E}_{q_\phi(z|s)}[\log p(s|z)] + \beta D_{\text{KL}}(q_\phi(z|s)\|p(z)). \quad (9)$$

4.3.3 Conditional Variational Auto-Encoders

Instead of a generative model that learns to generate the dataset distribution, one might instead desire a more structured generative model that can generate samples based on structured input. One example of this is a conditional variational auto-encoder (CVAE) that conditions the output on some input variable c and samples from $p(x|c)$ (Sohn et al., 2015). For example, to train a model that generates images of digits given the desired digit, the input variable c might be a one-hot encoded vector of the desired digit.

A CVAE trains $q_\phi(z|s, c)$ and $q_\psi(s|z, c)$, where both the encoder and decoder has access to the input variable c . The CVAE then minimizes:

$$\mathcal{L}_{\text{CVAE}} = -\mathbb{E}_{q_\phi(z|s,c)}[\log p(s|z, c)] + \beta D_{\text{KL}}(q_\phi(z|s, c)\|p(z)). \quad (10)$$

Samples are generated by first sampling a latent $z \sim p(z)$. Based on c , we can then decode z with $q_\psi(s|z, c)$ and visualize the output, which is in our case an image. In our framework $c = s_0$.

4.3.4 Reinforcement Learning with Imagined Goals

To learn skills from raw observations in a self-supervised manner, reinforcement learning with imagined goals (RIG) proposed to use representation learning combined with goal-conditioned RL (A. Nair et al., 2018b). The aim of RIG is to choose actions in order to make the state s_t reach a goal image s_g at test time. RIG first collects an interaction dataset $\{s_t\}$ and learns a latent representation by training a VAE on this data. Then, a goal-conditioned policy is trained to act in the environment in order to reach a given goal z_g . Exploration data is collected by rolling out the policy with goals that are “imagined” from the VAE prior; at test time, the policy can take in a goal image as input, encode

the image to the latent space, and act to reach the goal latent. A key limitation of this method is that sampling goals from the VAE prior during training time assumes that every state in the dataset is reachable at any time. However, in general, this assumption may not be true.

4.4 SELF-SUPERVISED LEARNING WITH CONTEXT-CONDITIONED REPRESENTATIONS

In this work, our goal is to enable the learning of flexible goal-conditioned policies that can be used to successfully perform a variety of tasks in a variety of contexts – e.g., with different objects in the scene. Such policies must learn from large amounts of experience, and although it is in principle possible to use random exploration to collect this experience, this quickly becomes impractical in the real world. It is, therefore, necessary for the robot to set its own goals during self-supervised training, to collect meaningful experience. However, in diverse settings, many randomly generated goals may not be feasible – e.g., the robot cannot push a red puck if the red puck is not present in the scene. We propose to extend off-policy goal-conditioned reinforcement learning with a conditional goal setting model, which proposes only those goals that are currently feasible. This enables a learning regime with imagined goals that is more realistic for real-world robotic systems that must generalize effectively to a range of objects and settings.

4.4.1 Context-Conditioned VAEs

To train a generative model that can improve the generation of feasible goals in varied scenes, we use a modified CVAE that uses the initial state s_0 in a rollout as the input c , which we call the “context” for that rollout. The modified CVAE, which we call a context-conditioned VAE (CC-VAE), is shown in Figure 26. While most CVAE applications use a one-hot vector as the input, we use an image s_0 . This image is encoded with a convolutional encoder e_0 into a compact representation z_c . Note that by design, e and e_0 do not share weights, as they are intended to encode different factors of variation in the images. The context z_c is used to output the latent representation z_t , as well as the reconstruction of the state \hat{s}_t . In addition, z_c is used alone to (deterministically) decode $\hat{s}_0 = d_0(z_c)$. The objective is given by

$$\mathcal{L}_{\text{CC-VAE}} = \mathcal{L}_{\text{CVAE}} + \log p(s_0|z_c). \quad (11)$$

Algorithm 3 Context-Conditioned RIG

Require: Encoders $\mu(s_t, s_0)$, $e_0(s_0)$, policy $\pi_\theta(\bar{z}, \bar{z}_g)$, goal-conditioned value function $Q_w(\bar{z}, \bar{z}_g)$, dataset $\mathcal{D} = \{\tau^{(i)}\}$ of trajectories.

- 1: Train CC-VAE on \mathcal{D} by optimizing equation 11.
- 2: **for** $n = 0, \dots, N - 1$ episodes **do**
- 3: Sample latent goal $z_g \sim p(z)$, $\bar{z}_g = (z_g, z_c)$.
- 4: Sample initial state $s_0 \sim p(s_0)$.
- 5: Encode $z_c = e_0(s_0)$
- 6: **for** $t = 0, \dots, H - 1$ steps **do**
- 7: Observe s_t and encode $\bar{z}_t = (\mu(s_t, s_0), z_c)$
- 8: Get action $a_t = \pi_\theta(\bar{z}_t, \bar{z}_g) + \text{noise}$.
- 9: Get next state $s_{t+1} \sim p(\cdot | s_t, a_t)$.
- 10: Store $(\bar{z}_t, a_t, \bar{z}_{t+1}, \bar{z}_g)$ into replay buffer \mathcal{R} .
- 11: Sample transition $(\bar{z}, a, \bar{z}', \bar{z}_g) \sim \mathcal{R}$.
- 12: Compute new reward $r = -\|\bar{z}' - \bar{z}_g\|$.
- 13: Minimize equation 35 using $(\bar{z}, a, \bar{z}', \bar{z}_g, r)$.
- 14: **end for**
- 15: **end for**=o

Due to the information bottleneck on z_t , this loss function penalizes information passing through z_t but allows for unrestricted information flow from z_c . Therefore, the optimal solution would encode as much information as possible in z_c , while only including the state information that changes within a trajectory in the latent variable z_t . These are precisely the features of most interest for control.

4.4.2 Context-Conditioned Reinforcement Learning with Imagined Goals

We propose to use our context-conditioned VAE in the RIG framework to learn policies over environments with visual diversity, where each episode might involve interacting with a different scene and different objects. We first collect a dataset of trajectories $\mathcal{D} = \{\tau^{(i)}\}$ by executing random actions in the environment. We then learn a CC-VAE, as detailed in Section 4.4.1, to learn a factored representation of the image observations. To use the CC-VAE for self-supervised learning, we save the first image s_0 when starting a rollout. We compute the encoding of s_0 , $z_c = e_0(s_0)$. Let \bar{z} denote the context concatenated vector (z, z_c) , and let $\mu(s, s_0)$ denote the mean of $q_\phi(z|s, s_0)$. We then use RIG in the \bar{z} latent space by encoding observations with μ , meaning that we train a goal-conditioned policy $\pi(\bar{z}, \bar{z}_g)$ and a goal-conditioned Q-function $Q(\bar{z}, \bar{z}_g)$.

To collect data, we sample a latent goal for each rollout from the prior $z_g \sim N(0, I)$, as in RIG. For every observation s_t , we compute the mean encoding $\mu_t(s_t, s_0)$. We then obtain a rollout of the policy by executing $\pi(\bar{z}, \bar{z}_g)$. The reward at each timestep is the

latent distance $\|\bar{\mu}_t - \bar{z}_g\|$.

The policy and Q-function can be trained with any off-policy reinforcement learning algorithm. We use TD3 in our implementation (Fujimoto et al., 2018b). Our policy and Q-function are goal-conditioned, and we take advantage of being able to relabel the goals for each transition to improve sample efficiency (Andrychowicz et al., 2017b; A. Nair et al., 2018b; Vitchyr H. Pong et al., 2019). However, when relabeling a goal \bar{z}_g with a random goal from the environment, the context-conditioning is still preserved. That is, if $z'_g \sim N(0, 1)$ is the new sampled goal, we use $\bar{z}'_g = (z'_g, z_c)$. This ensures that the relabeled goal is compatible with the scene for the corresponding transition.

After training, we can use the learned policy π to reach a visually indicated goal. Given a goal image s_g , we encode it into a latent goal $z_g = \mu(s_g, s_0)$. Then, we execute the policy with the latent goal \bar{z}_g , just as during the training phase. The complete algorithm is presented in Algorithm 3.

4.5 EXPERIMENTS

In our experiments, we aim to answer the following questions:

1. How does our method compare to prior work at learning self-supervised skills in visually diverse environments?
2. Do context-conditioned VAEs learn an image representation that produces coherent and diverse goals that are suitable for the current scene?
3. Can our proposed context-conditioned RIG method handle diverse real-world data and learn effective policies under visual variation in the real world?

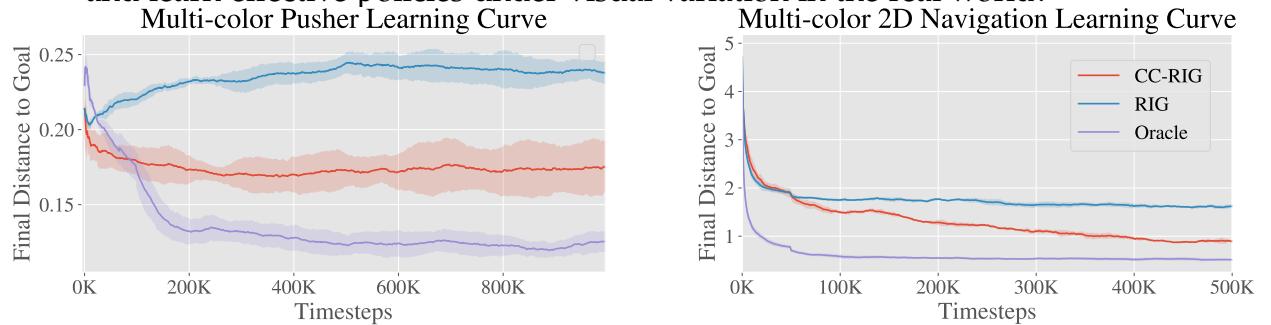


Figure 18: Self-supervised learning results in visually varied simulated environments. CC-RIG significantly outperforms RIG and is competitive with the oracle method that has direct access to ground truth states. The simulated pusher environment (left) is shown in Figure 19 and the navigation environment is shown in Figure ??.

4.5.1 Self-Supervised Learning in Simulation

In simulation, we can conduct controlled experiments and evaluate against known underlying state values to measure the performance of our approach and prior methods. As a simulation test-bed, we use a multi-color pusher environment simulated in MuJoCo ([Todorov et al., 2012](#)). In this environment (Figure 19(left)), a simulated Sawyer arm is tasked with pushing a circular puck to a target position, specified by a goal image at test time. On each rollout, the puck color is set to a random RGB value. Therefore, the goal proposals for each method must adequately account for the color of the puck – a goal that requires moving a red puck to a given location is impossible if only a blue puck is present in the scene.

We compare the following algorithms: **CC-RIG**. Our method using a CC-VAE for representation learning, as described in Section 4.4.2. **RIG**. Reinforcement learning with imagined goals ([A. Nair et al., 2018b](#)) using a standard VAE, as described in Section 4.3.4. **Oracle**. The oracle agent runs goal-conditioned RL with direct access to state information. Achieving performance similar to the oracle indicates that an algorithm loses little from using raw image observations over ground truth state.

Learning curves comparing these methods are presented in the plot on the left in Figure 18. CC-RIG outperforms RIG significantly, and standard RIG is not able to improve beyond the initial random policy. The performance of CC-RIG approaches that of the oracle policy, which has access to the true state. This suggests that, in visually varied environments, self-supervised learning is possible so long as the visual complexity is factored out with representation learning, and the proposed goals are consistent with the appearance of the current scene.

4.5.2 Generalizing to Varying Appearance and Dynamics with Self-Supervised Learning

In this experiment, to study changing both visual appearance and physical dynamics, we study how well our method can generalize when the environment dynamics change. We use a simulated 2D navigation task, where the goal is to navigate a point robot around an obstacle. The arrangement of the obstacles is chosen from a set of 15 possible configurations, and the color of the point robot is generated from a random RGB value. Learning curves obtained by training the different methods above in this environment are presented in Figure 18. CC-RIG requires more samples to learn, but eventually approaches the oracle performance. RIG, in comparison, plateaus with poor performance. This environment is explained further in the supplementary, in Section ?? and Figure ??.

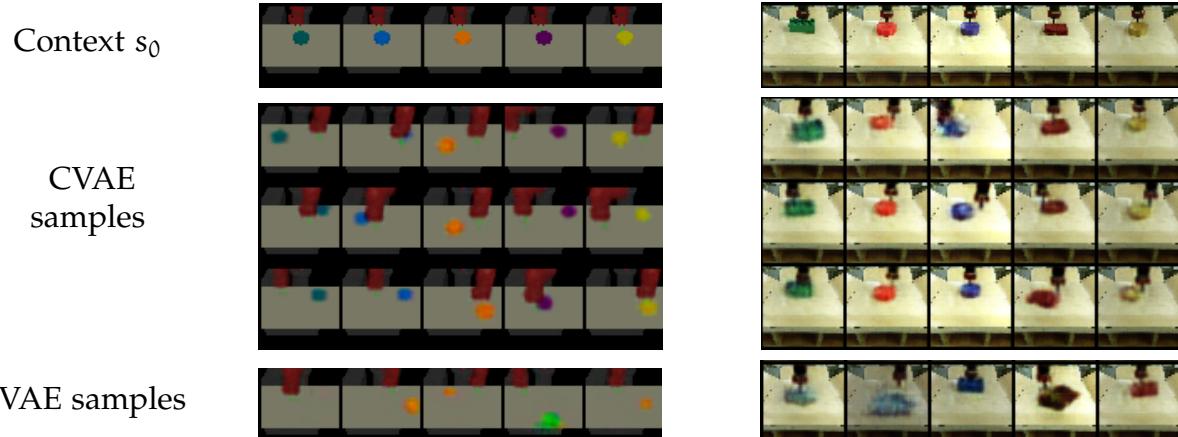


Figure 19: Comparing samples from our CC-VAE model to a standard VAE. The initial image s_0 is shown on the top row, and samples conditioned on s_0 are shown below. Our model coherently maintains object color and geometry in its samples, suggesting that the context conditioned model can successfully factor out the scene-specific object identity from the variable object position. This enables the use of the CC-VAE for goal proposals in visually diverse scenes.

4.5.3 Context-Conditioned VAE Goal Sampling

To better understand why CC-RIG outperforms RIG, we compare the samples from our CC-VAE to a standard VAE. Samples from both models are shown in Figure 19. The quality of the samples reveals why the CC-VAE provides better goal setting for self-supervised learning. In all environments, the samples from the CC-VAE maintain the background, object shape, and object color from the initial state. Therefore, the goals are more meaningful in the CC-VAE latent space.

This kind of visualization is a good indicator for the suitability of the representation for self-supervised learning. Diverse, coherent samples indicate that the latent space captures the appropriate factors of change in the environment and can be useful for self-supervised policy learning. Good samples also suggest that the latent space is well-structured, and therefore distances in the latent space should provide a good reward function for goal-reaching. In practice, we also look at the quality of the reconstructions. Good reconstructions confirm that the latent variables capture sufficient information about the image to be used in place of the image itself as a state representation.

4.5.4 Real-World Robotic Evaluation

In this experiment, we evaluate whether our method can handle manipulating visually varied objects in the real world. We use CC-RIG to train a Sawyer robot to manipulate a variety of objects, placed one at a time in a bin in front of the robot. As before, the training phase is self-supervised, and the robot must match a given goal image at test time. The robot setup is shown in Figure 26.

We first collect a large dataset with random actions and train a CC-VAE on the data. Samples from the model are shown in Figure 20. The CC-VAE learns to generate goals with the correct object. To handle varying brightness at different times of the day, we added data augmentation by applying a color jitter filter to (s_0, s_t) pairs. As seen in the figure, the model is robust to this factor of variation. Each sample contains the same type of object, brightness level, and background as the initial state that it is conditioned on. However, crucially, these factors of variation are not present in z_t , as evidenced by the fact they do not vary within each column of Figure 20, but the object position does.

Next, we run CC-RIG with the trained CC-VAE to learn to reach visually indicated goals in a self-supervised manner. We first conduct fully off-policy training using the same dataset as was used to train the CC-VAE, consisting of 50,000 samples (about 3 hours) of total interaction with 20 objects. Then, we collect a small amount of additional on-policy data to finetune the policy, analogous to recent work on large-scale vision-based robotic reinforcement learning (Kalashnikov et al., 2018c). The robot learns to push objects to target locations, indicated by a goal image. The real-world results are presented in Figure 20. Because it is difficult to automatically detect the positions of objects, we show some representative rollout examples, and we compute several distance metrics between the final state of a rollout and the goal: **CVAE distance**. CC-VAE latent space distance between final image and goal. **VAE distance**. VAE latent space distance between final image and goal. **Pixel distance**. We manually label the center of mass of the object in the final image and goal image, and compute the distance between them. **Object distance**. We measure the distance between the physical goal position of the object and the final position. In each metric, CC-RIG outperforms RIG.

At training time, the dataset consists of interaction with 20 objects. The result of running CC-RIG on novel objects that were not included in the dataset are shown in the table as “CC-RIG, novel” and in the rollouts in Figure 20. These results show that our method can also generalize its experience to push novel objects it has not seen before.

Real-World Pushing Results

	CVAE distance	VAE distance	Pixel distance	Object distance (cm)
RIG	2.37 ± 0.97	2.41 ± 0.93	93.9 ± 41.7	17.1 ± 8.2
CC-RIG	1.66 ± 0.63	2.17 ± 0.88	56.8 ± 34.5	14.0 ± 6.9
CC-RIG, novel	1.51 ± 0.71	1.91 ± 0.87	53.1 ± 24.9	11.5 ± 2.9

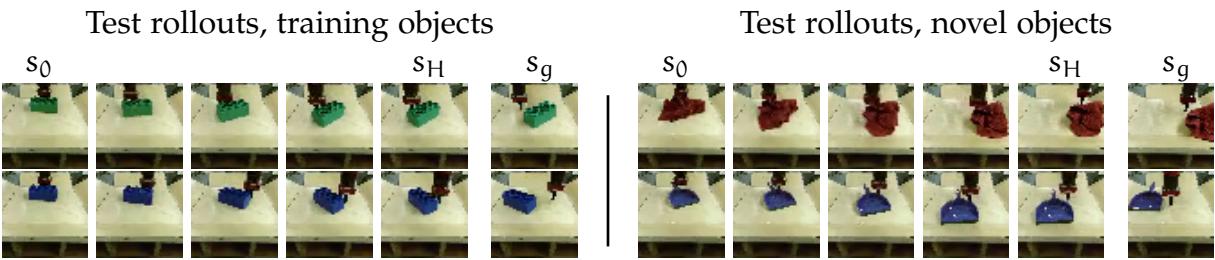


Figure 20: The table above shows the performance of our method in the real-world, evaluated with four different evaluation metrics¹. CC-RIG outperforms RIG in each one, even when tested on novel objects that it has not been trained on. Test rollouts of our method are shown on training objects on the left and unseen novel objects on the right. Successful rollouts where the object is pushed to the goal location are shown in top row, and failure modes are shown in the bottom row.

4.6 CONCLUSION

We presented a method for sample-efficient, flexible self-supervised task learning for environments with visual diversity. Our method can learn effective behavior without external supervision in simulated environments with randomized colors and layout, and in a real-world pushing task with differently colored pucks. Each environment contains an axis of visual variation that requires our algorithm to utilize an intelligent goal-setting strategy, to ensure that the self-proposed goals are consistent with the tasks and feasible in the current scene.

The main idea behind our method is to devise a context-conditioned goal proposal mechanism, allowing our self-supervised reinforcement learning algorithm to propose goals for itself that are feasible to reach. This context-conditioned VAE model factors out the unchanging context of a rollout, such as which objects are present in the scene, from the controllable aspects, such as the object positions to construct a more generalizable goal proposal model.

We believe this contribution will enable scalable learning in the real world. An agent manipulating objects in the real world must handle many forms of variation: different manipulation skills to learn, objects to manipulate, as well as variation in lighting, textures, etc. Methods that learn from data must be able to represent these variations while at the same time taking advantage of common structure across objects and tasks in order to achieve practical sample efficiency. Future work will address the remaining challenges to achieve this vision.

¹ The first three metrics are computed on 40 trajectories per method, and we report mean \pm standard deviation. Object distance is computed on 10 trajectories per method, and we report median \pm standard deviation.

Part II

UTILIZING PRIOR DATA FOR CONTROL

5

OVERCOMING EXPLORATION IN REINFORCEMENT LEARNING WITH DEMONSTRATIONS

5.1 INTRODUCTION

RL has found significant success in decision making for solving games, so what makes it more challenging to apply in robotics? A key difference is the difficulty of exploration, which comes from the choice of reward function and complicated environment dynamics. In games, the reward function is usually given and can be directly optimized. In robotics, we often desire behavior to achieve some binary objective (e.g., move an object to a desired location or achieve a certain state of the system) which naturally induces a sparse reward. Sparse reward functions are easier to specify and recent work suggests that learning with a sparse reward results in learned policies that perform the desired objective instead of getting stuck in local optima (Andrychowicz et al., 2017b; Večerík et al., 2017). However, exploration in an environment with sparse reward is difficult since with random exploration, the agent rarely sees a reward signal.

The difficulty posed by a sparse reward is exacerbated by the complicated environment dynamics in robotics. For example, system dynamics around contacts are difficult to model and induce a sensitivity in the system to small errors. Many robotics tasks also require executing multiple steps successfully over a long horizon, involve high dimensional control, and require generalization to varying task instances. These conditions further result in a situation where the agent so rarely sees a reward initially that it is not able to learn at all.

All of the above means that random exploration is not a tenable solution. Instead, in this work we show that we can use demonstrations as a guide for our exploration. To test our method, we solve the problem of stacking several blocks at a given location from a random initial state. Stacking blocks has been studied before in the literature (Marc Peter Deisenroth et al., 2011a; Duan et al., 2017) and exhibits many of the difficulties

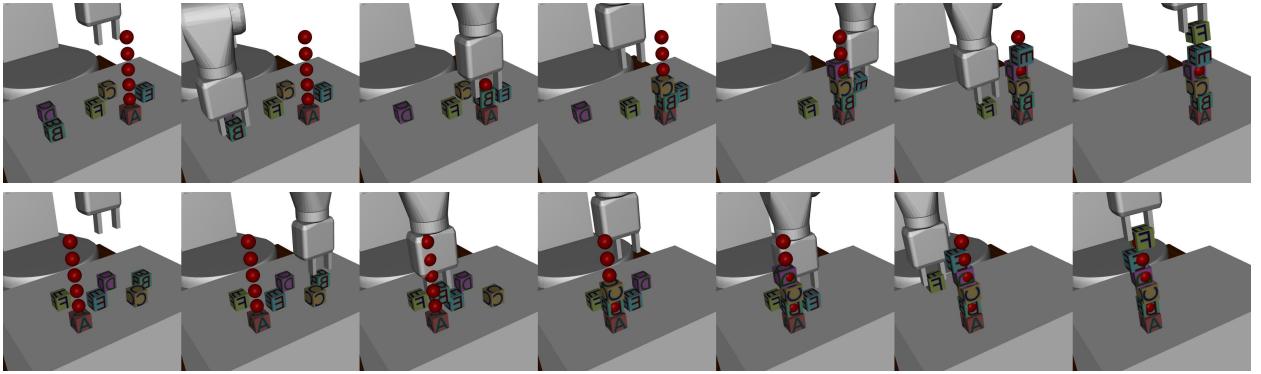


Figure 21: We present a method using reinforcement learning to solve the task of block stacking shown above. The robot starts with 6 blocks labelled A through F on a table in random positions and a target position for each block. The task is to move each block to its target position. The targets are marked in the above visualization with red spheres which do not interact with the environment. These targets are placed in order on top of block A so that the robot forms a tower of blocks. This is a complex, multi-step task where the agent needs to successfully manage multiple contacts to succeed. Frames from rollouts of the learned policy are shown. A video of our experiments can be found at: <http://ashvin.me/demoddpg-website>

mentioned: long horizons, contacts, and requires generalizing to each instance of the task. We limit ourselves to 100 human demonstrations collected via teleoperation in virtual reality. Using these demonstrations, we are able to solve a complex robotics task in simulation that is beyond the capability of both reinforcement learning and imitation learning.

The primary contribution of this paper is to show that demonstrations can be used with reinforcement learning to solve complex tasks where exploration is difficult. We introduce a simple auxiliary objective on demonstrations, a method of annealing away the effect of the demonstrations when the learned policy is better than the demonstrations, and a method of resetting from demonstration states that significantly improves and speeds up training policies. By effectively incorporating demonstrations into RL, we short-circuit the random exploration phase of RL and reach nonzero rewards and a reasonable policy early on in training. Finally, we extensively evaluate our method against other commonly used methods, such as initialization with learning from demonstrations and fine-tuning with RL, and show that our method significantly outperforms them.

5.2 RELATED WORK

Learning methods for decision making problems such as robotics largely divide into two classes: imitation learning and reinforcement learning (RL). In imitation learning (also called learning from demonstrations) the agent receives behavior examples from an expert and attempts to solve a task by copying the expert’s behavior. In RL, an agent attempts to maximize expected reward through interaction with the environment. Our work combines aspects of both to solve complex tasks.

Imitation Learning: Perhaps the most common form of imitation learning is behavior cloning (BC), which learns a policy through supervised learning on demonstration state-action pairs. BC has seen success in autonomous driving (Pomerleau, 1989; Bojarski et al., 2016), quadcopter navigation (Giusti et al., 2015), locomotion (Nakanishi et al., 2004; Kalakrishnan et al., 2009). BC struggles outside the manifold of demonstration data. Dataset Aggregation (DAGGER) augments the dataset by interleaving the learned and expert policy to address this problem of accumulating errors (Ross et al., 2011). However, DAGGER is difficult to use in practice as it requires access to an expert during all of training, instead of just a set of demonstrations.

Fundamentally, BC approaches are limited because they do not take into account the task or environment. Inverse reinforcement learning (IRL) (A. Ng and Russell, 2000) is another form of imitation learning where a reward function is inferred from the demonstrations. Among other tasks, IRL has been applied to navigation (Ziebart et al., 2008), autonomous helicopter flight (Abbeel and Andrew Y Ng, 2004), and manipulation (Finn et al., 2016a). Since our work assumes knowledge of a reward function, we omit comparisons to IRL approaches.

Reinforcement Learning: Reinforcement learning methods have been harder to apply in robotics, but are heavily investigated because of the autonomy they could enable. Through RL, robots have learned to play table tennis (Peters et al., 2010), swing up a cartpole, and balance a unicycle (Marc Peter Deisenroth and Rasmussen, 2011). A renewal of interest in RL cascaded from success in games (Mnih et al., 2015; D. Silver et al., 2016b), especially because of the ability of RL with large function approximators (ie. deep RL) to learn control from raw pixels. Robotics has been more challenging in general but there has been significant progress. Deep RL has been applied to manipulation tasks (LevineFDA15), grasping (Pinto and Abhinav Gupta, 2016; Levine et al., 2016b), opening a door (S. Gu et al., 2017), and locomotion (Lillicrap et al., 2016; Mnih et al., 2016; Schulman et al., 2015). However, results have been attained predominantly in simulation per high sample complexity, typically caused by exploration challenges.

Robotic Block Stacking: Block stacking has been studied from the early days of AI and robotics as a task that encapsulates many difficulties of more complicated tasks we want to solve, including multi-step planning and complex contacts. SHRDLU (Winograd, 1972) was one of the pioneering works, but studied block arrangements only in terms of logic and natural language understanding. More recent work on task and motion planning considers both logical and physical aspects of the task (Leslie Pack Kaelbling and Lozano-Perez, 2011; Kavraki et al., 1996; Srivastava et al., 2014), but requires domain-specific engineering. In this work we study how an agent can learn this task without the need of domain-specific engineering.

One RL method, PILCO (Marc Peter Deisenroth and Rasmussen, 2011) has been applied to a simple version of stacking blocks where the task is to place a block on a tower (Marc Peter Deisenroth et al., 2011a). Methods such as PILCO based on learning forward models naturally have trouble modelling the sharply discontinuous dynamics of contacts; although they can learn to place a block, it is a much harder problem to grasp the block in the first place. One-shot Imitation (Duan et al., 2017) learns to stack blocks in a way that generalizes to new target configurations, but uses more than 100,000 demonstrations to train the system. A heavily shaped reward can be used to learn to stack a Lego block on another with RL (Popov et al., 2017). In contrast, our method can succeed from fully sparse rewards and handle stacking several blocks.

Combining RL and Imitation Learning: Previous work has combined reinforcement learning with demonstrations. Demonstrations have been used to accelerate learning on classical tasks such as cart-pole swing-up and balance (Schaal, 1997). This work initialized policies and (in model-based methods) initialized forward models with demonstrations. Initializing policies from demonstrations for RL has been used for learning to hit a baseball (Peters and Schaal, 2008c) and for underactuated swing-up (Kober and Peter, 2008). Beyond initialization, we show how to extract more knowledge from demonstrations by using them effectively throughout the entire training process.

Our method is closest to two recent approaches — Deep Q-Learning From Demonstrations (DQfD) (Hester et al., 2018) and DDPG From Demonstrations (DDPGfD) (Večerík et al., 2017) which combine demonstrations with reinforcement learning. DQfD improves learning speed on Atari, including a margin loss which encourages the expert actions to have higher Q-values than all other actions. This loss can make improving upon the demonstrator policy impossible which is not the case for our method. Prior work has previously explored improving beyond the demonstrator policy in simple environments by introducing slack variables (Kim et al., 2013), but our method uses a learned value to actively inform the improvement. DDPGfD solves simple robotics tasks akin to peg inser-

tion using DDPG with demonstrations in the replay buffer. In contrast to this prior work, the tasks we consider exhibit additional difficulties that are of key interest in robotics: multi-step behaviours, and generalization to varying goal states. While previous work focuses on speeding up already solvable tasks, we show that we can extend the state of the art in RL with demonstrations by introducing new methods to incorporate demonstrations.

5.3 BACKGROUND

5.3.1 Reinforcement Learning

We consider the standard Markov Decision Process framework for picking optimal actions to maximize rewards over discrete timesteps in an environment E . We assume that the environment is fully observable. At every timestep t , an agent is in a state x_t , takes an action a_t , receives a reward r_t , and E evolves to state x_{t+1} . In reinforcement learning, the agent must learn a policy $a_t = \pi(x_t)$ to maximize expected returns. We denote the return by $R_t = \sum_{i=t}^T \gamma^{(i-t)} r_i$ where T is the horizon that the agent optimizes over and γ is a discount factor for future rewards. The agent's objective is to maximize expected return from the start distribution $J = \mathbb{E}_{r_i, s_i \sim E, a_i \sim \pi}[R_0]$.

A variety of reinforcement learning algorithms have been developed to solve this problem. Many involve constructing an estimate of the expected return from a given state after taking an action:

$$Q^\pi(s_t, a_t) = \mathbb{E}_{r_i, s_i \sim E, a_i \sim \pi}[R_t | s_t, a_t] \quad (12)$$

$$= \mathbb{E}_{r_t, s_{t+1} \sim E}[r_t + \gamma \mathbb{E}_{a_{t+1} \sim \pi}[Q^\pi(s_{t+1}, a_{t+1})]] \quad (13)$$

We call Q^π the action-value function. Equation 25 is a recursive version of equation 24, and is known as the Bellman equation. The Bellman equation allows for methods to estimate Q that resemble dynamic programming.

5.3.2 DDPG

Our method combines demonstrations with one such method: Deep Deterministic Policy Gradients (DDPG) (Lillicrap et al., 2016). DDPG is an off-policy model-free reinforcement learning algorithm for continuous control which can utilize large function approximators such as neural networks. DDPG is an actor-critic method, which bridges

the gap between policy gradient methods and value approximation methods for RL. At a high level, DDPG learns an action-value function (critic) by minimizing the Bellman error, while simultaneously learning a policy (actor) by directly maximizing the estimated action-value function with respect to the parameters of the policy.

Concretely, DDPG maintains an actor function $\pi(s)$ with parameters θ_π , a critic function $Q(s, a)$ with parameters θ_Q , and a replay buffer R as a set of tuples (s_t, a_t, r_t, s_{t+1}) for each transition experienced. DDPG alternates between running the policy to collect experience and updating the parameters. Training rollouts are collected with extra noise for exploration: $a_t = \pi(s) + \mathcal{N}$, where \mathcal{N} is a noise process.

During each training step, DDPG samples a minibatch consisting of N tuples from R to update the actor and critic networks. DDPG minimizes the following loss L w.r.t. θ_Q to update the critic:

$$y_i = r_i + \gamma Q(s_{i+1}, \pi(s_{i+1})) \quad (14)$$

$$L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i | \theta_Q))^2 \quad (15)$$

The actor parameters θ_π are updated using the policy gradient:

$$\nabla_{\theta_\pi} J = \frac{1}{N} \sum_i \nabla_a Q(s, a | \theta_Q)|_{s=s_i, a=\pi(s)} \nabla_{\theta_\pi} \pi(s | \theta_\pi)|_{s_i} \quad (16)$$

To stabilize learning, the Q value in equation 14 is usually computed using a separate network (called the *target* network) whose weights are an exponential average over time of the critic network. This results in smoother target values.

Note that DDPG is a natural fit for using demonstrations. Since DDPG can be trained off-policy, we can use demonstration data as off-policy training data. We also take advantage of the action-value function $Q(s, a)$ learned by DDPG to better use demonstrations.

5.3.3 Multi-Goal RL

Instead of the standard RL setting, we train agents with parametrized goals, which lead to more general policies (Schaul et al., 2015b) and have recently been shown to make learning with sparse rewards easier (Andrychowicz et al., 2017b). Goals describe the task we expect the agent to perform in the given episode, in our case they specify the desired positions of all objects. We sample the goal g at the beginning of every episode. The function approximators, here π and Q , take the current goal as an additional input.

5.3.4 Hindsight Experience Replay (HER)

To handle varying task instances and parametrized goals, we use Hindsight Experience Replay (HER) ([Andrychowicz et al., 2017b](#)). The key insight of HER is that even in failed rollouts where no reward was obtained, the agent can transform them into successful ones by assuming that a state it saw in the rollout was the actual goal. HER can be used with any off-policy RL algorithm assuming that for every state we can find a goal corresponding to this state (i.e. a goal which leads to a positive reward in this state).

For every episode the agent experiences, we store it in the replay buffer twice: once with the original goal pursued in the episode and once with the goal corresponding to the final state achieved in the episode, as if the agent intended on reaching this state from the very beginning.

5.4 METHOD

Our method combines DDPG and demonstrations in several ways to maximally use demonstrations to improve learning. We describe our method below and evaluate these ideas in our experiments.

5.4.1 Demonstration Buffer

First, we maintain a second replay buffer R_D where we store our demonstration data in the same format as R . In each minibatch, we draw an extra N_D examples from R_D to use as off-policy replay data for the update step. These examples are included in both the actor and critic update. This idea has been introduced in ([Večerík et al., 2017](#)).

5.4.2 Behavior Cloning Loss

Second, we introduce a new loss computed only on the demonstration examples for training the actor.

$$L_{BC} = \sum_{i=1}^{N_D} \pi(s_i | \theta_\pi) - a_i^2 \quad (17)$$

This loss is a standard loss in imitation learning, but we show that using it as an auxiliary loss for RL improves learning significantly. The gradient applied to the actor parameters

θ_π is:

$$\lambda_1 \nabla_{\theta_\pi} J - \lambda_2 \nabla_{\theta_\pi} L_{BC} \quad (18)$$

(Note that we maximize J and minimize L_{BC} .) Using this loss directly prevents the learned policy from improving significantly beyond the demonstration policy, as the actor is always tied back to the demonstrations. Next, we show how to account for sub-optimal demonstrations using the learned action-value function.

5.4.3 Q-Filter

We account for the possibility that demonstrations can be suboptimal by applying the behavior cloning loss only to states where the critic $Q(s, a)$ determines that the demonstrator action is better than the actor action:

$$L_{BC} = \sum_{i=1}^{N_D} \pi(s_i | \theta_\pi) - a_i^2 \mathbb{1}_{Q(s_i, a_i) > Q(s_i, \pi(s_i))} \quad (19)$$

The gradient applied to the actor parameters is as in equation 18. We label this method using the behavior cloning loss and Q-filter “Ours” in the following experiments.

5.4.4 Resets to demonstration states

To overcome the problem of sparse rewards in very long horizon tasks, we reset some training episodes using states and goals from demonstration episodes. Restarts from within demonstrations expose the agent to higher reward states during training. This method makes the additional assumption that we can restart episodes from a given state, as is true in simulation.

To reset to a demonstration state, we first sample a demonstration $D = (x_0, u_0, x_1, u_1, \dots, x_N, u_N)$ from the set of demonstrations. We then uniformly sample a state x_i from D . As in HER, we use the final state achieved in the demonstration as the goal. We roll out the trajectory with the given initial state and goal for the usual number of timesteps. At evaluation time, we do not use this procedure.

We label our method with the behavior cloning loss, Q-filter, and resets from demonstration states as “Ours, Resets” in the following experiments.

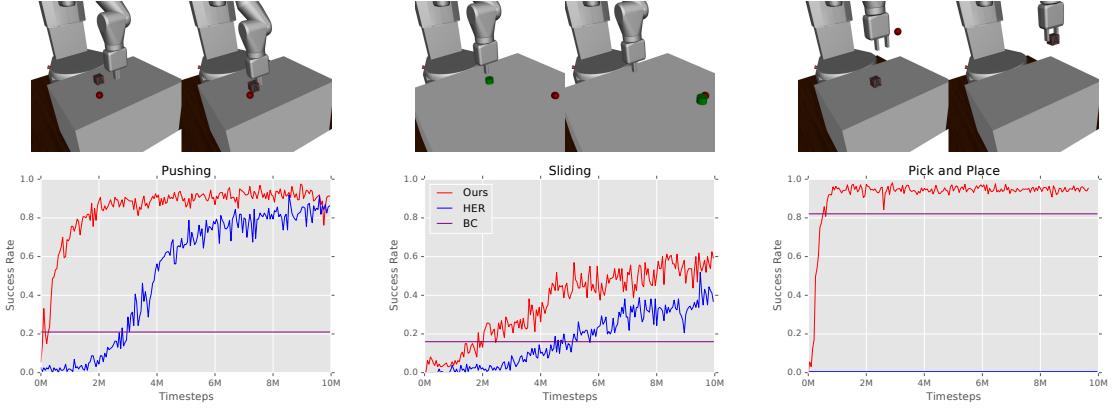


Figure 22: Baseline comparisons on tasks from (Andrychowicz et al., 2017b). Frames from the learned policy are shown above each task. Our method significantly outperforms the baselines. On the right plot, the HER baseline always fails.

5.5 EXPERIMENTAL SETUP

5.5.1 Environments

We evaluate our method on several simulated MuJoCo (Todorov et al., 2012) environments. In all experiments, we use a simulated 7-DOF Fetch Robotics arm with parallel grippers to manipulate one or more objects placed on a table in front of the robot.

The agent receives the positions of the relevant objects on the table as its observations. The control for the agent is continuous and 4-dimensional: 3 dimensions that specify the desired end-effector position¹ and 1 dimension that specifies the desired distance between the robot fingers. The agent is controlled at 50Hz frequency.

We collect demonstrations in a virtual reality environment. The demonstrator sees a rendering of the same observations as the agent, and records actions through a HTC Vive interface at the same frequency as the agent. We have the option to accept or reject a demonstration; we only accept demonstrations we judge to be mostly correct. The demonstrations are not optimal. The most extreme example is the “sliding” task, where only 7 of the 100 demonstrations are successful, but the agent still sees rewards for these demonstrations with HER.

¹ In the 10cm x 10cm x 10cm cube around the current gripper position

5.5.2 Training Details

To train our models, we use Adam (D. Kingma and Ba, 2015) as the optimizer with learning rate 10^{-3} . We use $N = 1024$, $N_D = 128$, $\lambda_1 = 10^{-3}$, $\lambda_2 = 1.0/N_D$. The discount factor γ is 0.98. We use 100 demonstrations to initialize R_D . The function approximators π and Q are deep neural networks with ReLU activations and L2 regularization with the coefficient 5×10^{-3} . The final activation function for π is tanh, and the output value is scaled to the range of each action dimension. To explore during training, we sample random actions uniformly within the action space with probability 0.1 at every step, and the noise process \mathcal{N} is uniform over $\pm 10\%$ of the maximum value of each action dimension. Task-specific information, including network architectures, are provided in the next section.

5.5.3 Overview of Experiments

We perform three sets of experiments. In Sec. 5.6, we provide a comparison to previous work. In Sec. 5.7 we solve block stacking, a difficult multi-step task with complex contacts that the baselines struggle to solve. In Sec. 5.8 we do ablations of our own method to show the effect of individual components.

5.6 COMPARISON WITH PRIOR WORK

5.6.1 Tasks

We first show the results of our method on the simulated tasks presented in the Hind-sight Experience Replay paper (Andrychowicz et al., 2017b). We apply our method to three tasks:

1. *Pushing*. A block placed randomly on the table must be moved to a target location on the table by the robot (fingers are blocked to avoid grasping).
2. *Sliding*. A puck placed randomly on the table must be moved to a given target location. The target is outside the robot's reach so it must apply enough force that the puck reaches the target and stops due to friction.
3. *Pick-and-place*. A block placed randomly on the table must be moved to a target location in the air. Note that the original paper used a form of initializing from favorable states to solve this task. We omit this for our experiment but discuss and evaluate the initialization idea in an ablation.

As in the prior work, we use a fully sparse reward for this task. The agent is penalized if the object is not at its goal position:

$$r_t = \begin{cases} 0, & \text{if } \|x_i - g_i\| < \delta \\ -1, & \text{otherwise} \end{cases} \quad (20)$$

where the threshold δ is 5cm.

5.6.2 Results

Fig. 22 compares our method to HER without demonstrations and behavior cloning. Our method is significantly faster at learning these tasks than HER, and achieves significantly better policies than behavior cloning does. Measuring the number of timesteps to get to convergence, we exhibit a 4x speedup over HER in pushing, a 2x speedup over HER in sliding, and our method solves the pick-and-place task while HER baseline cannot solve it at all.

The pick-and-place task showcases the shortcoming of RL in sparse reward settings, even with HER. In pick-and-place, the key action is to grasp the block. If the robot could manage to grasp it a small fraction of the time, HER discovers how to achieve goals in the air and reinforces the grasping behavior. However, grasping the block with random actions is extremely unlikely. Our method pushes the policy towards demonstration actions, which are more likely to succeed.

In the HER paper, HER solves the pick-and-place task by initializing half of the rollouts with the gripper grasping the block. With this addition, pick-and-place becomes the easiest of the three tasks tested. This initialization is similar in spirit to our initialization idea, but takes advantage of the fact that pick-and-place with any goal can be solved starting from a block grasped at a certain location. This is not always true (for example, if there are multiple objects to be moved) and finding such a keyframe for other tasks would be difficult, requiring some engineering and sacrificing autonomy. Instead, our method guides the exploration towards grasping the block through demonstrations. Providing demonstrations does not require expert knowledge of the learning system, which makes it a more compelling way to provide prior information.

5.7 MULTI-STEP EXPERIMENTS

5.7.1 Block Stacking Task

To show that our method can solve more complex tasks with longer horizon and sparser reward, we study the task of block stacking in a simulated environment as shown in Fig. 26 with the same physical properties as the previous experiments. Our experiments show that our approach can solve the task in full and learn a policy to stack 6 blocks with demonstrations and RL. To measure and communicate various properties of our method, we also show experiments on stacking fewer blocks, a subset of the full task.

We initialize the task with blocks at 6 random locations $x_1 \dots x_6$. We also provide 6 goal locations $g_1 \dots g_6$. To form a tower of blocks, we let $g_1 = x_1$ and $g_i = g_{i-1} + (0, 0, 5\text{cm})$ for $i \in 2, 3, 4, 5$.

By stacking N blocks, we mean N blocks reach their target locations. Since the target locations are always on top of x_1 , we start with the first block already in position. So stacking N blocks involves $N - 1$ pick-and-place actions. To solve stacking N , we allow the agent $50 * (N - 1)$ timesteps. This means that to stack 6 blocks, the robot executes 250 actions or 5 seconds.

We recorded 100 demonstrations to stack 6 blocks, and use subsets of these demonstrations as demonstrations for stacking fewer blocks. The demonstrations are not perfect; they include occasionally dropping blocks, but our method can handle suboptimal demonstrations. We still rejected more than half the demonstrations and excluded them from the demonstration data because we knocked down the tower of blocks when releasing a block.

5.7.2 Rewards

Two different reward functions are used. To test the performance of our method under fully sparse reward, we reward the agent only if all blocks are at their goal positions:

$$r_t = \min_i \mathbb{1}_{\|x_i - g_i\| < \delta} \quad (21)$$

The threshold δ is the size of a block, 5cm. Throughout the paper we call this the “sparse” reward.

To enable solving the longer horizon tasks of stacking 4 or more blocks, we use the “step” reward :

$$r_t = -1 + \sum_i \mathbb{1}_{\|x_i - g_i\| < \delta} \quad (22)$$

Task	Ours	Ours, Resets	BC	HER	BC+ HER
Stack 2, Sparse	99%	97%	65%	0%	65%
Stack 3, Sparse	99%	89%	1%	0%	1%
Stack 4, Sparse	1%	54%	-	-	-
Stack 4, Step	91%	73%	0%	0%	0%
Stack 5, Step	49%	50%	-	-	-
Stack 6, Step	4%	32%	-	-	-

Figure 23: Comparison of our method against baselines. The value reported is the median of the best performance (success rate) of all randomly seeded runs of each method.

Note the step reward is still very sparse; the robot only sees the reward change when it moves a block into its target location. We subtract 1 only to make the reward more interpretable, as in the initial state the first block is already at its target.

Regardless of the reward type, an episode is considered successful for computing success rate if all blocks are at their goal position in their final state.

5.7.3 Network architectures

We use a 4 layer networks with 256 hidden units per layer for π and Q for the HER tasks and stacking 3 or fewer blocks. For stacking 4 blocks or more, we use an attention mechanism (Bahdanau et al., 2015) for the actor and a larger network. The attention mechanism uses a 3 layer network with 128 hidden units per layer to query the states and goals with one shared head. Once a state and goal is extracted, we use a 5 layer network with 256 hidden units per layer after the attention mechanism. Attention speeds up training slightly but does not change training outcomes.

5.7.4 Baselines

We include the following methods to compare our method to baselines on stacking 2 to 6 blocks.²

² Because of computational constraints, we were limited to 5 random seeds per method for stacking 3 blocks, 2 random seeds per method for stacking 4 and 5 blocks, and 1 random seed per method for stacking 6 blocks. Although we are careful to draw conclusions from few random seeds, the results are consistent with our collective experience training these models. We report the median of the random seeds

Ours: Refers to our method as described in section 5.4.3.

Ours, Resets: Refers to our method as described in section 5.4.3 with resets from demonstration states (Sec. 5.4.4).

BC: This method uses behavior cloning to learn a policy. Given the set of demonstration transitions R_D , we train the policy π by supervised learning. Behavior cloning requires much less computation than RL. For a fairer comparison, we performed a large hyperparameter sweep over various networks sizes, attention hyperparameters, and learning rates and report the success rate achieved by the best policy found.

HER: This method is exactly the one described in Hindsight Experience Replay ([Andrychowicz et al., 2017b](#)), using HER and DDPG.

BC+HER: This method first initializes a policy (actor) with BC, then finetunes the policy with RL as described above.

5.7.5 Results

We are able to learn much longer horizon tasks than the other methods, as shown in Fig. 23. The stacking task is extremely difficult using HER without demonstrations because the chance of grasping an object using random actions is close to 0. Initializing a policy with demonstrations and then running RL also fails since the actor updates depend on a reasonable critic and although the actor is pretrained, the critic is not. The pretrained actor weights are therefore destroyed in the very first epoch, and the result is no better than BC alone. We attempted variants of this method where initially the critic was trained from replay data. However, this also fails without seeing on-policy data.

The results with sparse rewards are very encouraging. We are able to stack 3 blocks with a fully sparse reward without resetting to the states from demonstrations, and 4 blocks with a fully sparse reward if we use resetting. With resets from demonstration states and the step reward, we are able to learn a policy to stack 6 blocks.

5.8 ABLATION EXPERIMENTS

In this section we perform a series of ablation experiments to measure the importance of various components of our method. We evaluate our method on stacking 3 to 6 blocks.

We perform the following ablations on the best performing of our models on each task:

everywhere applicable.

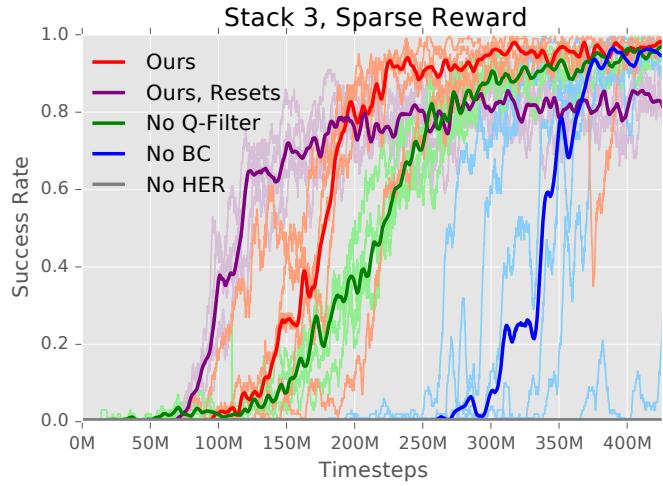


Figure 24: Ablation results on stacking 3 blocks with a fully sparse reward. We run each method 5 times with random seeds. The bold line shows the median of the 5 runs while each training run is plotted in a lighter color. Note “No HER” is always at 0% success rate. Our method without resets learns faster than the ablations. Our method with resets initially learns faster but converges to a worse success rate.

No BC Loss: This method does not apply the behavior cloning gradient during training. It still has access to demonstrations through the demonstration replay buffer.

No Q-Filter: This method uses standard behavioral cloning loss instead of the loss from equation Eq. 19, which means that the actor tries to mimic the demonstrator’s behaviour regardless of the critic.

No HER: Hindsight Experience Replay is not used.

5.8.1 Behavior Cloning Loss

Without the behavior cloning loss, the method is significantly worse in every task we try. Fig. 24 shows the training curve for learning to stack 3 blocks with a fully sparse reward. Without the behavior cloning loss, the system is about 2x slower to learn. On longer horizon tasks, we do not achieve any success without this loss.

To see why, consider the training curves for stacking 4 blocks shown in Fig. 25. The “No BC” policy learns to stack only one additional block. Without the behavior cloning loss, the agent only has access to the demonstrations through the demonstration replay buffer. This allows it to view high-reward states and incentivizes the agent to stack more blocks, but there is a stronger disincentive: stacking the tower higher is risky and could

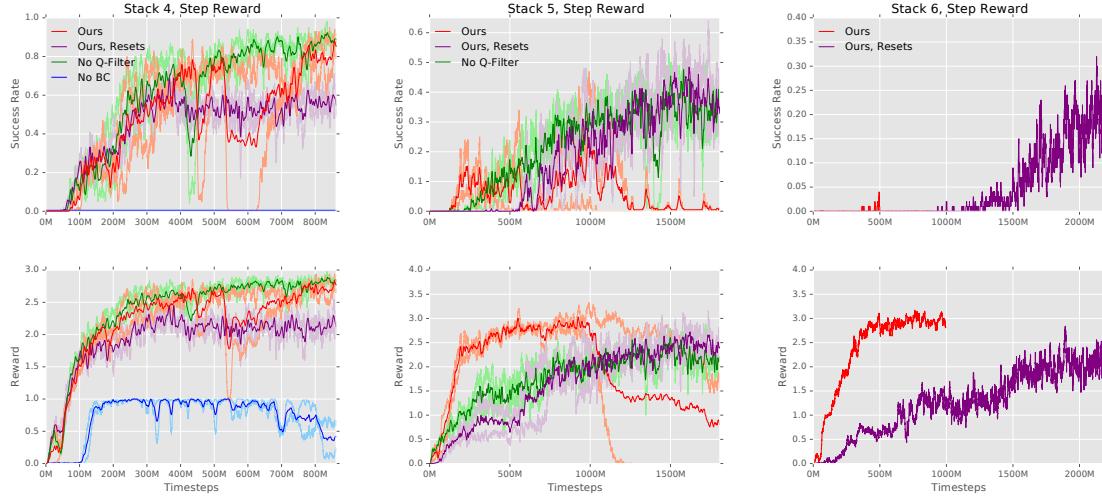


Figure 25: Ablation results on longer horizon tasks with a step reward. The upper row shows the success rate while the lower row shows the average reward at the final step of each episode obtained by different algorithms. For stacking 4 and 5 blocks, we use 2 random seeds per method. The median of the runs is shown in bold and each training run is plotted in a lighter color. Note that for stacking 4 blocks, the “No BC” method is always at 0% success rate. As the number of blocks increases, resets from demonstrations becomes more important to learn the task.

result in lower reward if the agent knocks over a block that is already correctly placed. Because of this risk, which is fundamentally just another instance of the agent finding a local optimum in a shaped reward, the agent learns the safer behavior of pausing after achieving a certain reward. Explicitly weighting behavior cloning steps into gradient updates forces the policy to continue the task.

5.8.2 Q-Filter

The Q-Filter is effective in accelerating learning and achieving optimal performance. Fig. 24 shows that the method without filtering is slower to learn. One issue with the behavior cloning loss is that if the demonstrations are suboptimal, the learned policy will also be suboptimal. Filtering by Q-value gives a natural way to anneal the effect of the demonstrations as it automatically disables the BC loss when a better action is found. However, it gives mixed results on the longer horizon tasks. One explanation is that in the step reward case, learning relies less on the demonstrations because the reward signal is stronger. Therefore, the training is less affected by suboptimal demonstrations.

5.8.3 Resets From Demonstrations

We find that initializing rollouts from within demonstration states greatly helps to learn to stack 5 and 6 blocks but hurts training with fewer blocks, as shown in Fig. 25. Note that even where resets from demonstration states helps the final success rate, learning takes off faster when this technique is not used. However, since stacking the tower higher is risky, the agent learns the safer behavior of stopping after achieving a certain reward. Resetting from demonstration states alleviates this problem because the agent regularly experiences higher rewards.

This method changes the sampled state distribution, biasing it towards later states. It also inflates the Q values unrealistically. Therefore, on tasks where the RL algorithm does not get stuck in solving a subset of the full problem, it could hurt performance.

5.9 DISCUSSION AND FUTURE WORK

We present a system to utilize demonstrations along with reinforcement learning to solve complicated multi-step tasks. We believe this can accelerate learning of many tasks, especially those with sparse rewards or other difficulties in exploration. Our method is very general, and can be applied on any continuous control task where a success criterion can be specified and demonstrations obtained.

An exciting future direction is to train policies directly on a physical robot. Fig. 22 shows that learning the pick-and-place task takes about 1 million timesteps, which is about 6 hours of real world interaction time. This can realistically be trained on a physical robot, short-cutting the simulation-reality gap entirely. Many automation tasks found in factories and warehouses are similar to pick-and-place but without the variation in initial and goal states, so the samples required could be much lower. With our method, no expert needs to be in the loop to train these systems: demonstrations can be collected by users without knowledge about machine learning or robotics and rewards could be directly obtained from human feedback.

A major limitation of this work is sample efficiency on solving harder tasks. While we could not solve these tasks with other learning methods, our method requires a large amount of experience which is impractical outside of simulation. To run these tasks on physical robots, the sample efficiency will have to improved considerably. We also require demonstrations which are not easy to collect for all tasks. If demonstrations are not available but the environment can be reset to arbitrary states, one way to learn goal-reaching but avoid using demonstrations is to reuse successful rollouts as in (Florensa

et al., 2018b).

Finally, our method of resets from demonstration states requires the ability to reset to arbitrary states. Although we can solve many long-horizon tasks without this ability, it is very effective for the hardest tasks. Resetting from demonstration rollouts resembles curriculum learning: we solve a hard task by first solving easier tasks. If the environment does not afford setting arbitrary states, then other curriculum methods will have to be used.

6

RESIDUAL REINFORCEMENT LEARNING FOR ROBOT CONTROL

6.1 INTRODUCTION

Robots in today’s manufacturing environments typically perform repetitive tasks, and often lack the ability to handle variability and uncertainty. Commonly used control algorithms, such as PID regulators and the computed torque method, usually follow predefined trajectories with little adaptive behavior. Many manufacturing tasks require some degree of adaptability or feedback to the environment, but significant engineering effort and expertise is required to design feedback control algorithms for these industrial robots. The engineering time for fine tuning such a controller might be similar in cost to the robot hardware itself. Being able to quickly and easily design feedback controllers for industrial robots would significantly broaden the space of manufacturing tasks that can be automated by robots.

Why is designing a feedback controller for many tasks hard with classical methods? While conventional feedback control methods can solve tasks such as path following efficiently, applications that involve contacts between the robot and its environment are difficult to approach with conventional control methods. Identifying and characterizing contacts and friction is difficult—even if a physical model provides reasonable contact behavior, identifying the physical parameters of a contact interaction accurately is very hard. Hence, it is often difficult to achieve adaptable yet robust control behavior, and significant control tuning effort is required as soon as these elements are introduced. Another drawback of conventional control methods is their lack of behavior generalization. Thus, all possible system behaviors must be considered a priori at design time.

Reinforcement learning (RL) methods hold the promise of solving these challenges because they allow agents to learn behaviors through interaction with their surrounding environments and ideally generalize to new scenarios that differ from the specifications at the control design stage. Moreover, RL can handle control problems that are difficult

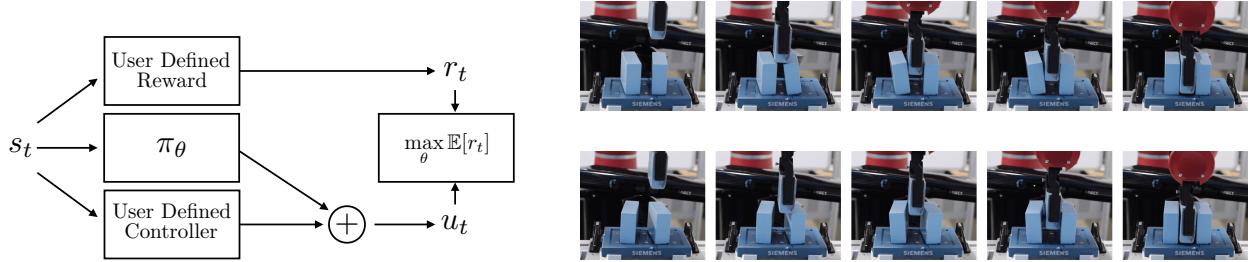


Figure 26: We train an agent directly in the real world to solve a model assembly task involving contacts and unstable objects. An outline of our method, which consists of combining hand-engineered controllers with a residual RL controller, is shown on the left. Rollouts of residual RL solving the block insertion task are shown on the right. Residual RL is capable of learning a feedback controller that adapts to variations in the orientations of the standing blocks and successfully completes the task of inserting a block between them. Videos are available at residualrl.github.io

to approach with conventional controllers because the control goal can be specified indirectly as a term in a reward function and not explicitly as the result of a control action. All of these aspects are considered enablers for truly autonomous manufacturing systems and important for fully flexible lot-size one manufacturing (Kannengiesser et al., 2017). However, standard RL methods require the robot learn through interaction, which can be unsafe initially, and collecting the amount of interaction that is needed to learn a complex skill from scratch can be time consuming.

In this paper, we study control problems that are difficult to approach with conventional feedback control methods. However, the problems possess structure that can be partially handled with conventional feedback control, e.g. with impedance control. The residual part of the control task, which is the part that must consider contacts and external object dynamics, is solved with RL. The outputs of the conventional controller and RL are superposed to form the commanded control. The main contribution of this paper is a methodology that combines conventional feedback control with deep RL methods and is illustrated in Fig. 26. Our main motivation is a control approach that is suitable for real-world control problems in manufacturing, where the exploratory behavior of RL is a safety concern and the data requirements of deep RL can be expensive. We provide a thorough evaluation of our method on a block assembly task in simulation and on physical hardware. When the initial orientation of the blocks is noisy, our hand-designed controller fails to solve the task, while residual RL successfully learns to perform the task in under 3 hours. This suggests that we could usefully apply our method to practical manufacturing problems.

6.2 PRELIMINARIES

In this section, we set up our problem and summarize the foundations of classical control and reinforcement learning that we build on in our approach.

6.2.1 Problem Statement - System Theoretic Interpretation

The class of control problems that we are dealing with in this paper can be viewed from a dynamical systems point of view as follows. Consider a dynamical system that consists of a fully actuated robot and underactuated objects in the robot's environment. The robot and the objects in its environment are described by their states s_m and s_o , respectively. The robot can be controlled through the control input u while the objects cannot be directly controlled. However, the robot's states are coupled with the objects' states so that indirect control of s_o is possible through u . This is for example the case if the agent has large inertia and is interacting with small parts as is common in manufacturing. The states of agent and objects can either be fully observable or they can be estimated from measurements.

The time-discrete equations of motion of the overall dynamical system comprise the robot and objects and can be stated as

$$s_{t+1} = \begin{bmatrix} s_{m,t+1} \\ s_{o,t+1} \end{bmatrix} = \begin{bmatrix} A(s_{m,t}) & 0 \\ B(s_{m,t}, s_{o,t}) & C(s_{o,t}) \end{bmatrix} \begin{bmatrix} s_{m,t} \\ s_{o,t} \end{bmatrix} + D \begin{bmatrix} u_t \\ 0 \end{bmatrix}, \quad (23)$$

where the states can also be subject to algebraic constraints, which we do not state explicitly here.

The type of control objectives that we are interested in can be summarized as controlling the agent in order to manipulate the objects while also fulfilling a geometric objective such as trajectory following. It is difficult to solve the control problem directly with conventional feedback control approaches, which compute the difference between a desired and a measured state variable. In order to achieve best system performance feedback control methods require well understood and modeled state transition dynamics. Finding the optimal control parameters can be difficult or even impossible if the system dynamics are not fully known.

In equation 23 the state transition matrices although $A(s_m)$ and $C(s_o)$ are usually known to a certain extent, because they represent rigid body dynamics, the coupling matrix $B(s_m, s_o)$ is usually not known. Physical interactions such as contacts and friction

forces are the dominant effects that $B(s_m, s_o)$ needs to capture, which also applies to algebraic constraints, which are functions of s_m and s_o as well. Hence, conventional feedback control synthesis for determining u to control s_o is very difficult, and requires trial and error in practice. Another difficulty for directly designing feedback controllers is due to the fact that, for many control objectives, the states s_o need to fulfill conditions that cannot be expressed as deviations (errors) from desired states. This is often the case when we only know the final goal rather than a full trajectory.

Instead of directly designing a feedback control system, we can instead specify the goal via a reward function. These reward functions can depend on both s_m and s_o , where the terms that depend on s_m are position related objectives.

6.2.2 Interpretation as a Reinforcement Learning Problem

In reinforcement learning, we consider the standard Markov decision process framework for picking optimal actions to maximize rewards over discrete timesteps in an environment E . At every timestep t , an agent is in a state s_t , takes an action u_t , receives a reward r_t , and E evolves to state s_{t+1} . In reinforcement learning, the agent must learn a policy $u_t = \pi(s_t)$ to maximize expected returns. We denote the return by $R_t = \sum_{i=t}^T \gamma^{(i-t)} r_i$, where T is the horizon that the agent optimizes over and γ is a discount factor for future rewards. The agent's objective is to maximize expected return from the start distribution $J = \mathbb{E}_{r_i, s_i \sim E, a_i \sim \pi}[R_0]$.

Unlike the previous section, RL does not attempt to model the unknown coupled dynamics of the agent and the object. Instead, it finds actions that maximizes rewards, without making any assumptions about the system dynamics. In this paper, we use value-based RL methods. These methods estimate the state-action value function:

$$Q^\pi(s_t, u_t) = \mathbb{E}_{r_i, s_i \sim E, u_i \sim \pi}[R_t | s_t, u_t] \quad (24)$$

$$= \mathbb{E}_{r_t, s_{t+1} \sim E}[r_t + \gamma \mathbb{E}_{u_{t+1} \sim \pi}[Q^\pi(s_{t+1}, u_{t+1})]] \quad (25)$$

Equation 25 is a recursive version of Equation 24, and is known as the Bellman equation. The Bellman equation allows us to estimate Q via approximate dynamic programming. Value-based methods can be learned off-policy, making them very sample efficient, which is vital for real-world robot learning.

6.3 METHOD

Based on the analysis in Sec. 10.3, we introduce a control system that consists of two parts. The first part is based on conventional feedback control theory and maximizes all reward terms that are functions of s_m . An RL method is superposed and maximizes the reward terms that are functions of s_o .

6.3.1 Residual Reinforcement Learning

In most robotics tasks, we consider rewards of the form:

$$r_t = f(s_m) + g(s_o). \quad (26)$$

The term $f(s_m)$ is assumed to be a function, which represents a geometric relationship of robot states, such as a Euclidean distance or a desired trajectory. The second term of the sum $g(s_o)$ can be a general class of functions. Concretely, in our model assembly task, $f(s_m)$ is the reward for moving the robot gripper between the standing blocks, while $g(s_o)$ is the reward for keeping the standing blocks upright and in their original positions.

The key insight of residual RL is that in many tasks, $f(s_m)$ can be easily optimized a priori of any environment interaction by conventional controllers, while $g(s_o)$ may be easier to learn with RL which can learn fine-grained hand-engineered feedback controllers even with friction and contacts. To take advantage of the efficiency of conventional controllers but also the flexibility of RL, we choose:

$$u = \pi_H(s_m) + \pi_\theta(s_m, s_o) \quad (27)$$

as the control action, where $\pi_H(s_m)$ is the human-designed controller and $\pi_\theta(s_m, s_o)$ is a learned policy parametrized by θ and optimized by an RL algorithm to maximize expected returns on the task.

Inserting equation 32 into equation 23 one can see that a properly designed feedback control law for $\pi_H(s_m)$ is able to provide exponentially stable error dynamics of s_m if the learned controller π_θ is neglected and the sub statespace is stabilizable. This is equivalent to maximizing equation 26 for the case f represents errors between actual and desired states.

The residual controller $\pi_\theta(s_m, s_o)$ can now be used to maximize the reward term $g(s_o)$ in equation 26. Since the control sequence equation 32 enters equation 23 through the

dynamics of s_m and s_m is in fact the control input to the dynamics of s_o , we cannot simply use the a-priori hand-engineered feedback controller to achieve zero error of s_m and independently achieve the control objective on s_o . Through the coupling of states we need to perform an overall optimization of equation 32, whereby the hand-engineered feedback controller provides internal structures and eases the optimization related to the reward term $f(s_m)$.

Algorithm 4 Residual reinforcement learning

Require: policy π_θ , hand-engineered controller π_H .

- 1: **for** $n = 0, \dots, N - 1$ episodes **do**
 - 2: Initialize random process \mathcal{N} for exploration
 - 3: Sample initial state $s_0 \sim E$.
 - 4: **for** $t = 0, \dots, H - 1$ steps **do**
 - 5: Get policy action $u_t = \pi_\theta(s_t) + \mathcal{N}_t$.
 - 6: Get action to execute $u'_t = u_t + \pi_H(s_t)$.
 - 7: Get next state $s_{t+1} \sim p(\cdot | s_t, u'_t)$.
 - 8: Store (s_t, u_t, s_{t+1}) into replay buffer \mathcal{R} .
 - 9: Sample set of transitions $(s, u, s') \sim \mathcal{R}$.
 - 10: Optimize θ using RL with sampled transitions.
 - 11: **end for**
 - 12: **end for**
-

6.3.2 Method Summary

Our method is summarized in Algorithm 5. The key idea is to combine the flexibility of RL with the efficiency of conventional controllers by additively combining a learnable parametrized policy with a fixed hand-engineered controller.

As our underlying RL algorithm, we use a variant of twin delayed deep deterministic policy gradients (TD3) as described in Fujimoto et al., 2018b. TD3 is a value-based RL algorithm for continuous control based off of the deep deterministic policy gradient (DDPG) algorithm Lillicrap et al., 2016. We have found that TD3 is stable, sample-efficient, and requires little manual tuning compared to DDPG. We used the publicly available `rlkit` implementation of TD3 V. Pong et al., 2018. Our method is independent of the choice of RL algorithm, and we could apply residual RL to any other RL algorithm.

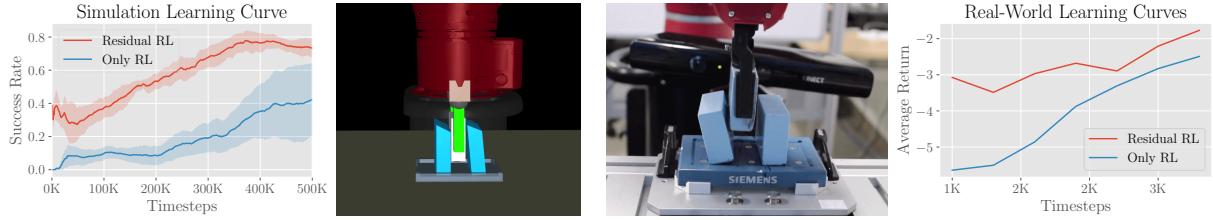


Figure 27: Block assembly task in simulation (left) and real-world (right). The task is to insert a block between the two blocks on the table without moving the blocks or tipping them over. In the learning curves, we compare our method with RL without any hand-engineered controller¹. In both simulation and real-world experiments, we see that residual RL learns faster than RL alone, while achieving better performance than the hand-engineered controller.

6.4 EXPERIMENTAL SETUP

We evaluate our method on the task shown in Fig. 27, both in simulation and in the real world. This section introduces the details of the experimental setup and provides an overview of the experiments.

6.4.1 Simulated Environment

We use MuJoCo (Todorov et al., 2012), a full-featured simulator for model-based optimization considering body contacts, to evaluate our method in simulation. This environment consists of a simulated Sawyer robot arm with seven degrees of freedom and a parallel gripper. We command the robot with a Cartesian-space position controller.

6.4.2 Real-World Environment

The real-world environment is largely the same as the simulated environment, except for the controller, rewards, and observations. We command the robot with a compliant joint-space impedance controller we have developed to be smooth and tolerant of contacts. The positioning of the block being inserted is similar to the simulation but the observation is estimated from a camera-based tracking system as we do not have access to ground truth position information.

¹ In all simulation plots, we use 10 random seeds and report a 95% confidence interval for the mean.

6.4.3 Overview of Experiments

In our experiments we evaluate the following research questions:

1. Does incorporating a hand-designed controller improve the performance and sample-efficiency of RL algorithms, while still being able to recover from an imperfect hand-designed controller?
2. Can our method allow robots to be more tolerant of variation in the environment?
3. Can our method successfully control noisy systems, compared to classical control methods?

6.5 EXPERIMENTS

6.5.1 Sample Efficiency of Residual RL

In this section, we compare our residual RL method with the human controller alone and RL alone. The following methods are compared:

1. Only RL: using the same underlying RL algorithm as our method but without adding a hand-engineered policy
2. Residual RL: our method which trains a superposition of the hand-engineered controller and a neural network policy, with RL

6.5.2 Effect of Environment Variation

In automation, environments can be subject to noise and solving manufacturing tasks become more difficult as variability in the environment increases. It is difficult to manually design feedback controllers that are robust to environment variation, as it might require significant human expertise and tuning. In this experiment, we vary the initial orientation of the blocks during each episode and demonstrate that residual RL can still solve the task. We compare its performance to that of the hand-engineered controller.

To introduce variability in simulation, on every reset we sampled the rotation of each block independently from a uniform distribution $U[-r, r], r \in \{0, 0.05, 0.1, 0.15, 0.2, 0.25, 0.3\}$.

Similarly, in the real world experiments, on every reset we randomly rotated each block to one of three orientations: straight, tilted clockwise, or tilted counterclockwise (tilt was $\pm 20^\circ$ from original position).

6.5.3 Recovering from Control Noise

Due to a host of issues, such as defective hardware or poorly tuned controller parameters, feedback controllers might have induced noise. Conventional feedback control policies are determined a priori and do not adapt their behavior from data. However, RL methods are known for their ability to cope with shifting noise distributions and are capable of recovering from such issues.

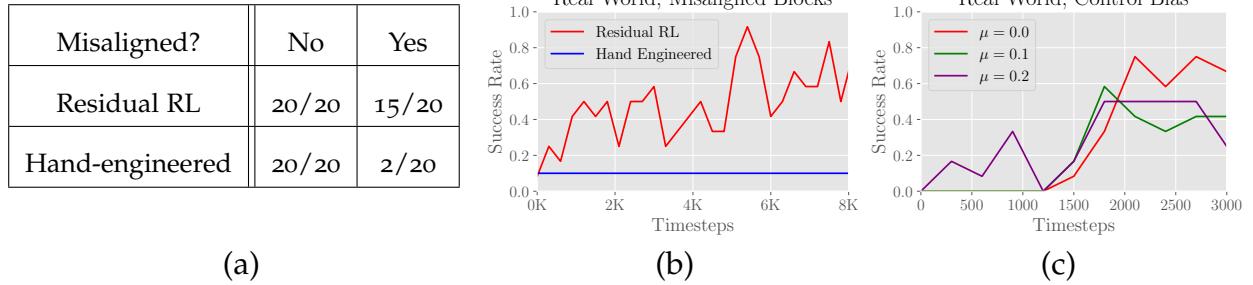


Figure 28: Outcome of our residual RL method in different experiments during the block assembly task in the real-world. Success rate is recorded manually by human judgment of whether the blocks stayed upright and ended in the correct position. Plot (a) compares the insertion success of residual RL and hand-designed controller depending on the block orientation during run time. Plot (b) shows the success rate of the insertion process during training, where on every reset the blocks are randomly rotated: straight, tilted clockwise, or tilted counterclockwise ($\pm 20^\circ$) and plot (c) shows the increasing success rate of our method for biased controllers as well even as control bias increases.

In this experiment, we introduce a control noise, including biased control noise, and demonstrate that residual RL can still successfully solve the task, while a hand-engineered controller cannot. The control noise follows a normal distribution and is added to the control output of the system at every step:

$$u'_t = u_t + \mathcal{N}(\mu, \sigma^2) \quad (28)$$

To test tolerance to control noise, we set $\mu = 0$ and vary $\sigma \in [0.01, 0.1]$. In theory, RL could adapt to a noisy controller by learning more robust solutions to the task which are less sensitive to perturbations.

Furthermore, to test tolerance to a biased controller, we set $\sigma = 0.05$ and vary $\mu \in [0, 0.2]$. To optimize the task reward, RL can learn to simply counteract the bias.

6.5.4 Sim-to-Real with Residual RL

As an alternative to analytic solutions of real-world control problems, we can often instead model the forward dynamics of the problem (ie. a simulator). With access to such a model, we can first find a solution to the problem with our possibly inaccurate model, and then use residual RL to find a realistic control solution in the real world.

In this experiment, we attempt the block insertion task with the side blocks fixed in place. The hand-engineered policy π_H in this case comes from training a parametric policy in simulation of the same scenario (with deep RL). We then use this policy as initialization for residual RL in the real world.

6.6 RESULTS

We trained our method to optimize the insertion process in simulation as well as on physical hardware. This section provides the results of our discussed experiments and shows the functionality of our method.

6.6.1 Sample Efficiency of Residual RL

First, we compare residual RL and pure RL without a hand-engineered controller on the insertion task. Fig. 27 shows in simulation and real-world that residual RL achieves a better final performance and requires less samples than RL alone, both in simulation and on physical hardware. Unlike residual RL, the pure RL approach needs to learn the structure of the position control problem from scratch, which explains the difference in sample efficiency. As samples are expensive to collect in the real world, residual RL is better suited for solving real-world tasks. Moreover, RL shows a broader spatial variance during training and needs to explore a wider set of states compared to residual RL, which can be potentially dangerous in hardware deployments.

6.6.2 Effect of Environment Variation

In previous set of experiments, both standing blocks were placed in their initial position without any position or orientation error. In this case, the hand-engineered controller performs well, as both blocks are placed such that there is a sufficiently large defined gap for insertion. However, once the initial orientation of the blocks is randomized, the gap between the blocks and the goal position does not afford directly inserting from

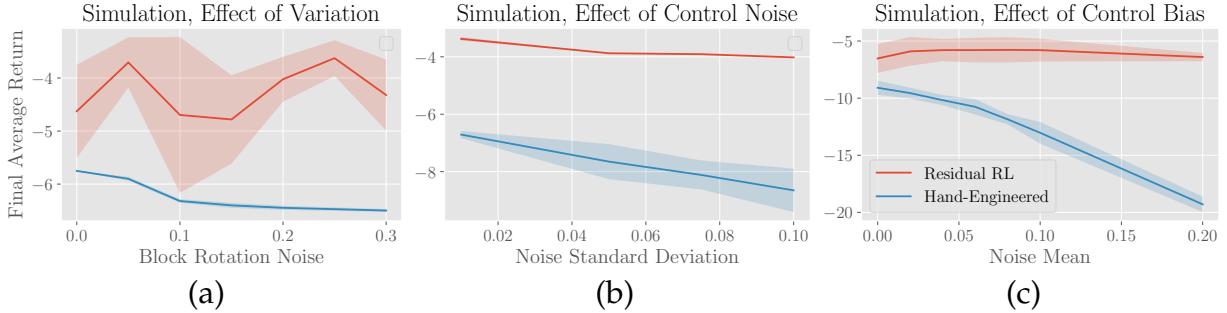


Figure 29: Simulation results for different experiments. In each plot, the final average return obtained by running the method for various settings of a parameter is shown. Plot (a) shows that residual RL can adjust to noise in the environment caused by rotation of the blocks in a range of 0 to 0.3 rad. In plot (b), residual RL finds robust strategies in order to reduce the effect of control noise, as the final average return is not greatly affected by the magnitude of noise. Plot (c) shows that residual RL can compensate for biased controllers and maintains good performance as control bias increases, while the performance of the hand-designed controller dramatically deteriorates with higher control bias.

above. Therefore, the hand-engineered controller struggles to solve the insertion task, succeeding in only 2/20 trials, while residual RL still succeeds in 15/20 trials. These results are summarized in Fig. 28 (a). Rollouts from the learned policy are included in 26. In this experiment, the agent demonstrably learns consistent small corrective feedback behaviors in order to slightly nudge the blocks in the right direction without tipping them over, a behavior that is very difficult to manually specify.

The result of this experiment showcases the strength of residual RL. Since the human controller specifies the general trajectory of the optimal policy, environment samples are required only to learn this corrective feedback behavior. The real-world learning curve for the experiment in Fig. 28 (b) shows that this behavior is gradually acquired over the course of eight thousand samples, which is only about three hours of real-world training time.

We further studied the effect of the block orientation changing after every reset in simulation. The results are shown in Fig. 29 (a). The simulation results show that the performance of the hand-engineered controller decreases as the block rotation angle increases, whereas our control method maintains a constant average performance over different variations.

6.6.3 Recovering from Control Noise

In this experiment, we observe that residual RL is able to cope with actuator noise, including biased actuator noise. Quantitative results for simulation are shown in Fig. 29 (b) and (c). In Fig. 29 (c) our method keeps the average return constant and correct for biased controllers even as control bias increases, whereas the hand-engineered controller cannot compensate biased input and its performance deteriorates as control bias increases. The same applies for adding control noise to the control output as shown in Fig. 29 (b).

For the hardware experiments, only biased actuator noise is investigated. These results are shown in Fig. 28 (c). These learning curves show that even as more control bias is introduced, training in the real world proceeds without significant issues. This result suggests the potential for RL to address practical issues in automation such as sensor drift.

6.6.4 Sim-to-Real with Residual RL

The result of the sim-to-real experiment is shown in Fig. 30. In this experiment, each setting was run with three random seeds. Adding policy initialization from simulation significantly speeds up both RL and residual RL. In particular, residual RL with policy initialization from simulation successfully solves the task extremely quickly: in under one thousand timesteps of environment interaction. This method poses a highly sample efficient, practical way to solve robotics problems with difficult contact dynamics.

6.7 RELATED WORK

Reinforcement learning for robotics holds the promise of greater autonomy and reliability, which could be vital to improving our manufacturing processes beyond its current limitations. RL methods have been difficult to apply in robotics because of sample efficiency, safety, and stability issues. Still, RL has been used to allow robots to learn tasks such as playing table tennis (Peters et al., 2010), swinging up a cartpole and balancing a unicycle (Marc Peter Deisenroth and Rasmussen, 2011), grasping (Pinto et al., 2017; Levine et al., 2017), opening a door (S. Gu et al., 2017), and general manipulation tasks (Levine et al., 2016a; Haarnoja et al., 2018a). RL, particularly deep RL, tends to be data-hungry; even learning simple tasks can require many hours of interaction. To bring these methods into factories and warehouses, they must be able to consistently solve complex tasks, multi-step tasks. One way to enable these methods to solve these complex tasks is

to introduce prior human knowledge into the learning system, as our method does.

Prior work in RL has incorporated human prior knowledge for solving tasks in various ways. One such way is reward shaping (Andrew Y. Ng et al., 1999), where additional rewards auxiliary to the real objective are included in order to guide the agent towards the desired behavior. Reward shaping can effectively encode a policy. For example, to train an agent to perform block stacking, each step can be encoded into the reward (Popov et al., 2017). Often, intensive reward shaping is key for RL to succeed at a task and prior work has even considered reward shaping as part of the learning system (Daniel et al., 2014). Reward shaping in order to tune agent behavior is a very manual process and recovering a good policy with reward shaping can be as difficult as specifying the policy itself. Hence, in our method we allow for human specification of both rewards and policies—whichever might be more practical for a particular task.

Further work has incorporated more specialized human knowledge into RL systems. One approach is to use trajectory planners in RL in order to solve robotics tasks (G. Thomas et al., 2018). However, since the method optimizes trajectory following instead of the task reward, generalization can be difficult when aspects of the environment change. Other work has focused on human feedback (Loftin et al., 2014; Saunders et al., 2018; Torrey and Taylor, 2013; Frank et al., 2008; Christiano et al., 2017) to inform the agent about rewards or to encourage safety. However, in many robotics tasks, providing enough information about the task through incremental human feedback is difficult.

Another way to include prior knowledge in RL is through demonstrations (Peters and Schaal, 2008c; Kober and Peter, 2008; Rajeswaran et al., 2018; Hester et al., 2018; Večerík et al., 2017; A. Nair et al., 2018a). Demonstrations can substantially simplify the exploration problem as the agent begins training having already received examples of high-reward transitions and therefore knows where to explore (Subramanian et al., 2016). However, providing demonstrations requires humans to be able to teleoperate the robot to perform the task. In contrast, our method only requires a conventional controller for

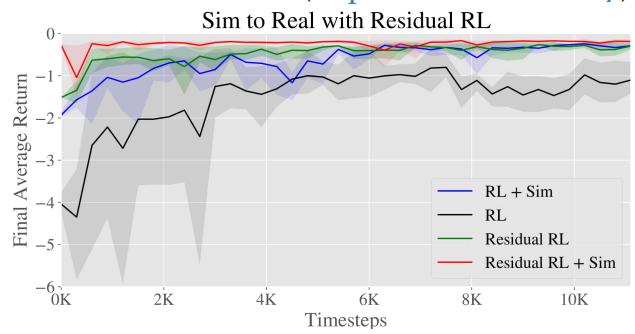


Figure 30: Real-world block insertion results using residual RL for sim-to-real transfer. “Sim” indicates that the real-world policy was initialized by reinforcement learning in simulation. Residual RL with simulation initialization successfully solves the task with little environment experience required.

motion, which ships with most robots.

Prior knowledge can also be induced through neural network architecture choices. Deep residual networks with additive residual blocks achieved state of the art results in many computer vision tasks by enabling training of extremely deep networks (He et al., 2016). In RL, structured control nets showed improvement on several tasks by splitting the policy into a linear module and a non-linear module (Srouji et al., 2018). Prior work in adaptive flight control has also considered compensating linearized controllers with neural networks (Johnson and Calise, 2000). Most closely related to our work, residual policy learning concurrently and independently explores training a residual policy in the context of simulated long-horizon, sparse-reward tasks with environment variation and sensor noise (T. Silver et al., 2018). Our work instead focuses on achieving practical real-world training of contact-intensive tasks.

6.8 CONCLUSION

In this paper we study the combination of conventional feedback control methods with deep RL. We presented a control method that utilizes conventional feedback control along with RL to solve complicated manipulation tasks involving friction and contacts with unstable objects. We believe this approach can accelerate learning of many tasks, especially those where the control problem can be solved in large part by prior knowledge but requires some model-free reasoning to solve perfectly. Our results demonstrate that the combination of conventional feedback control and RL can circumvent the disadvantages of both and provide a sample efficient controller that can cope with contact dynamics.

7

DEEP REINFORCEMENT LEARNING FOR INDUSTRIAL INSERTION TASKS WITH VISUAL INPUTS AND NATURAL REWARDS

7.1 INTRODUCTION

Many industrial tasks on the edge of automation require a degree of adaptability that is difficult to achieve with conventional robotic automation techniques. While standard control methods, such as PID controllers, are heavily employed to automate many tasks in the context of positioning, tasks that require significant adaptability or tight visual perception-control loops are often beyond the capabilities of such methods, and therefore are typically performed manually. Standard control methods can struggle in presence of complex dynamical phenomena that are hard to model analytically, such as complex contacts. Reinforcement learning (RL) offers a different solution, relying on trial and error learning instead of accurate modeling to construct an effective controller. RL with expressive function approximation, i.e. deep RL, has further shown to automatically handle high dimensional inputs such as images (Mnih et al., 2013).

However, deep RL has thus far not seen wide adoption in the automation community due to several practical obstacles. Sample efficiency is one obstacle: tasks must be completed without excessive interaction time or wear and tear on the robot. Progress in recent years on developing better RL algorithms has led to significantly better sample efficiency, even in dynamically complicated tasks (Haarnoja et al., 2018a; Hessel et al., 2018), but remains a challenge for deploying RL in real-world robotics contexts. Another major, often underappreciated, obstacle is goal specification: while prior work in RL assumes a reward signal to optimize, it is often carefully shaped to allow the system to learn (Andrew Y. Ng et al., 1999; Popov et al., 2017; Daniel et al., 2014). Obtaining such dense reward signals can be a significant challenge, as one must additionally build a perception system that allows computing dense rewards on state representations. Shaping a

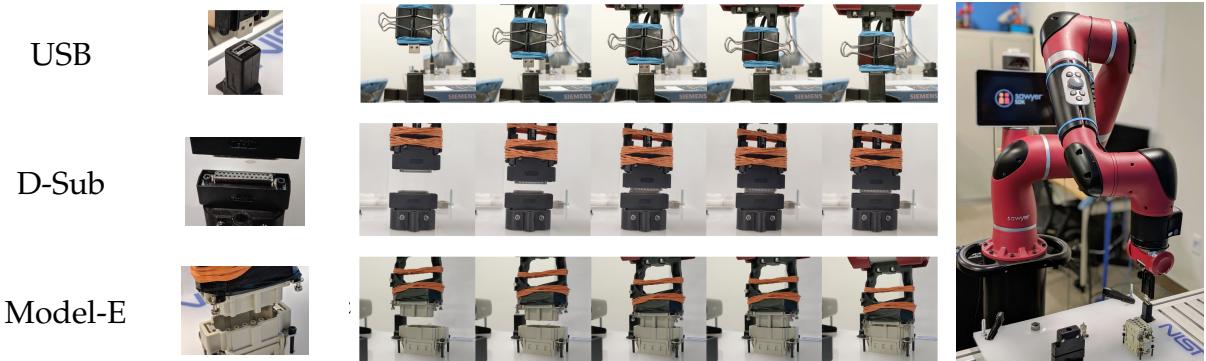


Figure 31: We train policies directly in the real world to solve connector insertion tasks from raw pixel input and without access to ground-truth state information for reward functions. Left: top-down views of the connectors. Middle: a rollout from a learned policy that successfully completes the insertion task for each connector is shown. Right: a full view of the robot setup. Videos of the results are available at industrial-insertion-rl.github.io

reward function so that an agent can learn from it is also a manual process that requires considerable manual effort. An ideal RL system would learn from rewards that are natural and easy to specify. How can we enable robots to autonomously perform complex tasks without significant engineering effort to design perception and reward systems?

We first consider an end-to-end approach that learns a policy from images, where the images serve as both the state representation and the goal specification. Using goal images is not fully general, but can successfully represent tasks when the task is to reach a final desired state (A. Nair et al., 2018b). Specifying goals via goal images is convenient, and makes it possible to specify goals with minimal manual effort. Using images as the state representation also allows a robot to learn behaviors that utilize direct visual feedback, which provides some robustness to sensor and actuator noise.

Secondly, we consider learning from simple and sparse reward signals. Sparse rewards can often be obtained conveniently, for instance from human-provided labels or simple instrumentation. In many electronic assembly tasks, which we consider here, we can directly detect whether the electronics are functional, and use that signal as a reward. Learning from sparse rewards poses a challenge, as exploration with sparse reward signals is difficult, but by using sufficient prior information about the task, one can overcome this challenge. To handle this challenge, we extend the residual RL approach Johannink et al., 2019; T. Silver et al., 2018, which learns a parametric policy on top of a fixed, hand-specified controller, to the setting of vision-based manipulation.

In our experiments, we show that we can successfully complete real-world tight tolerance assembly tasks, such as inserting USB connectors, using RL from images with

reward signals that are convenient for users to specify. We can learn from only a sparse reward based on the electrical connection for a USB adapter plug, and we demonstrate learning insertion skills with rewards based only on goal images. These reward signals require no extra engineering and are easy to specify for many tasks. Beyond showing the feasibility of RL for solving these tasks, we evaluate multiple RL algorithms across three tasks and study their robustness to imprecise positioning and noise.

7.2 RELATED WORK

Learning has been applied previously in a variety of robotics contexts. Different forms of learning have enabled autonomous driving (Pomerleau, 1989), biped locomotion (Nakanishi et al., 2004), block stacking (Marc Peter Deisenroth et al., 2011b), grasping (Pinto and Abhinav Gupta, 2016), and navigation (Giusti et al., 2015; Pathak et al., 2018). Among these methods, many involve reinforcement learning, where an agent learns to perform a task by maximizing a reward signal. Reinforcement learning algorithms have been developed and applied to teach robots to perform tasks such as balancing a robot (Marc Peter Deisenroth and Rasmussen, 2011), playing ping-pong (Peters et al., 2010) and baseball (Peters and Schaal, 2008c). The use of large function approximators, such as neural networks, in RL has further broadened the generality of RL (Mnih et al., 2013). Such techniques, called “deep” RL, have further allowed robots to be trained directly in the real world to perform fine-grained manipulation tasks from vision (Levine et al., 2016a), open doors (S. Gu et al., 2016), play hockey (Chebotar et al., 2017a), stack Lego blocks (M. Zhang et al., 2019), use dexterous hands (Zhu et al., 2019), and grasp objects (Kalashnikov et al., 2018b). In this work we further explore solving real-world robotics tasks using RL.

Many RL algorithms introduce prior information about the specific task to be solved. One common method is reward shaping (Andrew Y. Ng et al., 1999), but reward shaping can become arbitrarily difficult as the complexity of the task increases. Other methods incorporate a trajectory planner (G. Thomas et al., 2018) but for complex assembly tasks, trajectory planners require a host of information about objects and geometries which can be difficult to provide.

Another body of work on incorporating prior information studies using demonstrations either to initialize a policy (Peters and Schaal, 2008c; Kober and Peter, 2008), infer reward functions using inverse reinforcement learning (Finn et al., 2016a; Ziebart et al., 2008) or to improve the policy throughout the learning procedure (Hester et al., 2018; A. Nair et al., 2018a; Rajeswaran et al., 2018). These methods require multiple demonstra-

tions, which can be difficult to collect, especially for assembly tasks, although learning a reward function by classifying goal states (Singh et al., 2019) may partially alleviate this issue. More recently, manually specifying a policy and learning the residual task has been proposed (Johannink et al., 2019; T. Silver et al., 2018). In this work we evaluate both residual RL and combining RL with learning from demonstrations.

Previous work has also tackled high precision assembly tasks, especially insertion-type tasks. One line of work focuses on obtaining high dimensional observations, including geometry, forces, joint positions and velocities (R. Li et al., 2014; Tamar et al., 2017; Inoue et al., 2017; Luo et al., 2019), but this information is not easily procured, increasing complexity of the experiments and the supervision required. Other work relies on external trajectory planning or very high precision control (Inoue et al., 2017; Tamar et al., 2017), but this can be brittle to error in other components of the system, such as perception. We show how our method not only solves insertion tasks with much less information about the environment, but also does so under noisy conditions.

7.3 ELECTRIC CONNECTOR PLUG INSERTION TASKS

In this work, we empirically evaluate learning methods on a set of electric connector assembly tasks, pictured in Fig. 31. Connector plug insertions are difficult for two reasons. First, the robot must be very precise in lining up the plug with its socket. As we show in our experiments, errors as small as ± 1 mm can lead to consistent failure.

Second, there is significant friction when the connector plug touches the socket, and the robot must learn to apply sufficient force in order to insert the plug. Image sequences of successful insertions are shown in Fig. 31, where it is also possible to see details of the gripper setup that we used to ensure a failure free, fully automated training process. In our experiments, we use a 7 degrees of freedom Sawyer robot with end-effector control, meaning that the action signal u_t can be interpreted as the relative end-effector movement in Cartesian coordinates. The robot’s underlying internal control pipeline is

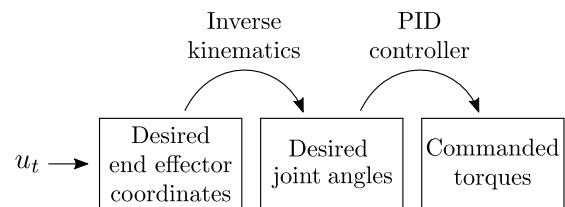


Figure 32: Illustration of the robot’s cascade control scheme. The actions u_t are computed at a frequency of up to 10 Hz, desired joint angles are obtained by inverse kinematics, and a joint-space impedance controller with anti-windup PID control commands actuator torques at 1000 Hz.

illustrated in Figure 32.

To comprehensively evaluate connector assembly tasks, we experiment on a variety of connectors. Each connector offers a different challenge in terms of required precision and force to overcome friction. We chose to benchmark the controllers performance on the insertion of a USB connector, a U-Sub connector, and a waterproof Model-E connector manufactured by MISUMI. All the explored use cases were part of the IROS 2017 Robotic Grasping and Manipulation Competition (Falco et al., 2018), included as part of a task board developed by NIST to benchmark the performance of assembly robots.

7.3.1 *Adapters*

In the following we describe the used adapters, USB, D-Sub, and Model-E. The observed difficulty of the insertion increases in that order.

USB. The USB connector is a ubiquitous, widely-used connector and offers a challenging insertion task. Because the adapter becomes smoother and therefore easier to insert over time due to wear and tear, we periodically replace the adapter. Of the three tested adapters, the USB adapter is the easiest.

D-sub. Inserting this adapter requires aligning several pins correctly, and is therefore more sensitive than inserting the USB adapter. It also requires more downward force due to a tighter fit.

Model-E. This adapter is the most difficult of the three tested connectors as it contains several edges and grooves to align and requires significant downward force to successfully insert the part.

7.3.2 *Experimental Settings*

We consider three settings in our experiments in order to evaluate how plausible it is to solve these tasks with more convenient state representations and reward functions and to evaluate the performance of different algorithms changes as the setting is modified.

3.2.1 Visual. In this experiment, we evaluate whether the RL algorithms can learn to perform the connector assembly tasks from vision without having access to state information. The state provided to the learned policy is a 32×32 grayscale image, such as shown in Fig. 33.

For goal specification, we use a goal image, avoiding the need for state information to compute rewards. The reward is the pixelwise L1 distance to the given goal image. Being able to learn from such a setup is compelling as it does not require any extra state estimation and many tasks can be specified easily by a goal image.

3.2.2. Sparse. In this experiment, the reward is obtained by directly measuring whether the connection is alive and transmitting:

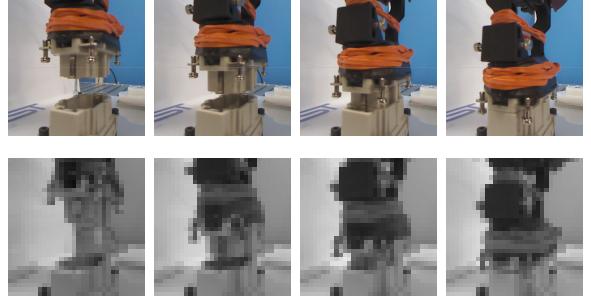


Figure 33: Successful insertion on the Model-E connector. The image-based RL algorithms receives only receives the 32×32 grayscale image as the observation.

$$r = \begin{cases} 1, & \text{if insertion signal detected} \\ 0, & \text{else.} \end{cases} \quad (29)$$

This is the exact true reward for the task of connecting a cable, and can be naturally obtained in many manufacturing systems. As state, the robot is given the Cartesian coordinates of the end-effector x_t and the vertical force f_z that is acting on the end-effector. We could only automatically detect the USB connection, so we only include the USB adapter for the sparse experiments.

3.2.3. Dense. In this experiment, the robot receives a manually shaped reward based on the distance to the target location x^* . We use the reward function

$$r_t = -\alpha \cdot \|x_t - x^*\|_1 - \frac{\beta}{(\|x_t - x^*\|_2 + \varepsilon)} - \varphi \cdot f_z, \quad (30)$$

where $0 < \varepsilon \ll 1$. The hyperparameters are set to $\alpha = 100$, $\beta = 0.002$, and $\varphi = 0.1$. When an insertion is indicated through a distance measurement, the sign of the force term flips, so that $\varphi = -0.1$ when the connector is inserted. This rewards the agent for pressing down after an insertion and showed to improve the learning process. The force measurements are calibrated before each rollout to account for measurement bias and to decouple the measurements from the robot pose.

7.4 METHODS

To solve the connector insertion tasks, we consider and evaluate a variety of RL algorithms.

7.4.1 Preliminaries

In a Markov decision process (MDP), an agent at every time step is at state $s_t \in \mathcal{S}$, takes actions $u_t \in \mathcal{U}$, receives a reward $r_t \in \mathbb{R}$, and the state evolves according to environment transition dynamics $p(s_{t+1}|s_t, u_t)$. The goal of reinforcement learning is to choose actions $u_t \sim \pi(u_t|s_t)$ to maximize the expected returns $\mathbb{E}[\sum_{t=0}^H \gamma^t r_t]$ where H is the horizon and γ is a discount factor. The policy $\pi(u_t|s_t)$ is often chosen to be an expressive parametric function approximator, such as a neural network, as we use in this work.

7.4.2 Efficient Off-Policy Reinforcement Learning

One class of RL methods additionally estimates the expected discounted return after taking action u from state s , the Q-value $Q(s, u)$. Q-values can be recursively defined with the Bellman equation:

$$Q(s_t, u_t) = \mathbb{E}_{s_{t+1}}[r_t + \gamma \max_{u_{t+1}} Q(s_{t+1}, u_{t+1})] \quad (31)$$

and learned from off-policy transitions (s_t, u_t, r_t, s_{t+1}) . Because we are interested in sample-efficient real-world learning, we use such RL algorithms that can take advantage of off-policy data.

For control with continuous actions, computing the required maximum in the Bellman equation is difficult. Continuous control algorithms such as deep deterministic policy gradients (DDPG) (Lillicrap et al., 2016) additionally learn a policy $\pi_\theta(u_t|s_t)$ to approximately choose the maximizing action. In this paper we specifically consider two related reinforcement learning algorithms that lend themselves well to real-world learning as they are sample efficient, stable, and require little hyperparameter tuning.

Twin Delayed Deep Deterministic Policy Gradients (TD3). Like DDPG, TD3 optimizes a deterministic policy (Fujimoto et al., 2018b) but uses two Q-function approximators to reduce value overestimation Van Hasselt et al., 2016 and delayed policy updates to stabilize training.

Soft Actor Critic (SAC). SAC is an off-policy value-based reinforcement learning

Algorithm 5 Residual reinforcement learning

Require: policy π_θ , hand-engineered controller π_H .

- 1: **for** $n = 0, \dots, N - 1$ episodes **do**
- 2: Sample initial state $s_0 \sim E$.
- 3: **for** $t = 0, \dots, H - 1$ steps **do**
- 4: Get policy action $u_t \sim \pi_\theta(u_t|s_t)$.
- 5: Get action to execute $u'_t = u_t + \pi_H(s_t)$.
- 6: Get next state $s_{t+1} \sim p(\cdot | s_t, u'_t)$.
- 7: Store (s_t, u_t, s_{t+1}) into replay buffer \mathcal{R} .
- 8: Sample set of transitions $(s, u, s') \sim \mathcal{R}$.
- 9: Optimize θ using RL with transitions.
- 10: **end for**
- 11: **end for**

method based on the maximum entropy reinforcement learning framework with a stochastic policy (Haarnoja et al., 2018a).

We used the implementation of these RL algorithms publicly available at `rlkit` (V. Pong et al., 2018).

7.4.3 Residual Reinforcement Learning

Instead of randomly exploring from scratch, we can inject prior information into an RL algorithm in order to speed up the training process, as well as to minimize unsafe exploration behavior. In residual RL, actions u_t are chosen by additively combining a fixed policy $\pi_H(s_t)$ with a parametric policy $\pi_\theta(u_t|s_t)$:

$$u_t = \pi_H(s_t) + \pi_\theta(u_t|s_t). \quad (32)$$

The parameters θ can be learned using any RL algorithm. In this work, we evaluate both SAC and TD3, explained in the previous section. The residual RL implementation that we use in our experiments is summarized in Algorithm 5.

A simple P-controller serves as the hand-designed controller π_H of our experiments. The P-controller operates in Cartesian space and calculates the current control action by

$$\pi_H(s_t) = -k_p \cdot (x_t - x^*), \quad (33)$$

where x^* denotes the commanded goal location. As control gains we use $k_p = [1, 1, 0.3]$. This P-controller quickly centers the end-effector above the goal position and reaches the goal after about 10 time steps from the reset position, which is located 5cm above the goal.

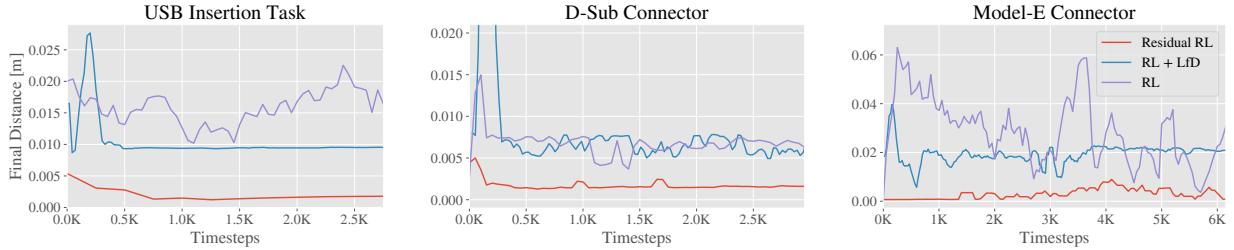


Figure 34: Resulting final mean distance during the vision-based training. The comparison includes RL, residual RL, and RL with learning from demonstrations. Only residual RL manages to deal with the high-dimensional input and consistently solve all the tasks after the given amount of training. The other methods learn to move downwards, but often get stuck in the beginning of the insertion and fail to recover from unsuccessful attempts.

7.4.4 Learning from Demonstrations

Another method to incorporate prior information is to use demonstrations from an expert policy to guide exploration during RL. We first collected demonstrations with a joystick controller. Then, we add a behavior cloning loss while performing RL that pushes the policy towards the demonstrator actions, as previously considered in (A. Nair et al., 2018a). Instead of DDPG, the underlying algorithm RL algorithm used is TD3.

7.5 EXPERIMENTS

We evaluate our method, which combines residual RL with easy-to-obtain reward signals, on a variety of connector assembly tasks performed on a real robot. In this section, we consider two types of natural rewards that are intuitive for users to provide: an image directly specifying a goal and a binary sparse reward indicating success. For both cases, we report success rates on tasks they solve. We aim to answer the following questions: (1) Can such trained policies provide comparable performance to policies that are trained with densely-shaped rewards? (2) Are these trained policies robust to small variations and noise?

5.1 Vision-based Learning. For the vision-based learning experiments, we use only raw image observations and ℓ_1 distance between the current image and goal image as the reward signal. Sample images that the robot received are shown in Fig. 33. We evaluate this type of reward on all three connectors. In our experiments, we use 32×32 grayscale images.

5.2 Learning from Sparse Rewards. In the sparse reward experiment, we use the binary signal of the connector being electrically connected as the reward signal. This experiment is most applicable to electronic manufacturing settings where the electrical connection between connectors can be directly measured. We only evaluate the sparse reward setting on the USB connector, as it was straightforward to obtain the electrical connection signal.

5.3 Perfect State Information. After evaluating the tasks in the above settings, we further evaluate with full state information with a dense and carefully shaped reward signal, given in Eq. 30, that incorporates distance to the goal and force information. Evaluating in this setting gives us an “oracle” that can be compared to the previous experiments in order to understand how much of a challenge sparse or image rewards pose for various algorithms.

5.4 Robustness. For safe and reliable future usage, it is required that the insertion controller is robust against small measurement or calibration errors that can occur when disassembling and reassembling a mechanical system. In this experiment, small goal perturbations are introduced in order to uncover the required setup precision of our algorithms.

5.5 Exploration Comparison. One advantage of using reinforcement learning is the exploratory behavior that allows the controller to adapt from new experiences unlike a deterministic control law. The two RL algorithms we consider in this paper, SAC and TD3, explore differently. SAC maintains a stochastic policy, and the algorithm also adapts the stochasticity through training. TD3 has a deterministic policy, but uses another noise process (in our case Gaussian) to inject exploratory behavior during training time. We compare the two algorithms, as well as when they are used in conjunction with residual RL, in order to evaluate the effect of the different exploration schemes.

7.6 RESULTS

We analyze the performance of policies learned with residual RL, as well as other methods, based on their ability to achieve the task goal, as well as the distance of the final object location to the goal pose over the course of training. To study the robustness of

D-Sub Connector		Goal		Model-E Connector		Goal	
		Perfect	Noisy			Perfect	Noisy
Pure RL	Dense Images, SAC	16%	0%	Pure RL	Dense Images, SAC	0%	0%
	Images, TD3	0%	0%		Images, TD3	0%	0%
	Images	12%	12%		Images	20%	20%
RL + LfD	Images	52%	52%	Residual RL	Dense Images, SAC	100%	76%
Residual RL	Dense Images, SAC	100%	64%		Images, TD3	100%	76%
	Images, TD3	52%	52%		Human	P-Controller	0%
Human	P-Controller	100%	44%		P-Controller	52%	24%

Table 1: We report average success out of 25 policy executions after training is finished for each method. For noisy goals, noise is added in form of ± 1 mm perturbations of the goal location. Residual RL, particularly with SAC, tends to be the best performing method across all three connectors. For the Model-E connector, only residual RL solves the task in the given amount of training time.

the learned policies, we also evaluate them in conditions where the goal connector position is perturbed, in order to understand the tolerance of RL policies to imprecise object placement.

7.6.1 Vision-based Learning

The results of the vision-based experiment are shown in Fig. 34. Our experiments show that a successful and consistent vision-based insertion policy can be learned from relatively few samples using residual RL.

This result suggests that goal-specification through images is a practical way to solve these types of industrial tasks. Although image-based rewards are often very sparse and hard to learn from, in this case the distance between images corresponds to a relatively dense reward signal which is sufficient to distinguish the different stages of the insertion process.

Interestingly, during training with standard RL, the policy would sometimes learn to “hack” the reward signal by moving down in the image in front of or behind the socket. In contrast, the stabilizing human-engineered controller in residual RL provides sufficient horizontal control to prevent this. The initial controller also scaffolds the learning process, by providing a very strong initialization that requires the reinforcement learning algorithm to only learn the final phase of the insertion. This produces substantially

better performance in conjunction with vision-based rewards.

USB Connector		Goal	
		Perfect	Noisy
Pure RL	Dense	28%	20%
	Sparse, SAC	16%	8%
	Sparse, TD3	44%	28%
	Images, SAC	36%	32%
	Images, TD3	28%	28%
RL + LfD	Sparse	100%	32%
	Images	88%	60%
Residual RL	Dense	100%	84%
	Sparse, SAC	88%	84%
	Sparse, TD3	100%	36%
	Images, SAC	100%	80%
	Images, TD3	0%	0%
Human	P-Controller	100%	60%

Table 2: Learning curves for solving the USB insertion task with a sparse reward. Final distance to goal is shown; lower is better. Residual RL and RL with learning from demonstrations both solve the task relatively quickly, while RL alone takes about twice as long to solve the task at the same performance.

7.6.2 Learning From Sparse Rewards

In this experiment, we compare these methods on the USB insertion task with sparse rewards. The results are reported in Fig. 35. All methods are able to achieve very high success rates in the sparse setting. This result shows that we can learn precise industrial insertion tasks in sparse-reward settings, which can often be obtained much more easily than a dense, shaped reward. In fact, prior work has found that the final policy for sparse rewards can outperform the final policy for dense rewards as it does not suffer from a misspecified objective (Andrychowicz et al., 2017b).

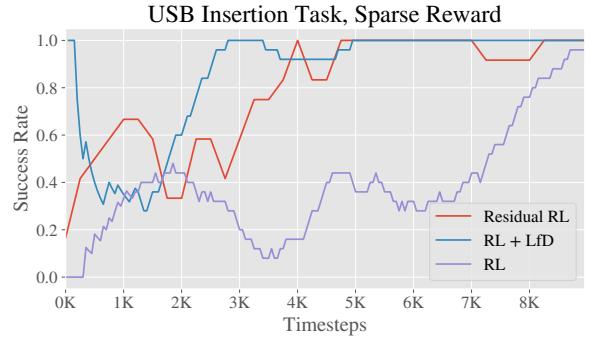


Figure 35: Learning curves for solving the USB insertion task with a sparse reward. Final distance to goal is shown; lower is better. Residual RL and RL with learning from demonstrations both solve the task relatively quickly, while RL alone takes about twice as long to solve the task at the same performance.

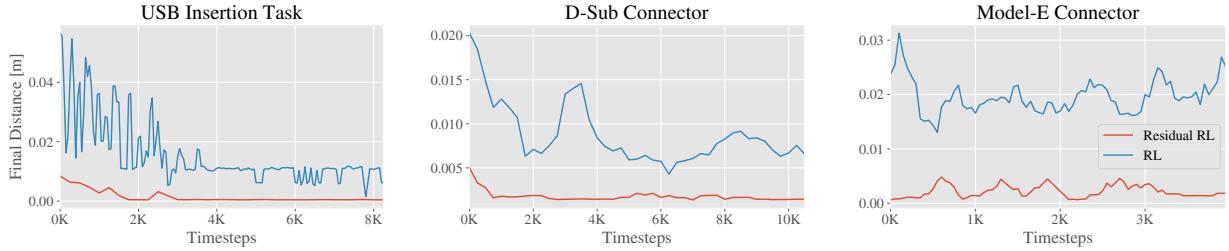


Figure 36: Plots of the final mean distance to the goal during the state-based training. Final distances greater than 0.01 m indicate unsuccessful insertions. Here, the residual RL approach performs noticeably better than pure RL and is often able to solve the task during the exploration in the early stages of the training.

7.6.3 Perfect State Information

The results of the experiment with perfect state information and dense rewards is shown in Fig. 36. In this case, residual RL still outperforms standard RL, though the better-shaped reward enables standard RL to make more initial progress than with the other reward signals. However, the hand-designed shaped reward function makes it harder for the policy to actually perform the full insertion, potentially because the more complex reward landscape provides other competing goals to the policy. The final performance with sparse rewards on the USB insertion task is substantially better.

7.6.4 Robustness

In the previous set of experiments, the goal locations were known exactly. In this case, the hand-engineered controller performs well. However, once noise is added to the goal location, the deterministic P-controller struggles. To test robustness, a goal perturbation is created artificially, and the controllers are tested under this condition. All results of our robustness evaluations are listed in Fig. 1 and Fig. ???. In the presence of a $\pm 1\text{mm}$ perturbation, the residual RL controller succeeds more often on the USB and D-Sub tasks, and significantly more often on the Model-E task. Unlike the hand-engineered controller, residual RL consistently solved this task and overcame goal perturbations in 16/25 trials. The agent demonstrably learns small but consistent corrective feedback behaviors in order to move in the right direction during the descent motion, a behavior that is very difficult to specify manually. This behavior illustrates the strength of residual RL. Since the human controller already specifies the general trajectory of the optimal policy, environment samples are only required to learn this corrective feedback behavior.

7.6.5 Exploration Comparison

All experiments were also performed using TD3 instead of SAC. The final success rates of these experiments are included in Fig. 1. When combined with residual RL, SAC and TD3 perform comparably. However, TD3 is often substantially less robust. These results are likely explained by the exploration strategy of the two algorithms. TD3 has a deterministic policy and fixed noise during training, so once it observes some high-reward states, it quickly learns to repeat that trajectory. SAC adapts the noise to the correct scale, helping SAC stay robust to small perturbations, and because SAC learns the value function for a stochastic policy, it is able to handle some degree of additive noise effectively. We found that the outputted action of TD3 approaches the extreme values at the edge of the allowed action space, while SAC executed less extreme actions, which likely further improved robustness.

7.7 CONCLUSION

In this paper, we studied deep reinforcement learning in a practical setting, and demonstrated that deep RL can solve complex industrial assembly tasks with tight tolerances. We showed that we can learn insertion policies with raw image observations with either binary outcome-based rewards, or rewards based on goal images. We conducted a series of experiments for various connector type assemblies, and demonstrated the feasibility of our method under challenging conditions, such as noisy goal specification and complex connector geometries. Reinforcement learning algorithms that can automatically learn complex assembly tasks with easy-to-specify reward functions have the potential to automate a wide range of assembly tasks, making this technology a promising direction forward for flexible and capable robotic manipulators.

There remains significant challenges for applying these techniques in more complex environments. One practical direction for future work is focusing on multi-stage assembly tasks through vision. This would pose a challenge to the goal-based policies as the background would be visually more complex. Moreover, multi-step tasks involve adapting to previous mistakes or inaccuracies, which could be difficult but should be able to be handled by RL. Extending the presented approach to multi-stage assembly tasks will pave the road to a higher robot autonomy in flexible manufacturing.

8

AWAC: ACCELERATING ONLINE REINFORCEMENT LEARNING WITH OFFLINE DATASETS

8.1 INTRODUCTION

Learning models that generalize effectively to complex open-world settings, from image recognition (Krizhevsky et al., 2012) to natural language processing (Devlin et al., 2019), relies on large, high-capacity models as well as large, diverse, and representative datasets. Leveraging this recipe of pre-training from large-scale offline datasets has the potential to provide significant benefits for reinforcement learning (RL) as well, both in terms of generalization and sample complexity. But most existing RL algorithms collect data online from scratch every time a new policy is learned, which can quickly become impractical in domains like robotics where physical data collection has a non-trivial cost. In the same way that powerful models in computer vision and NLP are often pre-trained on large, general-purpose datasets and then fine-tuned on task-specific data, practical instantiations of reinforcement learning for real world robotics problems will need to be able to incorporate large amounts of prior data effectively into the learning process, while still collecting additional data online for the task at hand. Doing so effectively will make the online data collection process much more practical while still allowing robots operating in the real world to continue improving their behavior.

For data-driven reinforcement learning, offline datasets consist of trajectories of states, actions and associated rewards. This data can potentially come from demonstrations for the desired task (Schaal, 1997; Atkeson and Schaal, 1997), suboptimal policies (Gao et al., 2018), demonstrations for related tasks (Zhou et al., 2019), or even just random exploration in the environment. Depending on the quality of the data that is provided, useful knowledge can be extracted about the dynamics of the world, about the task being solved, or both. Effective data-driven methods for deep reinforcement learning should be able to use this data to pre-train offline while improving with online fine-tuning.

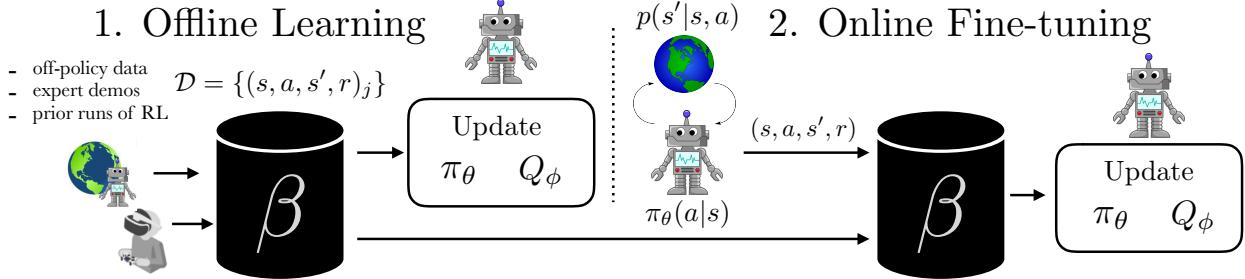


Figure 37: We study learning policies by offline learning on a prior dataset \mathcal{D} and then fine-tuning with online interaction. The prior data could be obtained via prior runs of RL, expert demonstrations, or any other source of transitions. Our method, advantage weighted actor critic (AWAC) is able to learn effectively from offline data and fine-tune in order to reach expert-level performance after collecting a limited amount of interaction data. Videos and data are available at awacrl.github.io

Since this prior data can come from a variety of sources, we would like to design an algorithm that does not utilize different types of data in any privileged way. For example, prior methods that incorporate demonstrations into RL directly aim to mimic these demonstrations (A. Nair et al., 2018a), which is desirable when the demonstrations are known to be optimal, but imposes strict requirements on the type of offline data, and can cause undesirable bias when the prior data is not optimal. While prior methods for fully offline RL provide a mechanism for utilizing offline data (Fujimoto et al., 2019a; Kumar et al., 2019a), as we will show in our experiments, such methods generally are not effective for fine-tuning with online data as they are often too conservative. In effect, prior methods require us to choose: Do we assume prior data is optimal or not? Do we use only offline data, or only online data? To make it feasible to learn policies for open-world settings, we need algorithms that learn successfully in any of these cases.

In this work, we study how to build RL algorithms that are effective for pre-training from off-policy datasets, but also well suited to continuous improvement with online data collection. We systematically analyze the challenges with using standard off-policy RL algorithms (Haarnoja et al., 2018a; Kumar et al., 2019a; Abdolmaleki et al., 2018) for this problem, and introduce a simple actor critic algorithm that elegantly bridges data-driven pre-training from offline data and improvement with online data collection. Our method, which uses dynamic programming to train a critic but a supervised learning style update to train a constrained actor, combines the best of supervised learning and actor-critic algorithms. Dynamic programming can leverage off-policy data and enable sample-efficient learning. The simple supervised actor update implicitly enforces a constraint that mitigates the effects of distribution shift when learning from offline



Figure 38: Utilizing prior data for online learning allows us to solve challenging real-world robotics tasks, such as this dexterous manipulation task where the learned policy must control a 4-fingered hand to reposition an object.

data (Fujimoto et al., 2019a; Kumar et al., 2019a), while avoiding overly conservative updates.

We evaluate our algorithm on a wide variety of robotic control tasks, using a set of simulated dexterous manipulation problems as well as three separate real-world robots: drawer opening with a 7-DoF robotic arm, picking up an object with a multi-fingered hand, and rotating a valve with a 3-fingered claw. Our algorithm, Advantage Weighted Actor Critic (AWAC), is able to quickly learn successful policies for these challenging tasks, in spite of high dimensional action spaces and uninformative, sparse reward signals. We show that AWAC finetunes much more efficiently after offline pretraining as compared to prior methods and, given a fixed time budget, attains significantly better performance on the real-world tasks. Moreover, AWAC can utilize different types of prior data without any algorithmic changes: demonstrations, suboptimal data, or random exploration data. The contribution of this work is not just another RL algorithm, but a systematic study of what makes offline pre-training with online fine-tuning unique compared to the standard RL paradigm, which then directly motivates a simple algorithm, AWAC, to address these challenges. We additionally discuss the design decisions required for applying AWAC as a practical tool for real-world robotic skill learning.

8.2 PRELIMINARIES

We use the standard reinforcement learning notation, with states \mathbf{s} , actions \mathbf{a} , policy $\pi(\mathbf{a}|\mathbf{s})$, rewards $r(\mathbf{s}, \mathbf{a})$, and dynamics $p(\mathbf{s}'|\mathbf{s}, \mathbf{a})$. The discounted return is defined as $R_t = \sum_{i=t}^T \gamma^i r(s_i, a_i)$, for a discount factor γ and horizon T which may be infinite. The objective

of an RL agent is to maximize the expected discounted return $J(\pi) = \mathbb{E}_{p_{\pi}(\tau)}[R_0]$ under the distribution induced by the policy. The optimal policy can be learned directly (e.g., using policy gradient to estimate $\nabla J(\pi)$ (R. J. Williams, 1992)), but this is often ineffective due to high variance of the estimator. Many algorithms attempt to reduce this variance by making use of the value function $V^{\pi}(s) = \mathbb{E}_{p_{\pi}(\tau)}[R_t|s]$, action-value function $Q^{\pi}(s, a) = \mathbb{E}_{p_{\pi}(\tau)}[R_t|s, a]$, or advantage $A^{\pi}(s, a) = Q^{\pi}(s, a) - V^{\pi}(s)$. The action-value function for a policy can be written recursively via the Bellman equation:

$$Q^{\pi}(s, a) = r(s, a) + \gamma \mathbb{E}_{p(s'|s, a)}[V^{\pi}(s')] \quad (34)$$

$$= r(s, a) + \gamma \mathbb{E}_{p(s'|s, a)}[\mathbb{E}_{\pi(a'|s')}[Q^{\pi}(s', a')]]. \quad (35)$$

Instead of estimating policy gradients directly, actor-critic algorithms maximize returns by alternating between two phases (Konda and Tsitsiklis, 2000): policy evaluation and policy improvement. During the policy evaluation phase, the critic $Q^{\pi}(s, a)$ is estimated for the current policy π . This can be accomplished by repeatedly applying the Bellman operator \mathcal{B} , corresponding to the right-hand side of Equation 35, as defined below:

$$\mathcal{B}^{\pi}Q(s, a) = r(s, a) + \gamma \mathbb{E}_{p(s'|s, a)}[\mathbb{E}_{\pi(a'|s')}[Q^{\pi}(s', a')]]. \quad (36)$$

By iterating according to $Q^{k+1} = \mathcal{B}^{\pi}Q^k$, Q^k converges to Q^{π} (Sutton and Barto, 1998). With function approximation, we cannot apply the Bellman operator exactly, and instead minimize the Bellman error with respect to Q -function parameters ϕ_k :

$$\phi_k = \arg \min_{\phi} \mathbb{E}_{\mathcal{D}}[(Q_{\phi}(s, a) - y)^2], \quad (37)$$

$$y = r(s, a) + \gamma \mathbb{E}_{s', a'}[Q_{\phi_{k-1}}(s', a')]. \quad (38)$$

During policy improvement, the actor π is typically updated based on the current estimate of Q^{π} . A commonly used technique (Lillicrap et al., 2016; Fujimoto et al., 2018b; Haarnoja et al., 2018a) is to update the actor $\pi_{\theta_k}(a|s)$ via likelihood ratio or pathwise derivatives to optimize the following objective, such that the expected value of the Q -function Q^{π} is maximized:

$$\theta_k = \arg \max_{\theta} \mathbb{E}_{s \sim \mathcal{D}}[\mathbb{E}_{\pi_{\theta}(a|s)}[Q_{\phi_k}(s, a)]] \quad (39)$$

Actor-critic algorithms are widely used in deep RL (Mnih et al., 2016; Lillicrap et al., 2016; Haarnoja et al., 2018a; Fujimoto et al., 2018b). With a Q -function estimator, they

can in principle utilize off-policy data when used with a replay buffer for storing prior transition tuples, which we will denote β , to sample previous transitions, although we show that this by itself is insufficient for our problem setting.

8.3 CHALLENGES IN OFFLINE RL WITH ONLINE FINE-TUNING

In this section, we study the unique challenges that exist when pre-training using offline data, followed by fine-tuning with online data collection. We first describe the problem, and then analyze what makes this problem difficult for prior methods.

8.3.1 Problem Definition

A static dataset of transitions, $\mathcal{D} = \{(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)_j\}$, is provided to the algorithm at the beginning of training. This dataset can be sampled from an arbitrary policy or mixture of policies, and may even be collected by a human expert. This definition is general and encompasses many scenarios: learning from demonstrations, random data, prior RL experiments, or even from multi-task data. Given the dataset \mathcal{D} , our goal is to leverage \mathcal{D} for pre-training and use a small amount of online interaction to learn the optimal policy $\pi^*(\mathbf{a}|\mathbf{s})$, as depicted in Fig 37. This setting is representative of many real-world RL settings, where prior data is available and the aim is to learn new skills efficiently. We first study existing algorithms empirically in this setting on the HalfCheetah-v2 Gym environment¹. The prior dataset consists of 15 demonstrations from an expert policy and 100 suboptimal trajectories sampled from a behavioral clone of these demonstrations. All methods for the remainder of this paper incorporate the prior dataset, unless explicitly labeled “scratch”.

8.3.2 Data Efficiency

One of the simplest ways to utilize prior data such as demonstrations for RL is to pre-train a policy with imitation learning, and fine-tune with on-policy RL (Abhishek Gupta et al., 2019a; Rajeswaran et al., 2018). This approach has two drawbacks: (1) prior data may not be optimal; (2) on-policy fine-tuning is data inefficient as it does not reuse the prior data in the RL stage. In our setting, data efficiency is vital. To this end, we require

¹ We use this environment for analysis because it helps understand and accentuate the differences between different algorithms. More challenging environments like the ones shown in Fig 40 are too hard to solve to analyze variants of different methods.

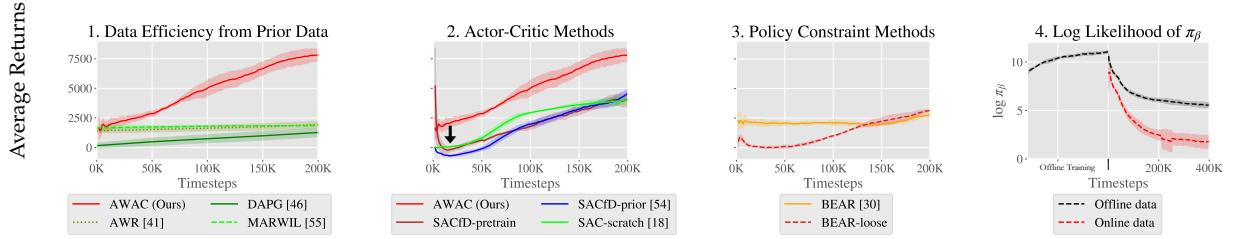


Figure 39: Analysis of prior methods on HalfCheetah-v2 using offline RL with online fine-tuning.

(1) On-policy methods (DAPG, AWR, MARWIL) learn relatively slowly, even with access to prior data. We present our method, AWAC, as an example of how off-policy RL methods can learn much faster. (2) Variants of soft actor-critic (SAC) with offline training (performed before timestep 0) and fine-tuning. We see a “dip” in the initial performance, even if the policy is pretrained with behavioral cloning. (3) Offline RL method BEAR (Kumar et al., 2019a) on offline training and fine-tuning, including a “loose” variant of BEAR with a weakened constraint. Standard offline RL methods fine-tune slowly, while the “loose” BEAR variant experiences a similar dip as SAC. (4) We show that the fit of the behavior models $\hat{\pi}_\beta$ used by these offline methods degrades as new data is added to the buffer during fine-tuning, potentially explaining their poor fine-tuning performance.

algorithms that are able to reuse arbitrary off-policy data during online RL for data-efficient fine-tuning. We find that algorithms that use on-policy fine-tuning (Rajeswaran et al., 2018; Abhishek Gupta et al., 2019a), or Monte-Carlo return estimation (Peters and Schaal, 2007a; Q. Wang et al., 2018a; Peng et al., 2019a) are generally much less efficient than off-policy actor-critic algorithms, which iterate between improving π and estimating Q^π via Bellman backups. This can be seen from the results in Figure 39 plot 1, where on-policy methods like DAPG (Rajeswaran et al., 2018) and Monte-Carlo return methods like AWR (Peng et al., 2019a) and MARWIL (Q. Wang et al., 2018a) are an order of magnitude slower than off-policy actor-critic methods. Actor-critic methods, shown in Figure 39 plot 2, can in principle use off-policy data. However, as we will discuss next, naively applying these algorithms to our problem suffers from a different set of challenges.

8.3.3 Bootstrap Error in Offline Learning with Actor-Critic Methods

When standard off-policy actor-critic methods are applied to this problem setting, they perform poorly, as shown in the second plot in Figure 39: despite having a prior dataset in the replay buffer, these algorithms do not benefit significantly from offline training. We evaluate soft actor critic (Haarnoja et al., 2018a), a state-of-the-art actor-critic algo-

rithm for continuous control. Note that “SAC-scratch,” which does not receive the prior data, performs similarly to “SACfD-prior,” which does have access to the prior data, indicating that the off-policy RL algorithm is not actually able to make use of the off-policy data for pre-training. Moreover, even if the SAC policy is pre-trained by behavior cloning, labeled “SACfD-pretrain”, we still observe an initial decrease in performance, and performance similar to learning from scratch.

This challenge can be attributed to off-policy bootstrapping error accumulation, as observed in several prior works (Sutton and Barto, 1998; Kumar et al., 2019a; Yifan Wu et al., 2020; Levine et al., n.d.; Fujimoto et al., 2019a). In actor-critic algorithms, the target value $Q(s', a')$, with $a' \sim \pi$, is used to update $Q(s, a)$. When a' is outside of the data distribution, $Q(s', a')$ will be inaccurate, leading to accumulation of error on static datasets.

Offline RL algorithms (Fujimoto et al., 2019a; Kumar et al., 2019a; Yifan Wu et al., 2020) propose to address this issue by explicitly adding constraints on the policy improvement update (Equation 39) to avoid bootstrapping on out-of-distribution actions, leading to a policy update of this form:

$$\arg \max_{\theta} \mathbb{E}_{s \sim \mathcal{D}} [\mathbb{E}_{\pi_\theta(a|s)} [Q_{\phi_k}(s, a)]] \text{ s.t. } D(\pi_\theta, \pi_\beta) \leq \epsilon. \quad (40)$$

Here, π_θ is the actor being updated, and $\pi_\beta(a|s)$ represents the (potentially unknown) distribution from which all of the data seen so far (both offline data and online data) was generated. In the case of a replay buffer, π_β corresponds to a mixture distribution over all past policies. Typically, π_β is not known, especially for offline data, and must be estimated from the data itself. Many offline RL algorithms (Kumar et al., 2019a; Fujimoto et al., 2019a; Noah Y. Siegel et al., 2020b) explicitly fit a parametric model to samples for the distribution π_β via maximum likelihood estimation, where samples from π_β are obtained simply by sampling uniformly from the data seen thus far: $\hat{\pi}_\beta = \max_{\hat{\pi}_\beta} \mathbb{E}_{s, a \sim \pi_\beta} [\log \hat{\pi}_\beta(a|s)]$. After estimating $\hat{\pi}_\beta$, prior methods implement the constraint given in Equation 40 in various ways, including penalties on the policy update (Kumar et al., 2019a; Yifan Wu et al., 2020) or architecture choices for sampling actions for policy training (Fujimoto et al., 2019a; Noah Y. Siegel et al., 2020b). As we will see next, the requirement for accurate estimation of $\hat{\pi}_\beta$ makes these methods difficult to use with online fine-tuning.

8.3.4 Excessively Conservative Online Learning

While offline RL algorithms with constraints (Kumar et al., 2019a; Fujimoto et al., 2019a; Yifan Wu et al., 2020) perform well offline, they struggle to improve with fine-tuning, as shown in the third plot in Figure 39. We see that the purely offline RL performance (at “oK” in Fig. 39) is much better than the standard off-policy methods shown in Section 8.3.3. However, with additional iterations of online fine-tuning, the performance increases very slowly (as seen from the slope of the BEAR curve in Fig 39). What causes this phenomenon?

This can be attributed to challenges in fitting an accurate behavior model as data is collected online during fine-tuning. In the offline setting, behavior models must only be trained once via maximum likelihood, but in the online setting, the behavior model must be updated online to track incoming data. Training density models online (in the “streaming” setting) is a challenging research problem (Ramapuram et al., 2017), made more difficult by a potentially complex multi-modal behavior distribution induced by the mixture of online and offline data. To understand this, we plot the log likelihood of learned behavior models on the dataset during online and offline training for the HalfCheetah task. As we can see in the plot, the accuracy of the behavior models ($\log \pi_\beta$ on the y-axis) reduces during online fine-tuning, indicating that it is not fitting the new data well during online training. When the behavior models are inaccurate or unable to model new data well, constrained optimization becomes too conservative, resulting in limited improvement with fine-tuning. This analysis suggests that, in order to address our problem setting, we require an off-policy RL algorithm that constrains the policy to prevent offline instability and error accumulation, but not so conservatively that it prevents online fine-tuning due to imperfect behavior modeling. Our proposed algorithm, which we discuss in the next section, accomplishes this by employing an *implicit* constraint, which does not require *any* explicit modeling of the behavior policy.

8.4 ADVANTAGE WEIGHTED ACTOR CRITIC: A SIMPLE ALGORITHM FOR FINE-TUNING FROM OFFLINE DATASETS

In this section, we will describe the advantage weighted actor-critic (AWAC) algorithm, which trains an off-policy critic and an actor with an *implicit* policy constraint. We will show AWAC mitigates the challenges outlined in Section 8.3. AWAC follows the design for actor-critic algorithms as described in Section 10.3, with a policy evaluation step to learn Q^π and a policy improvement step to update π . AWAC uses off-policy temporal-

difference learning to estimate Q^π in the policy evaluation step, and a policy improvement update that is able to obtain the benefits of offline RL algorithms at training from prior datasets, while avoiding the overly conservative behavior described in Section 8.3.4. We describe the policy improvement step in AWAC below, and then summarize the entire algorithm.

Policy improvement for AWAC proceeds by learning a policy that maximizes the value of the critic learned in the policy evaluation step via TD bootstrapping. If done naively, this can lead to the issues described in Section 8.3.4, but we can avoid the challenges of bootstrap error accumulation by restricting the policy distribution to stay close to the data observed thus far during the actor update, while maximizing the value of the critic. At iteration k , AWAC therefore optimizes the policy to maximize the estimated Q-function $Q^{\pi_k}(\mathbf{s}, \mathbf{a})$ at every state, while constraining it to stay close to the actions observed in the data, similar to prior offline RL methods, though this constraint will be enforced differently. Note from the definition of the advantage in Section 10.3 that optimizing $Q^{\pi_k}(\mathbf{s}, \mathbf{a})$ is equivalent to optimizing $A^{\pi_k}(\mathbf{s}, \mathbf{a})$. We can therefore write this optimization as:

$$\pi_{k+1} = \arg \max_{\pi \in \Pi} \mathbb{E}_{\mathbf{a} \sim \pi(\cdot | \mathbf{s})} [A^{\pi_k}(\mathbf{s}, \mathbf{a})] \quad (41)$$

$$\text{s.t. } D_{KL}(\pi(\cdot | \mathbf{s}) \| \pi_\beta(\cdot | \mathbf{s})) \leq \epsilon. \quad (42)$$

As we saw in Section 8.3.3, enforcing the constraint by incorporating an explicit learned behavior model (Kumar et al., 2019a; Fujimoto et al., 2019a; Yifan Wu et al., 2020; Noah Y. Siegel et al., 2020b) leads to poor fine-tuning performance. Instead, we enforce the constraint *implicitly*, without learning a behavior model. We first derive the solution to the constrained optimization in Equation 41 to obtain a non-parametric closed form for the actor. This solution is then projected onto the parametric policy class *without* any explicit behavior model. The analytic solution to Equation 41 can be obtained by enforcing the KKT conditions (Peters and Schaal, 2007a; Peters et al., 2010; Peng et al., 2019a). The Lagrangian is:

$$\mathcal{L}(\pi, \lambda) = \mathbb{E}_{\mathbf{a} \sim \pi(\cdot | \mathbf{s})} [A^{\pi_k}(\mathbf{s}, \mathbf{a})] + \lambda(\epsilon - D_{KL}(\pi(\cdot | \mathbf{s}) \| \pi_\beta(\cdot | \mathbf{s}))), \quad (43)$$

and the closed form solution to this problem is $\pi^*(\mathbf{a} | \mathbf{s}) \propto \pi_\beta(\mathbf{a} | \mathbf{s}) \exp\left(\frac{1}{\lambda} A^{\pi_k}(\mathbf{s}, \mathbf{a})\right)$. When using function approximators, such as deep neural networks as we do, we need to project the non-parametric solution into our policy space. For a policy π_θ with parameters θ , this can be done by minimizing the KL divergence of π_θ from the optimal non-parametric

solution π^* under the data distribution $\rho_{\pi_\beta}(\mathbf{s})$:

$$\arg \min_{\theta} \rho_{\pi_\beta}(\mathbf{s}) [\mathcal{D}_{\text{KL}}(\pi^*(\cdot|\mathbf{s})||\pi_\theta(\cdot|\mathbf{s}))] \quad (44)$$

$$= \arg \min_{\theta} \rho_{\pi_\beta}(\mathbf{s}) [\pi^*(\cdot|\mathbf{s})[-\log \pi_\theta(\cdot|\mathbf{s})]] \quad (45)$$

Note that the parametric policy could be projected with either direction of KL divergence. Choosing the reverse KL results in explicit penalty methods (Yifan Wu et al., 2020) that rely on evaluating the density of a learned behavior model. Instead, by using forward KL, we can compute the policy update by sampling directly from β :

$$\theta_{k+1} = \arg \max_{\theta} \mathbb{E}_{\mathbf{s}, \mathbf{a} \sim \beta} \left[\log \pi_\theta(\mathbf{a}|\mathbf{s}) \exp \left(\frac{1}{\lambda} A^{\pi_k}(\mathbf{s}, \mathbf{a}) \right) \right]. \quad (46)$$

This actor update amounts to weighted maximum likelihood (i.e., supervised learning), where the targets are obtained by re-weighting the state-action pairs observed in the current dataset by the predicted advantages from the learned critic, *without* explicitly learning any parametric behavior model, simply sampling (\mathbf{s}, \mathbf{a}) from the replay buffer β . See Appendix ?? for a more detailed derivation and Appendix ?? for specific implementation details.

Avoiding explicit behavior modeling. Note that the update in Equation 46 completely avoids any modeling of the previously observed data β with a parametric model. By avoiding any explicit learning of the behavior model AWAC is far less conservative than methods which fit a model $\hat{\pi}_\beta$ explicitly, and better incorporates new data during online fine-tuning, as seen from our results in Section 8.6. This derivation is related to AWR (Peng et al., 2019a), with the main difference that AWAC uses an off-policy Q-function Q^π to estimate the advantage, which greatly improves efficiency and even final performance (see results in Section 8.6.1.1). The update also resembles ABM-MPO, but ABM-MPO *does* require modeling the behavior policy which, as discussed in Section 8.3.4, can lead to poor fine-tuning. In Section 8.6.1.1, AWAC outperforms ABM-MPO on a range of challenging tasks.

Policy evaluation. During policy evaluation, we estimate the action-value $Q^\pi(\mathbf{s}, \mathbf{a})$ for the current policy π , as described in Section 10.3. We utilize a temporal difference learning scheme for policy evaluation (Haarnoja et al., 2018a; Fujimoto et al., 2018b), minimizing the Bellman error as described in Equation 36. This enables us to learn very efficiently from off-policy data. This is particularly important in our problem setting to effectively use the offline dataset, and allows us to significantly outperform alternatives

using Monte-Carlo evaluation or TD(λ) to estimate returns (Peng et al., 2019a).

Algorithm 6 Advantage Weighted Actor Critic (AWAC)

```

1: Dataset  $\mathcal{D} = \{(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)\}$ 
2: Initialize buffer  $\beta = \mathcal{D}$ 
3: Initialize  $\pi_\theta, Q_\phi$ 
4: for iteration  $i = 1, 2, \dots$  do
5:   Sample batch  $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r) \sim \beta$ 
6:    $y = r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{s}', \mathbf{a}'}[Q_{\phi_{k-1}}(\mathbf{s}', \mathbf{a}')]$ 
7:    $\phi \leftarrow \arg \min_{\phi} \mathbb{E}_{\mathcal{D}}[(Q_\phi(\mathbf{s}, \mathbf{a}) - y)^2]$ 
8:    $\theta \leftarrow \arg \max_{\theta} \mathbb{E}_{\mathbf{s}, \mathbf{a} \sim \beta} [\log \pi_\theta(\mathbf{a} | \mathbf{s}) \exp(\frac{1}{\lambda} A^{\pi_k}(\mathbf{s}, \mathbf{a}))]$ 
9:   if  $i > \text{num\_offline\_steps}$  then
10:     $\tau_1, \dots, \tau_K \sim p_{\pi_\theta}(\tau)$ 
11:     $\beta \leftarrow \beta \cup \{\tau_1, \dots, \tau_K\}$ 
12:   end if
13: end for=o

```

Algorithm summary. The full AWAC algorithm for offline RL with online fine-tuning is summarized in Algorithm 6. In a practical implementation, we can parameterize the actor and the critic by neural networks and perform SGD updates from Eqn. 46 and Eqn. 37. Specific details are provided in Appendix ???. AWAC ensures data efficiency with off-policy critic estimation via bootstrapping, and avoids offline bootstrap error with a constrained actor update. By avoiding explicit modeling of the behavior policy, AWAC avoids overly conservative updates.

While AWAC is related to several prior RL algorithms, we note that there are key differences that make it particularly amenable to the problem setting we are considering – offline RL with online fine-tuning – that other methods are unable to tackle. As we show in our experimental analysis with direct comparisons to prior work, every one of the design decisions being made in this work are important for algorithm performance. As compared to AWR (Peng et al., 2019a), AWAC uses TD bootstrapping for significantly more efficient and even asymptotically better performance. As compared to offline RL techniques like ABM (Noah Y. Siegel et al., 2020b), MPO (Abdolmaleki et al., 2018), BEAR (Kumar et al., 2019a) or BCQ (Fujimoto et al., 2019a) this work is able to avoid the need for any behavior modeling, thereby enabling the *online* fine-tuning part of the problem much better. As shown in Fig 40, when these seemingly ablations are made to AWAC, the algorithm performs significantly worse.

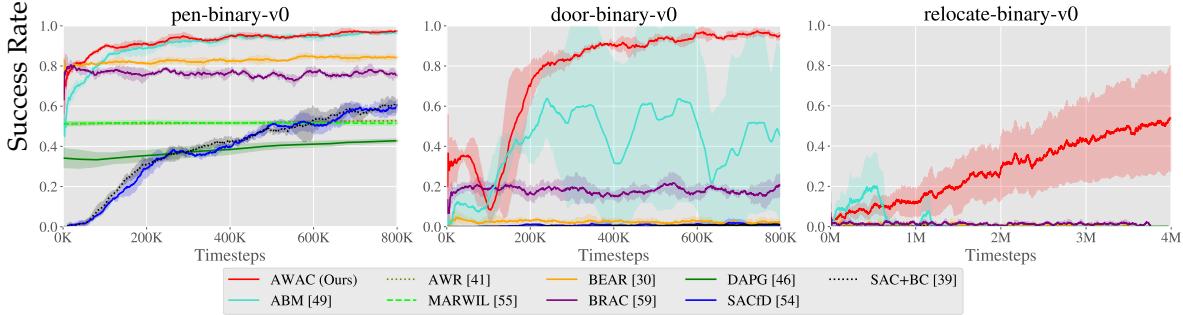


Figure 40: Comparative evaluation on the dexterous manipulation tasks. These tasks are difficult due to their high action dimensionality and reward sparsity. We see that AWAC is able to learn these tasks with little online data collection required (100K samples \approx 16 minutes of equivalent real-world interaction time). Meanwhile, most prior methods are not able to solve the harder two tasks: door opening and object relocation.

8.5 RELATED WORK

Off-policy RL algorithms are designed to reuse off-policy data during training, and have been studied extensively (Konda and Tsitsiklis, 2000; Degris et al., 2012; Mnih et al., 2016; Haarnoja et al., 2018a; Fujimoto et al., 2018b; Bhatnagar et al., 2009; Peters and Schaal, 2008a; S. Zhang et al., 2019; Wawrzynski, 2009; Balduzzi and Ghifary, 2015). While standard off-policy methods are able to benefit from including data seen *during* a training run, as we show in Section 8.3.3 they struggle when training from previously collected offline data from other policies, due to error accumulation with distribution shift (Fujimoto et al., 2019a; Kumar et al., 2019a). Offline RL methods aim to address this issue, often by constraining the actor updates to avoid excessive deviation from the data distribution (Lange et al., 2012b; P. S. Thomas and Brunskill, 2016; Hallak et al., 2015; Hallak et al., 2016; Hallak and Mannor, 2017; Agarwal et al., 2019; Kumar et al., 2019a; Fujimoto et al., 2019a; Fakoor et al., 2019; Nachum et al., 2019; Noah Y. Siegel et al., 2020b; Levine et al., n.d.; Ruiyi Zhang et al., 2020). One class of these methods utilize importance sampling (P. S. Thomas and Brunskill, 2016; Ruiyi Zhang et al., 2020; Nachum et al., 2019; Degris et al., 2012; Jiang and L. Li, 2016; Hallak and Mannor, 2017). Another class of methods perform offline reinforcement learning via dynamic programming, with an explicit constraint to prevent deviation from the data distribution (Lange et al., 2012b; Kumar et al., 2019a; Fujimoto et al., 2019a; Yifan Wu et al., 2020; Jaques et al., 2019). While these algorithms perform well in the purely offline settings, we show in Section 8.3.4 that such methods tend to be overly conservative, and therefore may not learn efficiently when fine-tuning with online data collection. In contrast, our algorithm AWAC is compara-

ble to these algorithms for offline pre-training, but learns much more efficiently during subsequent fine-tuning.

Prior work has also considered the special case of learning from *demonstration* data. One class of algorithms initializes the policy via behavioral cloning from demonstrations, and then fine-tunes with reinforcement learning (Peters and Schaal, 2008c; Ijspeert et al., 2002; Theodorou et al., 2010; Kim et al., 2013; Rajeswaran et al., 2018; Abhishek Gupta et al., 2019a; Zhu et al., 2019). Most such methods use on-policy fine-tuning, which is less sample-efficient than off-policy methods that perform value function estimation. Other prior works have incorporated demonstration data into the replay buffer using off-policy RL methods (Večerík et al., 2017; A. Nair et al., 2017). We show in Section 8.3.3 that these strategies can result in a large dip in performance during online fine-tuning, due to the inability to pre-train an effective value function from offline data. In contrast, our work shows that using supervised learning style policy updates can allow for better bootstrapping from demonstrations as compared to Večerík et al. (2017) and A. Nair et al. (2017).

Our method builds on algorithms that implement a maximum likelihood objective for the actor, based on an expectation-maximization formulation of RL (Peters and Schaal, 2007a; Neumann and Peters, 2008; Theodorou et al., 2010; Peters et al., 2010; Peng et al., 2019a; Abdolmaleki et al., 2018; Q. Wang et al., 2018a). Most closely related to our method in this respect are the algorithms proposed by Peng et al. (2019a) (AWR) and Noah Y. Siegel et al. (2020b) (ABM). Unlike AWR, which estimates the value function of the *behavior* policy, V^{π_β} via Monte-Carlo estimation or TD- λ , our algorithm estimates the Q-function of the *current* policy Q^π via bootstrapping, enabling much more efficient learning, as shown in our experiments. Unlike ABM, our method does not require learning a separate function approximator to model the behavior policy π_β , and instead directly samples the dataset. As we discussed in Section 8.3.4, modeling π_β can be a major challenge for online fine-tuning. While these distinctions may seem somewhat subtle, they are important and we show in our experiments that they result in a large difference in algorithm performance. Finally, our work goes beyond the analysis in prior work, by studying the issues associated with pre-training and fine-tuning in Section 8.3. Closely to our work, Z. Wang et al. (2020b) proposed critic regularized regression for offline RL, which uses off-policy Q-learning and an equivalent policy update. In contrast to this concurrent work, we specifically study the offline pretraining online fine-tuning problem, which this prior work does not address, analyze why other methods are ineffective in this setting, and show that our approach enables strong fine-tuning results on challenging dexterous manipulation tasks and real-world robotic systems.

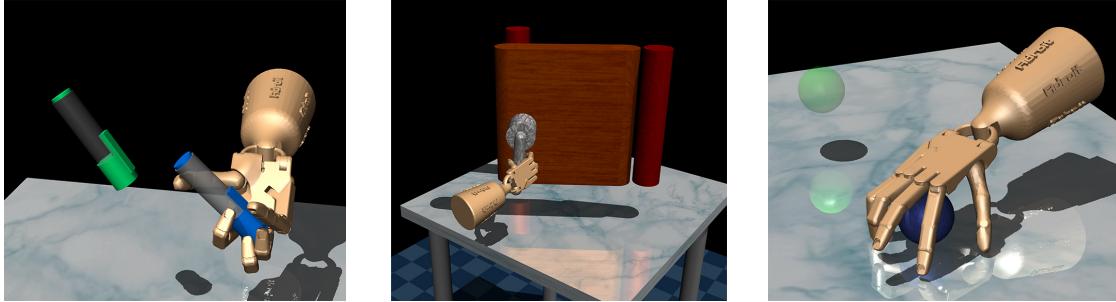


Figure 41: Illustration of dexterous manipulation tasks in simulation. These tasks exhibit sparse rewards, high-dimensional control, and complex contact physics.

The idea of bootstrapping learning from prior data for real-world robotic learning is not a new one; in fact, it has been extensively explored in the context of providing initial rollouts to bootstrap policy search (Kober and Peter, 2008; Peters and Schaal, 2008c; Kormushev et al., 2010), initializing dynamic motion primitives (Bentivegna et al., 2003; Kormushev et al., 2010; Mülling et al., 2013) in the context of on-policy reinforcement learning algorithms (Rajeswaran et al., 2018; Zhu et al., 2019), inferring reward shaping (Yuchen Wu et al., 2020) and even for inferring reward functions (Ziebart et al., 2008; Abbeel and Andrew Y Ng, 2004). Our work shows how we can generalize the idea of bootstrapping robotic learning from prior data to include arbitrary sub-optimal data rather than just demonstration data and shows the ability to continue improving beyond this data as well.

8.6 EXPERIMENTAL EVALUATION

In our experimental evaluation we aim to answer the following question:

1. Does AWAC effectively combine prior data with online experience to learn complex robotic control tasks more efficiently than prior methods?
2. Is AWAC able to learn from sub-optimal or random data?
3. Does AWAC provide a practical way to bootstrap real-world robotic reinforcement learning?

In the following sections, we study these questions using several challenging and high-dimensional simulated robotic tasks, as well as three separate real-world robotic platforms. Videos of all experiments are available at awacrl.github.io

8.6.1 Simulated Experiments

We study the first two questions in challenging simulation environments.

8.6.1.1 Comparative Evaluation When Bootstrapping From Prior Data

We study tasks in simulation that have significant exploration challenges, where offline learning and online fine-tuning are likely to be effective. We begin our analysis with a set of challenging sparse reward dexterous manipulation tasks proposed by [Rajeswaran et al. \(2018\)](#) in simulation. These tasks involve complex manipulation skills using a 28-DoF five-fingered hand in the MuJoCo simulator ([Todorov et al., 2012](#)) shown in Figure 40: in-hand rotation of a pen, opening a door by unlatching the handle, and picking up a sphere and relocating it to a target location. The reward functions in these environments are binary 0-1 rewards for task completion.² [Rajeswaran et al. \(2018\)](#) provide 25 human demonstrations for each task, which are not fully optimal but do solve the task. Since this dataset is small, we generated another 500 trajectories of interaction data by constructing a behavioral cloned policy, and then sampling from this policy.

First, we compare our method on these dexterous manipulation tasks against prior methods for off-policy learning, offline learning, and bootstrapping from demonstrations. Specific implementation details are discussed in Appendix ???. The results are shown in Fig. 40. Our method is able to leverage the prior data to quickly attain good performance, and the efficient off-policy actor-critic component of our approach fine-tunes much more quickly than demonstration augmented policy gradient (DAPG), the method proposed by [Rajeswaran et al. \(2018\)](#). For example, our method solves the pen task in 120K timesteps, the equivalent of just 20 minutes of online interaction. While the baseline comparisons and ablations make some amount of progress on the pen task, alternative off-policy RL and offline RL algorithms are largely unable to solve the door and relocate task in the time-frame considered. We find that the design decisions to use off-policy critic estimation allow AWAC to outperform AWR ([Peng et al., 2019a](#)) while the implicit behavior modeling allows AWAC to significantly outperform ABM ([Noah Y. Siegel et al., 2020b](#)), although ABM does make some progress. [Rajeswaran et al. \(2018\)](#) show that DAPG can solve variants of these tasks with more well-shaped rewards, but requires considerably more samples.

² [Rajeswaran et al. \(2018\)](#) use a combination of task completion factors as the sparse reward. For instance, in the door task, the sparse reward as a function of the door position d was $r = 10\mathbb{1}_{d>1.35} + 8\mathbb{1}_{d>1.0} + 2\mathbb{1}_{d>1.2} - 0.1\|d - 1.57\|_2$. We only use the fully sparse success measure $r = \mathbb{1}_{d>1.4}$, which is substantially more difficult.

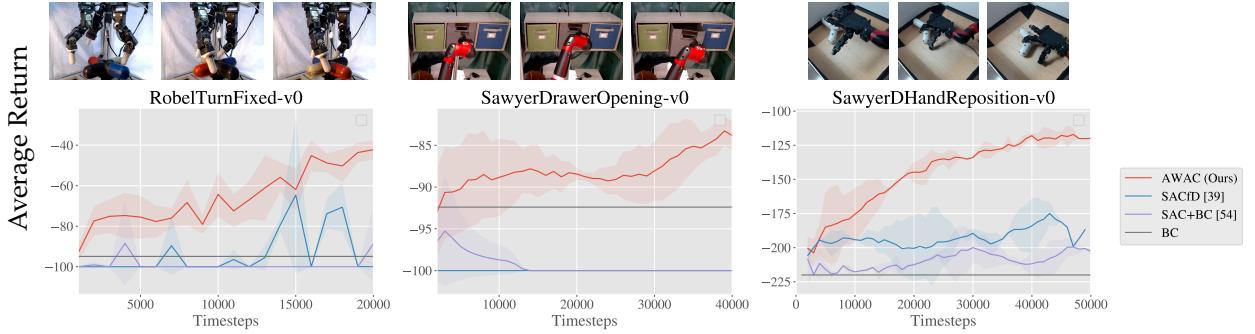


Figure 42: Algorithm comparison on three real-world robotic environments. Images of real world robotic tasks are pictured above. Left: a three fingered D’claw must rotate a valve 180° . Middle: a Sawyer robot must slide open a drawer using a hook attachment. Right: a dexterous hand attached to a Sawyer robot must re-position an object to the center of the table. On each task, AWAC trained offline achieves reasonable performance (shown at timestep 0) and then steadily improves from online interaction. Other methods, which also all have access to prior data, fail to utilize the prior data effectively offline and therefore exhibit slow or no online improvement. Videos of all experiments are available at awacrl.github.io

Additionally, we evaluated all methods on the Gym MuJoCo locomotion benchmarks, similarly providing demonstrations as offline data. The results plots for these experiments are included in Appendix ?? in the supplementary materials. These tasks are substantially easier than the sparse reward manipulation tasks described above, and a number of prior methods also perform well. SAC+BC and BRAC perform on par with our method on the HalfCheetah task, and ABM performs on par with our method on the Ant task, while our method outperforms all others on the Walker2D task. However, our method matches or exceeds the best prior method in all cases, whereas no other single prior method attains good performance on all tasks.

8.6.1.2 Fine-Tuning from Random Policy Data

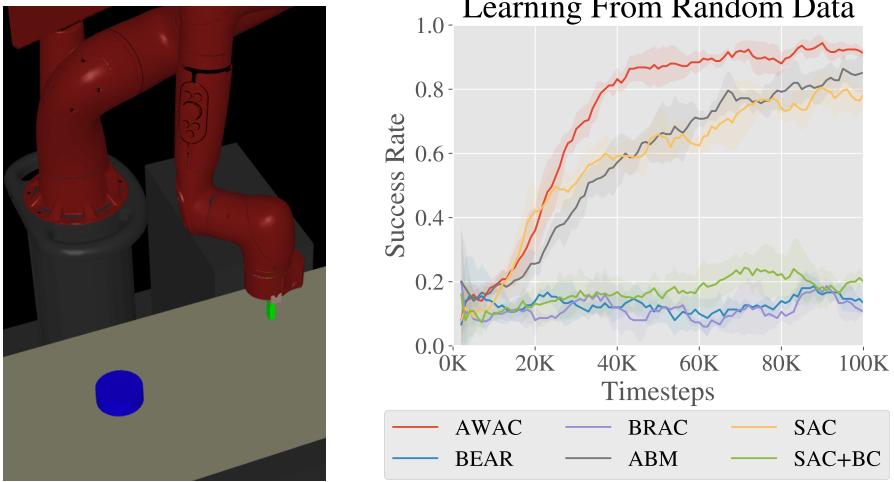


Figure 43: Comparison of fine-tuning from an initial dataset of suboptimal data on a Sawyer robot pushing task.

An advantage of using off-policy RL for reinforcement learning is that we can also incorporate suboptimal data, rather than demonstrations. In this experiment, we evaluate on a simulated tabletop pushing environment with a Sawyer robot pictured in Fig 40 and described further in Appendix ???. To study the potential to learn from suboptimal data, we use an off-policy dataset of 500 trajectories generated by a random process. The task is to push an object to a target location in a 40cm x 20cm goal space. The results are shown in Figure 43. We see that while many methods begin at the same initial performance, AWAC learns the fastest online and is actually able to make use of the offline dataset effectively.

8.6.2 Real-World Robot Learning with Prior Data

We next evaluate AWAC and several baselines on a range of real-world robotic systems, shown in the top row of Fig 42. We study the following tasks: rotating a valve with a 3-fingered claw, repositioning an object with a 4-fingered hand, and opening a drawer with a Sawyer robotic arm. The dexterous manipulation tasks involve fine finger coordination to properly reorient and reposition objects, as well as high dimensional state and action spaces. The Sawyer drawer opening task requires accurate arm movements to properly hook the end-effector into the handle of the drawer. To ensure continuous operation,

all environments are fitted with an automated reset mechanism that executes before each trajectory is collected, allowing us to run real-world experiments without human supervision. Since real-world experiments are significantly more time-consuming, we could not compare to the full range of prior methods in the real world, but we include comparisons with the following methods: direct behavioral cloning (BC) of the provided data (which is reasonable in these settings, since the prior data includes demonstrations), off-policy RL with soft actor-critic (SAC) (Haarnoja et al., 2018a), where the prior data is included in the replay buffer and used to pretrain the policy (which refer to as SACfD), and a modified version of SAC that includes an added behavioral cloning loss (SAC+BC), which is analogous to A. Nair et al. (2018a) or an off-policy version of Rajeswaran et al. (2018). Further implementation details of these algorithms are provided in Appendix ?? in the supplementary materials.

Next, we describe the experimental setup for hardware experiments. Precise details of the hardware setup can be found in Appendix ?? in the supplementary materials.

Dexterous Manipulation with a 3-Fingered Claw. This task requires controlling a 3-fingered, 9 DoF robotic hand, introduced by Ahn et al. (2019), to rotate a 4-pronged valve object by 180 degrees. To properly perform this task, multiple fingers need to coordinate to stably and efficiently rotate the valve into the desired orientation. The state space of the system consists of the joint angles of all the 9 joints in the claw, and the action space consists of the joint positions of the fingers, which are followed by the robot using a low-level PID controller. The reward for this task is sparse: -1 if the valve is rotated within 0.25 radians of the target, and 0 otherwise. Note that this reward function is significantly more difficult than the dense, well-shaped reward function typically used in prior work (Ahn et al., 2019). The prior data consists of 10 trajectories collected using kinesthetic teaching, combined this with 200 trajectories obtained through executing a policy trained via imitation learning in the environment.

Drawer Opening with a Sawyer Arm. This task requires controlling a Sawyer arm to slide open a drawer. The robot uses 3-dimensional end-effector control, and is equipped with a hook attachment to make the drawer opening possible. The state space is 4-dimensional, consisting of the position of the robot end-effector and the linear position of the drawer, measured using an encoder. The reward is sparse: -1 if the drawer is open beyond a threshold and 0 otherwise. For this task, the prior data consists of 10 demonstration trajectories collected using via teleoperation with a 3D mouse, as well as 500 trajectories obtained through executing a policy trained via imitation learning in the environment. This task is challenging because it requires very precise insertion of the hook attachment into the opening, as pictured in Fig 42, before the robot can open

the drawer. Due to the sparse reward, making learning progress on this task requires utilizing prior data to construct an initial policy that at least sometimes succeeds.

Dexterous Manipulation with a Robotic Hand. This task requires controlling a 4-fingered robotic hand mounted on a Sawyer robotic arm to reposition an object (Abhishek Gupta et al., 2021). The task requires careful coordination between the hand and the arm to manipulate the object accurately. The reward for this task is a combination of the negative distance between the hand and the object and the negative distance between the object and the target. The actions are 19-dimensional, consisting of 16-dimensional finger control and 3-dimensional end effector control of the arm. For this task, the prior data of 19 trajectories were collected using kinesthetic teaching and combined with 50 trajectories obtained by executing a policy trained with imitation learning on this data.

The results on these tasks are shown in Figure 42. We first see that AWAC attains performance that is comparable to the best prior method from offline training alone, as indicated by the value at time step 0 (which corresponds to the beginning of online finetuning). This means that, during online interaction, AWAC collects data that is of higher quality, and improves more quickly. The prior methods struggle to improve from online training on these tasks, likely because the sparse reward function and challenging dynamics make progress very difficult from a bad initialization. These results suggest that AWAC is effectively able to leverage prior data to bootstrap online reinforcement learning in the real world, even on tasks with difficult and uninformative reward functions.

8.7 DISCUSSION AND FUTURE WORK

We have discussed the challenges existing RL methods face when fine-tuning from prior datasets, and proposed an algorithm, AWAC, for this setting. The key insight in AWAC is that an implicitly constrained actor-critic algorithm is able to both train offline and continue to improve with more experience. We provide detailed empirical analysis of the design decisions behind AWAC, showing the importance of off-policy learning, bootstrapping and the particular form of implicit constraint enforcement. To validate these ideas, we evaluate on a variety of simulated and real world robotic problems.

While AWAC is able to effectively leverage prior data for significantly accelerating learning, it does run into some limitations. Firstly, it can be challenging to choose the appropriate threshold for constrained optimization. Resolving this would involve exploring adaptive threshold tuning schemes. Secondly, while AWAC is able to avoid over-conservative behavior empirically, in future work, we hope to analyze theoretical factors that go into building a good finetuning algorithm. And lastly, in the future we plan on

applying AWAC to more broadly incorporate data across different robots, labs and tasks rather than just on isolated setups. By doing so, we hope to enable an even wider array of robotic applications.

9

OFFLINE REINFORCEMENT LEARNING WITH IMPLICIT Q-LEARNING

Offline reinforcement learning (RL) addresses the problem of learning effective policies entirely from previously collected data, without online interaction (Fujimoto et al., 2019b; Lange et al., 2012a). This is very appealing in a range of real-world domains, from robotics to logistics and operations research, where real-world exploration with untrained policies is costly or dangerous, but prior data is available. However, this also carries with it major challenges: improving the policy beyond the level of the behavior policy that collected the data requires estimating values for actions other than those that were seen in the dataset, and this, in turn, requires trading off policy improvement against distributional shift, since the values of actions that are too different from those in the data are unlikely to be estimated accurately. Prior methods generally address this by either constraining the policy to limit how far it deviates from the behavior policy (Fujimoto et al., 2019b; Yifan Wu et al., 2019; Fujimoto and S. S. Gu, 2021; Kumar et al., 2019b; A. Nair et al., 2020; Z. Wang et al., 2020a), or by regularizing the learned value functions to assign low values to out-of-distribution actions (Kumar et al., 2020; Kostrikov et al., 2021a). Nevertheless, this imposes a trade-off between how much the policy improves and how vulnerable it is to misestimation due to distributional shift. Can we devise an offline RL method that avoids this issue by never needing to directly query or estimate values for actions that were not seen in the data?

In this work, we start from an observation that *in-distribution* constraints widely used in prior work might not be sufficient to avoid value function extrapolation, and we ask whether it is possible to learn an optimal policy with *in-sample learning*, without *ever* querying the values of any unseen actions. The key idea in our method is to approximate an upper expectile of the distribution over values with respect to the distribution of dataset actions for each state. We alternate between fitting this value function with expectile regression, and then using it to compute Bellman backups for training the Q-function.

We show that we can do this simply by modifying the loss function in a SARSA-style TD backup, without ever using out-of-sample actions in the target value. Once this Q-function has converged, we extract the corresponding policy using advantage-weighted behavioral cloning. This approach does not require explicit constraints or explicit regularization of out-of-distribution actions during value function training, though our policy extraction step does implicitly enforce a constraint, as discussed in prior work on advantage-weighted regression (Peters and Schaal, 2007b; Peng et al., 2019b; A. Nair et al., 2020; Z. Wang et al., 2020a).

Our main contribution is implicit Q-learning (IQL), a new offline RL algorithm that avoids ever querying values of unseen actions while still being able to perform multi-step dynamic programming updates. Our method is easy to implement by making a small change to the loss function in a simple SARSA-like TD update and is computationally very efficient. Furthermore, our approach demonstrates the state-of-the-art performance on D4RL, a popular benchmark for offline reinforcement learning. In particular, our approach significantly improves over the prior state-of-the-art on challenging Ant Maze tasks that require to “stitch” several sub-optimal trajectories. Finally, we demonstrate that our approach is suitable for finetuning; after initialization from offline RL, IQL is capable of improving policy performance utilizing additional interactions.

9.1 RELATED WORK

A significant portion of recently proposed offline RL methods are based on either constrained or regularized approximate dynamic programming (e.g., Q-learning or actor-critic methods), with the constraint or regularizer serving to limit deviation from the behavior policy. We will refer to these methods as “multi-step dynamic programming” algorithms, since they perform true dynamic programming for multiple iterations, and therefore can in principle recover the optimal policy if provided with high-coverage data. The constraints can be implemented via an explicit density model (Yifan Wu et al., 2019; Fujimoto et al., 2019b; Kumar et al., 2019b; Ghasemipour et al., 2021), implicit divergence constraints (A. Nair et al., 2020; Z. Wang et al., 2020a; Peters and Schaal, 2007b; Peng et al., 2019b; Noah Y Siegel et al., 2020a), or by adding a supervised learning term to the policy improvement objective (Fujimoto and S. S. Gu, 2021). Several works have also proposed to directly regularize the Q-function to produce low values for out-of-distribution actions (Kostrikov et al., 2021a; Kumar et al., 2020; Fakoor et al., 2021). Our method is also a multi-step dynamic programming algorithm. However, in contrast to prior works, our method completely avoids directly querying the learned Q-function with unseen actions during training, removing the need for any constraint during this stage, though the

subsequent policy extraction, which is based on advantage-weighted regression (Peng et al., 2019b; A. Nair et al., 2020), does apply an implicit constraint. However, this policy does not actually influence value function training.

In contrast to multi-step dynamic programming methods, several recent works have proposed methods that rely either on a single step of policy iteration, fitting the value function or Q-function of the behavior policy and then extracting the corresponding greedy policy (Peng et al., 2019b; Brandfonbrener et al., 2021; Gulcehre et al., 2021), or else avoid value functions completely and utilize behavioral cloning-style objectives (L. Chen et al., 2021). We collectively refer to these as “single-step” approaches. These methods avoid needing to query unseen actions as well, since they either use no value function at all, or learn the value function of the behavior policy. Although these methods are simple to implement and effective on the MuJoCo locomotion tasks in D4RL, we show that such single-step methods perform very poorly on more complex datasets in D4RL, which require combining parts of suboptimal trajectories (“stitching”). Prior multi-step dynamic programming methods perform much better in such settings, as does our method. We discuss this distinction in more detail in Section 9.4.1. Our method also shares the simplicity and computational efficiency of single-step approaches, providing an appealing combination of the strengths of both types of methods.

Our method is based on estimating the characteristics of a random variable. Several recent works involve approximating statistical quantities of the value function distribution. In particular, quantile regression (Koenker and Hallock, 2001) has been previously used in reinforcement learning to estimate the quantile function of a state-action value function (Dabney et al., 2018b; Dabney et al., 2018a; Kuznetsov et al., 2020). Although our method is related, in that we perform expectile regression, our aim is not to estimate the distribution of values that results from stochastic transitions, but rather estimate expectiles of the state value function with respect to random actions. This is a very different statistic: our aim is not to determine how the Q-value can vary with different future outcomes, but how the Q-value can vary with different actions *while averaging together future outcomes due to stochastic dynamics*. While prior work on distributional RL can also be used for offline RL, it would suffer from the same action extrapolation issues as other methods, and would require similar constraints or regularization, while our method does not.

9.2 PRELIMINARIES

The RL problem is formulated in the context of a Markov decision process (MDP) $(\mathbf{S}, p_0(s), p(s'|s, a), r(s, a), \gamma)$, where \mathbf{S} is a state space, a is an action space, $p_0(s)$ is a distribution of initial states, $p(s'|s, a)$ is the environment dynamics, $r(s, a)$ is a reward function, and γ is a discount factor. The agent interacts with the MDP according to a policy $\pi(a|s)$. The goal is to obtain a policy that maximizes the cumulative discounted returns:

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{\pi} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 \sim p_0(\cdot), a_t \sim \pi(\cdot|s_t), s_{t+1} \sim p(\cdot|s_t, a_t) \right].$$

Off-policy RL methods based on approximate dynamic programming typically utilize a state-action value function (Q-function), referred to as $Q(s, a)$, which corresponds to the discounted returns obtained by starting from the state s and action a , and then following the policy π .

OFFLINE REINFORCEMENT LEARNING. In contrast to online (on-policy or off-policy) RL methods, offline RL uses previously collected data without any additional data collection. Like many recent offline RL methods, our work builds on approximate dynamic programming methods that minimize temporal difference error, according to the following loss:

$$L_{TD}(\theta) = \mathbb{E}_{(s, a, s') \sim [(\mathcal{D})]} [(r(s, a) + \gamma \max_{a'} Q_{\hat{\theta}}(s', a') - Q_{\theta}(s, a))^2], \quad (47)$$

where \mathcal{D} is the dataset, $Q_{\theta}(s, a)$ is a parameterized Q-function, $Q_{\hat{\theta}}(s, a)$ is a target network (e.g., with soft parameters updates defined via Polyak averaging), and the policy is defined as $\pi(s) = \arg \max_a Q_{\theta}(s, a)$. Most recent offline RL methods modify either the value function loss (above) to regularize the value function in a way that keeps the resulting policy close to the data, or constrain the arg max policy directly. This is important because out-of-distribution actions a' can produce erroneous values for $Q_{\hat{\theta}}(s', a')$ in the above objective, often leading to *overestimation* as the policy is defined to maximize the (estimated) Q-value.

9.3 IMPLICIT Q-LEARNING

In this work, we aim to entirely avoid querying *out-of-sample* (unseen) actions in our TD loss. Although the goal of this work is to approximate the optimal Q-function, we start by considering fitted Q evaluation with a SARSA-style objective which has been considered in prior work on Offline Reinforcement Learning ([Brandfonbrener et al., 2021](#); [Gulcehre](#)

et al., 2021). This objective aims to learn the value of the dataset policy π_β (also called the behavior policy):

$$L(\theta) = \mathbb{E}_{(s,a,s',a') \sim [(\mathbf{r}(s, a) + \gamma Q_{\hat{\theta}}(s', a') - Q_\theta(s, a))^2]}. \quad (48)$$

This objective never queries values for out-of-sample actions, in contrast to Equation (47). One specific property of this objective that is important for this work is that it uses mean squared error (MSE) that fits $Q_\theta(s, a)$ to predict the mean statistics of the TD targets. Thus, if we assume unlimited capacity and no sampling error, the optimal parameters should satisfy

$$Q_{\theta^*}(s, a) \approx \mathbf{r}(s, a) + \gamma \mathbb{E}_{\substack{s' \sim p(\cdot|s, a) \\ a' \sim \pi_\beta(\cdot|s)}} [Q_{\hat{\theta}}(s', a')]. \quad (49)$$

Prior work (Brandfonbrener et al., 2021; Gulcehre et al., 2021; Peng et al., 2019b) has proposed directly using this objective to learn Q^{π_β} , and then train the policy π_ψ to maximize Q^{π_β} . This avoids any issues with out-of-distribution actions, since the TD loss only uses dataset actions. However, while this procedure works well empirically on simple MuJoCo locomotion tasks in D4RL, we will show that it performs very poorly on more complex tasks that benefit from multi-step dynamic programming. In our method, which we derive next, we retain the benefits of using this SARSA-like objective, but modify it so that it allows us to perform multi-step dynamic programming and learn a near-optimal Q -function.

Our method will perform a Q -function update similar to Equation (48), but we will aim to estimate the maximum Q -value over actions that are in the support of the data distribution. Crucially, we will show that it is possible to do this *without ever querying the learned Q -function on out-of-sample actions* by utilizing expectile regression. Formally, the value function we aim to learn is given by:

$$L(\theta) = \mathbb{E}_{(s,a,s') \sim [(\mathbf{r}(s, a) + \gamma \max_{\substack{a' \in \\ \text{s.t. } \pi_\beta(a'|s') > 0}} Q_{\hat{\theta}}(s', a') - Q_\theta(s, a))^2]}. \quad (50)$$

Our algorithm, implicit Q-Learning (IQL), aims to estimate this objective while evaluating the Q -function only on the state-action pairs in the dataset. To this end, we propose to fit $Q_\theta(s, a)$ to estimate state-conditional expectiles of the target values, and show that specific expectiles approximate the maximization defined above. In Section 9.3.4 we show that this approach performs multi-step dynamic programming in theory, and in Section 9.4.1 we show that it does so in practice.

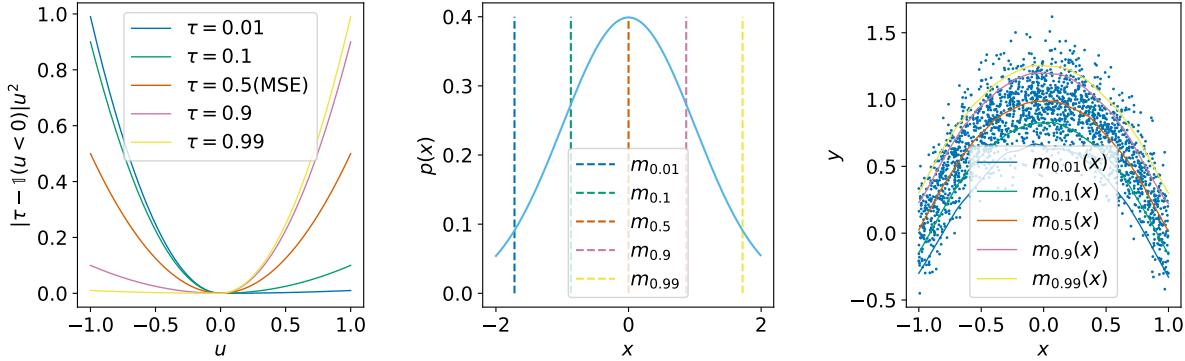


Figure 44: **Left:** The asymmetric squared loss used for expectile regression. $\tau = 0.5$ corresponds to the standard mean squared error loss, while $\tau = 0.9$ gives more weight to positive differences. **Center:** Expectiles of a normal distribution. **Right:** an example of estimating state conditional expectiles of a two-dimensional random variable. Each x corresponds to a distribution over y . We can approximate a maximum of this random variable with expectile regression: $\tau = 0.5$ corresponds to the conditional mean statistics of the distribution, while $\tau \approx 1$ approximates the maximum operator over in-support values of y .

9.3.1 Expectile Regression

Practical methods for estimating various statistics of a random variable have been thoroughly studied in applied statistics and econometrics. The $\tau \in (0, 1)$ expectile of some random variable X is defined as a solution to the asymmetric least squares problem:

$$\arg \min_{m_\tau} \mathbb{E}_{x \sim X} [L_2^\tau(x - m_\tau)], \text{ where } L_2^\tau(u) = |\tau - 1(u < 0)|u^2.$$

That is, for $\tau > 0.5$, this asymmetric loss function downweights the contributions of x values smaller than m_τ while giving more weights to larger values (see Figure 44, left). Expectile regression is closely related to quantile regression (Koenker and Hallock, 2001), which is a popular technique for estimating quantiles of a distribution widely used in reinforcement learning (Dabney et al., 2018b; Dabney et al., 2018a)¹. The quantile regression loss is defined as an asymmetric ℓ_1 loss.

¹ Our method could also be derived with quantiles, but since we are not interested in learning all of the expectiles/quantiles, unlike prior work (Dabney et al., 2018b; Dabney et al., 2018a), it is more convenient to estimate a single expectile because this involves a simple modification to the MSE loss that is already used in standard RL methods. We found it to work somewhat better than quantile regression with its corresponding ℓ_1 loss.

We can also use this formulation to predict expectiles of a conditional distribution:

$$\arg \min_{m_\tau(x)} \mathbb{E}_{(x,y) \sim [L_2^\tau(y - m_\tau(x))]}.$$

Figure 44 (right) illustrates conditional expectile regression on a simple two-dimensional distribution. Note that we can optimize this objective with stochastic gradient descent. It provides unbiased gradients and is easy to implement with standard machine learning libraries.

9.3.2 Learning the Value Function with Expectile Regression

Expectile regression provides us with a powerful framework to estimate statistics of a random variable beyond mean regression. We can use expectile regression to modify the policy evaluation objective in Equation (48) to predict an upper expectile of the TD targets that approximates the maximum of $r(s, a) + \gamma Q_{\hat{\theta}}(s', a')$ over actions a' constrained to the dataset actions, as in Equation (50). This leads to the following expectile regression objective:

$$L(\theta) = \mathbb{E}_{(s,a,s',a') \sim [L_2^\tau(r(s, a) + \gamma Q_{\hat{\theta}}(s', a') - Q_\theta(s, a))]}.$$

However, this formulation has a significant drawback. Instead of estimating expectiles just with respect to the actions in the support of the data, it also incorporates stochasticity that comes from the environment dynamics $s' \sim p(\cdot|s, a)$. Therefore, a large target value might not necessarily reflect the existence of a single action that achieves that value, but rather a “lucky” sample that happened to have transitioned into a good state. We resolve this by introducing a separate value function that approximates an expectile only with respect to the action distribution, leading to the following loss:

$$L_V(\psi) = \mathbb{E}_{(s,a) \sim [L_2^\tau(Q_{\hat{\theta}}(s, a) - V_\psi(s))]}.$$
 (51)

We can then use this estimate to update the Q-functions with the MSE loss, which averages over the stochasticity from the transitions and avoids the “lucky” sample issue mentioned above:

$$L_Q(\theta) = \mathbb{E}_{(s,a,s') \sim [(r(s, a) + \gamma V_\psi(s') - Q_\theta(s, a))^2]}.$$
 (52)

Note that these losses do not use any explicit policy, and only utilize actions from the dataset for both objectives, similarly to SARSA-style policy evaluation. In Section 9.3.4, we will show that this procedure recovers the optimal Q-function under some assumptions. Also, even though only one action is available for every state in the dataset for continuous action spaces, due to neural network generalization, the expectile regression does not result in SARSA-style policy evaluation as shown in Section 9.4.2.

9.3.3 Policy Extraction and Algorithm Summary

While our modified TD learning procedure learns an approximation to the optimal Q-function, it does not explicitly represent the corresponding policy, and therefore requires a separate policy extraction step. While one can consider any technique for policy extraction that constrains the learned policy to stay close to the dataset actions, we aim for a simple method for policy extraction. As before, we aim to avoid using out-of-samples actions. Therefore, we extract the policy with advantage-weighted regression (Peters and Schaal, 2007b; Peng et al., 2019b) previously successfully used for policy extraction in Offline RL (Q. Wang et al., 2018b; A. Nair et al., 2020; Brandfonbrener et al., 2021):

$$L_\pi(\phi) = \mathbb{E}_{(s,a) \sim [\exp(\beta(Q_{\hat{\theta}}(s, a) - V_\psi(s))) \log \pi_\phi(a|s)]}, \quad (53)$$

where $\beta \in [0, \infty)$ is an inverse temperature. Note that this objective does not clone all actions from the dataset but, as shown in prior work, this objective learns a policy that maximizes the Q-values subject to a distribution constraint (Peters and Schaal, 2007b; Peng et al., 2019b; A. Nair et al., 2020). This step can be seen as selecting and cloning the most optimal actions in the dataset.

Our final algorithm consists of two stages. First, we fit the value function and Q, performing a number of gradient updates alternating between Eqn. (51) and (52). Second, we perform stochastic gradient descent on Equation (53). For both steps, we use a version of clipped double Q-learning (Fujimoto et al., 2018a), taking a minimum of two Q-functions for V-function and policy updates. We summarize our final method in Algorithm 7. Note that the policy does not influence the value function in any way, and

Algorithm 7 Implicit Q-learning

```

o: Initialize parameters  $\psi, \theta, \hat{\theta}, \phi$ .
o: TD learning (IQL):
o: for each gradient step do
o:    $\psi \leftarrow \psi - \lambda_V \nabla_\psi L_V(\psi)$ 
o:    $\theta \leftarrow \theta - \lambda_Q \nabla_\theta L_Q(\theta)$ 
o:    $\hat{\theta} \leftarrow (1 - \alpha)\hat{\theta} + \alpha\theta$ 
o: Policy extraction (AWR):
o: for each gradient step do
o:    $\phi \leftarrow \phi - \lambda_\pi \nabla_\phi L_\pi(\phi)$ 
=0

```

therefore extraction could be performed either concurrently or after TD learning. Concurrent learning provides a way to use IQL with online finetuning, as we discuss in Section 9.4.3.

9.3.4 Analysis

In this section, we will show that IQL can recover the optimal value function under the dataset support constraints. First, we prove a simple lemma that we will then use to show how our approach can enable learning the optimal value function.

Lemma 3. *Let X be a real-valued random variable with a bounded support and supremum of the support is x^* . Then,*

$$\lim_{\tau \rightarrow 1} m_\tau = x^*$$

One can show that expectiles of a random variable have the same supremum x^* . Moreover, for all τ_1 and τ_2 such that $\tau_1 < \tau_2$, we get $m_{\tau_1} \leq m_{\tau_2}$. Therefore, the limit follows from the properties of bounded monotonically non-decreasing functions.

In the following theorems, we show that under certain assumptions, our method indeed approximates the optimal state-action value Q^* and performs multi-step dynamical programming. We first prove a technical lemma relating different expectiles of the Q -function, and then derive our main result regarding the optimality of our method.

For the sake of simplicity, we introduce the following notation for our analysis. Let $\mathbb{E}_{x \sim X}^\tau[x]$ be a τ^{th} expectile of X (e.g., $\mathbb{E}^{0.5}$ corresponds to the standard expectation). Then, we define $V_\tau(s)$ and $Q_\tau(s, a)$, which correspond to optimal solutions of Eqn. 51 and 52 correspondingly, recursively as:

$$\begin{aligned} V_\tau(s) &= \mathbb{E}_{a \sim \pi_\beta(\cdot|s)}^\tau [Q_\tau(s, a)], \\ Q_\tau(s, a) &= r(s, a) + \gamma \mathbb{E}_{s' \sim p(\cdot|s, a)} [V_\tau(s')]. \end{aligned}$$

Lemma 4. *For all s, τ_1 and τ_2 such that $\tau_1 < \tau_2$ we get $V_{\tau_1}(s) \leq V_{\tau_2}(s)$.*

Proof. The proof follows the policy improvement proof (Sutton and Barto, 2018). See ??.

Corollary 1. For any τ and s we have $V_\tau(s) \leq \max_{\substack{a \in \\ s.t. \pi_\beta(a|s) > 0}} Q^*(s, a)$ where $V_\tau(s)$ is defined as above and $Q^*(s, a)$ is an optimal state-action value function constrained to the dataset and defined as

$$Q^*(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim p(\cdot|s, a)} \left[\max_{\substack{a' \in \\ s.t. \pi_\beta(a'|s') > 0}} Q^*(s', a') \right].$$

Proof. The proof follows from the observation that convex combination is smaller than maximum. \square

Theorem 1.

$$\lim_{\tau \rightarrow 1} V_\tau(s) = \max_{\substack{a \in \\ s.t. \pi_\beta(a|s) > 0}} Q^*(s, a).$$

Proof. Follows from combining Lemma 3 and Corollary 1. \square

Therefore, for a larger value of $\tau < 1$, we get a better approximation of the maximum. On the other hand, it also becomes a more challenging optimization problem. Thus, we treat τ as a hyperparameter. Due to the property discussed in Theorem 1 we dub our method implicit Q-learning (IQL). We also emphasize that our value learning method defines the entire spectrum of methods between SARSA ($\tau = 0.5$) and Q-Learning ($\tau \rightarrow 1$). Note that in contrast to other multi-step methods, IQL absorbs the policy improvement step into value learning. Therefore, fitting Q-function corresponds to the policy evaluation step, while fitting the value function with IQL corresponds to *implicit* policy improvement.

9.4 EXPERIMENTAL EVALUATION

Our experiments aim to evaluate our method comparatively, in contrast to prior offline RL methods, and in particular to understand how our approach compares both to single-step methods and multi-step dynamic programming approaches. We will first demonstrate the benefits of multi-step dynamic programming methods, such as ours, in contrast to single-step methods, showing that on some problems this difference can be extremely large. We will then compare IQL with state-of-the-art single-step and multi-step algorithms on the D4RL (Fu et al., 2020) benchmark tasks, studying the degree to which we can learn effective policies using only the actions in the dataset. We examine domains that contain near-optimal trajectories, where single-step methods perform well, as well

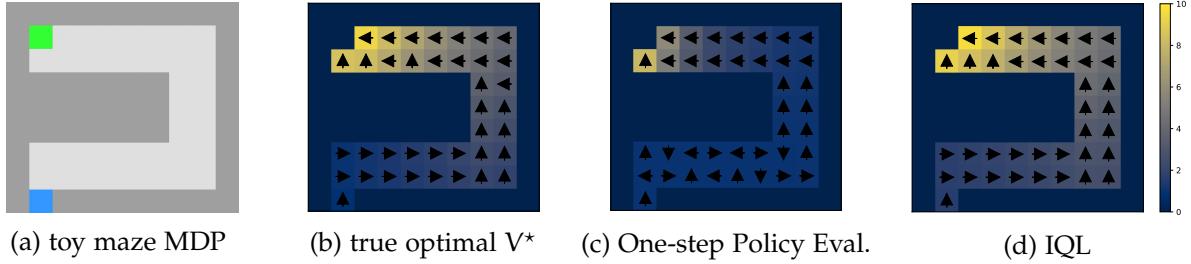


Figure 45: Evaluation of our algorithm on a toy umaze environment (a). When the static dataset is heavily corrupted by suboptimal actions, one-step policy evaluation results in a value function that degrades to zero far from the rewarding states too quickly (c). Our algorithm aims to learn a near-optimal value function, combining the best properties of SARSA-style evaluation with the ability to perform multi-step dynamic programming, leading to value functions that are much closer to optimality (shown in (b)) and producing a much better policy (d).

as domains with no optimal trajectories at all, which require multi-step dynamic programming. Finally, we will study how IQL compares to prior methods when finetuning with online RL starting from an offline RL initialization.

9.4.1 The Difference Between One-Step Policy Improvement and IQL

We first use a simple maze environment to illustrate the importance of multi-step dynamic programming for offline RL. The maze has a u-shape, a single start state, and a single goal state (see Figure 45a). The agent receives a reward of 10 for entering the goal state and zero reward for all other transitions. With a probability of 0.25, the agent transitions to a random state, and otherwise to the commanded state. The dataset consists of 1 optimal trajectory and 99 trajectories with uniform random actions. Due to a short horizon of the problem, we use $\gamma = 0.9$.

Figure 45 (c, d) illustrates the difference between single-step methods which fit $Q^{\pi}(s, a)$ via SARSA-style objective, in this case represented by One step RL (Brandfonbrener et al., 2021; Q. Wang et al., 2018b; Gulcehre et al., 2021) and IQL with $\tau = 0.95$. Note that these methods represent a special case of our method with $\tau = 0.5$. Although states closer to the high reward state will still have higher values, these values decay much faster as we move further away than they would for the optimal value function, and the resulting policy is highly suboptimal. Since IQL (d) performs iterative dynamic programming, it correctly propagates the signal, and the values are no longer dominated by noise. The resulting value function closely matches the true optimal value function (b).

9.4.2 Comparisons on Offline RL Benchmarks

Next, we evaluate our approach on the D4RL benchmark in comparison to prior methods (see Table 3). The MuJoCo tasks in D4RL consist of the Gym locomotion tasks, the Ant Maze tasks, and the Adroit and Kitchen robotic manipulation environments. Some prior works, particularly those proposing one-step methods, focus entirely on the Gym locomotion tasks. However, these tasks include a significant fraction of near-optimal trajectories in the dataset. In contrast, the Ant Maze tasks, especially the medium and large ones, contain very few or no near-optimal trajectories, making them very challenging for one-step methods. These domains require “stitching” parts of suboptimal trajectories

that travel between different states to find a path from the start to the goal of the maze (Fu et al., 2020). As we will show, multi-step dynamic programming is essential in these domains. The Adroit and Kitchen tasks are comparatively less discriminating, and we found that most RL methods perform similarly to imitation learning in these domains (Florence et al., 2021). We therefore focus our analysis on the Gym locomotion and Ant Maze domains, but include full Adroit and Kitchen results in ?? for completeness.

COMPARISONS AND BASELINES. We compare to methods that are representative of both multi-step dynamic programming and one-step approaches. In the former category, we compare to CQL (Kumar et al., 2020), TD3+BC (Fujimoto and S. S. Gu, 2021), and AWAC (A. Nair et al., 2020). In the latter category, we compare to Onestep RL (Brandfonbrener et al., 2021) and Decision Transformers (L. Chen et al., 2021). We obtained the Decision Transformers results on Ant Maze subsets of D4RL tasks using the author-provided implementation² and following authors instructions communicated over email. We obtained results for TD3+BC and Onestep RL (Exp. Weight) directly from the authors. Note that L. Chen et al. (2021) and Brandfonbrener et al. (2021) incorrectly report results for some prior methods, such as CQL, using the “-vo” environments. These gen-

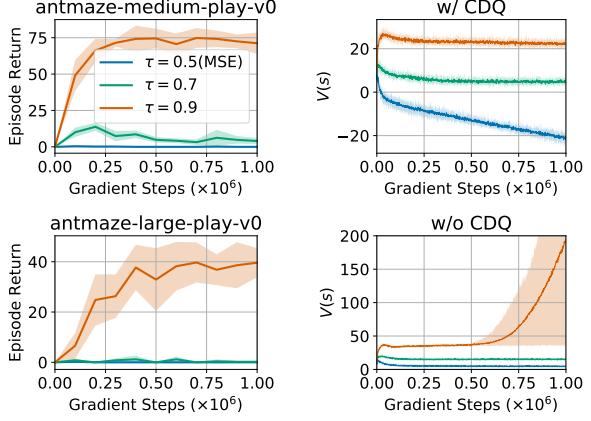


Figure 46: **Left:** Estimating a larger expectile τ is crucial for antmaze tasks that require dynamical programming (“stitching”). **Right:** Clipped double Q-Learning (CDQ) is crucial for learning values for $\tau = 0.9$.

² <https://github.com/kzl/decision-transformer>

erally produce lower scores than the “-v2” environments that these papers use for their own methods. We use the “-v2” environments for all methods to ensure a fair comparison, resulting in higher values for CQL. Because of this fix, our reported CQL scores are higher than all other prior methods. We obtained results for “-v2” datasets using an author-suggested implementation.³ On the Gym locomotion tasks (halfcheetah, hopper, walker2d), we find that IQL performs comparably to the best performing prior method, CQL. On the more challenging Ant Maze task, IQL outperforms CQL, and outperforms the one-step methods by a very large margin.

RUNTIME. Our approach is also computationally faster than the baselines (see Table 3). For the baselines, we measure runtime for our reimplementations of the methods in JAX (Bradbury et al., 2018) built on top of JAXRL (Kostrikov, 2021), which are typically faster than the original implementations. For example, the original implementation of CQL takes more than 4 hours to perform 1M updates, while ours takes only 80 minutes. Even so, IQL still requires about 4x less time than our reimplementation of CQL on average, and is comparable to the fastest prior one-step methods. We did not reimplement Decision Transformers due to their complexity and report runtime of the original implementation.

EFFECT OF τ HYPERPARAMETER. We also demonstrate that it is crucial to compute a larger expectile on tasks that require “stitching” (see Figure 46). We provide complete results in ???. With larger values of τ , our method approximates Q-learning better, leading to better performance on the Ant Maze tasks. Moreover, due to neural network generalization, values learned with expectile regression increase with a larger τ and do not degrade to behavior policy values ($\tau = 0.5$). Finally, clipped double Q-Learning is crucial for estimating values for a larger $\tau = 0.9$.

9.4.3 *Online Fine-tuning after Offline RL*

The policies obtained by offline RL can often be improved with a small amount of online interaction. IQL is well-suited for online fine-tuning for two reasons. First, IQL has strong offline performance, as shown in the previous section, which provides a good initialization. Second, IQL implements a weighted behavioral cloning policy extraction step, which has previously been shown to allow for better online policy improvement compared to other types of offline constraints (A. Nair et al., 2020). To evaluate the finetuning capability of various RL algorithms, we first run offline RL on each dataset, then run 1M steps of online RL, and then report the final performance. We compare to

³ <https://github.com/young-geng/CQL>

AWAC ([A. Nair et al., 2020](#)), which has been proposed specifically for online finetuning, and CQL ([Kumar et al., 2020](#)), which showed the best performance among prior methods in our experiments in the previous section. Exact experimental details are provided in Appendix ???. We use the challenging Ant Maze D4RL domains ([Fu et al., 2020](#)), as well as the high-dimensional dexterous manipulation environments from [Rajeswaran et al. \(2018\)](#), which [A. Nair et al. \(2020\)](#) propose to use to study online adaptation with AWAC. Results are shown in Table 4. On the Ant Maze domains, IQL significantly outperforms both prior methods after online finetuning. CQL attains the second best score, while AWAC performs comparatively worse due to much weaker offline initialization. On the dexterous hand tasks, IQL performs significantly better than AWAC on relocate-binary-vo, comparably on door-binary-vo, and slightly worse on pen-binary-vo, with the best overall score.

9.5 CONCLUSION

We presented implicit Q-Learning (IQL), a general algorithm for offline RL that completely avoids any queries to values of out-of-sample actions during training while still enabling multi-step dynamic programming. To our knowledge, this is the first method that combines both of these features. This has a number of important benefits. First, our algorithm is computationally efficient: we can perform 1M updates on one GTX1080 GPU in less than 20 minutes. Second, it is simple to implement, requiring only minor modifications over a standard SARSA-like TD algorithm, and performing policy extraction with a simple weighted behavioral cloning procedure resembling supervised learning. Finally, despite the simplicity and efficiency of this method, we show that it attains excellent performance across all of the tasks in the D4RL benchmark, matching the best prior methods on the MuJoCo locomotion tasks, and exceeding the state-of-the-art performance on the challenging ant maze environments, where multi-step dynamic programming is essential for good performance.

Table 3: Averaged normalized scores on MuJoCo locomotion and Ant Maze tasks. Our method outperforms prior methods on the challenging Ant Maze tasks, which require dynamic programming, and is competitive with the best prior methods on the locomotion tasks.

Dataset	BC	10%BC	BCQ	DT	ABM	AWAC	Oonestep RL	TD ₃ +BC	CQL	IQL (Ours)
halfcheetah-m-v2	42.6	42.5	47.0	42.6 ± 0.1	53.6	43.5	48.4 \pm 0.1	48.3 \pm 0.3	44.0 ± 5.4	47.4 \pm 0.2
hopper-m-v2	52.9	56.9	56.7	67.6 \pm 1.0	0.7	57.0	59.6 \pm 2.5	59.3 ± 4.2	58.5 ± 2.1	66.2 \pm 5.7
walker2d-m-v2	75.3	75.0	72.6	74.0 ± 1.4	0.5	72.4	81.8 \pm 2.2	83.7 ± 2.1	72.5 ± 0.8	78.3 \pm 8.7
halfcheetah-m-r-v2	36.6	40.6	40.4	36.6 ± 0.8	50.5	40.5	38.1 ± 1.3	44.6 \pm 0.5	45.5 \pm 0.5	44.2 \pm 1.2
hopper-m-r-v2	18.1	75.9	53.3	82.7 ± 7.0	49.6	37.2	97.5 \pm 0.7	60.9 ± 18.8	95.0 \pm 6.4	94.7 \pm 8.6
walker2d-m-r-v2	26.0	62.5	52.1	66.6 ± 3.0	53.8	27.0	49.5 ± 12.0	81.8 \pm 5.5	77.2 ± 5.5	73.8 ± 7.1
halfcheetah-m-e-v2	55.2	92.9	89.1	86.8 ± 1.3	18.5	42.8	93.4 \pm 1.6	90.7 \pm 4.3	91.6 \pm 2.8	86.7 ± 5.3
hopper-m-e-v2	52.5	110.9	81.8	107.6 \pm 1.8	0.7	55.8	103.3 ± 1.9	98.0 ± 9.4	105.4 \pm 6.8	91.5 ± 14.3
walker2d-m-e-v2	107.5	109.0	109.5	108.1 \pm 0.2	3.5	74.5	113.0 \pm 0.4	110.1 \pm 0.5	108.8 \pm 0.7	109.6 \pm 1.0
locomotion-v2 total	466.7	666.2	602.5	672.6 \pm 16.6	231.4	450.7	684.6 \pm 22.7	677.4 \pm 44.5	698.5 \pm 31.0	692.4 \pm 52.1
antmaze-u-vo	54.6	62.8	89.8	59.2	59.9	56.7	64.3	78.6	74.0	87.5 \pm 2.6
antmaze-u-d-vo	45.6	50.2	83.0	53.0	48.7	49.3	60.7	71.4	84.0	62.2 ± 13.8
antmaze-m-p-vo	0.0	5.4	15.0	0.0	0.0	0.0	0.3	10.6	61.2	71.2 \pm 7.3
antmaze-m-d-vo	0.0	9.8	0.0	0.0	0.5	0.7	0.0	3.0	53.7	70.0 \pm 10.9
antmaze-l-p-vo	0.0	0.0	0.0	0.0	0.	0.0	0.0	0.2	15.8	39.6 \pm 5.8
antmaze-l-d-vo	0.0	6.0	0.0	0.0	0.0	1.0	0.0	0.0	14.9	47.5 \pm 9.5
antmaze-vo total	100.2	134.2	187.8	112.2	109.1	107.7	125.3	163.8	303.6	378.0 \pm 49.9
total	566.9	800.4	790.3	784.8	340.5	558.4	809.9	841.2	1002.1	1070.4 \pm 102.0
kitchen-vo total	154.5	-	-	-	-	-	-	-	144.6	159.8 \pm 22.6
adroit-vo total	104.5	-	-	-	-	-	-	-	93.6	118.1 \pm 30.7
total+kitchen+adroit	825.9	-	-	-	-	-	-	-	1240.3	1348.3 \pm 155.3
runtime	10m	10m		960m		20m	20m*	20m	80m	20m

*: Note that it is challenging to compare one-step and multi-step methods directly. Also, [Brandfonbrener et al., 2021](#) reports results for a set of hyperparameters, such as batch and network size, that is significantly different from other methods. We report results for the original hyperparameters and runtime for a comparable set of hyperparameters.

Dataset	AWAC	CQL	IQL (Ours)
antmaze-umaze-vo	56.7 → 59.0	70.1 → 99.4	88.0 → 96.3
antmaze-umaze-diverse-vo	49.3 → 49.0	31.1 → 99.4	67.0 → 49.0
antmaze-medium-play-vo	0.0 → 0.0	23.0 → 0.0	69.0 → 89.2
antmaze-medium-diverse-vo	0.7 → 0.3	23.0 → 32.3	71.8 → 91.4
antmaze-large-play-vo	0.0 → 0.0	1.0 → 0.0	36.8 → 51.8
antmaze-large-diverse-vo	1.0 → 0.0	1.0 → 0.0	42.2 → 59.8
antmaze-vo total	107.7 → 108.3	151.5 → 231.1	374.8 → 437.5
pen-binary-vo	44.6 → 70.3	31.2 → 9.9	37.4 → 60.7
door-binary-vo	1.3 → 30.1	0.2 → 0.0	0.7 → 32.3
relocate-binary-vo	0.8 → 2.7	0.1 → 0.0	0.0 → 31.0
hand-vo total	46.7 → 103.1	31.5 → 9.9	38.1 → 124.0
total	154.4 → 211.4	182.8 → 241.0	412.9 → 561.5

Table 4: Online finetuning results showing the initial performance after offline RL, and performance after 1M steps of online RL. In all tasks, IQL is able to finetune to a significantly higher performance than the offline initialization, with final performance that is comparable to or better than the best of either AWAC or CQL on all tasks except pen-binary-vo.

Part III

AFFORDANCE LEARNING IN UNSEEN ENVIRONMENTS

10

LEARNING NEW SKILLS BY IMAGINING VISUAL AFFORDANCES

10.1 INTRODUCTION

Suppose that you need to learn to open a new kind of drawer in a kitchen. While this new skill might demand some amount of trial and error, you would likely be able to use your mental model and past experience to *imagine* the drawer in the open position, and perhaps even imagine likely intermediate steps, such as grasping the handle, even if you do not yet know precisely how to perform the task. Borrowing the terminology put forward by Gibson (Gibson, 1979), the drawer presents the *affordance* of being “openable,” and you are aware of this affordance from your past experience with other similar objects. In fact, infants learn about affordances such as movability, suckability, graspability, and digestibility through interaction and sensory feedback (Berger, 2014). Learning and utilizing affordances through interaction allows an agent to acquire diverse, meaningful experiences even in unfamiliar situations. However, this way of learning new skills differs markedly from the approach taken by most robotic learning algorithms: the most widely used exploration methods are generally *undirected*, and focus more on seeking out novelty and surprise (Houthooft et al., 2016; Tang et al., 2017; Pathak et al., 2017), rather than familiar and previously seen outcomes. In this paper, we study how robots operating entirely from pixel input can learn about affordances and, when faced with a new and unfamiliar environment, can utilize a previously trained model of possible outcomes to propose potential *goals* that they can practice in this new environment, so as to explore and update their policy efficiently.

We study affordance learning through the framework of self-supervised goal-conditioned reinforcement learning (RL). Learning a new skill in this framework requires generalization in terms of goal setting (*affordances*) and generalization in terms of goal reaching (*behavior*). Prior methods for goal-conditioned RL learn a policy to reach a goal state without the need for an externally provided reward function, and are able to mas-

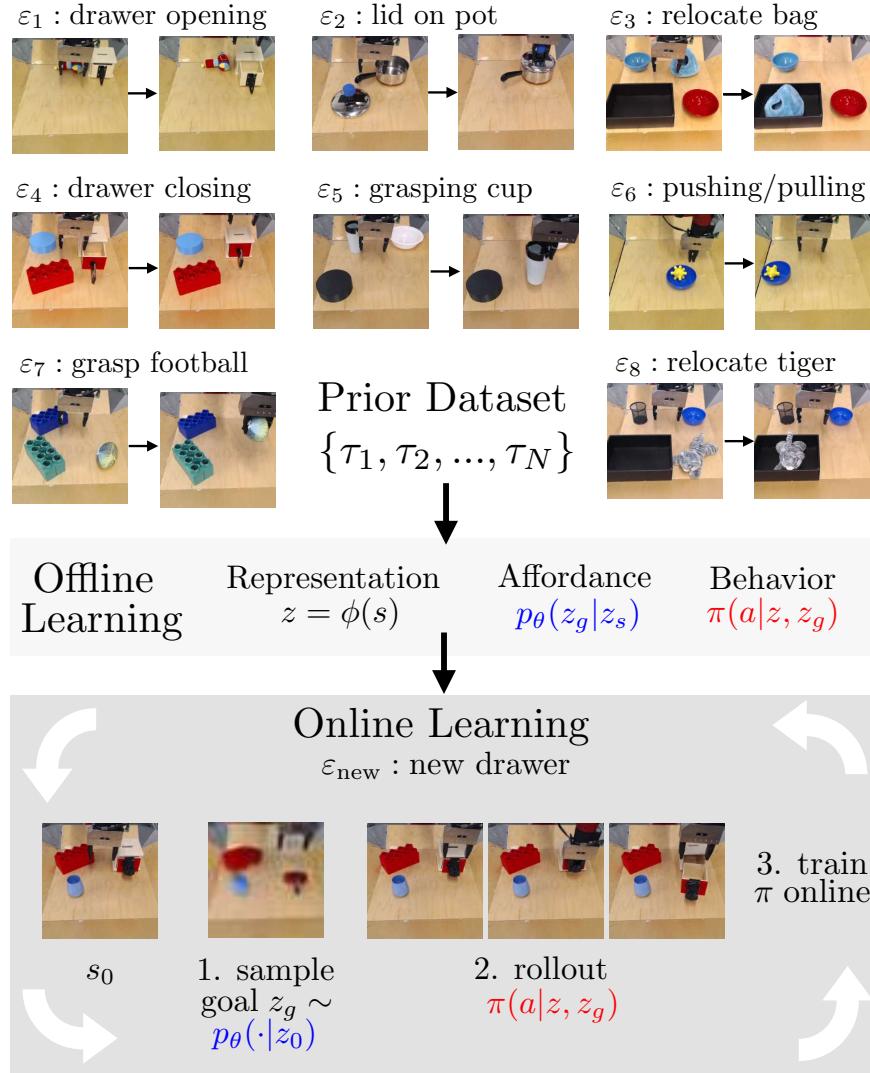


Figure 47: We propose a system for efficient self-supervised robot learning in an unseen environment \mathcal{E}_{new} by utilizing prior data \mathcal{D} of trajectories from related similar environments $\mathcal{E}_i \sim p(\mathcal{E})$. Tasks are specified via a target goal image. From prior data, we learn an encoder $\mathbf{z} = \phi(\mathbf{s})$ of images (representation) for compressing observations and self-generating rewards, a model of what tasks might be tested in the new environment (affordance), and goal-conditioned policy to accomplish a given task (behavior). While this provides reasonable performance in some test environments, perfecting the policy in test environments may require additional interaction in the test environment. Dropped in a test environment \mathcal{E}_{new} without a given goal, we run RL online in order to practice potential tasks, with goals sampled from the affordance model. Online behavior learning allows us to improve the policy π for \mathcal{E}_{new} even when it contains new and unseen objects.

ter skills such as pushing and grasping objects from image observations (Agrawal et al., 2016; A. Nair et al., 2018b; Lynch et al., 2019; A. Nair et al., 2019a). They learn skills by setting goals to explore, and learning a policy to reach them. However, while these prior works have studied how to learn goal-conditioned policies in individual environments, they do not consider what happens when the robot enters a *new* environment where the policy does not simply generalize zero-shot and needs to be trained further. Our approach to solving tasks in this setting is to learn affordances, represented by a generative model of possible outcomes, and then sample possible affordances in a new environment to explore the new environment efficiently.

Our method, visuomotor affordance learning (VAL), uses expressive conditional models for learning generalizable affordances, along with off-policy RL to learn generalizable behaviors. First, we propose to use more expressive generative models for RL to learn compressed representations of images that can reconstruct unseen objects and help the policy generalize to them. Second, we propose to learn an expressive conditional generative model that generalizes past data to set meaningful goals for novel environments, enabling an understanding of object affordances. Third, we demonstrate how we can use off-policy RL with all prior data to initialize a general-purpose goal-conditioned policy, then fine-tune the policy with additional data to master a skill in a new environment. Combining RL and representation learning, we show that we are able to learn skills with a small amount of online exploration on novel objects in real-world robotic scenarios such as object grasping, relocation, and drawer opening and closing.

The main contribution of this work is to present a learning system that can learn robotic skills entirely from image pixels in novel environments by utilizing prior data to generate goals and generalize behavior in new settings. We demonstrate that our method can learn complex manipulation skills including grasping, drawer opening, pushing, and object re-positioning for a diverse set of objects in simulation. We also demonstrate our method in the real world on a Sawyer robot, where our method is able to learn tasks such as grasping and placing unseen objects and opening and closing unseen drawers after only five minutes of online interaction.

10.2 RELATED WORK

RL has been applied previously to robotic manipulation (Kober and Peter, 2008; Peters et al., 2010; Levine et al., 2016a), and also various other applications from playing games (Mnih et al., 2013; D. Silver et al., 2016a) to locomotion (Benbrahim and Franklin, 1997; Kohl and Stone, 2004; Marc Peter Deisenroth and Rasmussen, 2011; G. Williams

et al., 2017). Such approaches require an external reward function, but obtaining this reward function itself poses a challenge to exploring in novel uninstrumented environments as we consider in this paper. Thus, in this work we focus on self-supervised RL methods that do not assume externally provided reward functions.

When learning without an externally provided reward function, one common idea is to use novelty-based intrinsic reward functions (Chentanez et al., 2005; Lopes et al., 2012; M. Bellemare et al., 2016; Houthooft et al., 2016; Stadie et al., 2016; Pathak et al., 2017). State novelty-based methods eventually visit all possible states, but do not necessarily learn a useful policy from purely optimizing the exploration objective. An alternative exploration framework that learns a useful policy even solely from the intrinsic objective is goal reaching (L P Kaelbling, 1993; Schaul et al., 2015b; Adrien Baranes and Pierre-Yves Oudeyer, 2012b; Andrychowicz et al., 2017b; A. Nair et al., 2018b; Nachum et al., 2018; Held et al., 2018; Pérez et al., 2018c; Warde-Farley et al., 2019; Vitchyr H Pong et al., 2020a): by picking a distance measure between states, setting goals, and attempting to reach them, an agent can discover all potential goals in its environment. We refer the reader to the survey of Colas et al. (Colas et al., 2021) for a full classification and discussion of these methods. However, these prior methods do not study the question of how to set goals in new environments, which is vital for collecting coherent experience when faced with a new task. Our method utilizes representation learning and off-policy RL to generalize prior experience to set exploration goals in new settings.

Most similar to our work, context-conditioned reinforcement learning with imagined goals (CCRG) learns a conditional variational auto-encoder (CVAE) (Sohn et al., 2015) that generates goals conditional on the current scene (A. Nair et al., 2019a). CCRG was able to learn pushing skills that generalized mainly to object color and partially to object geometry. Our work differs from CCRG in a number of ways. First, we learn expressive generative models that are able to generate goals in scenes with significantly more visual diversity. Second, we learn a diverse set of skills (e.g., grasping, drawer opening, object placing) that require the goal generation to understand affordances of the environment. Finally, we show that we can use off-policy RL on prior experience, in addition to fine-tuning further on a single specific task to learn new skills. These differences allow our method to better operate in real-world scenarios, as borne out in our experiments.

Another line of work explores the use of affordances in RL, robotics, and control, historically through the lens of perception (Zech et al., 2017; Hassanin et al., 2018; Yamanobe et al., 2018). Affordances have also been discussed previously in reinforcement learning in order to accelerate planning in model-based RL by planning over only a subset of relevant actions (Abel et al., 2014; Khetarpal et al., 2020; Xu et al., 2021). Our work is more

related to the view of affordances in developmental robotics, where affordances were hypothesized to be useful for learning general manipulation skills (Hart and Grupen, 2010; Min et al., 2016). In our work, we show how goal-conditioned RL utilizing from prior data can put these ideas into practice on real-world robotics systems.

10.3 PRELIMINARIES

In this section, we cover preliminaries on RL, goal-conditioned RL, and self-supervised visual RL.

Goal-conditioned reinforcement learning. In goal-conditioned RL, we augment the standard Markov decision process (MDP), which is defined in terms of states $s_t \in \mathcal{S}$, actions $a_t \in \mathcal{A}$, and environment dynamics $s_{t+1} \sim p(\cdot|s_t, a_t)$, with goals $g \in \mathcal{G}$ that represent the agent’s intention to perform one of a variety of tasks drawn from the task family $p(g)$. The reward function is also goal-conditioned, and given by some function $r(s, g)$. The discounted return is defined as $R_t = \sum_{i=0}^H \gamma^i r(s_i, g)$, where γ is a discount factor and H is the horizon, which may be infinite. The aim of the agent is to optimize a policy $\pi(a_t|s_t, g)$ to maximize the expected discounted return $J(\pi) = \mathbb{E}_g[R_0]$. Efficient off-policy RL algorithms have been proposed to learn goal-conditioned policies (Schaul et al., 2015b; Andrychowicz et al., 2017b).

Self-supervised visual reinforcement learning. For scalable robot learning, we cannot always assume known shaped reward functions for tasks. In the absence of such reward functions, Andrychowicz et al. propose goal-state reaching as a natural objective (Andrychowicz et al., 2017b): tasks are defined by state outcomes, where the goal space $\mathcal{G} = \mathcal{S}$, the state space, and the reward is a goal-reaching objective $r(s, g) = -1_{\|s-g\|>\epsilon}$. The task distribution $p(g)$ is chosen to be the feasible states of the robot and objects it interacts with.

But when states are high-dimensional (e.g., images), two issues arise: we do not know the task distribution $p(g)$, and exact reaching of a target goal state is impractical. Reinforcement learning with imagined goals (RIG) (A. Nair et al., 2018b) addresses these issues with a generative model, which is used to learn a latent space of observations and similarity metric on images. Specifically, a variational auto-encoder (D. P. Kingma and Welling, 2014) with encoder $\phi(z_t|s_t)$ and prior $p(z)$ is learned. At training time, the robot sets goals for itself in latent space by sampling a goal latent $z_g \sim p(z)$ and learns a policy $\pi(z_t, z_g)$ to reach latent goals. At test time, the robot can be tasked with a goal image g and execute the learned policy with goal latent $z_g \sim \phi(\cdot|g)$ to match the goal image. In this way, goal-conditioned RL with generative models enables self-supervised learning

in a single environment \mathcal{E} where the task distribution $p(\mathbf{g})$ is not known apriori.

10.4 PROBLEM SETTING

In this paper, we now consider acting in a distribution of environments $p(\mathcal{E})$ with shared structure.¹ Each environment \mathcal{E}_i has its own task distribution $p_i(\mathbf{g})$. As before, tasks are defined by goal states, so $p_i(\mathbf{g})$ represents the potential outcomes of interest in that environment. We assume that the outcomes of interest depends only on the appearance of the environment. When \mathcal{E} is fully observed, this means $p(\mathbf{g}|s_0)$ is shared across environments. As prior data, the robot has access to a training dataset $\mathcal{D} = \{\tau_1, \tau_2, \dots, \tau_N\}$ of trajectories from prior environments, where the trajectories achieve a final outcome $s_T \sim p_i(\mathbf{g})$ – that is, they succeed on tasks from the underlying task distribution for that environment. Now, the robot is placed in a new environment $\mathcal{E}_{\text{new}} \sim p(\mathcal{E})$, and must learn to solve tasks in the new environment through self-supervised practice at training time, such that it can accomplish tasks sampled from $p_i(\mathbf{g})$ at test time. The environments we consider vary visually and dynamically in terms of the objects that are present and potential tasks they afford, such as being able to lift various objects and open different drawers; such real-world variation is presented in Figure 47.

Thus, the agent in this new setting must generalize its prior experience \mathcal{D} to practice potential skills it may be asked to perform at test time in the new environment efficiently, even when it encounters novel objects. At test time, the robot is evaluated in terms of its ability to accomplish a task in \mathcal{E}_{new} specified by a goal image. To perform well in this setting, a method must attempt to infer $p_{\text{new}}(\mathbf{g})$, which should be possible since $p(\mathbf{g}|s_0)$ is common across environments and \mathcal{D} contains trajectories that achieve goals from the same distribution. Note that given observations from \mathcal{E}_{new} , the agent may still have to practice the various behaviors the environment affords if there are multiple, since the test task distribution is unknown at training time.

10.5 VISUOMOTOR AFFORDANCE LEARNING

In this section, we present visuomotor affordance learning (VAL), our method for self-supervised learning in novel environments utilizing prior data from related environ-

¹ Meta-learning approaches have also studied learning a new task quickly, given experience on a set of related MDPs (Duan et al., 2016; Finn et al., 2017; Rakelly et al., 2019). The aims of our method are related to meta-learning, in that we also aim to learn in new environments more quickly, but we do not assume being given user-specified tasks or rewards to solve in the new environment.

Algorithm 8 Visual Affordance Learning

Require: Dataset \mathcal{D} , policy $\pi(\mathbf{a}|\mathbf{z}, \mathbf{z}_g)$, Q-function $Q(\mathbf{z}, \mathbf{a}, \mathbf{z}_g)$, RL algorithm \mathcal{A} , replay buffer \mathcal{R} , relabeling strategy $p_{RS}(\mathbf{z})$, environment family $p(\mathcal{E})$.

- 1: Learn encoder $\phi(\mathbf{z}|\mathbf{s})$ by generative model of \mathcal{D}
- 2: Learn affordances $p(\mathbf{z}_t|\mathbf{z}_0)$ by generative model of \mathcal{D}
- 3: Add latent encoding of \mathcal{D} to the replay buffer
- 4: Initialize π and Q by running \mathcal{A} offline
- 5: Sample $\mathcal{E}_{\text{new}} \sim p(\mathcal{E})$, $\mathcal{E}_{\text{new}} = (p_{\text{new}}(\mathbf{s}_0), \pi_{\text{new}}(\mathbf{s}_{t+1}|\mathbf{s}, \mathbf{a}))$
- 6: **for** $1, \dots, N_{\text{episodes}}$ **do**
- 7: Sample initial state $\mathbf{s}_0 \sim p_{\text{new}}(\mathbf{s}_0)$.
- 8: Sample goal $\mathbf{z}_g \sim p(\mathbf{z}_t|\mathbf{z}_0)$
- 9: **for** $t = 0, \dots, H$ **do**
- 10: Sample $\mathbf{a}_t \sim \pi(\cdot|\mathbf{z}_t, \mathbf{z}_g)$
- 11: Sample $\mathbf{s}_{t+1} \sim p_{\text{new}}(\cdot|\mathbf{s}_t, \mathbf{a}_t)$
- 12: **end for**
- 13: Store trajectory $(\mathbf{z}_1, \mathbf{a}_1, \dots, \mathbf{z}_H)$ in replay buffer \mathcal{R} .
- 14: **for** $1, \dots, N_{\text{train_steps}}$ **do**
- 15: Sample transition $(\mathbf{z}_t, \mathbf{a}_t, \mathbf{z}_{t+1}, \mathbf{z}_g)$
- 16: Relabel with $\mathbf{z}'_g \sim p_{RS}(\mathbf{z}_g)$ and recompute reward
- 17: Update π and Q with relabeled transition using \mathcal{A}
- 18: **end for**
- 19: **end for**

ments. VAL consists of three learning phases: (A) an affordance learning phase to learn affordances from the prior data, (B) an offline behavior learning phase to learn behaviors from the prior data, and (C) an online behavior learning phase where the agent actively interacts with the test environment using affordances and learns potential behaviors in the new environment. The overall method is summarized in Algorithm 8.

10.5.1 Affordance Learning

The affordance learning system must generate goals that induce coherent exploration trajectories even in new environments. We would like to learn a model that, given an observation s_0 in a new environment, generates a potential goal state the agent might be tasked to reach. With high-dimensional image observations, attempting to sample such goal states directly is a difficult generative modeling problem. Instead, we first learn a lower-dimensional latent space $p(\mathbf{z}_t|\mathbf{s}_t)$ by training a generative model through

image reconstruction. With sufficiently expressive models and enough data, the generative model represents images in a manner that it can reconstruct even unseen objects. Given such a latent space, we can then learn affordances by training a conditional model $p(\mathbf{z}_t|\mathbf{z}_0)$ to generate goals that are plausible outcomes of an initial state, even for unseen environments.

To instantiate these two models, we must choose a class of latent variable generative models for $p(\mathbf{s}_t|\mathbf{z}_t)$ and conditional affordance models $p(\mathbf{z}_t|\mathbf{z}_0)$. For the first, while many choices would be suitable, including models such as variational auto-encoders (VAEs) (D. P. Kingma and Welling, 2014) and generative adversarial networks (GANs) (Goodfellow et al., 2014; Donahue et al., 2017), in our implementation we use the VQVAE model (Aaron van den Oord et al., 2017). The VQVAE is expressive enough to represent very diverse datasets, and be able to reconstruct even unseen objects with a high level of detail. We do not require image reconstructions for our method, but the ability to partially reconstruct unseen objects suggests that model expressively represents geometry and color information that may be important for learning affordances and behaviors. In the VQVAE case, ϕ is deterministic, so we will use the shorthand for the latent embedding $\mathbf{z}_t = \phi(\mathbf{s}_t)$, where \mathbf{z}_t is the continuous latent resulting after quantization. In our experiments, we compare this choice of model to other expressive models: VAE (D. P. Kingma and Welling, 2014), CVAE (Sohn et al., 2015), and BiGAN (Donahue et al., 2017).

Next, given a latent space, we need to sample potential goals from this space that is predictive of which goals might be tasked at test time in the new environment \mathcal{E}_{new} . We use a conditional PixelCNN model (Aaron van den Oord et al., 2016) in the latent space to do so, conditioned on the initial state s_0 . The PixelCNN model is trained to maximize $\mathbb{E}_{\tau \sim \mathcal{D}, (\mathbf{s}_0, \mathbf{s}_t) \sim \tau, z \sim \phi(\cdot | \mathbf{s})} [\log p_\theta(\mathbf{z}_t | \mathbf{z}_0)]$, where θ is the parameters of the PixelCNN density model. To generate exploration goals, we sample $z_g \sim p_\theta(\cdot | \mathbf{z}_0)$, where $\mathbf{z}_0 = \phi(\mathbf{s}_0)$ is the encoding of an image of the current setting. The PixelCNN model in VAL being conditional allows us to generate meaningful goals that might be achievable with a novel object. As we show in Section 10.6.1, the conditional model allows us to sample affordances such as opening drawers and lifting objects, even for new environments that are visually complex with variation in the identity, color, geometry, and functionality of objects.

10.5.2 Offline Behavior Learning

Given learned affordances of what behaviors the agent may be tasked with, the agent needs to learn how to actually accomplish those behaviors. In this phase, we wish to learn a reasonable goal-conditioned policy with offline RL. While trained on a limited fixed offline dataset, the hope is that the policy can generalize learning from offline data to accomplish desired behaviors in a new environment, allowing us to either perform tasks successfully zero-shot without further learning, or collect meaningful exploration data in the next phase (Section 10.5.3).

To learn with offline RL while allowing the possibility of quickly fine-tuning in a new environment, we use advantage weighted actor critic (AWAC) (A. Nair et al., 2020) as the underlying RL algorithm. AWAC is an off-policy RL algorithm that has shown strong performance in utilizing prior data for offline pretraining while still being amenable to online fine-tuning. The aim of behavior learning is to optimize a goal-conditioned policy $\pi(a|z, z_g)$ which is able to solve any task in the task distribution $p(z_g)$. We do not assume an external reward function, so we require a reward function to optimize. Following prior work (Andrychowicz et al., 2017b; A. Nair et al., 2018b), we optimize a goal reaching objective: to maximize the density $p_{Z_g}(z = z_g)$; in practice we use the sparse reward function $r(z, z_g) = -\mathbb{1}_{\|z-z_g\|>\epsilon}$, where ϵ is a fixed threshold as it encourages fully solving tasks in a binary fashion. For a particular transition in the replay buffer $(z_t, a_t, z_{t+1}, z_g, r)$, we can also relabel the goal with a new goal z'_g and the recompute the reward. In practice we keep z_g with 20% probability, future hindsight experience replay with 40% probability, and sample $z'_g \sim p(z_t|z_0)$ with 40% probability. Importantly, due to using a compressed representation, we can expect that z_t in any new environment will be semantically similar to past experience. Thus, in this phase, we can obtain a policy that generalizes partially to a new environment.

10.5.3 Online Behavior Learning

Guaranteeing zero-shot generalization for every possible new environment is in general impossible. Instead, we will discuss how to finetune in a specific environment \mathcal{E}_{new} using affordances. In this phase, we utilize the learned affordances which inform *what* tasks to perform, and the offline learned behaviors that inform *how* to perform those tasks.

At online training time, the new task distribution $p_{\text{new}}(g)$ is unknown, so we use the affordance model to sample potential tasks. Thus, to collect coherent exploration data, we sample goals from the affordance module $z_g \sim p_\theta(\cdot|z_0)$ and roll out the goal-

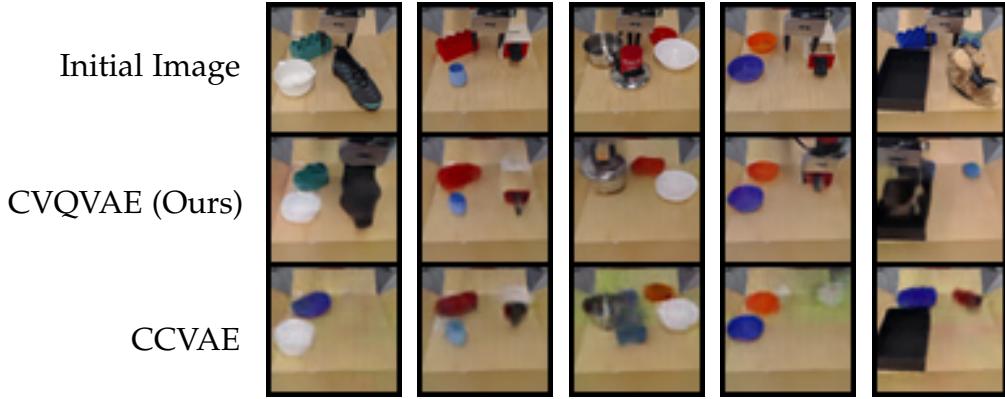


Figure 48: Samples on unseen objects in the real world. In each column, the top image is the conditioning image s_0 and the images below are conditionally sampled images from the corresponding generative model. Our model, CVQVAE, generates clear and diverse samples.

conditioned policy $\pi(\mathbf{a}|\mathbf{z}, \mathbf{z}_g)$. We then iterate between improving the policy with off-policy RL, and collecting exploration data and appending it to the replay buffer. To learn a new task, exploration data in the new environment for fine-tuning the policy is extremely valuable, so this iteration allows us to quickly fine-tune. The online learning process is illustrated in the bottom box in Figure 47.

In summary, VAL enables self-supervised learning in novel environments utilizing prior data. We first learn a generative model for learning a latent space and affordances, and learn how to accomplish tasks with off-policy RL. Then, both are used in an online behavior learning phase to perfect potential behaviors in a new environment. The overall method is summarized in Algorithm 8.

10.6 REAL-WORLD EXPERIMENTAL EVALUATION

We first evaluate our method in a real-world manipulation setting with a Sawyer robot and diverse objects and tasks. This setting is very challenging due to the variety of objects, scenes, and tasks that the robot interacts with, and being limited to using image inputs.

Real-world setting. The real-world setting is shown in Figure 47. A Sawyer robot is controlled at 5Hz with 4 degrees of control: 3 dimensions of end-effector velocity control in Euclidean space and one dimension to open and close the gripper. The environments span 10 drawer handles, 10 pot handles, 40 toys, and 60 distractor objects. To create a new environment \mathcal{E} in the real world, we randomly sample one interaction object as well as 2-

Task	VAL (Ours) Offline → Online	CCRIG Offline → Online	VAL training			VAL testing		
	s_0	s_T	$d(z_g)$	s_0	s_T	s_g		
(1) Pickup shoe	12.5% → 50%	0% → 16.6%	(1)					
(2) Drawer closing	25% → 100%	0% → 12.5%	(2)					
(3) Drawer opening	62.5% → 100%	0% → 0%	(3)					
(4) Place object in tray	25% → 75%	0% → 0%	(4)					
(5) Lid on pot	37.5% → 87.5%	0% → 0%	(5)					

Figure 49: Real-world results. Left, success rates per method for the five tasks tested. We report the offline performance, followed by the performance after five minutes of online fine-tuning. With VAL, we see reasonable initial offline performance followed by significant improvement on all tasks. Meanwhile CCRIG fails to succeed any of the tasks. Right, film strips of VAL during training (left, with decoded affordance proposals) and testing (right, with goal images) are visualized. Videos are available at <https://sites.google.com/view/val-rl>

3 distractor objects. The behaviors that the environments afford therefore span grasping and relocating toys, opening and closing drawers, as well as covering and uncovering pots. Each test environment contains an unseen interaction object and a random set of 2-3 distractor objects, with all positions randomized.

Our experiments aim to answer the following questions:

1. Does the learned generative model sample plausible affordances in new scenes?
2. Is VAL able to accelerate learning of real-world manipulation skills online in new settings?

10.6.1 Generative Models for Affordance Learning

Expressive generative models enable us to propose desired outcomes in new environments even when the current policy cannot yet achieve them. We can preview this capability by inspecting the training procedure of the models, which also gives insights into which models will tend to perform well when used for interactive learning. Specifically, we inspect model samples to see if the model can actually output plausible candidate affordances. We evaluate our model, which combines a VQVAE with a conditional pixel CNN, and prior models that have been used in self-supervised RL as well as other expressive generative models. We compare (1) VAE (D. P. Kingma and Welling, 2014), (2) CVAE (Sohn et al., 2015), (3) BiGAN (Donahue et al., 2017), (4) Conditional BiGAN, (5)

Ours, a VQVAE with a conditional PixelCNN.

Sampled affordances are shown in Figure 48. Our model is able to sample coherent tasks to perform where other methods produce indistinct or uninterrupted samples. For example, the VQVAE model turns an unseen drawer (with a wide handle) into one with a thin handle, which exists in the prior data. This illustrates its ability to utilize prior data for generating conducive goals for online RL. CCVAE samples are usually less coherent, often missing the object completely and not capturing the geometry of unseen objects. The CBGAN also struggles to produce realistic images, while the VAE fails completely as it is not a conditional model cannot represent the wide diversity of potential images.

10.6.2 Real-World Visuomotor Affordance Learning

Next, we investigate whether VAL can handle real-world visual complexity and object diversity to learn control policies for varied tasks on a Sawyer robot. The setup is shown in Figure 47: the robot is tasked with fine-tuning in a particular environment, beginning with about 1,000 trajectories of prior data for affordance learning and offline behavior learning. The protocol for collecting prior data, as well as examples of images from the data are shown in Appendix ??.

After pretraining affordances, a policy, and a Q function on the prior data, the robot is dropped in a new test environment. We evaluate five test environments which each contain distractor objects as well as objects that may be interacted with such a shoe that can be picked up, drawers than can be opened or closed, and a lid which may be placed on a pot. These behaviors are demonstrated on similar objects in the prior dataset, but the objects during test time are previously unseen - for instance, the drawer has a different handle. The agent can interact with the environment without supervision to collect more data and improve the policy; then to evaluate, the agent is tasked with a specific goal image that corresponds to a task such as opening a closed drawer.

Results running VAL in the real world are shown in the table in Figure 49, reporting the success rates on five test tasks for our method and CCRIG both offline and after one epoch of online training, which includes 10 interactive trials, amounting to less than five minutes of real-world interaction time. On all five tasks, VAL shows nontrivial offline performance followed by strong online improvement. Qualitatively, our method is able to learn behaviors such as recovering from a missed grasp by returning to the grasp. Film strips of our method are shown on the right side of Figure 49. In contrast, CCRIG struggles to make progress on all five tasks. CCRIG fails for two reasons. First,

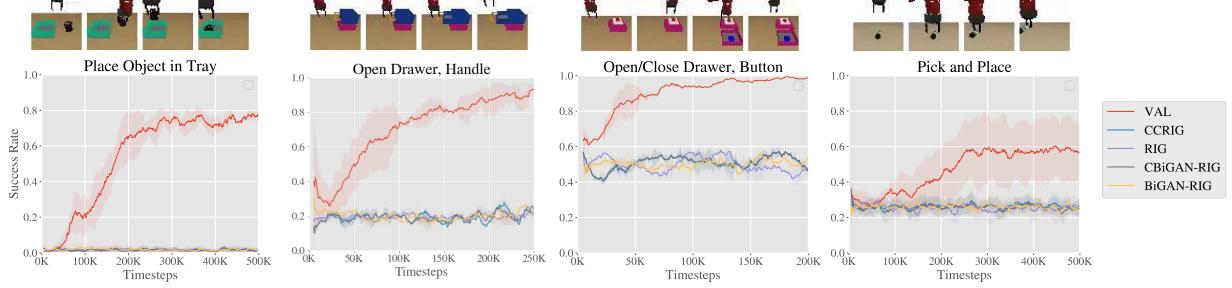


Figure 50: Learning curves for simulation experiments, fine-tuning on an unseen environment. Our method is able to learn these tasks online, while none of the baselines or prior methods are able to make meaningful learning progress in this setting. A successful rollout of each task in a test environment is shown above the corresponding plots.

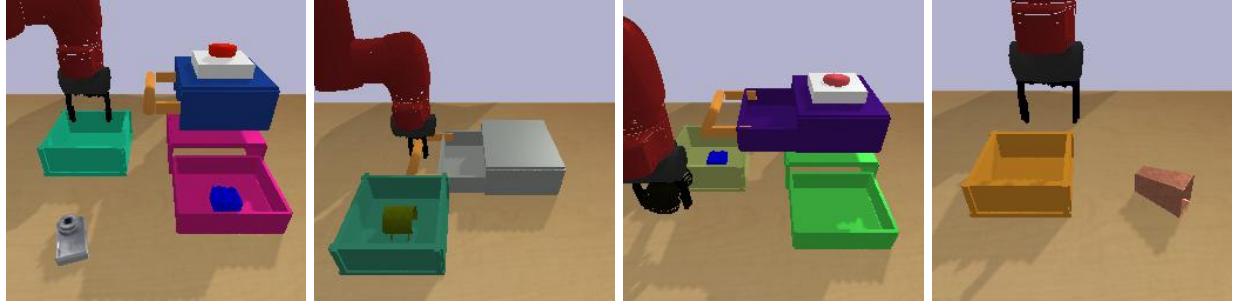


Figure 51: Randomly sampled scenes from our simulated multi-task environment. To practice in a sampled scene, the agent must infer the potential behaviors that the scene affords.

the quality of the sampled affordances are significantly poorer. Second, as can be seen in videos, CCRIG sometimes comes close to solving the task, but does not fully solve it, preventing improvement after subsequent training.

10.7 EXPERIMENTAL EVALUATION IN SIMULATION

In order to further study and understand VAL, we carry out more experiments in simulation, where we can better control the quantity of data and ablate parts of the method. These experiments aim to answer the following questions:

1. Does VAL outperform prior self-supervised RL methods in accelerating learning a new task from prior data?
2. Can VAL scale with additional data for affordance and behavior learning?

Simulated setting. Randomly generated workspaces in our simulated multi-task environment are shown in Figure 53. In this PyBullet simulation (Coulmans and Bai, 2016–

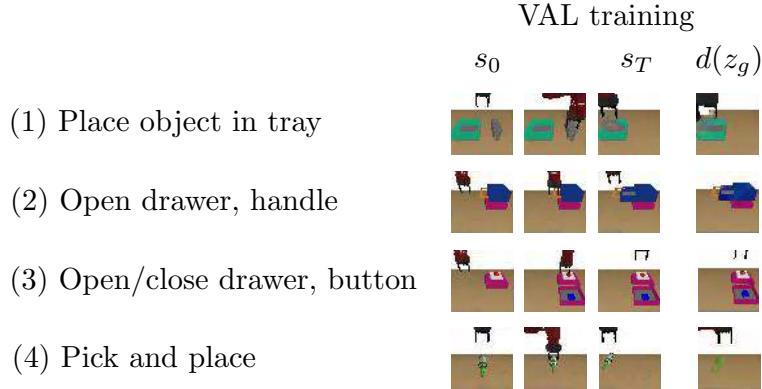


Figure 52: Training rollouts from VAL. For each environment, the reconstruction of a sampled affordance is shown on the rightmost column, and frames from a trajectory attempting to achieve that affordance is shown on the left.

2021), the robot is faced with multiple potential tasks in each environment based on which objects are present: opening and closing a drawer by the handle, opening and closing a different drawer by pressing a button, grasping objects, and moving objects into drawers or a box. To sample a new environment \mathcal{E} , we randomize the existence, position, color, and orientation of the following: two drawers, a box, a button, and an object. If an object is present it is chosen from a set of 84 object geometries. To successfully learn in a test environment, the method must be able to explore in the environment based on the behaviors that environment affords.

10.7.1 Self-Supervised Online Fine-Tuning from Prior Data

We first compare VAL against prior methods in this simulated setting. The robot receives a prior dataset \mathcal{D} of trajectories in the pre-training environments, which is utilized for learning affordances and offline RL. The details of this dataset are explained further in Appendix ???. After the pre-training phase, the robot is placed in a test setting and begins online fine-tuning. We evaluate the policy on goal images in the new environment, sampled from expert trajectories, and report whether the final state of the policy’s trajectory matches the state of the goal image within a chosen threshold.

Learning curves of the online training phase for various objects are shown in Figure 50. We see that for each task, VAL outperforms prior methods which do not make progress on learning these tasks at all. On the drawer opening task, offline RL achieves nontrivial performance of about 60%, but online interaction allows fine-tuning to over 90% success rate.

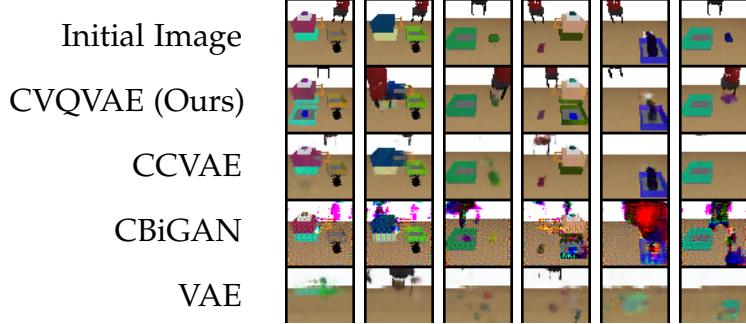


Figure 53: Samples on test environments in simulation. In each column, the top image is the conditioning image s_0 and the images below are conditionally sampled images. The left four columns are relatively successful samples for our affordance model, each showing the potential outcome of a behavior. The right two columns are failure modes, lacking diversity or altering an object’s geometry or color.

How is VAL able to outperform prior methods so significantly? One major reason is the quality of sampled goals. Samples for the affordance model for the simulated domain are shown in Figure 53. We compare (1) VAE (D. P. Kingma and Welling, 2014), (2) CVAE (Sohn et al., 2015), (3) BiGAN (Donahue et al., 2017), (4) Conditional BiGAN, (5) Ours, a VQVAE with a conditional PixelCNN. Our model produces diverse, coherent samples of possible outcomes (i.e., affordances) in new scenes with novel objects. In comparison to our model, conditional VAE samples tend to be blurry and do not capture the geometry of unseen objects well.

10.7.2 Scalable Robot Learning with VAL

For general-purpose robot learning, we would ideally like to learn diverse skills continuously, using prior experience to perpetually improve at learning new skills. To study this possibility, we examine whether solving tasks in a new environment can be sped up by collecting larger amounts of prior experience. In this experiment conducted in simulation in order to carefully control the data quantity, the robot receives a prior dataset \mathcal{D} of K trajectories for running VAL to grasp an unseen object in a new environment. We vary K to observe whether the method can benefit from larger amounts of prior data.

Learning curves of the online training phase, averaged over five test objects, are shown in Figure 54. First, we see from the starting point of each curve that the offline policy already generalizes to some level for grasping objects, but the average success rate is only around 35%. Importantly, note that training on more data only slightly improves generalization after offline training (at timestep 0). Then, fine-tuning results in rapid policy

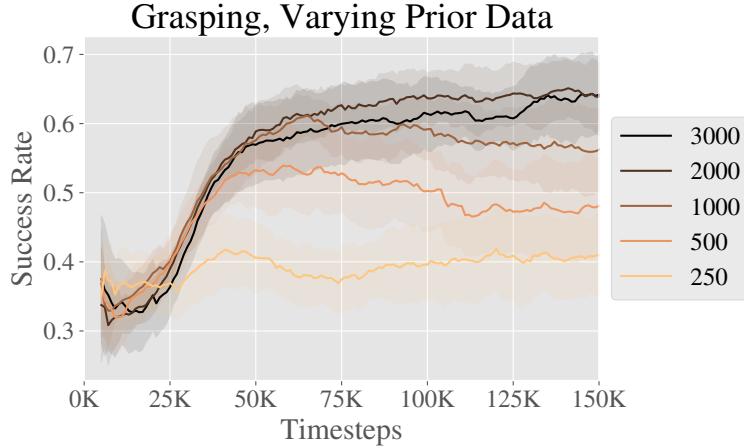


Figure 54: Learning curves for simulated grasping of novel objects with VAL, using data from an increasing number of training objects for offline RL. We collect 50 trajectories per training object, and each line is labeled with the total number of training trajectories.

improvement up to around 65% success rate after only 150,000 timesteps when utilizing the most prior data, compared to 40% with the least prior data. Thus, more prior data significantly accelerates learning in new environments even when the initial performance is comparable. This suggests that VAL can be deployed in a continual learning setting, with each new task being learned faster as it benefits from the increasing dataset size.

10.8 CONCLUSION

We present visuomotor affordance learning (VAL), a method for learning tasks online in a new environment without supervision, utilizing trajectories from other related environments. VAL uses expressive generative models to learn visual affordances, combines these affordances with off-policy goal-conditioned RL to learn skills offline, and then fine-tunes in a new environment online. Like deep learning in domains such as computer vision (Krizhevsky et al., 2012) and natural language processing (Devlin et al., 2019) which have been driven by large datasets and generalization, robotics will likely require learning from a similar scale of data. In future work, VAL could enable such systems by allowing autonomous collection of coherent exploration data in diverse real-world settings.

11

PLANNING TO PRACTICE: EFFICIENT ONLINE FINE-TUNING BY COMPOSING GOALS IN LATENT SPACE

11.1 INTRODUCTION

Developing controllers that can solve a range of tasks and generalize broadly in real-world settings is a long-standing challenge in robotics. Such generalization in other domains such as computer vision and natural language processing has been attributed to training large machine learning models on massive datasets (Krizhevsky et al., 2012). Consequently, one promising approach to robustly handle a wide variety of potential situations that a robot might encounter is to train on a large dataset of robot behaviors. Prior work in robotics has demonstrated effectively generalization from training on large datasets for individual tasks, such as grasping (Levine et al., 2017; Yu et al., 2021). However, a general-purpose robot should be able to perform a wide range of skills, and should also be *taskable*. That is, it must be able to complete a specific task when specified by a human, including temporally extended tasks that require sequencing many skills together to complete. This topic has been studied in prior work on goal-conditioned reinforcement learning, where a robot aims to perform a task given a desired end state (Khazatsky et al., 2021; Chebotar et al., 2021; Kalashnikov et al., 2021). What remains to make these methods widely applicable for real-world robotics?

While goal-conditioned policies can be trained effectively for relatively short-horizon tasks, temporally extended multi-stage can pose a significant challenge for current methods. These tasks present a major exploration challenge during online learning, and a major challenge for credit assignment during offline learning. In this paper, we aim to address these challenges by combining two ideas. The first is that long-horizon goal-reaching tasks can be decomposed into shorter-horizon tasks consisting of subgoals. The second is that these subgoals can be used to *fine-tune* a goal-conditioned policy online, even if its performance from offline data is poor. The first idea enables us to address

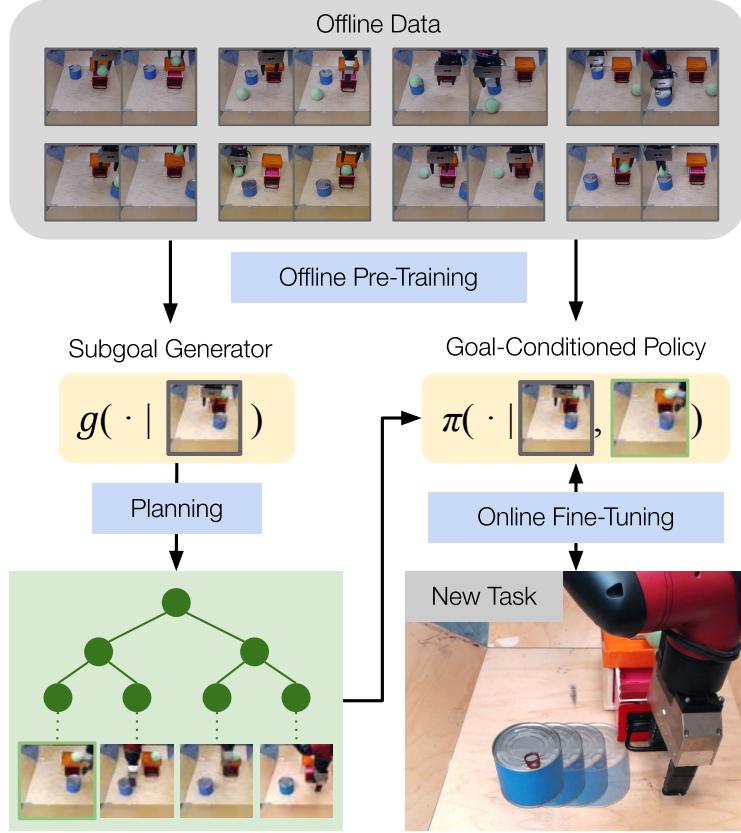


Figure 55: Our method, Plan to Practice (PTP), solves long-horizon goal-conditioned tasks by combining planning and fine-tuning. We begin with an offline dataset containing a variety of behaviors, and train a subgoal generator and goal-conditioned policy on this data. Then, to learn a more complex multi-stage tasks, we optimize over subgoals using the subgoal generator, which corresponds to a planning procedure over (visual) subgoals, and fine-tune the policy with online RL by practicing these subgoals. This enables the robot to solve multi-stage tasks directly from images.

the exploration challenge, by automatically generating intermediate subgoals that can be “practiced” on the way to a longer-horizon final goal. In the framework of goal-conditioned RL, solving long-horizon tasks can be reduced to the problem of optimization over a sequence of subgoals for the goal-conditioned policy, and this optimization over subgoals can be regarded as a kind of high-level planning, where the optimizer selects waypoints for achieving a distant goal. The high-level planner itself can use a learned high-level model.

However, if we rely entirely on offline data, credit assignment challenges make it difficult to perform longer-horizon tasks even with subgoal planning. Even if the offline

RL policy performs well on each individual skill, there may be errors from stitching skills together because the initial states of each stage diverge from the offline data when they are composed together. In practice, this leads to poor performance when using only offline training. Therefore, the second key idea in our work is to utilize subgoal planning not merely to *perform* a multi-stage task, but also to make it possible to *practice* that task to finetune it online. While online training for temporally extended tasks is ordinarily difficult, by addressing the exploration challenge with subgoal planning, we make it possible for the robot to practice a series of relatively short-horizon tasks, which makes this kind of finetuning feasible. Thus, the planner acts both as a higher level policy when performing the task, and as a scaffolding curriculum for finetuning the lower-level goal-conditioned policy. By collecting data actively in a specific environment, we can directly experience the distribution shift and can use reinforcement learning to improve performance under this shift.

To this end, we propose Planning to Practice (PTP), an approach that efficiently trains a goal-conditioned policy to solve multi-step tasks by setting subgoals to exploit the compositional structure of the offline data. An outline is shown in Fig. 55. Our approach is based around a planner that composes generated subgoals to guide the goal-conditioned policy during an online fine-tuning phase. To propose diverse and reachable subgoals to form the candidate plans, we design a conditional subgoal generator based on conditional variational autoencoder (CVAE) (Sohn et al., 2015). Through training on the offline dataset, the conditional subgoal generator captures the distribution of reachable subgoals from a given state and generates sequences of subgoals from the learned latent space in a recursive manner. Our subgoal planning algorithm hierarchically searches for subgoals in a coarse-to-fine manner using multiple conditional subgoal generators that trained to generate goals at different temporal resolutions. Both the goal-conditioned policy and the conditional subgoal generators are pre-trained on the offline data, and the policy is fine-tuned on the novel target task.

Our main contribution is a system for learning to solve long-horizon goal-reaching tasks by fine-tuning the goal-conditioned policy with subgoal planning in a learned latent space. We evaluate our approach on multi-stage robotic manipulation tasks with raw image observations and image goals in both simulation and the real world. After being pre-trained on short demonstrations of primitive interactions, our approach is able to find feasible subgoal sequences as plans for unseen final goals by recursively generating subgoals with the learned conditional subgoal generators. By comparing our approach with both model-free methods and prior approaches that optimize over subgoals, we demonstrate that the produced plans significantly improve the learning efficiency and

the resultant success rates during the online fine-tuning.

11.2 RELATED WORK

We propose to use a combination of optimization-based planning and fine-tuning with goal-conditioned reinforcement learning from prior data in order to allow robots to learn temporally extended skills. In this section, we cover prior methods in offline RL, planning, goal-conditioned RL, and how they relate to our method.

Learning from prior data. Offline reinforcement learning methods learn from prior data (Lange et al., 2012a; Fujimoto et al., 2019b; Kumar et al., 2019b; C. Zhang et al., 2021; Kumar et al., 2020; Fujimoto and S. S. Gu, 2021; Singh et al., 2020), and can also finetune through online interaction (A. Nair et al., 2020; Villaflor et al., 2020; Lu et al., 2021; Khazatsky et al., 2021; S. Lee et al., 2021; Meng et al., 2021). Such methods have been used in a variety of robotic settings (Kalashnikov et al., 2018a; Cabi et al., 2019; Kalashnikov et al., 2021; Lu et al., 2021). Our focus is not on introducing new offline RL methods. Rather, our work shows that planning over subgoals for a goal-conditioned policy that is pretrained offline can enable finetuning for temporally extended skills that would otherwise be very difficult to learn.

Goal-conditioned reinforcement learning. The aim of goal-conditioned reinforcement learning (GCRL) is to control the agent to efficiently reach specified goal states (Leslie Pack Kaelbling, 1993; Schaul et al., 2015a; Eysenbach et al., 2021). Compared to policies that are trained to solve a fixed task, the same goal-conditioned policy can perform a variety of tasks when it is commanded with different goals. Such flexibility allows GCRL to better share knowledge across different tasks and make use of goal relabeling techniques to improve the sample efficiency without meticulous reward engineering (Andrychowicz et al., 2017a; Vitchyr H. Pong et al., 2020b; M. Fang et al., 2019; Ding et al., 2019; Abhishek Gupta et al., 2019b; Sun et al., 2019; Eysenbach et al., 2020; Ghosh et al., 2021). Prior has explored various strategies for proposing goals for exploration (A. Nair et al., 2018b; A. Nair et al., 2019b; Khazatsky et al., 2021; Chane-Sane et al., 2021), and studied goal-conditioned RL from offline data (Chebotar et al., 2021). However, such works generally aim to learn short-horizon behaviors, and learning to reach goals that require multiple stages (e.g., several manipulation primitives) is very difficult, as shown in our experiments. Our work aims to extend model-free goal-conditioned RL methods by incorporating elements of planning to enable effective finetuning for multi-stage tasks.

Planning. A wide range of methods have been developed for planning in robotics. At the most abstract level, symbolic task planning searches over discrete logical formulas

to accomplish abstract goals (Fikes and Nilsson, 1971). Motion planning methods solve the geometric problem of reaching a goal configuration with dynamics and collision constraints (Kavraki et al., 1996; Koenig and Likhachev, 2002; Karaman and Frazzoli, 2011; Zucker et al., 2013; Kalakrishnan et al., 2011a). Prior methods have also considered task and motion planning as a combined problem (Srivastava et al., 2014). These methods generally assume high-level structured representations of environments and tasks, which can be difficult to actualize in real-world environments. Since in our setting we only have image inputs and not structured scene representations, we focus on methods that can handle raw images for observations and task specification.

Combining goal-conditioned RL and planning. A number of recent works have sought to integrate concepts from planning with goal-conditioned policies in order to plan sequences of subgoals for longer-horizon tasks (Nasiriany et al., 2019; Eysenbach et al., 2019; K. Fang et al., 2019; Charlesworth and Montana, 2020; Pertsch et al., 2020; Sharma et al., 2021; T. Zhang et al., 2021). These prior methods either propose subgoals from the set of previously seen states, or directly optimize over subgoals, often by utilizing a latent variable model to obtain a concise representation of image-based states (A. Nair et al., 2018b; Ichter et al., 2018; S. Nair and Finn, 2019; Nasiriany et al., 2019; Pertsch et al., 2020; Khazatsky et al., 2021; Chane-Sane et al., 2021). The method we employ is most closely related conceptually to the method proposed by Pertsch et al. (Pertsch et al., 2020), which also employs a hierarchical subgoal optimization, and the method proposed by Nasiriany et al. (Nasiriany et al., 2019), which also optimizes over sequences of latent vectors from a generative model. Our approach makes a number of low-level improvements, including the use of a conditional generative model (A. Nair et al., 2019b), which we show leads to significantly better performance. More importantly, our method differs conceptually from these prior works in that our focus is specifically on utilizing subgoal optimization as a way to enable finetuning goal-conditioned policies for longer-horizon tasks. We show that it is in fact this capacity to enable effective finetuning that enables our method to solve more complex multi-stage tasks in our experiments.

11.3 PROBLEM STATEMENT

In this paper, we consider the problem of learning to complete a long-horizon task specified by a goal image. The robot learns over a variety of initial configurations and goal distributions, which cover a range of behaviors such as opening or closing a drawer, and picking, placing, or pushing an object. As prior data, the robot has access to an offline dataset of trajectories $\mathcal{D}_{\text{offline}} = \{\tau_1, \tau_2, \dots, \tau_N\}$ for offline pre-training. In each trajectory,

the robot is controlled by a human tele-operator or a scripted policy to achieve one of the goals the environment affords. A goal-conditioned policy is pre-trained on this dataset using offline RL algorithms.

After offline pre-training, the robot is placed in a particular environment that it has online access to interact in. Even though the initial configuration of this environment may have been included in the set of training environments, the goal distribution for this environment at test time requires sequencing multiple skills together, which is not present in the offline data. For instance, as shown in Fig. 55, the robot would need to first slide away the can that blocks the drawer, then reaches the handle of the drawer, and finally opens the drawer.

Naïvely running offline RL may not solve the long-horizon test tasks for two reasons. First, the robot is given a test goal distribution that is long horizon but offline dataset consists of individual skills. The method needs to somehow compose these individual skills autonomously in order to succeed at goals drawn from the test distribution. Second, offline RL may not solve the task due to distribution shift. Distribution shift appears in two forms: distribution shift between transitions in the prior data and transitions obtained by the actively rolling out the policy, and the distribution shift introduced when performing tasks sequentially. If the robot may actively interact in the new environment to improve its policy, how can the robot further practice and improve its performance?

11.4 PRELIMINARIES

We consider a goal-conditioned Markov Decision Process (MDP) denoted by a tuple $M = (\mathcal{S}, \mathcal{A}, \rho, P, \mathcal{G}, \gamma)$ with state space \mathcal{S} , action space \mathcal{A} , initial state probability ρ , transition probability P , a goal space \mathcal{G} , and discount factor γ . In each episode, a desired goal $s_g \in \mathcal{G}$ is sampled for the robot to reach. At each time step t , a goal-conditioned policy $\pi(a_t|s_t, s_g)$ selects an action $a_t \in \mathcal{A}$ conditioned on the current state s_t and goal s_g . After each step, the robot receives the goal-reaching reward $r_t(s_{t+1}, s_g)$. The robot aims to reach the goal by maximizing the average cumulative reward $\mathbb{E}[\sum_t \gamma^t r_t]$. Our approach learns a goal-conditioned policy π for solving the target task specified by a desired final goal s_g . The goal-conditioned policy is pre-trained on a previously collected offline dataset $\mathcal{D}_{\text{offline}}$ and then fine-tuned to reach s_g by accumulating data into an online replay buffer $\mathcal{D}_{\text{online}}$. $\mathcal{D}_{\text{offline}}$ contains diverse short-horizon interactions with objects in the environment. During online fine-tuning, we would like the policy to learn to improve and compose these short-horizon behavior for multi-stage tasks specified by s_g .

Defining informative goal-reaching rewards and extracting useful state representations

from high-dimensional raw observations such as images can be challenging. Following the practice in prior work (A. Nair et al., 2018b; Khazatsky et al., 2021), we pre-train a state encoder $h = \phi(s)$ to extract the latent state representation h . By encoding the states and goals to the latent space, we can obtain an informative goal-reaching reward function $r_t = R(h_{t+1}, h_g)$ by computing $h_{t+1} = \phi(s_{t+1})$ and $h_g = \phi(s_g)$. Specifically, $R(h_{t+1}, h_g)$ returns 0 when $\|h_{t+1}, h_g\| < \epsilon$ and -1 otherwise, where ϵ is a selected threshold. In addition, we also use $\phi(s_{t+1})$ as the backbone feature extractor in all of our models that take s as an input. For simplicity, we directly use s to denote h in the rest of the paper. The details of the state encoder are explained in Sec. 11.6.2.

11.5 PLANNING TO PRACTICE

We propose Planning to Practice (PTP), an approach that efficiently fine-tunes a goal-conditioned policy to solve novel tasks. To enable the robot to efficiently learn to solve the target task, we propose to use subgoals to facilitate the online fine-tuning of the goal-conditioned policy. Given the initial state s_0 and the goal state s_g , we search for a sequence of K subgoals $\hat{s}_1 : K = \hat{s}_1, \dots, \hat{s}_K$ to guide the robot to reach s_g . Such subgoals will inform the goal-conditioned policy $\pi(a|s, s_g)$ what is the immediate next step on the path to s_g and provide the policy more dense reward signals compared to directly using the final goal. We choose the sequence of subgoals at the beginning of each episode and feed the first subgoal in the sequence to the goal-conditioned policy. The policy will switch to the next subgoal in the sequence when the current subgoal is reached or the time budget assigned for the current subgoal runs out.

The main challenge is to search for a sequence of subgoals that can lead to the desired final goals while ensuring each subgoal is a valid state that can be reached from the previous subgoal. Particularly when the states correspond to full images, most vectors will not actually represent valid states, and indeed naïvely optimizing over image pixels may simply result in out-of-distribution inputs that lead to erroneous results when input into the goal-conditioned policy.

As outlined in Fig. 55, we devise a method to effectively propose and select valid subgoal sequences to guide online fine-tuning by means of a generative model. At the heart of our approach is a conditional subgoal generator $g(\cdot|s_0)$ that recursively produces candidate subgoals in a hierarchical manner conditioned on the initial state s_0 . To find the optimal sequence of subgoals $\hat{s}_{1:K}^*$, we first sample N candidate sequences $\hat{s}_{1:K}^1, \dots, \hat{s}_{1:K}^N$ from the state space using the conditional subgoal generator. Then we rank the candidate sequences using a cost function $c(s_0, \hat{s}_{1:K}, s_g)$. The sequence that corresponds to the

Algorithm 9 Planning To Practice (PTP)

Require: set of final goals \mathcal{G} , time horizon T , offline data $\mathcal{D}_{\text{offline}}$, number of subgoals K .

- o: Train $\pi(a|s, s_g)$ and $g(s, z)$ on $\mathcal{D}_{\text{offline}}$.
- o: Initialize the online replay buffer $\mathcal{D}_{\text{online}} \leftarrow \emptyset$.
- o: **while** not converged **do**
- o: Reset the environment and observe s_0 .
- o: Sample s_g from \mathcal{G} .
- o: Plan for the subgoals $\hat{s}_{1:K}$.
- o: $k \leftarrow 1$
- o: **for** $t = 1, \dots, T$ **do**
- o: Compute the action $a_t \leftarrow \pi(a_t|s_t, \hat{s}_K)$
- o: Observe the state s_{t+1} and the reward r_t
- o: $\mathcal{D}_{\text{online}} \leftarrow \mathcal{D}_{\text{online}} \cup (s_t, a_t, r_t, s_{t+1})$.
- o: **if** $t \pmod{\Delta t} == 0$ **or** $\|s_{t+1} - \hat{s}_K\| < \epsilon$ **then**
- o: $k \leftarrow \min(k + 1, K)$
- o: Train π on batches sampled from $\mathcal{D}_{\text{offline}}$ and $\mathcal{D}_{\text{online}}$.

=0

lowest cost will be selected as $\hat{s}_{1:K}^*$ for the goal-conditioned policy. Through this sampling-based planning procedure, we choose the subgoal for guiding the goal-conditioned policy π during online fine-tuning. The overall algorithm is summarized in Algorithm 9. Next we describe the design of each module in details.

11.5.1 Conditional Subgoal Generation

The effectiveness of our planner relies on the generation of diverse and feasible sequences of subgoals as candidates. Specifically, we would like to generate the candidates by sampling from the distribution of suitable subgoal sequences $p(\hat{s}_1, \dots, \hat{s}_K | s_0)$ conditioned on the initial state s_0 . Most existing methods independently sample the subgoal at each step from a learned prior distribution (Pertsch et al., 2020) or a replay buffer (Eysenbach et al., 2019), which is unlikely to propose useful plans for tasks with large, combinatorial state spaces (i.e., with multiple objects).

We propose to break down $p(\hat{s}_1, \dots, \hat{s}_K | s_0)$ into $p(\hat{s}_1 | s_0) \prod_{i=1}^k p(\hat{s}_i | \hat{s}_{i-1})$ through modeling the conditional distribution $p(s' | s)$ of the reachable next subgoal s' . By utilizing temporal compositionality, the conditional subgoal generation paradigm improves generalization and enables generation of sequences of arbitrary lengths.

We use a conditional variational encoder (CVAE) (Sohn et al., 2015) to capture the

distribution of reachable goals $p(s'|s)$. In the CVAE, we define the decoder as $g(s, z)$ and the encoder as $q(z|s, s')$, where z is the learned latent representation of the transitions and it is sampled from a prior probability $p(z)$. To propose a sequence of subgoals, we use $g(s, z)$ as the conditional subgoal generator. Conditioned on the initial state s_0 , the first subgoal \hat{s}_1 can be generated as $\hat{s}_1 = g(s_0, z_1)$ given the sampled z_1 . Then the i_{th} subgoal can be recursively generated by sampling $z_i \sim p(z)$ and computing $\hat{s}_i = g(\hat{s}_{i-1}, z_i)$ given the previous subgoal \hat{s}_{i-1} . In this way, we could sample a sequence of i.i.d. latent representations z_1, \dots, z_K and recursively generate $\hat{s}_1, \dots, \hat{s}_K$ conditioned on the initial state s_0 using the conditional subgoal generator.

The CVAE is trained to minimize the evidence lower bound (ELBO) (D. P. Kingma and Welling, 2014) of $p(s'|s)$ given the offline dataset \mathcal{D} . During training, we sample transitions (s_t, s_τ) from the offline dataset to form the minibatches, where $\tau = t + \Delta t$ is a future step that is Δt steps ahead. Instead of using a fixed Δt , we sample Δt from a range for each transition to provide richer data. To encourage the trained model to be robust to compounding errors, we sample sequences composed of multiple states and use the subgoal reconstructed at the previous step as the context in the next step. Therefore, the objective for training the conditional subgoal generator is:

$$\mathbb{E}_{q(z|s_t, s_\tau)} \|s_\tau - g(s_t, z)\|^2 + D_{KL}[q(z|s_t, s_\tau) || p(z)] \quad (54)$$

where $D_{KL}[\cdot || \cdot]$ indicates the KL-Divergence.

11.5.2 Efficient Planning in the Latent Space

We build a planner that efficiently searches for sequences of subgoals in the latent space as shown in Algorithm 10. To tackle the large search space of candidate subgoal sequences, we design a hierarchical planning algorithm that searches for subgoals in a coarse-to-fine manner and re-use the previously selected subgoals as candidates in new episodes.

The hierarchical planning is conducted at L levels with different temporal resolutions $\Delta t_1, \dots, \Delta t_L$. The temporal resolution of each level is an integral multiple of that of the previous level, i.e., $\Delta t_i = M\Delta t_{i-1}$, where M is a scaling factor and is set to 2 in our experiments. We first plan for the subgoals $\hat{s}_1^1, \hat{s}_2^1, \dots$ on the first level. Then the subgoals $\hat{s}_{1:K}^1$ of finer temporal resolution are planned on each level l to connect the subgoals planned on the previous level $l-1$. Specifically, given the adjacent subgoals \hat{s}_i^{l-1} and \hat{s}_{i+1}^{l-1} produced on the previous level, we plan for a segment of M subgoals $\hat{s}_{i*M+1}^l, \dots, \hat{s}_{(i+1)*M}^l$ on the level

Algorithm 10 Plan(s_0, s_g, L, K, M, N)

Require: the initial state s_0 , the goal state s_g , number of subgoals K , number of levels L , multiplier M , number of samples N .

- o: Sample N latent action sequences $\{z_{1:K}^i\}_{i=1}^N$.
- o: Recursively generate subgoals $\{\hat{s}_{1:K}^i\}_{i=1}^N$ using $g(s, z)$.
- o: Select $z_{1:K}^*$ and $\hat{s}_{1:K}^*$ of the lowest cost.
- o: Update $z_{1:K}^*$ and $\hat{s}_{1:K}^*$ using MPPI.
- o: **if** $L = 1$ **then**
- o: **return** $\hat{s}_{1:K}^*$
- o: **else**
- o: Denote $\hat{s}_0^* \leftarrow s_0$.
- o: Initialize the plan $\hat{\mathcal{S}}$ as an empty list
- o: **for** $i = 1, \dots, K$ **do**
- o: Append Plan($\hat{s}_{i-1}^*, \hat{s}_i^*, L - 1, M, M, N$) to \mathcal{S}
- o: **return** $\hat{\mathcal{S}}$

=0

l , by treating \hat{s}_i^{l-1} and \hat{s}_{i+1}^{l-1} as the initial state and final goal state in Eqn. 56. The planned segments are returned to the previous level and concatenated as a more fine-grained plan. For this purpose, we train L conditional subgoal generators to propose subgoals that are $\Delta t_1, \dots, \Delta t_L$ steps away, respectively. In contrast to the prior work (Pertsch et al., 2020), the conditional subgoal generators enable us to plan for unseen goals that are beyond the temporal horizon of the demonstrations in the offline dataset by exploiting the compositional structure of the demonstrations. By recursively generating the subgoals across time at each level, we only need to enforce that the temporal resolution of the top level Δt_L is smaller than since the the conditional subgoal generator $f^l(s, z)$ needs to be trained on trajectories at least $\Delta t_L + 1$ steps in length.

We maintain a latent plan buffer for each level to further facilitate the planning with the conditional subgoal generator. After each episode, the selected latent representations on each level are appended to the corresponding latent plan buffer. In each target task, the subgoals are supposed to have the same semantic meaning. In spite of the variations of the initial and goals state in each episode, the optimal plans in the latent space can often be similar to each other. Therefore, we sample half of the latent representations from the prior distribution $p(z)$ and the other half from the latent plan buffer among the initial samples to enhance the chance of finding a close initial guess.

We build our planner upon the model predictive path integral (MPPI) (Gandhi et al., 2021), which iteratively optimizes the plan through importance sampling. In each

interaction, we perturb the chosen plan in the latent space with a small Gaussian noise as new candidates.

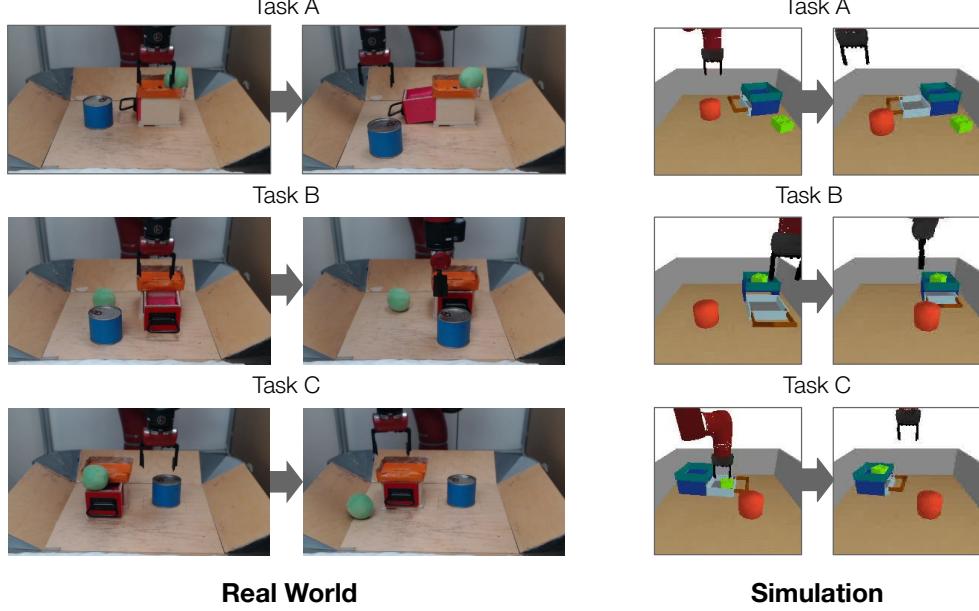


Figure 56: Target tasks. Three multi-stage tasks are designed for our experiments in the simulation and the real world respectively. In each target task, the robot needs to strategically interact with the environment (e.g. first takes out an object in the drawer then closes the drawer). The initial state and the desired goal state are shown for each task.

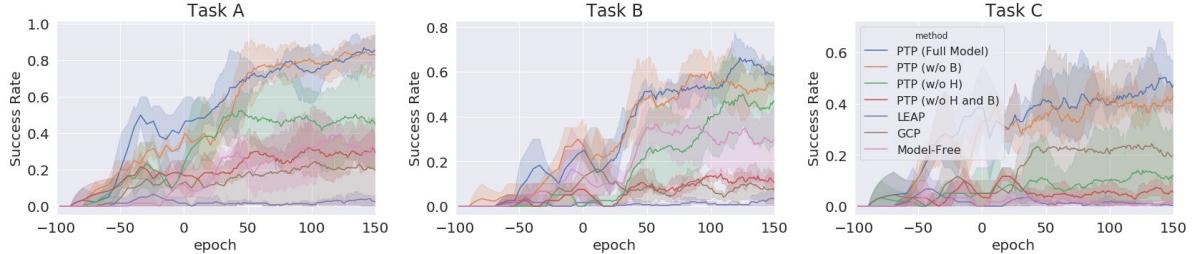


Figure 57: Quantitative comparison in simulation. The average success rate across 3 runs is shown with the shaded region indicating the standard deviation. The negative x-axis indicates the epochs of offline pre-training and positive x-axis indicates epochs of online fine-tuning. Using offline learning and planning, our method PTP is able to solve these tasks partially (at 0 epochs). Then with online finetuning the performance improves further. In contrast, prior methods have lower offline performance and do not fine-tune successfully in most cases, as they do not collect coherent online data.

11.5.3 Cost Function For Feasible Subgoals

To provide informative guidance to the policy $\pi(a|s, s_g)$, we would like that the final goal s_g can be reached at the end of the episode while encouraging the transition between each pair of subgoals to be feasible within a limited time budget. As explained in Sec. 11.4, the goal state is considered to be reached when the Euclidean distance between the last subgoal in the plan and the desired goal is less than a threshold δ in the learned latent space. The feasibility of each transition between adjacent subgoals can be measured using the goal-conditioned value function $V(s, s')$ trained by the reinforcement learning algorithm. Therefore, finding the subgoals $\hat{s}_{1:K}^*$ can be formulated as a constrained optimization problem:

$$\begin{aligned} & \text{minimize} && \|s_g - \hat{s}_K\| \\ & \text{subject to} && V(\hat{s}_i, \hat{s}_{i+1}) \geq \delta, \text{ for } i = 0, \dots, K-1 \end{aligned} \quad (55)$$

where we use $\hat{s}_0 = s_0$ to denote the initial state for convenience. By re-writing Eq. 55 as a Lagrangian, we obtain the cost function with a weight η :

$$c(s_0, \hat{s}_{1:K}, s_g) = \|s_g - \hat{s}_K\| + \eta \sum_{i=0}^{K-1} V(\hat{s}_i, \hat{s}_{i+1}) \quad (56)$$

The details of our method are explained in Sec. 11.6.2.

11.6 EXPERIMENTS

In our experiments, we aim to answer the following questions: 1) Can PTP propose and select feasible subgoals as plans for real-world robotic manipulation tasks? 2) Can the subgoals planned by PTP facilitate online fine-tuning of the goal-conditioned policies to solve target tasks unseen in the offline dataset? 3) How does each design option affect the performance of PTP? Videos of our experimental results are available on the project website: sites.google.com/view/planning-to-practice

11.6.1 Experimental Setup

Environment. As shown in Fig. 56, our experiments are conducted in a table-top manipulation environment with a Sawyer robot. At the beginning of each episode, a fixed

drawer and two movable objects are randomly placed on the table. The robot can change the state of the environment by opening/closing the drawer, sliding the objects, and picking and placing objects into different destinations, . At each time step, the robot receives a 48×48 RGB image via a Logitech C920 camera as the observation and takes a 5-dimensional continuous action to change the gripper status through position control. The action dictates the change of the coordinates along the three axes, the change of the rotation, and the status of the fingers. We use PyBullet ([Coumans and Bai, 2016–2021](#)) for our simulated experiments.

Prior data. The prior data consists of varied demonstrations for different primitive tasks. In each demonstrated trajectory, we randomly initialize the environment and perform primitive interactions such as opening the drawer and poking the object. These trajectories are collected using teleoperation in the real world, and a scripted policy that uses privileged information of the environment (e.g., the object pose and the status of the drawer) in simulation. The trajectories vary in length from 5 to 150 time steps, with 2,344 trajectories in the real world and 4,000 in simulation.

Target tasks. In each target task, a desired goal state is specified by a 48×48 RGB image (same dimension with the observation). The robot is tasked to reach the goal state by interacting with the objects on the table. Task success for our evaluation is determined based on the object positions at the end of each episode (this metric is not used for learning). As shown in Fig. 56, we design three target tasks that require multi-stage interactions with the environment to complete. These target tasks are designed with temporal dependencies between stages (e.g. the robot needs to first move away a can that blocks the drawer before opening the drawer). The transitions from the initial state to the goal state are unseen in the offline data. The episode length is 400 steps in simulation and 125 steps in the real world, which are much longer than the time horizon of the demonstrations.

Baselines and ablations. We compare PTP with 3 baselines and 3 ablations. **Model-Free** uses a policy directly conditioned on the final goal and conducts online fine-tuning without using any subgoals. **LEAP** ([Nasiriany et al., 2019](#)) learns a variational auto-encoder (VAE) ([D. P. Kingma and Welling, 2014](#)) to capture the prior distribution of states and plans for subgoals without conditioning on any context. **GCP** ([Pertsch et al., 2020](#)) learns a goal predictor that hierarchically generates intermediate subgoals between the initial state and the goal state. To analyze the design options in PTP, we also compare with variations of our method by removing the latent plan buffer (**PTP (w/o B)**), the hierarchical planning algorithm (**PTP (w/o H)**), and both of these two designs (**PTP (w/o H and B)**). All methods use the same neural network architecture in the goal-conditioned

policy and are pre-trained on the same offline dataset.

11.6.2 Implementation Details

Following Khazatsky et al., 2021, we use a vector quantized variational autoencoder (VQ-VAE) (Aäron van den Oord et al., 2017) as the state encoder, which encodes a $48 \times 48 \times 3$ image to a 720-dimensional encoding. The conditional subgoal generator is implemented with a U-Net architecture (Ronneberger et al., 2015) and decodes the subgoal from a 8-dimensional latent representations conditioned on the encoding of the current state. In our planner, we use $L = 3$, $K = 8$, $M = 2$, $N = 1024$, and we run MPPI for 5 iteration on each level. g is trained to predict subgoals that are 15, 30, and 60 steps away. Implicit Q-Learning (IQL) (Kostrikov et al., 2021b) is used as the underlying RL algorithm for offline pre-training and online fine-tuning with default hyperparameters. We use the same network architectures for the policy and the value functions from (Khazatsky et al., 2021) for simulation experiments. For real-world experiments, we use a convolutional neural network instead. We use Adam optimizer with a learning rate of $3 \cdot 10^{-4}$ and a batch size of 1024. During training, we relabel the goal with future hindsight experience replay (Andrychowicz et al., 2017a) with 70% probability. We use $\epsilon = 3$ for the reward function defined in Sec. 11.4, $\eta = 0.01$ in Eqn. 56.

11.6.3 Quantitative Comparisons

We evaluate PTP and baselines on three unseen target tasks. We use simulated versions of these tasks for comparisons and ablations, and real-world tasks, where all pretraining and finetuning uses only real-world data, to evaluate the practical effectiveness of the method.

Simulation. We first pre-train the goal-conditioned policy on the offline dataset for 100 epochs and the run online fine-tuning for the target task for 150 epochs. Each epoch takes 2,000 simulation steps (only during fine-tuning) and 2,000 training iterations. We run online fine-tuning using each method with 3 different random seeds. After each epoch, we test the policy in the target task for 5 episodes. We report the average success rate across 3 runs in Fig. 57 where the negative x-axis indicates the offline pre-training epochs and positive x-axis indicates the online fine-tuning epochs.

As shown in Fig. 57, our full model consistently outperforms baselines with a large performance gap. The generated subgoals not only enables the pre-trained policy to achieve higher success rate by breaking down the hard problems into easier pieces, but

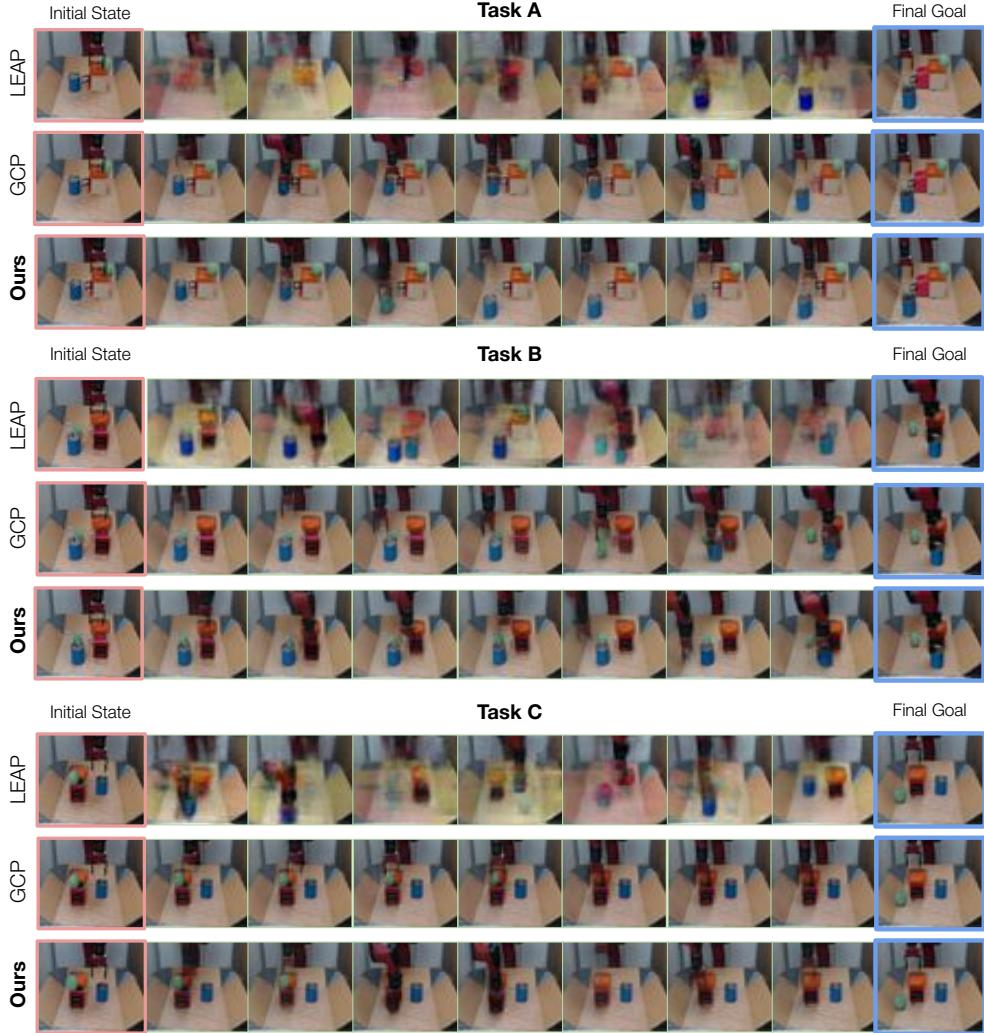


Figure 58: **Planned subgoal sequences.** Each row shows the sequence of subgoals produced by each method. The initial state and the final goal are shown at the two ends.

also introduces larger performance improvements during online fine-tuning. After fine-tuning for 150 epochs, the policy achieves the success rates of 84.9%, 59.9%, 49.3% in the three target tasks respectively. Compared to the policy pre-trained on the offline dataset, the performance is significantly improved (+31.6%, +37.8%, and +13.8%). When directly using the final goal or subgoals generated by baseline methods, the policy’s performance plateaus at around 0.0% to 30.0% and does not improve much during online fine-tuning.

We found that the hierarchical planner and the latent plan buffer are crucial for PTP’s performance. Without these two design options, the planner often suffers from the large

Table 5: The real-world success rates before and after online fine-tuning. The tasks are described in Sec. 11.6.1.

Task	PTP (Ours)	GCP
	Offline → Online	Offline → Online
Task A	12.5% → 62.5%	12.5% → 0.0%
Task B	75.0% → 100.0%	50.0% → 75.0%
Task C	25.0% → 50.0%	25.0% → 12.5%

search space of possible subgoal sequences and the resultant success rates decrease. The latent plan buffer significantly improves the performance of non-hierarchical PTP while it has a minor effect on hierarchical PTP.

Real-world evaluation. We pre-train the policy for 200 epochs and fine-tune it for 10 epochs. In each epoch, we run 10,000 training iterations and collect 1,000 steps in the real world. We train on three target tasks which are shown in Figure 56, and report the success rate of the goal-conditioned policy before and after online fine-tuning in Table 5. Planning enables the robot to succeed partially with just the offline initialized policy, achieving success rates of 12.5%, 75.0% and 25.0% on the three tasks. (When the offline policy is conditioned on only the final goal image without planning, the success rate is 0%.) Then in each task, we fine-tune to a significantly higher success rate.

Qualitatively, at the beginning of fine-tuning, the robot often fails, deviating from the planned subgoals or colliding with the environment. With the planned subgoals, the original long-horizon task is broken down to short snippets that are easier to complete. Even if a subgoal is not reached successfully at first, the data is useful to collect additional experience and fine-tune the policy. After fine-tuning for 4-5 epochs, we already observe that the robot’s performance reaching subgoals during training time significantly improves, collecting even more coherent and useful data. After 10 epochs, we achieve success rates of 62.5%, 100.0% and 50.0%. In comparison, GCP cannot provide useful guidance to the policy when the generated goals are noisy.

11.6.4 Generated Subgoals

In Fig. 58, we present qualitative results of the generated subgoals for each task in the real world. Each row shows a sequence of generated subgoals produced by the planner

in each method. In all the three target tasks, PTP successfully plans for a sequence of subgoals that can lead to the desired final goal. The transition between adjacent subgoals are feasible within a short period of time. By comparison, both of the baseline methods fail to generate reasonable plans. Without conditioning on the current state, LEAP ([Nasiriany et al., 2019](#)) can hardly produce any realistic images of the environment. Most of the generated subgoals are highly noisy images with duplicated robot arms and objects. The quality of the subgoals produced by GCP ([Pertsch et al., 2020](#)) is higher than that of LEAP but still much worse than ours. GCP cannot generalize well for the initial state and the goal state that are out of the distribution of the offline dataset, which contains only short snippets of demonstrations.

11.7 CONCLUSION AND DISCUSSION

We presented PTP, a method for real-world learning of temporally extended skills by utilizing planning and fine-tuning to stitch together skills from prior data. First, planning is used to convert a long-horizon task into achievable subgoals for a lower level goal-conditioned policy trained from prior data. Then, the goal-conditioned policy is further fine-tuned with active online interaction, mitigating the distribution shift between the offline data and actual states seen during rollouts. This procedure allows robots to extend their capabilities autonomously, composing previously seen data into more complicated and useful skills.

BIBLIOGRAPHY

- Abbeel, Pieter and Andrew Y Ng (2004). "Apprenticeship learning via inverse reinforcement learning." In: *International Conference on Machine Learning (ICML)*, p. 1.
- Abdolmaleki, Abbas, Jost Tobias Springenberg, Yuval Tassa, Remi Munos, Nicolas Heess, and Martin Riedmiller (2018). "Maximum a Posteriori Policy Optimisation." In: *International Conference on Learning Representations (ICLR)*, pp. 1–19.
- Abel, David, Gabriel Barth-Maron, James Macglashan, and Stefanie Tellex (2014). "Toward Affordance-Aware Planning." In: *RSS Workshop on Affordances in Vision for Cognitive Robotics*. URL: <https://vimeo.com/88689171>.
- Adolph, Karen E and John M Franchak (2017). "The development of motor behavior." In: *Wiley interdisciplinary reviews. Cognitive science*. doi: [10.1002/wcs.1430](https://doi.org/10.1002/wcs.1430) (cit. on p. 1).
- Agarwal, Rishabh, Dale Schuurmans, and Mohammad Norouzi (2019). "An Optimistic Perspective on Offline Reinforcement Learning." In: *International Conference on Machine Learning (ICML)*. arXiv: [1907.04543v2](https://arxiv.org/abs/1907.04543v2).
- Agrawal, Pulkit, Ashvin Nair, Pieter Abbeel, Jitendra Malik, and Sergey Levine (2016). "Learning to Poke by Poking: Experiential Learning of Intuitive Physics." In: *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv: [1606 . 07419](https://arxiv.org/abs/1606.07419). URL: <http://arxiv.org/abs/1606.07419> (cit. on pp. 7, 44, 141).
- Ahn, Michael, Henry Zhu, Kristian Hartikainen, Hugo Ponte, Abhishek Gupta, Sergey Levine, and Vikash Kumar (Sept. 2019). "ROBEL: Robotics Benchmarks for Learning with Low-Cost Robots." In: *Conference on Robot Learning (CoRL)*. arXiv. arXiv: [1909 . 11639](https://arxiv.org/abs/1909.11639). URL: <http://arxiv.org/abs/1909.11639>.
- Andrychowicz, Marcin, Dwight Crow, Alex Ray, Jonas Schneider, Rachel Fong, Peter Welinder, Bob McGrew, Joshua Tobin, P. Abbeel, and Wojciech Zaremba (2017a). "Hindsight Experience Replay." In: *Advances in Neural Information Processing Systems*.
- Andrychowicz, Marcin, Filip Wolski, Alex Ray, Jonas Schneider, Rachel Fong, Peter Welinder, Bob McGrew, Josh Tobin, Pieter Abbeel, and Wojciech Zaremba (2017b). "Hindsight Experience Replay." In: *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv: [1707 . 01495](https://arxiv.org/abs/1707.01495). URL: <https://arxiv.org/pdf/1707.01495.pdf> % 20[http://arxiv.org/abs/1707.01495](https://arxiv.org/abs/1707.01495) (cit. on pp. 8, 12, 15, 17, 22, 23, 30, 44, 49, 56, 61, 62, 64, 65, 69, 99, 142, 143, 147).

- Atkeson, Christopher G and Stefan Schaal (1997). "Robot Learning From Demonstration." In: *International Conference on Machine Learning (ICML)*. URL: <http://www.cc.gatech.edu/fac/fChris..>
- Bahdanau, Dzmitry, Kyunghyun Cho, and Yoshua Bengio (2015). "Neural Machine Translation by Jointly Learning to Align and Translate." In: *International Conference on Learning Representations (ICLR)*. URL: <https://arxiv.org/pdf/1409.0473.pdf>.
- Baldazzi, David and Muhammad Ghifary (2015). "Compatible Value Gradients for Reinforcement Learning of Continuous Deep Policies." In: *CoRR* abs/1509.03005. arXiv: <1509.03005>. URL: <http://arxiv.org/abs/1509.03005>.
- Baranes, A and P-Y Oudeyer (2012). "Active Learning of Inverse Models with Intrinsically Motivated Goal Exploration in Robots." In: *Robotics and Autonomous Systems* 61.1, pp. 49–73. DOI: <10.1016/j.robot.2012.05.008>. arXiv: <arXiv:1301.4862v1>. URL: <http://dx.doi.org/10.1016/j.robot.2012.05.008> (cit. on p. 8).
- Baranes, Adrien and Pierre-Yves Oudeyer (2012a). "Active Learning of Inverse Models with Intrinsically Motivated Goal Exploration in Robots." In: *Robotics and Autonomous Systems* 61.1, pp. 49–73. DOI: <10.1016/j.robot.2012.05.008>. arXiv: <arXiv:1301.4862v1>. URL: <http://dx.doi.org/10.1016/j.robot.2012.05.008>.
- Baranes, Adrien and Pierre-Yves Oudeyer (2012b). "Active Learning of Inverse Models with Intrinsically Motivated Goal Exploration in Robots." In: *Robotics and Autonomous Systems* 61.1, pp. 49–73. DOI: <10.1016/j.robot.2012.05.008>. arXiv: <arXiv:1301.4862v1>. URL: <http://dx.doi.org/10.1016/j.robot.2012.05.008>.
- Barber, David and Felix V Agakov (2004). "Information maximization in noisy channels: A variational approach." In: *Advances in Neural Information Processing Systems*, pp. 201–208.
- Bellemare, Marc, Sriram Srinivasan, Georg Ostrovski, Tom Schaul, David Saxton, and Remi Munos (2016). "Unifying count-based exploration and intrinsic motivation." In: *Advances in Neural Information Processing Systems (NeurIPS)*, pp. 1471–1479 (cit. on pp. [13](#), [30](#), [142](#)).
- Benbrahim, Hamid and Judy A. Franklin (Dec. 1997). "Biped dynamic walking using reinforcement learning." In: *Robotics and Autonomous Systems* 22.3-4, pp. 283–302. ISSN: 09218890. DOI: [10.1016/S0921-8890\(97\)00043-2](10.1016/S0921-8890(97)00043-2).
- Bentivegna, Darrin C., Gordon Cheng, and Christopher G. Atkeson (2003). "Learning from Observation and from Practice Using Behavioral Primitives." In: *Robotics Research, The Eleventh International Symposium, ISRR, October 19-22, 2003, Siena, Italy*. Ed. by Paolo Dario and Raja Chatila. Vol. 15. Springer Tracts in Advanced Robotics.

- Springer, pp. 551–560. DOI: [10.1007/11008941__59](https://doi.org/10.1007/11008941__59). URL: https://doi.org/10.1007/11008941__59.
- Berger, Kathleen (2014). *The Developing Person Through the Life Span*.
- Bhatnagar, Shalabh, Richard S. Sutton, Mohammad Ghavamzadeh, and Mark Lee (2009). “Natural actor-critic algorithms.” In: *Autom.* 45.11, pp. 2471–2482. DOI: [10.1016/j.automatica.2009.07.008](https://doi.org/10.1016/j.automatica.2009.07.008). URL: <https://doi.org/10.1016/j.automatica.2009.07.008>.
- Bohg, Jeannette and Danica Kragic (2010). “Learning grasping points with shape context.” In: *Robotics and Autonomous Systems* 58.4, pp. 362–377.
- Bojarski, Mariusz, Davide Del Testa, Daniel Dworakowski, Bernhard Firner, Beat Flepp, Prasoon Goyal, Lawrence D. Jackel, Mathew Monfort, Urs Muller, Jiakai Zhang, Xin Zhang, Jake Zhao, Karol Zieba, Davide Del Testa, Daniel Dworakowski, Bernhard Firner, Beat Flepp, Prasoon Goyal, Lawrence D. Jackel, Mathew Monfort, Urs Muller, Jiakai Zhang, Xin Zhang, Jake Zhao, and Karol Zieba (2016). “End to End Learning for Self-Driving Cars.” In: *CoRR* abs/1604.0, pp. 1–9. arXiv: [1604.07316](https://arxiv.org/pdf/1604.07316.pdf). URL: <https://arxiv.org/pdf/1604.07316.pdf> <http://arxiv.org/abs/1604.07316>.
- Bradbury, James, Roy Frostig, Peter Hawkins, Matthew James Johnson, Chris Leary, Dougal Maclaurin, George Necula, Adam Paszke, Jake VanderPlas, Skye Wanderman-Milne, and Qiao Zhang (2018). *JAX: composable transformations of Python+NumPy programs*. Version 0.2.5. URL: <http://github.com/google/jax>.
- Brandfonbrener, David, William F Whitney, Rajesh Ranganath, and Joan Bruna (2021). “Offline RL Without Off-Policy Evaluation.” In: *arXiv preprint arXiv:2106.08909*.
- Burda, Yuri, Harri Edwards, Deepak Pathak, Amos Storkey, Trevor Darrell, and Alexei A Efros (2019). “Large-scale study of curiosity-driven learning.” In: *International Conference on Learning Representations (ICLR)*.
- Burda, Yuri, Harrison Edwards, Amos Storkey, and Oleg Klimov (2018). “Exploration by random network distillation.” In: *arXiv preprint arXiv:1810.12894*.
- Cabi, Serkan, Sergio Gómez Colmenarejo, Alexander Novikov, Ksenia Konyushkova, Scott Reed, Rae Jeong, Konrad Zolna, Yusuf Aytar, David Budden, Mel Vecerik, et al. (2019). “Scaling data-driven robotics with reward sketching and batch reinforcement learning.” In: *arXiv preprint arXiv:1909.12200*.
- Chane-Sane, Elliot, Cordelia Schmid, and Ivan Laptev (2021). “Goal-Conditioned Reinforcement Learning with Imagined Subgoals.” In: *ICML*.
- Charlesworth, Henry and G. Montana (2020). “PlanGAN: Model-based Planning With Sparse Rewards and Multiple Goals.” In: *ArXiv* abs/2006.00900.

- Chebotar, Yevgen, Karol Hausman, Yao Lu, Ted Xiao, Dmitry Kalashnikov, Jake Varley, Alex Irpan, Benjamin Eysenbach, Ryan Julian, Chelsea Finn, and Sergey Levine (Apr. 2021). "Actionable Models: Unsupervised Offline Reinforcement Learning of Robotic Skills." In.
- Chebotar, Yevgen, Karol Hausman, Marvin Zhang, Gaurav Sukhatme, Stefan Schaal, and Sergey Levine (2017a). "Combining Model-Based and Model-Free Updates for Trajectory-Centric Reinforcement Learning." In: *International Conference on Machine Learning (ICML)*. URL: <https://arxiv.org/pdf/1703.03078.pdf>.
- Chebotar, Yevgen, Mrinal Kalakrishnan, Ali Yahya, Adrian Li, Stefan Schaal, and Sergey Levine (2017b). "Path integral guided policy search." In: *2017 IEEE International Conference on Robotics and Automation (ICRA)*. IEEE, pp. 3381–3388.
- Chen, Lili, Kevin Lu, Aravind Rajeswaran, Kimin Lee, Aditya Grover, Michael Laskin, Pieter Abbeel, Aravind Srinivas, and Igor Mordatch (2021). "Decision transformer: Reinforcement learning via sequence modeling." In: *arXiv preprint arXiv:2106.01345*.
- Chen, Xi, Yan Duan, Rein Houthooft, John Schulman, Ilya Sutskever, and Pieter Abbeel (2016). "Infogan: Interpretable representation learning by information maximizing generative adversarial nets." In: *Advances in Neural Information Processing Systems (NeurIPS)*, pp. 2172–2180. arXiv: [1606 . 03657](https://arxiv.org/abs/1606.03657). URL: <http://arxiv.org/abs/1606.03657> (cit. on p. 8).
- Chentanez, Nuttapong, Andrew G Barto, and Satinder P Singh (2005). "Intrinsically motivated reinforcement learning." In: *Advances in neural information processing systems*, pp. 1281–1288.
- Cheung, Brian, Jesse A Livezey, Arjun K Bansal, and Bruno A Olshausen (2014). "Discovering hidden factors of variation in deep networks." In: *arXiv preprint arXiv:1412.6583* (cit. on p. 8).
- Christiano, Paul F, Jan Leike, Tom B Brown, Miljan Martic, Shane Legg, and Dario Amodei (2017). "Deep reinforcement learning from human preferences." In: *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv: [1706 . 03741](https://arxiv.org/pdf/1706.03741.pdf). URL: <https://arxiv.org/pdf/1706.03741.pdf%20http://arxiv.org/abs/1706.03741>.
- Colas, Cédric, Pierre Fournier, Olivier Sigaud, and Pierre-Yves Oudeyer (2018a). "CURIOS: Intrinsically Motivated Multi-Task, Multi-Goal Reinforcement Learning." In: *CoRR abs/1810.06284*.
- Colas, Cédric, Tristan Karch, Olivier Sigaud, and Pierre-Yves Oudeyer (2021). *Intrinsically Motivated Goal-Conditioned Reinforcement Learning: a Short Survey*. Tech. rep. arXiv: [2012.09830v2](https://arxiv.org/abs/2012.09830v2).

- Colas, Cédric, Olivier Sigaud, and Pierre-Yves Oudeyer (2018b). "GEP-PG: Decoupling Exploration and Exploitation in Deep Reinforcement Learning Algorithms." In: *International Conference on Machine Learning (ICML)*.
- Coumans, Erwin and Yunfei Bai (2016–2021). *PyBullet, a Python module for physics simulation for games, robotics and machine learning*. <http://pybullet.org>.
- Dabney, Will, Georg Ostrovski, David Silver, and Rémi Munos (2018a). "Implicit quantile networks for distributional reinforcement learning." In: *International conference on machine learning*. PMLR, pp. 1096–1105.
- Dabney, Will, Mark Rowland, Marc G Bellemare, and Rémi Munos (2018b). "Distributional reinforcement learning with quantile regression." In: *Thirty-Second AAAI Conference on Artificial Intelligence*.
- Daniel, Christian, Malte Viering, Jan Metz, Oliver Kroemer, and Jan Peters (2014). "Active Reward Learning." In: *Robotics: Science and Systems (RSS)*. URL: <http://www.roboticsproceedings.org/rss10/p31.pdf>.
- Degris, Thomas, Martha White, and Richard S. Sutton (May 2012). "Off-Policy Actor-Critic." In: *International Conference on Machine Learning (ICML)*. arXiv: [1205 . 4839](https://arxiv.org/abs/1205.4839). URL: <http://arxiv.org/abs/1205.4839>.
- Deisenroth, M. P., P. Englert, J. Peters, and D. Fox (May 2014). "Multi-task policy search for robotics." In: *2014 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 3876–3881. DOI: [10.1109/ICRA.2014.6907421](https://doi.org/10.1109/ICRA.2014.6907421).
- Deisenroth, Marc Peter and Carl Edward Rasmussen (2011). "PILCO: A model-based and data-efficient approach to policy search." In: *International Conference on Machine Learning (ICML)*, pp. 465–472. URL: <http://mlg.eng.cam.ac.uk/pub/pdf/DeiRas11.pdf>.
- Deisenroth, Marc Peter, Carl Edward Rasmussen, and Dieter Fox (2011a). "Learning to Control a Low-Cost Manipulator using Data-Efficient Reinforcement Learning." In: *Robotics: Science and Systems VII*, pp. 57–64.
- Deisenroth, Marc Peter, Carl Edward Rasmussen, and Dieter Fox (2011b). "Learning to Control a Low-Cost Manipulator using Data-Efficient Reinforcement Learning." In: *Robotics: Science and Systems (RSS) VII*, pp. 57–64. URL: <http://www.roboticsproceedings.org/rss07/p08.pdf>.
- Desjardins, Guillaume, Aaron Courville, and Yoshua Bengio (2012). "Disentangling factors of variation via generative entangling." In: *CoRR* abs/1210.5 (cit. on p. 8).
- Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova (Oct. 2019). "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding." In:

- Association for Computational Linguistics (ACL).* arXiv: [1810.04805](#). URL: <http://arxiv.org/abs/1810.04805> (cit. on pp. [2](#), [102](#), [154](#)).
- Ding, Yiming, Carlos Florensa, Mariano Phiellipp, and P. Abbeel (2019). “Goal-conditioned Imitation Learning.” In: *NeurIPS*.
- Donahue, Jeff, Philipp Krähenbühl, and Trevor Darrell (May 2017). “Adversarial Feature Learning.” In: *International Conference on Learning Representations (ICLR)*. arXiv: [1605.09782](#). URL: <http://arxiv.org/abs/1605.09782>.
- Duan, Yan, Marcin Andrychowicz, Bradly C. Stadie, Jonathan Ho, Jonas Schneider, Ilya Sutskever, Pieter Abbeel, and Wojciech Zaremba (2017). “One-Shot Imitation Learning.” In: *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv: [1703.07326](#).
- Duan, Yan, John Schulman, Xi Chen, Peter L. Bartlett, Ilya Sutskever, and Pieter Abbeel (Nov. 2016). “RL\$
2\$: Fast Reinforcement Learning via Slow Reinforcement Learning.” In: arXiv: [1611.02779](#). URL: <http://arxiv.org/abs/1611.02779>.
- Ebert, Frederik, Sudeep Dasari, Alex X Lee, Sergey Levine, and Chelsea Finn (2018). “Robustness via Retrying: Closed-Loop Robotic Manipulation with Self-Supervised Learning.” In: *Conference on Robot Learning (CoRL)*. arXiv: [1810.03043v1](#). URL: <https://arxiv.org/pdf/1810.03043.pdf>.
- Ebert, Frederik, Chelsea Finn, Sudeep Dasari, Annie Xie, Alex Lee, and Sergey Levine (n.d.). “Visual Foresight: Model-Based Deep Reinforcement Learning for Vision-Based Robotic Control.” In: (). arXiv: [1812.00568v1](#). URL: <https://sites.google.com/view/visualforesight>.
- Ebert, Frederik, Chelsea Finn, Alex X Lee, and Sergey Levine (2017). “Self-Supervised Visual Planning with Temporal Skip Connections.” In: *Conference on Robot Learning (CoRL)*. URL: <https://128.84.21.199/pdf/1710.05268.pdf%20https://arxiv.org/pdf/1710.05268.pdf> (cit. on pp. [7](#), [44](#)).
- Ekvall, Staffan and Danica Kragic (2004). “Interactive grasp learning based on human demonstration.” In: *IEEE International Conference on Robotics and Automation, 2004. Proceedings. ICRA'04*. 2004. Vol. 4. IEEE, pp. 3519–3524.
- Eysenbach, Benjamin, Xinyang Geng, Sergey Levine, and Ruslan Salakhutdinov (2020). “Rewriting History with Inverse RL: Hindsight Inference for Policy Improvement.” In: *ArXiv* abs/2002.11089.
- Eysenbach, Benjamin, Abhishek Gupta, Julian Ibarz, and Sergey Levine (2018). “Diversity is All You Need: Learning Skills without a Reward Function.” In: *arXiv preprint arXiv:1802.06070* (cit. on pp. [13](#), [31](#)).

- Eysenbach, Benjamin, Ruslan Salakhutdinov, and Sergey Levine (2019). "Search on the Replay Buffer: Bridging Planning and Reinforcement Learning." In: *NeurIPS*.
- Eysenbach, Benjamin, Ruslan Salakhutdinov, and Sergey Levine (2021). "C-Learning: Learning to Achieve Goals via Recursive Classification." In: *ArXiv* abs/2011.08909.
- Fakoor, Rasool, Pratik Chaudhari, and Alexander J Smola (2019). "P₃O: Policy-on Policy-off Policy Optimization." In: *Conference on Uncertainty in Artificial Intelligence (UAI)*. arXiv: 1905.01756v2. URL: <https://github.com/rasoolfa/P3O..>
- Fakoor, Rasool, Jonas Mueller, Kavosh Asadi, Pratik Chaudhari, and Alexander J Smola (2021). "Continuous doubly constrained batch reinforcement learning." In: *arXiv preprint arXiv:2102.09225*.
- Falco, Joe, Yu Sun, and Maximo Roa (2018). "Robotic Grasping and Manipulation Competition: Competitor Feedback and Lessons Learned." In: *Robotic Grasping and Manipulation*. Ed. by Yu Sun and Joe Falco. Cham: Springer International Publishing, pp. 180–189. ISBN: 978-3-319-94568-2.
- Fang, Kuan, Yuke Zhu, Animesh Garg, Silvio Savarese, and Li Fei-Fei (2019). "Dynamics Learning with Cascaded Variational Inference for Multi-Step Manipulation." In: *Conference on Robot Learning (CoRL)*.
- Fang, Meng, Tianyi Zhou, Yali Du, Lei Han, and Zhengyou Zhang (2019). "Curriculum-guided Hindsight Experience Replay." In: *NeurIPS*.
- Fikes, Richard E and Nils J Nilsson (1971). "STRIPS: A New Approach to the Application of Theorem Proving to Problem Solving." In: *IJCAI*.
- Finn, Chelsea, Pieter Abbeel, and Sergey Levine (2017). "Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks." In: *International Conference on Machine Learning (ICML)*. URL: <https://arxiv.org/pdf/1703.03400.pdf>.
- Finn, Chelsea and Sergey Levine (2016). "Deep Visual Foresight for Planning Robot Motion." In: *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv: 1610.00696. URL: <https://arxiv.org/pdf/1610.00696.pdf%20http://arxiv.org/abs/1610.00696> (cit. on pp. 7, 44).
- Finn, Chelsea, Sergey Levine, and Pieter Abbeel (2016a). "Guided Cost Learning: Deep Inverse Optimal Control via Policy Optimization." In: *International Conference on Machine Learning (ICML)*. URL: <https://arxiv.org/pdf/1603.00448.pdf>.
- Finn, Chelsea, Xin Yu Tan, Yan Duan, Trevor Darrell, Sergey Levine, and Pieter Abbeel (2016b). "Deep spatial autoencoders for visuomotor learning." In: *IEEE International Conference on Robotics and Automation (ICRA)*. Vol. 2016-June. IEEE, pp. 512–519. ISBN: 9781467380263. DOI: [10.1109/ICRA.2016.7487173](https://doi.org/10.1109/ICRA.2016.7487173). arXiv: 1509.06113 (cit. on pp. 8, 15).

- Florence, Pete, Corey Lynch, Andy Zeng, Oscar Ramirez, Ayzaan Wahid, Laura Downs, Adrian Wong, Johnny Lee, Igor Mordatch, and Jonathan Tompson (2021). "Implicit Behavioral Cloning." In: *arXiv preprint arXiv:2109.00137*.
- Florensa, Carlos, Jonas Degrave, Nicolas Heess, Jost Tobias Springenberg, and Martin Riedmiller (2018a). "Self-supervised Learning of Image Embedding for Continuous Control." In: *Workshop on Inference to Control at NeurIPS*.
- Florensa, Carlos, Jonas Degrave, Nicolas Heess, Jost Tobias Springenberg, and Martin Riedmiller (2019). "Self-supervised Learning of Image Embedding for Continuous Control." In: *arXiv preprint arXiv:1901.00943*.
- Florensa, Carlos, Yan Duan, and Pieter Abbeel (2017). "Stochastic neural networks for hierarchical reinforcement learning." In: *arXiv preprint arXiv:1704.03012* (cit. on pp. 13, 31).
- Florensa, Carlos, David Held, Markus Wulfmeier, and Pieter Abbeel (2018b). "Reverse Curriculum Generation for Reinforcement Learning." In: *International Conference on Learning Representations (ICLR)*. arXiv: 1707. 05300. URL: <https://arxiv.org/pdf/1707.05300.pdf>%20<http://arxiv.org/abs/1707.05300>.
- Frank, Jordan, Shie Mannor, and Doina Precup (2008). "Reinforcement Learning in the Presence of Rare Events." In: *International Conference on Machine Learning (ICML)*. URL: <http://citeserx.ist.psu.edu/viewdoc/download?doi=10.1.1.149.4737%7B%5C%7Drep=rep1%7B%5C%7Dtype=pdf>.
- Fu, Justin, Aviral Kumar, Ofir Nachum, George Tucker, and Sergey Levine (Apr. 2020). "D4RL: Datasets for Deep Data-Driven Reinforcement Learning." In: arXiv: 2004 . 07219. URL: <http://arxiv.org/abs/2004.07219>.
- Fu, Justin, John D Co-Reyes, and Sergey Levine (2017). "EX 2 : Exploration with Exemplar Models for Deep Reinforcement Learning." In: *Advances in Neural Information Processing Systems (NeurIPS)*. URL: <https://papers.nips.cc/paper/6851-ex2-exploration-with-exemplar-models-for-deep-reinforcement-learning.pdf>%20<https://arxiv.org/pdf/1703.01260.pdf>.
- Fujimoto, Scott and Shixiang Shane Gu (2021). "A Minimalist Approach to Offline Reinforcement Learning." In: *arXiv preprint arXiv:2106.06860*.
- Fujimoto, Scott, Herke Hoof, and David Meger (2018a). "Addressing function approximation error in actor-critic methods." In: *International Conference on Machine Learning*. PMLR, pp. 1587–1596.
- Fujimoto, Scott, Herke van Hoof, and David Meger (2018b). "Addressing Function Approximation Error in Actor-Critic Methods." In: *International Conference on Machine Learning (ICML)* (cit. on pp. 9, 11, 13, 15, 48, 79, 94, 105, 111, 113).

- Fujimoto, Scott, David Meger, and Doina Precup (Dec. 2019a). "Off-Policy Deep Reinforcement Learning without Exploration." In: *International Conference on Machine Learning (ICML)*. arXiv: 1812.02900. URL: <http://arxiv.org/abs/1812.02900>.
- Fujimoto, Scott, David Meger, and Doina Precup (2019b). "Off-policy deep reinforcement learning without exploration." In: *International Conference on Machine Learning*. PMLR, pp. 2052–2062.
- Gandhi, Manan S., Bogdan Vlahov, Jason Gibson, Grady Williams, and Evangelos A. Theodorou (2021). "Robust Model Predictive Path Integral Control: Analysis and Performance Guarantees." In: *IEEE Robotics and Automation Letters* 6, pp. 1423–1430.
- Gao, Yang, Huazhe Xu, Ji Lin, Fisher Yu, Sergey Levine, and Trevor Darrell (2018). "Reinforcement Learning from Imperfect Demonstrations." In: *CoRR* abs/1802.05313. arXiv: 1802.05313. URL: <http://arxiv.org/abs/1802.05313>.
- Ghasemipour, Seyed Kamyar Seyed, Dale Schuurmans, and Shixiang Shane Gu (2021). "Emaq: Expected-max q-learning operator for simple yet effective offline and online rl." In: *International Conference on Machine Learning*. PMLR, pp. 3682–3691.
- Ghosh, Dibya, Abhishek Gupta, Ashwin Reddy, Justin Fu, Coline Devin, Benjamin Eysenbach, and Sergey Levine (2021). "Learning to Reach Goals via Iterated Supervised Learning." In: *arXiv: Learning*.
- Gibson, James (1979). *The Ecological Approach to Visual Perception*.
- Giusti, Alessandro, Jérôme Jerome Guzzi, Dan C Cirean, Fang-Lin He, Juan P Rodríguez, Flavio Fontana, Matthias Faessler, Christian Forster, Jürgen Jürgen Schmidhuber, Gianni Di Caro, Davide Scaramuzza, Luca M Gambardella, Dan C. Ciresan, Fang-Lin He, Juan P. Rodriguez, Flavio Fontana, Matthias Faessler, Christian Forster, Jürgen Jürgen Schmidhuber, Gianni Di Caro, Davide Scaramuzza, and Luca M Gambardella (2015). "A Machine Learning Approach to Visual Perception of Forest Trails for Mobile Robots." In: *IEEE Robotics and Automation Letters (RAL)*. Vol. 1. 2, pp. 2377–3766. ISBN: 9781467380256. DOI: 10.1109/LRA.2015.2509024. URL: <http://bit.ly/perceivingtrails.%20http://ieeexplore.ieee.org/document/7358076/>.
- Goodfellow, Ian J, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio (2014). "Generative Adversarial Nets." In: *Advances in Neural Information Processing Systems (NeurIPS)*. URL: <https://arxiv.org/pdf/1406.2661.pdf>.
- Gu, Shixiang, Ethan Holly, Timothy Lillicrap, and Sergey Levine (2017). "Deep Reinforcement Learning for Robotic Manipulation with Asynchronous Off-Policy Updates." In: *IEEE International Conference on Robotics and Automation (ICRA)*. ISSN: 0028-0836.

- DOI: [10.1038/nature20101](https://doi.org/10.1038/nature20101). arXiv: [1610.00633](https://arxiv.org/abs/1610.00633). URL: <https://arxiv.org/pdf/1610.00633.pdf>%20[http://arxiv.org/abs/1610.00633](https://arxiv.org/abs/1610.00633).
- Gu, Shixiang, Timothy Lillicrap, Ilya Sutskever, and Sergey Levine (2016). “Continuous Deep Q-Learning with Model-based Acceleration.” In: *International Conference on Machine Learning (ICML)*. ISBN: 3405062780. DOI: [10.3390/robotics2030122](https://doi.org/10.3390/robotics2030122). arXiv: [1603.00748](https://arxiv.org/abs/1603.00748). URL: <https://arxiv.org/pdf/1603.00748.pdf>%20[http://arxiv.org/abs/1603.00748](https://arxiv.org/abs/1603.00748).
- Gulcehre, Caglar, Sergio Gómez Colmenarejo, Ziyu Wang, Jakub Sygnowski, Thomas Paine, Konrad Zolna, Yutian Chen, Matthew Hoffman, Razvan Pascanu, and Nando de Freitas (2021). “Regularized behavior value estimation.” In: *arXiv preprint arXiv:2103.09575*.
- Gupta, Abhishek, Benjamin Eysenbach, Chelsea Finn, and Sergey Levine (2018a). “Unsupervised Meta-Learning for Reinforcement Learning.” In: *CoRR* abs:1806.04640.
- Gupta, Abhishek, Vikash Kumar, Corey Lynch, Sergey Levine, and Karol Hausman (Oct. 2019a). “Relay Policy Learning: Solving Long-Horizon Tasks via Imitation and Reinforcement Learning.” In: *Conference on Robot Learning (CoRL)*. arXiv: [1910.11956](https://arxiv.org/abs/1910.11956). URL: [http://arxiv.org/abs/1910.11956](https://arxiv.org/abs/1910.11956).
- Gupta, Abhishek, Vikash Kumar, Corey Lynch, Sergey Levine, and Karol Hausman (2019b). “Relay Policy Learning: Solving Long-Horizon Tasks via Imitation and Reinforcement Learning.” In: *CoRL*.
- Gupta, Abhishek, Russell Mendonca, Yuxuan Liu, Pieter Abbeel, and Sergey Levine (2018b). “Meta-Reinforcement Learning of Structured Exploration Strategies.” In: *Advances in Neural Information Processing Systems (NIPS)*. arXiv: [arXiv : 1802 . 07245v1](https://arxiv.org/abs/1802.07245v1). URL: <https://arxiv.org/pdf/1802.07245.pdf>.
- Gupta, Abhishek, Justin Yu, Tony Zhao, Vikash Kumar, Kelvin Xu, Thomas Devlin, Aaron Rovinsky, and Sergey Levine (2021). “Reset-Free Reinforcement Learning via Multi-Task Learning: Learning Dexterous Manipulation Behaviors without Human Intervention.” In: *International Conference on Robotics and Automation (ICRA)*.
- Ha, David and Jürgen Schmidhuber (2018). “World Models.” In: *arXiv preprint arXiv:1803.10122* (cit. on p. 8).
- Haarnoja, Tuomas, Aurick Zhou, Pieter Abbeel, and Sergey Levine (2018a). “Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor.” In: *International Conference on Machine Learning (ICML)*. arXiv: [arXiv : 1801 . 01290v2](https://arxiv.org/abs/1801.01290v2). URL: <https://arxiv.org/pdf/1801.01290.pdf>.

- Haarnoja, Tuomas, Aurick Zhou, Kristian Hartikainen, George Tucker, Sehoon Ha, Jie Tan, Vikash Kumar, Henry Zhu, Abhishek Gupta, Pieter Abbeel, and Sergey Levine (2018b). "Soft Actor-Critic Algorithms and Applications." In: *CoRR* abs/1812.05905.
- Hallak, Assaf and Shie Mannor (2017). "Consistent On-Line Off-Policy Evaluation." In: *International Conference on Machine Learning (ICML)*. arXiv: [1702.07121v1](#).
- Hallak, Assaf, Francois Schnitzler, Timothy Mann, and Shie Mannor (2015). "Off-policy Model-based Learning under Unknown Factored Dynamics." In: *International Conference on Machine Learning (ICML)*.
- Hallak, Assaf, Aviv Tamar, Rémi Munos, and Shie Mannor (2016). "Generalized Emphatic Temporal Difference Learning: Bias-Variance Analysis." In: *Association for the Advancement of Artificial Intelligence (AAAI)*. arXiv: [1509.05172v2](#).
- Hart, Stephen and Roderic Grupen (2010). "Learning Generalizable Control Programs." In: *IEEE Transactions on Autonomous Mental Development*.
- Hassanin, Mohammed, Salman Khan, and Murat Tahtali (2018). *Visual Affordance and Function Understanding: A Survey*. Tech. rep. 1. arXiv: [1807.06775v1](#).
- Hausman, Karol, Jost Tobias Springenberg, Ziyu Wang, Nicolas Heess, and Martin Riedmiller (2018). "Learning an Embedding Space for Transferable Robot Skills." In: *International Conference on Learning Representations (ICLR)*, pp. 1–16.
- Hazan, Elad, Sham M. Kakade, Karan Singh, and Abby Van Soest (2019). "Provably Efficient Maximum Entropy Exploration." In: *International Conference on Machine Learning (ICML)*.
- He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun (2016). "Deep Residual Learning for Image Recognition." In: *Conference on Computer Vision and Pattern Recognition (CVPR)*. arXiv: [1512.03385v1](#). URL: <http://image-net.org/challenges/LSVRC/2015/>.
- Held, David, Xinyang Geng, Carlos Florensa, and Pieter Abbeel (2018). "Automatic Goal Generation for Reinforcement Learning Agents." In: *International Conference on Machine Learning (ICML)*. URL: <https://arxiv.org/pdf/1705.06366.pdf>.
- Hessel, Matteo, Joseph Modayil, Hado van Hasselt, Tom Schaul, Georg Ostrovski, Will Dabney, Dan Horgan, Bilal Piot, Mohammad Azar, and David Silver (2018). *Rainbow: Combining Improvements in Deep Reinforcement Learning*. arXiv: [1710.02298v1](#). URL: www.aaai.org%20https://arxiv.org/pdf/1710.02298.pdf.
- Hester, Todd, Matej Vecerik, Olivier Pietquin, Marc Lanctot, Tom Schaul, Bilal Piot, Dan Horgan, John Quan, Andrew Sendonaris, Gabriel Dulac-Arnold, Ian Osband, John Agapiou, Joel Z Leibo, and Audrunas Gruslys (2018). "Learning from Demonstrations for Real World Reinforcement Learning." In: *AAAI Conference on Artificial Intel-*

- ligence.* arXiv: [1704.03732](#). URL: <https://arxiv.org/pdf/1704.03732.pdf> <http://arxiv.org/abs/1704.03732>.
- Higgins, Irina, Loic Matthey, Arka Pal, Christopher Burgess, Xavier Glorot, Matthew Botvinick, Shakir Mohamed, and Alexander Lerchner (2017a). “\$\$-VAE: Learning basic visual concepts with a constrained variational framework.” In: *International Conference on Learning Representations (ICLR)* (cit. on pp. [9](#), [30](#)).
- Higgins, Irina, Arka Pal, Andrei A Rusu, Loic Matthey, Christopher P Burgess, Alexander Pritzel, Matthew Botvinick, Charles Blundell, and Alexander Lerchner (2017b). “Darla: Improving zero-shot transfer in reinforcement learning.” In: *International Conference on Machine Learning (ICML)* (cit. on p. [8](#)).
- Houthooft, Rein, Xi Chen, Yan Duan, John Schulman, Filip De Turck, and Pieter Abbeel (2016). “Variational Information Maximizing Exploration.” In: *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv: [1605.09674](#).
- Ichter, Brian, James Harrison, and Marco Pavone (Sept. 2018). “Learning Sampling Distributions for Robot Motion Planning.” In: *ICRA*, pp. 7087–7094. ISSN: 10504729. DOI: [10.1109/ICRA.2018.8460730](#).
- Ijspeert, Auke Jan, Jun Nakanishi, and Stefan Schaal (2002). “Learning Attractor Landscapes for Learning Motor Primitives.” In: *Advances in Neural Information Processing Systems (NeurIPS)*, pp. 1547–1554. ISBN: 1049-5258. URL: <https://papers.nips.cc/paper/2140-learning-attractor-landscapes-for-learning-motor-primitives.pdf>.
- Inoue, Tadanobu, Giovanni De Magistris, Asim Munawar, Tsuyoshi Yokoya, and Ryuki Tachibana (2017). “Deep reinforcement learning for high precision assembly tasks.” In: *IROS*, pp. 819–825.
- Jaques, Natasha, Asma Ghandeharioun, Judy Hanwen Shen, Craig Ferguson, Ágata Lapedriza, Noah Jones, Shixiang Gu, and Rosalind W. Picard (2019). “Way Off-Policy Batch Deep Reinforcement Learning of Implicit Human Preferences in Dialog.” In: *CoRR* abs/1907.00456. arXiv: [1907.00456](#). URL: <http://arxiv.org/abs/1907.00456>.
- Jiang, Nan and Lihong Li (2016). “Doubly Robust Off-policy Value Evaluation for Reinforcement Learning.” In: *International Conference on Machine Learning (ICML)*. arXiv: [1511.03722v3](#).
- Johannink, Tobias, Shikhar Bahl, Ashvin Nair, Jianlan Luo, Avinash Kumar, Matthias Loskyll, Juan Aparicio Ojea, Eugen Solowjow, and Sergey Levine (2019). “Residual Reinforcement Learning for Robot Control.” In: *IEEE International Conference on Robotics and Automation (ICRA)*. arXiv: [1812.03201v2](#). URL: <https://arxiv.org/pdf/1812.03201.pdf>.

- Johnson, Eric and Anthony Calise (2000). "Pseudo-Control Hedging: A New Method For Adaptive Control." In: *Advances in Navigation Guidance and Control Technology Workshop*. URL: <https://www.researchgate.net/publication/2492500%7B%5C%7DPseudo-Control%7B%5C%7DHedging%7B%5C%7DA%7B%5C%7DNew%7B%5C%7DMethod%7B%5C%7DFor%7B%5C%7DAdaptive%7B%5C%7DControl>.
- Jonschkowski, Rico, Roland Hafner, Jonathan Scholz, and Martin Riedmiller (2017). "Pves: Position-velocity encoders for unsupervised learning of structured state representations." In: *arXiv preprint arXiv:1705.09805* (cit. on p. 8).
- Kaelbling, L P (1993). "Learning to achieve goals." In: *International Joint Conference on Artificial Intelligence (IJCAI)*. Vol. vol.2, pp. 1094–8 (cit. on pp. 7, 8, 12, 22, 23, 30, 44, 45, 142).
- Kaelbling, Leslie Pack (1993). "Learning to Achieve Goals." In: *IJCAI*.
- Kaelbling, Leslie Pack and Tomas Lozano-Perez (2011). "Hierarchical task and motion planning in the now." In: *IEEE International Conference on Robotics and Automation (ICRA)*, pp. 1470–1477. ISSN: 10504729. DOI: [10.1109/ICRA.2011.5980391](https://doi.org/10.1109/ICRA.2011.5980391). URL: <http://people.csail.mit.edu/lpk/papers/hpn2.pdf%20http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=5980391>.
- Kahn, Gregory, Adam Villaflor, Bosen Ding, Pieter Abbeel, and Sergey Levine (2018). "Self-supervised Deep Reinforcement Learning with Generalized Computation Graphs for Robot Navigation." In: *IEEE International Conference on Robotics and Automation (ICRA)*. arXiv: [1709.10489v3](https://arxiv.org/pdf/1709.10489v3.pdf). URL: <https://arxiv.org/pdf/1709.10489.pdf>.
- Kalakrishnan, Mrinal, Jonas Buchli, Peter Pastor, and Stefan Schaal (2009). "Learning Locomotion over Rough Terrain using Terrain Templates." In: *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. URL: <https://pdfs.semanticscholar.org/02db/2c0100ffb02592e8738d0ffcf454224f4b1b.pdf>.
- Kalakrishnan, Mrinal, Sachin Chitta, Evangelos Theodorou, Peter Pastor, and Stefan Schaal (2011a). "STOMP: Stochastic trajectory optimization for motion planning." In: *Proceedings - IEEE International Conference on Robotics and Automation*, pp. 4569–4574. ISSN: 10504729. DOI: [10.1109/ICRA.2011.5980280](https://doi.org/10.1109/ICRA.2011.5980280).
- Kalakrishnan, Mrinal, Ludovic Righetti, Peter Pastor, and Stefan Schaal (2011b). "Learning force control policies for compliant manipulation." In: *2011 IEEE/RSJ International Conference on Intelligent Robots and Systems*. IEEE, pp. 4639–4644.
- Kalashnikov, Dmitry, Alex Irpan, Peter Pastor, Julian Ibarz, Alexander Herzog, Eric Jang, Deirdre Quillen, Ethan Holly, Mrinal Kalakrishnan, Vincent Vanhoucke, et al. (2018a).

- “Scalable deep reinforcement learning for vision-based robotic manipulation.” In: *Conference on Robot Learning*. PMLR, pp. 651–673.
- Kalashnikov, Dmitry, Alex Irpan, Peter Pastor, Julian Ibarz, Alexander Herzog, Eric Jang, Deirdre Quillen, Ethan Holly, Mrinal Kalakrishnan, Vincent Vanhoucke, and Sergey Levine (2018b). “QT-Opt: Scalable Deep Reinforcement Learning for Vision-Based Robotic Manipulation.” In: *Conference on Robot Learning (CoRL)*. arXiv: [1806.10293v3](https://arxiv.org/abs/1806.10293). URL: <https://goo.gl/ykQn6g..>
- Kalashnikov, Dmitry, Alex Irpan, Peter Pastor, Julian Ibarz, Alexander Herzog, Eric Jang, Deirdre Quillen, Ethan Holly, Mrinal Kalakrishnan, Vincent Vanhoucke, and Sergey Levine (29–31 Oct 2018c). “Scalable Deep Reinforcement Learning for Vision-Based Robotic Manipulation.” In: *Proceedings of The 2nd Conference on Robot Learning*. Vol. 87. Proceedings of Machine Learning Research. PMLR, pp. 651–673.
- Kalashnikov, Dmitry, Jacob Varley, Yevgen Chebotar, Benjamin Swanson, Rico Jonschkowski, Chelsea Finn, Sergey Levine, and Karol Hausman (2021). “MT-Opt: Continuous Multi-Task Robotic Reinforcement Learning at Scale.” In.
- Kannengiesser, Udo, Richard Heininger, Lubomir Billy, Pavol Terpak, Matthias Neubauer, Chris Stary, Dennis Majoe, Alexandra Totter, and David Bonaldi (Jan. 2017). *Lot-Size One Production*, pp. 69–111. ISBN: 978-3-319-48465-5. DOI: [10.1007/978-3-319-48466-2_4](https://doi.org/10.1007/978-3-319-48466-2_4).
- Karaman, Sertac and Emilio Frazzoli (June 2011). “Sampling-based algorithms for optimal motion planning:” in: *IJRR* 30 (7), pp. 846–894. ISSN: 02783649. DOI: [10.1177/0278364911406761](https://doi.org/10.1177/0278364911406761).
- Kavraki, Lydia, Petr Svestka, J-C Latombe, and Mark Overmars (1996). “Probabilistic roadmaps for path planning in high-dimensional configuration spaces.” In: *IEEE Transactions on Robotics and Automation* 12.4, pp. 566–580.
- Khazatsky, Alexander, Ashvin Nair, Dan Jing, and Sergey Levine (2021). “What Can I Do Here? Learning New Skills by Imagining Visual Affordances.” In: *ICRA*, pp. 14291–14297.
- Khetarpal, Khimya, Zafarali Ahmed, Gheorghe Comanici, David Abel, and Doina Precup (2020). “What can I do here? A Theory of Affordances in Reinforcement Learning.” In: *International Conference on Machine Learning (ICML)*. arXiv: [2006.15085v1](https://arxiv.org/abs/2006.15085).
- Kim, Beomjoon, Amir-Massoud Farahmand, Joelle Pineau, and Doina Precup (2013). “Learning from Limited Demonstrations.” In: *Advances in Neural Information Processing Systems (NeurIPS)*. URL: <https://papers.nips.cc/paper/4918-learning-from-limited-demonstrations.pdf>.

- Kingma, Diederik and Jimmy Ba (2015). "Adam: A method for stochastic optimization." In: *International Conference on Learning Representations (ICLR)*.
- Kingma, Diederik P and Max Welling (2014). "Auto-Encoding Variational Bayes." In: *International Conference on Learning Representations (ICLR)*. URL: <https://arxiv.org/pdf/1312.6114.pdf> (cit. on pp. 6, 9, 45, 143, 146, 149, 153, 163, 168).
- Kober, Jens, J Andrew Bagnell, and Jan Peters (2013). "Reinforcement learning in robotics: A survey." In: *The International Journal of Robotics Research* 32.11, pp. 1238–1274.
- Kober, Jens and J. Peter (2008). "Policy search for motor primitives in robotics." In: *Advances in Neural Information Processing Systems (NeurIPS)*. Vol. 97, pp. 83–117. ISBN: 1099401052236. DOI: [10.1007/978-3-319-03194-1_4](https://doi.org/10.1007/978-3-319-03194-1_4). URL: <http://papers.nips.cc/paper/3545-policy-search-for-motor-primitives-in-robotics.pdf>.
- Koenig, Sven and Maxim Likhachev (2002). "D* Lite." In: *AAAI*.
- Koenker, Roger and Kevin F Hallock (2001). "Quantile regression." In: *Journal of economic perspectives* 15.4, pp. 143–156.
- Kohl, Nate and Peter Stone (2004). "Machine Learning for Fast Quadrupedal Locomotion." In: *AAAI Conference on Artificial Intelligence*, pp. 611–616. URL: <http://www.cs.utexas.edu/%7B%5C%7D7Bnate,pstone%7B%5C%7D7D>.
- Konda, Vijay R and John N Tsitsiklis (2000). "Actor-Critic Algorithms." In: *Advances in Neural Information Processing Systems (NeurIPS)*.
- Kormushev, Petar, Sylvain Calinon, and Darwin G. Caldwell (2010). "Robot motor skill coordination with EM-based Reinforcement Learning." In: *2010 IEEE/RSJ International Conference on Intelligent Robots and Systems, October 18-22, 2010, Taipei, Taiwan*. IEEE, pp. 3232–3237. DOI: [10.1109/IROS.2010.5649089](https://doi.org/10.1109/IROS.2010.5649089). URL: <https://doi.org/10.1109/IROS.2010.5649089>.
- Kostrikov, Ilya (Oct. 2021). *JAXRL: Implementations of Reinforcement Learning algorithms in JAX*. DOI: [10.5281/zenodo.5535154](https://doi.org/10.5281/zenodo.5535154). URL: <https://github.com/ikostrikov/jaxrl>.
- Kostrikov, Ilya, Rob Fergus, Jonathan Tompson, and Ofir Nachum (2021a). "Offline reinforcement learning with fisher divergence critic regularization." In: *International Conference on Machine Learning*. PMLR, pp. 5774–5783.
- Kostrikov, Ilya, Ashvin Nair, and Sergey Levine (2021b). "Offline reinforcement learning with implicit q-learning." In: *arXiv preprint arXiv:2110.06169*.
- Krainin, Michael, Brian Curless, and Dieter Fox (2011). "Autonomous generation of complete 3D object models using next best view manipulation planning." In: *2011 IEEE International Conference on Robotics and Automation*. IEEE, pp. 5031–5037.

- Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E Hinton (2012). "Imagenet classification with deep convolutional neural networks." In: *Advances in Neural Information Processing Systems (NeurIPS)*, pp. 1097–1105 (cit. on pp. 2, 102, 154, 155).
- Kroemer, Oliver, Renaud Detry, Justus Piater, and Jan Peters (2010). "Combining Active Learning and Reactive Control for Robot Grasping." In: *Robotics and Autonomous systems* 58.9, pp. 1105–1116. URL: <https://orbi.uliege.be//bitstream/2268/60643/1/Kroemer-2010-RAS-author-postprint.pdf>.
- Kumar, Aviral, Justin Fu, George Tucker, and Sergey Levine (June 2019a). "Stabilizing Off-Policy Q-Learning via Bootstrapping Error Reduction." In: *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv: 1906.00949. URL: <http://arxiv.org/abs/1906.00949>.
- Kumar, Aviral, Justin Fu, George Tucker, and Sergey Levine (2019b). "Stabilizing off-policy q-learning via bootstrapping error reduction." In: *arXiv preprint arXiv:1906.00949*.
- Kumar, Aviral, Aurick Zhou, George Tucker, and Sergey Levine (2020). "Conservative q-learning for offline reinforcement learning." In: *arXiv preprint arXiv:2006.04779*.
- Kuznetsov, Arsenii, Pavel Shvechikov, Alexander Grishin, and Dmitry Vetrov (2020). "Controlling overestimation bias with truncated mixture of continuous distributional quantile critics." In: *International Conference on Machine Learning*. PMLR, pp. 5556–5566.
- Lange, Sascha, Thomas Gabel, and Martin Riedmiller (2012a). "Batch reinforcement learning." In: *Reinforcement learning*. Springer, pp. 45–73.
- Lange, Sascha, Thomas Gabel, and Martin A. Riedmiller (2012b). "Batch Reinforcement Learning." In: *Reinforcement Learning*. Ed. by Marco Wiering and Martijn van Otterlo. Vol. 12. Adaptation, Learning, and Optimization. Springer, pp. 45–73. DOI: [10.1007/978-3-642-27645-3_2](https://doi.org/10.1007/978-3-642-27645-3_2). URL: https://doi.org/10.1007/978-3-642-27645-3%5C_2.
- Lange, Sascha, Martin Riedmiller, Arne Voigtlander, and Arne Voigtländer (2012c). "Autonomous reinforcement learning on raw visual input data in a real world application." In: *International Joint Conference on Neural Networks (IJCNN)*. June. IEEE, pp. 1–8. ISBN: 9781467314909. DOI: [10.1109/IJCNN.2012.6252823](https://doi.org/10.1109/IJCNN.2012.6252823) (cit. on pp. 7, 8).
- Lange, Sascha and Martin A Riedmiller (2010). "Deep learning of visual control policies." In: *European Symposium on Artificial Neural Networks (ESANN)*. Citeseer. URL: <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.226.6898%7B%5C&%7Drep=rep1%7B%5C&%7Dtype=pdf> (cit. on p. 15).

- Lee, Alex, Sergey Levine, and Pieter Abbeel (2017). "Learning Visual Servoing with Deep Features and Fitted Q-Iteration." In: *International Conference on Learning Representations (ICLR)*. URL: <https://arxiv.org/pdf/1703.11000.pdf> (cit. on p. 8).
- Lee, Alex X, Richard Zhang, Frederik Ebert, Pieter Abbeel, Chelsea Finn, and Sergey Levine (n.d.). "Stochastic Adversarial Video Prediction." In: arXiv: 1804 . 01523v1. URL: https://alexlee-gk.github.io/video%7B%5C_%7Dprediction.
- Lee, Seunghyun, Younggyo Seo, Kimin Lee, Pieter Abbeel, and Jinwoo Shin (2021). "Offline-to-Online Reinforcement Learning via Balanced Replay and Pessimistic Q-Ensemble." In.
- Levine, Sergey, Chelsea Finn, Trevor Darrell, and Pieter Abbeel (2016a). "End-to-End Training of Deep Visuomotor Policies." In: *Journal of Machine Learning Research (JMLR)* 17.1, pp. 1334–1373. ISSN: 15337928. DOI: [10 . 1007 / s13398 - 014 - 0173 - 7 . 2](https://doi.org/10.1007/s13398-014-0173-7). arXiv: 1504.00702. URL: <https://arxiv.org/pdf/1504.00702.pdf> (cit. on pp. 7, 15, 44, 85, 90, 141).
- Levine, Sergey, Aviral Kumar, George Tucker, and Justin Fu (n.d.). *Offline Reinforcement Learning: Tutorial, Review, and Perspectives on Open Problems*. Tech. rep. arXiv: 2005 . 01643v1.
- Levine, Sergey, Peter Pastor, Alex Krizhevsky, and Deirdre Quillen (2016b). "Learning Hand-Eye Coordination for Robotic Grasping with Deep Learning and Large-Scale Data Collection." In: *arXiv preprint arXiv:1603.02199*.
- Levine, Sergey, Peter Pastor, Alex Krizhevsky, and Deirdre Quillen (2017). "Learning Hand-Eye Coordination for Robotic Grasping with Deep Learning and Large-Scale Data Collection." In: *International Journal of Robotics Research* (cit. on pp. 7, 44, 85, 155).
- Levy, Andrew, Robert Platt, and Kate Saenko (2017). "Hierarchical Actor-Critic." In: *arXiv preprint arXiv:1712.00948* (cit. on p. 8).
- Li, Rui, Robert Platt, Wenzhen Yuan, Andreas Ten Pas, Nathan Roscup, Mandayam A Srinivasan, and Edward Adelson (2014). "Localization and Manipulation of Small Parts Using GelSight Tactile Sensing." In: *International Conference on Intelligent Robots and Systems (IROS)*. URL: <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.448.7123%7B%5C&%7Drep=rep1%7B%5C&%7Dtype=pdf>.
- Lillicrap, Timothy P, Jonathan J Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra (2016). "Continuous control with deep reinforcement learning." In: *International Conference on Learning Representations (ICLR)*. ISBN: 0-7803-3213-X. DOI: [10 . 1613 / jair . 301](https://doi.org/10.1613/jair.301). arXiv: 9605103 [cs]. URL: <https://arxiv.org/pdf/1509.02971.pdf> (cit. on pp. 7, 9, 23, 45, 58, 60, 79, 94, 105).

- Lin, Xingyu, Harjatin Singh Baweja, and David Held (2019). "Reinforcement Learning without Ground-Truth State." In: *arXiv preprint arXiv:1905.07866*. arXiv: [1905.07866v1](#). URL: <https://arxiv.org/pdf/1905.07866.pdf>.
- Loftin, Robert, James Macglashan, Bei Peng, Matthew E Taylor, Michael L Littman, Jeff Huang, and David L Roberts (2014). "A Strategy-Aware Technique for Learning Behaviors from Discrete Human Feedback." In: *AAAI Conference on Artificial Intelligence*, pp. 937–943. ISBN: 9781577356783.
- Lopes, Manuel, Tobias Lang, Marc Toussaint, and Pierre-Yves Oudeyer (2012). "Exploration in model-based reinforcement learning by empirically estimating learning progress." In: *Advances in Neural Information Processing Systems*, pp. 206–214.
- Lu, Yao, Karol Hausman, Yevgen Chebotar, Mengyuan Yan, Eric Jang, Alexander Herzog, Ted Xiao, Alex Irpan, Mohi Khansari, Dmitry Kalashnikov, and Sergey Levine (2021). "AW-Opt: Learning Robotic Skills with Imitation and Reinforcement at Scale." In.
- Luo, J, E Solowjow, C Wen, J Aparicio Ojea, A Agogino, A Tamar, and Abbeel P (2019). "Reinforcement Learning on Variable Impedance Controller for High-Precision Robotic Assembly." In: *ICRA*.
- Lynch, Corey, Mohi Khansari, Ted Xiao, Vikash Kumar, Jonathan Tompson, Sergey Levine, and Pierre Sermanet (Mar. 2019). "Learning Latent Plans from Play." In: *Conference on Robot Learning (CoRL)*. arXiv: [1903.01973](#). URL: <http://arxiv.org/abs/1903.01973>.
- Martinez, David, Guillem Alenya, Pablo Jimenez, Carme Torras, Jürgen Rossmann, Nils Wantia, Eren Erdal Aksoy, Simon Haller, and Justus Piater (2014). "Active learning of manipulation sequences." In: *2014 IEEE International Conference on Robotics and Automation (ICRA)*. IEEE, pp. 5671–5678.
- Meng, Linghui, Muning Wen, Yaodong Yang, Chenyang Le, Xiyun Li, Weinan Zhang, Ying Wen, Haifeng Zhang, Jun Wang, and Bo Xu (2021). "Offline Pre-trained Multi-Agent Decision Transformer: One Big Sequence Model Tackles All SMAC Tasks." In.
- Min, Huaqing, An Yi, Ronghua Luo, Jinhui Zhu, and Sheng Bi (2016). "Affordance Research in Developmental Robotics: A Survey." In: *IEEE Transactions on Cognitive and Developmental Systems* 8.4. doi: [10.1109/TCDS.2016.2614992](#). URL: http://www.ieee.org/publications%7B%5C_%7Dstandards/publications/rights/index.html.
- Mnih, Volodymyr, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin Riedmiller (2013). "Playing Atari with Deep Reinforcement Learning." In: *NIPS Workshop on Deep Learning*, pp. 1–9. ISBN: 1476-4687 (Electronic) 0028-0836 (Linking). doi: [10.1038/nature14236](#). arXiv: [1312.5602](#). URL: <https://doi.org/10.1038/nature14236>.

<https://arxiv.org/pdf/1312.5602.pdf> <https://arxiv.org/abs/1312.5602> <http://www.cs.toronto.edu/%7B~%7Dmnih/docs/dqn.pdf>.

Mnih, Volodymyr, Koray Kavukcuoglu, David Silver, Andrei A Rusu, Joel Veness, Marc G Bellemare, Alex Graves, Martin Riedmiller, Andreas K Fidjeland, Georg Ostrovski, et al. (2015). "Human-level control through deep reinforcement learning." In: *Nature* 518.7540, pp. 529–533 (cit. on pp. 2, 58).

Mnih, Volodymyr, Adrià Puigdomènech Badia, Mehdi Mirza, Alex Graves, Tim Harley, Timothy P Lillicrap, David Silver, Koray Kavukcuoglu, Korayk@google Com, and Google Deepmind (2016). "Asynchronous Methods for Deep Reinforcement Learning." In: *International Conference on Machine Learning (ICML)*. URL: <https://arxiv.org/pdf/1602.01783.pdf> (cit. on pp. 9, 58, 105, 113).

Mohamed, Shakir and Danilo Jimenez Rezende (2015). "Variational information maximisation for intrinsically motivated reinforcement learning." In: *Advances in neural information processing systems*, pp. 2125–2133.

Mülling, Katharina, Jens Kober, Oliver Kroemer, and Jan Peters (2013). "Learning to select and generalize striking movements in robot table tennis." In: *Int. J. Robotics Res.* 32.3, pp. 263–279. DOI: [10.1177/0278364912472380](https://doi.org/10.1177/0278364912472380). URL: <https://doi.org/10.1177/0278364912472380>.

Nachum, Ofir, Yinlam Chow, Bo Dai, and Lihong Li (June 2019). "DualDICE: Behavior-Agnostic Estimation of Discounted Stationary Distribution Corrections." In: *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv: [1906 . 04733](https://arxiv.org/abs/1906.04733). URL: <http://arxiv.org/abs/1906.04733>.

Nachum, Ofir, Shane Shixiang Gu, Honglak Lee, and Sergey Levine (2018). "Data-Efficient Hierarchical Reinforcement Learning." In: *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv: [arXiv:1805.08296v2](https://arxiv.org/abs/1805.08296v2). URL: <https://sites.google.com/view/efficient-hrl>.

Nair, Ashvin, Shikhar Bahl, Alexander Khazatsky, Vitchyr Pong, Glen Berseth, and Sergey Levine (Oct. 2019a). "Contextual Imagined Goals for Self-Supervised Robotic Learning." In: *Conference on Robot Learning (CoRL)*. arXiv: [1910 . 11670](https://arxiv.org/abs/1910.11670). URL: <http://arxiv.org/abs/1910.11670>.

Nair, Ashvin, Shikhar Bahl, Alexander Khazatsky, Vitchyr H. Pong, Glen Berseth, and Sergey Levine (2019b). "Contextual Imagined Goals for Self-Supervised Robotic Learning." In: *CoRL*.

Nair, Ashvin, Dian Chen, Pulkit Agrawal, Phillip Isola, Pieter Abbeel, Jitendra Malik, Sergey Levine, Dian Chen, Phillip Isola, Pieter Abbeel, Jitendra Malik, and Sergey Levine (2017). "Combining Self-Supervised Learning and Imitation for Vision-Based

- Rope Manipulation." In: *IEEE International Conference on Robotics and Automation (ICRA)*. ISBN: 9781509046331. DOI: [10.1109/ICRA.2017.7989247](https://doi.org/10.1109/ICRA.2017.7989247). arXiv: [1703.02018](https://arxiv.org/abs/1703.02018).
- Nair, Ashvin, Murtaza Dalal, Abhishek Gupta, and Sergey Levine (June 2020). "Accelerating Online Reinforcement Learning with Offline Datasets." In: arXiv: [2006.09359](https://arxiv.org/abs/2006.09359). URL: <http://arxiv.org/abs/2006.09359>.
- Nair, Ashvin, Bob McGrew, Marcin Andrychowicz, Wojciech Zaremba, and Pieter Abbeel (2018a). "Overcoming Exploration in Reinforcement Learning with Demonstrations." In: *IEEE International Conference on Robotics and Automation (ICRA)*. URL: <https://arxiv.org/pdf/1709.10089.pdf>.
- Nair, Ashvin, Vitchyr Pong, Murtaza Dalal, Shikhar Bahl, Steven Lin, and Sergey Levine (2018b). "Visual Reinforcement Learning with Imagined Goals." In: *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv: [arXiv:1807.04742v1](https://arxiv.org/abs/1807.04742v1). URL: <https://sites.google.com/site/>.
- Nair, Suraj and Chelsea Finn (2019). "Hierarchical foresight: Self-supervised learning of long-horizon tasks via visual subgoal generation." In: *arXiv preprint arXiv:1909.05829*.
- Nakanishi, Jun, Jun Morimoto, Gen Endo, Gordon Cheng, Stefan Schaal, and Mitsuo Kawato (2004). "Learning from demonstration and adaptation of biped locomotion." In: *Robotics and Autonomous Systems*. Vol. 47. 2-3, pp. 79–91. ISBN: 0921-8890. DOI: [10.1016/j.robot.2004.03.003](https://doi.org/10.1016/j.robot.2004.03.003). URL: <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.77.534%7B%5C%7Drep=rep1%7B%5C%7Dtype=pdf>.
- Nasiriany, Soroush, Vitchyr H. Pong, Steven Lin, and Sergey Levine (2019). "Planning with Goal-Conditioned Policies." In: *NeurIPS*.
- Neumann, Gerhard and Jan Peters (2008). "Fitted Q-iteration by Advantage Weighted Regression." In: *Advances in Neural Information Processing Systems (NeurIPS)*.
- Ng, Andrew and Stuart Russell (2000). "Algorithms for Inverse Reinforcement Learning." In: *International Conference on Machine Learning (ICML)*. URL: <http://ai.stanford.edu/%7B~%7Dang/papers/icml00-irl.pdf>.
- Ng, Andrew Y., Daishi Harada, and Stuart Russell (1999). "Policy invariance under reward transformations: Theory and application to reward shaping." In: *International Conference on Machine Learning (ICML)*. URL: <https://www-cs.stanford.edu/people/ang/papers/shaping-icml99.pdf>.
- Oh, Junhyuk, Xiaoxiao Guo, Honglak Lee, Richard Lewis, and Satinder Singh (2015). "Action-Conditional Video Prediction using Deep Networks in Atari Games." In: *Advances in Neural Information Processing Systems (NeurIPS)*. URL: <https://arxiv.org/pdf/1507.08750v1.pdf> (cit. on p. 7).

- Oord, Aaron van den, Nal Kalchbrenner, Oriol Vinyals, Lasse Espeholt, Alex Graves, and Koray Kavukcuoglu (June 2016). "Conditional Image Generation with PixelCNN Decoders." In: *Advances in Neural Information Processing Systems*. Neural information processing systems foundation, pp. 4797–4805. arXiv: [1606 . 05328](https://arxiv.org/abs/1606.05328). URL: <http://arxiv.org/abs/1606.05328>.
- Oord, Aaron van den, Oriol Vinyals, and Koray Kavukcuoglu (Nov. 2017). "Neural Discrete Representation Learning." In: *Advances in Neural Information Processing Systems*. Vol. 2017-Decem. Neural information processing systems foundation, pp. 6307–6316. arXiv: [1711.00937](https://arxiv.org/abs/1711.00937). URL: <http://arxiv.org/abs/1711.00937>.
- Oord, Aäron van den, Oriol Vinyals, and Koray Kavukcuoglu (2017). "Neural Discrete Representation Learning." In: *NeurIPS*.
- Ostrovski, Georg, Marc G Bellemare, Aäron Oord, and Rémi Munos (2017). "Count-Based Exploration with Neural Density Models." In: *International Conference on Machine Learning (ICML)*, pp. 2721–2730.
- Pathak, Deepak, Pulkit Agrawal, Alexei A. Efros, and Trevor Darrell (2017). "Curiosity-Driven Exploration by Self-Supervised Prediction." In: *International Conference on Machine Learning (ICML)*. IEEE, pp. 488–489. arXiv: [1705 . 05363](https://arxiv.org/abs/1705.05363) (cit. on pp. [13](#), [30](#), [139](#), [142](#)).
- Pathak, Deepak, Parsa Mahmoudieh, Guanghao Luo, Pulkit Agrawal, Dian Chen, Yide Shentu, Evan Shelhamer, Jitendra Malik, Alexei A Efros, and Trevor Darrell (2018). "Zero-Shot Visual Imitation." In: *International Conference on Learning Representations (ICLR)*. URL: <https://arxiv.org/pdf/1804.08606.pdf> (cit. on pp. [7](#), [90](#)).
- Peng, Xue Bin, Aviral Kumar, Grace Zhang, and Sergey Levine (Sept. 2019a). "Advantage-Weighted Regression: Simple and Scalable Off-Policy Reinforcement Learning." In: arXiv: [1910.00177](https://arxiv.org/abs/1910.00177). URL: <http://arxiv.org/abs/1910.00177>.
- Peng, Xue Bin, Aviral Kumar, Grace Zhang, and Sergey Levine (2019b). "Advantage-weighted regression: Simple and scalable off-policy reinforcement learning." In: *arXiv preprint arXiv:1910.00177*.
- Péré, Alexandre, Sébastien Forestier, Olivier Sigaud, and Pierre-Yves Oudeyer (2018a). "Unsupervised Learning of Goal Spaces for Intrinsically Motivated Goal Exploration." In: *International Conference on Learning Representations (ICLR)*. URL: <https://arxiv.org/pdf/1803.00781.pdf> (cit. on p. [8](#)).
- Péré, Alexandre, Sébastien Forestier, Olivier Sigaud, and Pierre-Yves Oudeyer (2018b). "Unsupervised Learning of Goal Spaces for Intrinsically Motivated Goal Exploration." In: *International Conference on Learning Representations (ICLR)*. URL: <https://arxiv.org/pdf/1803.00781.pdf>.

- Péré, Alexandre, Sébastien Forestier, Olivier Sigaud, and Pierre-Yves Oudeyer (2018c). "Unsupervised Learning of Goal Spaces for Intrinsically Motivated Goal Exploration." In: *International Conference on Learning Representations (ICLR)*. URL: <https://arxiv.org/pdf/1803.00781.pdf>.
- Pertsch, Karl, Oleh Rybkin, Frederik Ebert, Chelsea Finn, Dinesh Jayaraman, and Sergey Levine (2020). "Long-Horizon Visual Planning with Goal-Conditioned Hierarchical Predictors." In: *ArXiv* abs/2006.13205.
- Peters, Jan, Katharina Mülling, and Yasemin Altün (2010). "Relative Entropy Policy Search." In: *AAAI Conference on Artificial Intelligence*, pp. 1607–1612. URL: https://pdfs.semanticscholar.org/ff47/526838ce85d77a50197a0c5f6ee5095156aa.pdf%20http://www-clmc.usc.edu/publications/P/Peters%7B%5C_%7DPOTTNCOAIPGAT%7B%5C_%7D2010.pdf.
- Peters, Jan and Stefan Schaal (2007a). "Reinforcement Learning by Reward-weighted Regression for Operational Space Control." In: *International Conference on Machine Learning (ICML)*. URL: <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.79.6266%7B%5C&%7Drep=rep1%7B%5C&%7Dtype=pdf>.
- Peters, Jan and Stefan Schaal (2007b). "Reinforcement learning by reward-weighted regression for operational space control." In: *Proceedings of the 24th international conference on Machine learning*, pp. 745–750.
- Peters, Jan and Stefan Schaal (2008a). "Natural Actor-Critic." In: *Neurocomputing* 71.7-9, pp. 1180–1190. DOI: [10.1016/j.neucom.2007.11.026](https://doi.org/10.1016/j.neucom.2007.11.026). URL: <https://doi.org/10.1016/j.neucom.2007.11.026>.
- Peters, Jan and Stefan Schaal (2008b). "Reinforcement learning of motor skills with policy gradients." In: *Neural Networks* 21.4. Robotics and Neuroscience, pp. 682–697. ISSN: 0893-6080. DOI: <https://doi.org/10.1016/j.neunet.2008.02.003>. URL: <http://www.sciencedirect.com/science/article/pii/S0893608008000701>.
- Peters, Jan and Stefan Schaal (2008c). "Reinforcement learning of motor skills with policy gradients." In: *Neural Networks* 21.4, pp. 682–697. ISSN: 08936080. DOI: [10.1016/j.neunet.2008.02.003](https://doi.org/10.1016/j.neunet.2008.02.003). arXiv: [arXiv:1411.3159v1](https://arxiv.org/abs/1411.3159v1). URL: <https://pdfs.semanticscholar.org/eb5b/459c8a3e56064158fb3514eeab763486e437.pdf>.
- Pinto, Lerrel, Marcin Andrychowicz, Peter Welinder, Wojciech Zaremba, and Pieter Abbeel (2018). "Asymmetric Actor Critic for Image-Based Robot Learning." In: *Robotics: Science and Systems (RSS)* (cit. on p. 7).
- Pinto, Lerrel, James Davidson, Rahul Sukthankar, and Abhinav Gupta (2017). "Robust Adversarial Reinforcement Learning." In: *International Conference on Machine Learning (ICML)*. URL: <https://arxiv.org/pdf/1703.02702.pdf>.

- Pinto, Lerrel and Abhinav Gupta (2016). "Supersizing Self-supervision: Learning to Grasp from 50K Tries and 700 Robot Hours." In: *IEEE International Conference on Robotics and Automation (ICRA)* (cit. on pp. 7, 44, 58, 90).
- Pomerleau, Dean A (1989). "Alvinn: An autonomous land vehicle in a neural network." In: *Advances in Neural Information Processing Systems (NeurIPS)*, pp. 305–313. ISBN: 1-558-60015-9. URL: <http://repository.cmu.edu/compsci>.
- Pong, Vitchyr, Shixiang Gu, Murtaza Dalal, and Sergey Levine (2018). "Temporal Difference Models: Model-Free Deep RL For Model-Based Control." In: *International Conference on Learning Representations (ICLR)*. URL: <https://arxiv.org/pdf/1802.09081.pdf> (cit. on pp. 8, 23, 30, 79, 95).
- Pong, Vitchyr H, Murtaza Dalal, Steven Lin, Ashvin Nair, Shikhar Bahl, and Sergey Levine (2020a). "Skew-Fit: State-Covering Self-Supervised Reinforcement Learning." In: *International Conference on Machine Learning (ICML)*. arXiv: [1903 . 03698v2](https://arxiv.org/pdf/1903.03698v2.pdf). URL: <https://arxiv.org/pdf/1903.03698.pdf>.
- Pong, Vitchyr H., Murtaza Dalal, Steven Lin, Ashvin Nair, Shikhar Bahl, and Sergey Levine (2019). "Skew-Fit: State-Covering Self-Supervised Reinforcement Learning." In: *CoRR* abs/1903.03698. arXiv: [1903 . 03698](https://arxiv.org/abs/1903.03698). URL: [http://arxiv.org/abs/1903.03698](https://arxiv.org/abs/1903.03698).
- Pong, Vitchyr H., Murtaza Dalal, Steven Lin, Ashvin Nair, Shikhar Bahl, and Sergey Levine (2020b). "Skew-Fit: State-Covering Self-Supervised Reinforcement Learning." In: *ArXiv* abs/1903.03698.
- Ponomarenko, Nikolay, Lina Jin, Oleg Ieremeiev, Vladimir Lukin, Karen Egiazarian, Jaakko Astola, Benoit Vozel, Kacem Chehdi, Marco Carli, Federica Battisti, et al. (2015). "Image database TID2013: Peculiarities, results and perspectives." In: *Signal Processing: Image Communication* 30, pp. 57–77 (cit. on p. 5).
- Popov, Ivaylo, Nicolas Heess, Timothy Lillicrap, Roland Hafner, Gabriel Barth-maron, Matej Vecerik, Thomas Lampe, Yuval Tassa, Tom Erez, Martin Riedmiller, and Martin Riedmiller Deepmind (2017). "Data-efficient Deep Reinforcement Learning for Dexterous Manipulation." In: *CoRR* abs/1704.0. arXiv: [1704 . 03073](https://arxiv.org/pdf/1704.03073.pdf). URL: <https://arxiv.org/pdf/1704.03073.pdf>.
- Rajeswaran, Aravind, Vikash Kumar, Abhishek Gupta, John Schulman, Emanuel Todorov, and Sergey Levine (2018). "Learning Complex Dexterous Manipulation with Deep Reinforcement Learning and Demonstrations." In: *Robotics: Science and Systems*. URL: <https://arxiv.org/pdf/1709.10087.pdf>.
- Rakelly, Kate, Aurick Zhou, Deirdre Quillen, Chelsea Finn, and Sergey Levine (2019). "Efficient Off-Policy Meta-Reinforcement Learning via Probabilistic Context Variables."

- In: *International Conference on Machine Learning (ICML)*. URL: <https://github.com/katerakelly/oyster..>
- Ramapuram, Jason, Magda Gregorova, and Alexandros Kalousis (May 2017). "Lifelong Generative Modeling." In: *Neurocomputing*. arXiv: 1705.09847. URL: <http://arxiv.org/abs/1705.09847>.
- Rauber, Paulo, Filipe Mutz, and Juergen Jürgen Schmidhuber (2017). "Hindsight policy gradients." In: *CoRR*. Vol. abs/1711.0. URL: <https://arxiv.org/pdf/1711.06006.pdf> (cit. on p. 8).
- Reed, Scott, Kihyuk Sohn, Yuting Zhang, and Honglak Lee (2014). "Learning to disentangle factors of variation with manifold interaction." In: *International Conference on Machine Learning (ICML)*, pp. 1431–1439 (cit. on p. 8).
- Ronneberger, Olaf, Philipp Fischer, and Thomas Brox (2015). "U-Net: Convolutional Networks for Biomedical Image Segmentation." In: *MICCAI*.
- Ross, Stéphane, Geoffrey J Gordon, and J Andrew Bagnell (2011). "A Reduction of Imitation Learning and Structured Prediction to No-Regret Online Learning." In: *International Conference on Artificial Intelligence and Statistics (AISTATS)*. URL: <https://arxiv.org/pdf/1011.0686.pdf%20https://www.cs.cmu.edu/%7B~%7Dsross1/publications/Ross-AIStats11-NoRegret.pdf>.
- Rubin, Donald B (1988). "Using the SIR algorithm to simulate posterior distributions." In: *Bayesian statistics 3*, pp. 395–402.
- Rusu, Andrei A, Matej Vecerik, Thomas Rothörl, Nicolas Heess, Razvan Pascanu, and Raia Hadsell (2017). "Sim-to-real robot learning from pixels with progressive nets." In: *Conference on Robot Learning (CoRL)* (cit. on p. 18).
- Saunders, William, Girish Sastry, Andreas Stuhlmüller, and Owain Evans (2018). "Trial without Error: Towards Safe Reinforcement Learning via Human Intervention." In: *International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*. arXiv: arXiv:1707.05173v1. URL: <https://arxiv.org/pdf/1707.05173.pdf>.
- Savinov, Nikolay, Anton Raichuk, Raphaël Marinier, Damien Vincent, Marc Pollefeys, Timothy Lillicrap, and Sylvain Gelly (2018). "Episodic curiosity through reachability." In: *arXiv preprint arXiv:1810.02274*.
- Schaal, Stefan (1997). "Learning from demonstration." In: *Advances in Neural Information Processing Systems (NeurIPS)*. 9, pp. 1040–1046. ISBN: 1558604863. DOI: 10.1016/j.robot.2004.03.001. URL: <http://www.cc.gatech.edu/faculty/stefan/Schaal.pdf%20http://www.ira.uka.de/users/rogalla/WebOrdnerMaterial/ml-robotlearning.pdf%20http://www.cc.gatech.edu/faculty/stefan.Schaal>.

- Schaul, Tom, Dan Horgan, Karol Gregor, and David Silver (2015a). "Universal Value Function Approximators." In: *ICML*.
- Schaul, Tom, Daniel Horgan, Karol Gregor, and David Silver (2015b). "Universal Value Function Approximators." In: *International Conference on Machine Learning (ICML)*, pp. 1312–1320. ISBN: 9781510810587. URL: <http://proceedings.mlr.press/v37/schaul15.pdf> <http://jmlr.org/proceedings/papers/v37/schaul15.html> (cit. on pp. 7, 23, 30, 44, 45, 61, 142, 143).
- Schenck, Connor and Dieter Fox (2017). "Visual closed-loop control for pouring liquids." In: *2017 IEEE International Conference on Robotics and Automation (ICRA)*. IEEE, pp. 2629–2636.
- Schmidhuber, Jürgen (1992). "Learning factorial codes by predictability minimization." In: *Neural Computation* 4.6, pp. 863–879 (cit. on p. 8).
- Schulman, John, Sergey Levine, Philipp Moritz, Michael I. Jordan, and Pieter Abbeel (2015). "Trust Region Policy Optimization." In: *International Conference on Machine Learning (ICML)*. ISBN: 0375-9687. DOI: [10.1162/1063-14927398](https://doi.org/10.1162/1063-14927398). arXiv: [1502.05477](https://arxiv.org/pdf/1502.05477.pdf). URL: <https://arxiv.org/pdf/1502.05477.pdf> [http://arxiv.org/abs/1502.05477](https://arxiv.org/abs/1502.05477).
- Sermanet, Pierre, Corey Lynch, Yevgen Chebotar, Jasmine Hsu, Eric Jang, Stefan Schaal, and Sergey Levine (2017). "Time-Contrastive Networks: Self-Supervised Learning from Video." In: *arXiv preprint arXiv:1704.06888* (cit. on p. 8).
- Sharma, Archit, Abhishek Gupta, Sergey Levine, Karol Hausman, and Chelsea Finn (2021). "Autonomous Reinforcement Learning via Subgoal Curricula." In.
- Siegel, Noah Y, Jost Tobias Springenberg, Felix Berkenkamp, Abbas Abdolmaleki, Michael Neunert, Thomas Lampe, Roland Hafner, Nicolas Heess, and Martin Riedmiller (2020a). "Keep doing what worked: Behavioral modelling priors for offline reinforcement learning." In: *arXiv preprint arXiv:2002.08396*.
- Siegel, Noah Y., Jost Tobias Springenberg, Felix Berkenkamp, Abbas Abdolmaleki, Michael Neunert, Thomas Lampe, Roland Hafner, Nicolas Heess, and Martin Riedmiller (2020b). *Keep Doing What Worked: Behavioral Modelling Priors for Offline Reinforcement Learning*. arXiv: [2002.08396 \[cs.LG\]](https://arxiv.org/abs/2002.08396).
- Silver, David, Aja Huang, Chris J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis (Jan. 2016a). "Mastering the game of Go with deep neural networks and tree search." In: *Nature* 529.7587, pp. 484–489. ISSN: 0028-0836. DOI: [10.1038/nature16961](https://doi.org/10.1038/nature16961). arXiv: [1610.00633](https://arxiv.org/abs/1610.00633). URL: [http://dx.doi.org/10.1038/nature16961](https://dx.doi.org/10.1038/nature16961)

[nature16961%20http://10.0.4.14/nature16961%20http://www.nature.com/nature/journal/v529/n7587/abs/nature16961.html%7B%5C#%7Dsupplementary-information%20http://airesearch.com/wp-content/uploads/2016/01/deepmind-mastering-go.pdf](http://dx.doi.org/10.1038/nature16961) (cit. on pp. 2, 43, 141).

Silver, David, Aja Huang, Chris J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis (Jan. 2016b). "Mastering the game of Go with deep neural networks and tree search." In: *Nature* 529.7587, pp. 484–489. ISSN: 0028-0836. DOI: [10.1038/nature16961](https://doi.org/10.1038/nature16961). arXiv: [1610.00633](https://arxiv.org/abs/1610.00633). URL: [http://dx.doi.org/10.1038/nature16961%20http://10.0.4.14/nature16961%20http://www.nature.com/nature/journal/v529/n7587/abs/nature16961.html%7B%5C#%7Dsupplementary-information%20http://airesearch.com/wp-content/uploads/2016/01/deepmind-mastering-go.pdf](http://dx.doi.org/10.1038/nature16961).

Silver, Tom, Kelsey Allen, Josh Tenenbaum, and Leslie Kaelbling (2018). "Residual Policy Learning." In: *arXiv preprint arXiv:1812.06298*. arXiv: [1812.06298](https://arxiv.org/abs/1812.06298). URL: <http://arxiv.org/abs/1812.06298>.

Singh, Avi, Larry Yang, Kristian Hartikainen, Chelsea Finn, and Sergey Levine (Apr. 2019). "End-to-End Robotic Reinforcement Learning without Reward Engineering." In: *Robotics: Science and Systems (RSS)*. arXiv: [1904.07854](https://arxiv.org/abs/1904.07854). URL: <http://arxiv.org/abs/1904.07854>.

Singh, Avi, Albert Yu, Jonathan Yang, Jesse Zhang, Aviral Kumar, and Sergey Levine (2020). "COG: Connecting new skills to past experience with offline reinforcement learning." In: *CoRL*.

Smith, Linda and Michael Gasser (2005). "The development of embodied cognition: Six lessons from babies." In: *Artificial life* 11.1-2, pp. 13–29. URL: <https://www.cogsci.msu.edu/DSS/2010-2011/Smith/6lessons.pdf> (cit. on p. 5).

Sohn, Kihyuk, Xinchen Yan, and Honglak Lee (2015). "Learning Structured Output Representation using Deep Conditional Generative Models." In: *Advances in Neural Information Processing Systems (NeurIPS)*. URL: <https://papers.nips.cc/paper/5775-learning-structured-output-representation-using-deep-conditional-generative-models.pdf>.

Srinivas, Aravind, Allan Jabri, Pieter Abbeel, Sergey Levine, and Chelsea Finn (2018). "Universal Planning Networks." In: *International Conference on Machine Learning (ICML)*. arXiv: [arXiv:1804.00645v2](https://arxiv.org/abs/1804.00645v2). URL: <https://sites.google.com>. (cit. on p. 8).

- Srivastava, Siddharth, Eugene Fang, Lorenzo Riano, Rohan Chitnis, Stuart Russell, and Pieter Abbeel (2014). "Combined Task and Motion Planning Through an Extensible Planner-Independent Interface Layer." In: *IEEE International Conference on Robotics and Automation (ICRA)*. URL: <https://people.eecs.berkeley.edu/%7B~%7Drussell/papers/icra14-planrob.pdf>.
- Srouji, Mario, Jian Zhang, and Ruslan Salakhutdinov (2018). "Structured Control Nets for Deep Reinforcement Learning." In: *International Conference on Machine Learning (ICML)*. arXiv: 1802.08311. URL: <http://arxiv.org/abs/1802.08311>.
- Stadie, Bradly C, Sergey Levine, and Pieter Abbeel (2016). "Incentivizing Exploration In Reinforcement Learning With Deep Predictive Models." In: *International Conference on Learning Representations (ICLR)*. URL: <https://arxiv.org/pdf/1507.00814.pdf>.
- Subramanian, Kaushik, Charles L Isbell, and Andrea L Thomaz (2016). "Exploration from Demonstration for Interactive Reinforcement Learning." In: *International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*. URL: www.ifaamas.org.
- Sun, Hao, Zhizhong Li, Xiaotong Liu, Dahua Lin, and Bolei Zhou (2019). "Policy Continuation with Hindsight Inverse Dynamics." In: *NeurIPS*.
- Sutton, Richard S and Andrew G Barto (1998). *Reinforcement Learning: An Introduction*. URL: <http://incompleteideas.net/sutton/book/bookdraft2016sep.pdf%20https://webdocs.cs.ualberta.ca/%7B~%7Dsutton/book/bookdraft2016sep.pdf>.
- Sutton, Richard S and Andrew G Barto (2018). *Reinforcement learning: An introduction*. MIT press.
- Sutton, Richard S, Joseph Modayil, Michael Delp, Thomas Degris, Patrick M Pilarski, Adam White, and Doina Precup (2011). "Horde: A Scalable Real-time Architecture for Learning Knowledge from Unsupervised Sensorimotor Interaction." In: *International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*. Vol. 10, pp. 761–768. URL: <https://www.cs.swarthmore.edu/%7B~%7Dmeeden/DevelopmentalRobotics/horde1.pdf> (cit. on p. 7).
- Sutton, Richard S, Doina Precup, and Satinder Singh (1999). "Between MDPs and semi-MDPs: A framework for temporal abstraction in reinforcement learning." In: *Artificial intelligence* 112.1-2, pp. 181–211.
- Tamar, Aviv, Garrett Thomas, Tianhao Zhang, Sergey Levine, and Pieter Abbeel (2017). "Learning from the hindsight plan Episodic MPC improvement." In: *ICRA*, pp. 336–343.
- Tang, Haoran, Rein Houthooft, Davis Foote, Adam Stooke, Xi Chen, Yan Duan, John Schulman, Filip De Turck, and Pieter Abbeel (2017). "#Exploration: A Study of Count-Based Exploration for Deep Reinforcement Learning." In: *Advances in Neu-*

- ral Information Processing Systems (NeurIPS)*. arXiv: [arXiv:1611.04717v3](https://arxiv.org/pdf/1611.04717v3.pdf). URL: <https://arxiv.org/pdf/1611.04717.pdf>.
- Theodorou, Evangelos A, Jonas Buchli, Jonas@buchli Org, Stefan Schaal, and Sschaal@usc Edu (2010). "A Generalized Path Integral Control Approach to Reinforcement Learning." In: *Journal of Machine Learning Research (JMLR)* 11, pp. 3137–3181. URL: <http://www.jmlr.org/papers/volume11/theodorou10a/theodorou10a.pdf>.
- Thomas, Garrett, Melissa Chien, Aviv Tamar, Juan Aparicio Ojea, and Pieter Abbeel (2018). "Learning Robotic Assembly from CAD." In: *IEEE International Conference on Robotics and Automation (ICRA)*. URL: <https://arxiv.org/pdf/1803.07635.pdf>.
- Thomas, Philip S. and Emma Brunskill (2016). "Data-Efficient Off-Policy Policy Evaluation for Reinforcement Learning." In: *Proceedings of the 33nd International Conference on Machine Learning, ICML 2016, New York City, NY, USA, June 19-24, 2016*. Ed. by Maria-Florina Balcan and Kilian Q. Weinberger. Vol. 48. JMLR Workshop and Conference Proceedings. JMLR.org, pp. 2139–2148. URL: <http://proceedings.mlr.press/v48/thomas16.html>.
- Thomas, Valentin, Jules Pondard, Emmanuel Bengio, Marc Sarfati, Philippe Beaudoin, Marie-Jean Meurs, Joelle Pineau, Doina Precup, Yoshua Bengio, Valentin Thoma, Joelle Pineau, Doina Precup, and Yoshua Bengio (2017). "Independently Controllable Factors." In: *NIPS Workshop*. URL: <https://arxiv.org/pdf/1703.07718.pdf> <https://arxiv.org/pdf/1708.01289.pdf> (cit. on p. 8).
- Todorov, Emanuel, Tom Erez, and Yuval Tassa (2012). "MuJoCo: A physics engine for model-based control." In: *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pp. 5026–5033. ISBN: 9781467317375. DOI: [10 . 1109 / IROS . 2012 . 6386109](https://doi.org/10.1109/IROS.2012.6386109). URL: <https://homes.cs.washington.edu/%7B~%7Dtodorov/papers/TodorovIROS12.pdf> (cit. on pp. 14, 50, 64, 80, 116).
- Torrey, Lisa and Matthew E Taylor (2013). "Teaching on a Budget: Agents Advising Agents in Reinforcement Learning." In: *International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*. URL: www.ifaamas.org.
- Van Hasselt, Hado, Arthur Guez, and David Silver (2016). "Deep Reinforcement Learning with Double Q-learning." In: *Association for the Advancement of Artificial Intelligence (AAAI)*. ISBN: 1509.06461v3. arXiv: [1509.06461v3](https://arxiv.org/pdf/1509.06461v3.pdf). URL: www.aaai.org.
- Večerík, Matej, Todd Hester, Jonathan Scholz, Fumin Wang, Olivier Pietquin, Bilal Piot, Nicolas Heess, Thomas Rothörl, Thomas Lampe, and Martin Riedmiller (2017). "Leveraging Demonstrations for Deep Reinforcement Learning on Robotics Problems with Sparse Rewards." In: *CoRR* abs/1707.0. URL: <https://arxiv.org/pdf/1707.08817.pdf>.

- Veeriah, Vivek, Junhyuk Oh, and Satinder Singh (2018). "Many-goals reinforcement learning." In: *arXiv preprint arXiv:1806.09605*.
- Villaflor, Adam, John Dolan, and Jeff Schneider (2020). "Fine-tuning Offline Reinforcement Learning With Model-Based Policy Optimization." In.
- Wang, Qing, Jiechao Xiong, Lei Han, Peng Sun, Han Liu, and Tong Zhang (2018a). "Exponentially Weighted Imitation Learning for Batched Historical Data." In: *Neural Information Processing Systems (NeurIPS)*.
- Wang, Qing, Jiechao Xiong, Lei Han, Peng Sun, Han Liu, and Tong Zhang (2018b). "Exponentially Weighted Imitation Learning for Batched Historical Data." In: *NeurIPS*, pp. 6291–6300.
- Wang, Ziyu, Alexander Novikov, Konrad Zolna, Jost Tobias Springenberg, Scott Reed, Bobak Shahriari, Noah Siegel, Josh Merel, Caglar Gulcehre, Nicolas Heess, et al. (2020a). "Critic regularized regression." In: *arXiv preprint arXiv:2006.15134*.
- Wang, Ziyu, Alexander Novikov, Konrad Zolna, Jost Tobias Springenberg, Scott Reed, Bobak Shahriari, Noah Siegel, Josh Merel, Caglar Gulcehre, Nicolas Heess, and Nando De Freitas (2020b). "Critic Regularized Regression." In: arXiv: [2006.15134v1](#).
- Warde-Farley, David, Tom Van De Wiele, Tejas Kulkarni, Catalin Ionescu, Steven Hansen, and Mnih Volodymyr (2019). "Unsupervised Control Through Non-Parametric Discriminative Rewards." In: *International Conference on Learning Representations (ICLR)*. ISBN: [1811.11359v1](#). arXiv: [1811.11359v1](#). URL: <https://arxiv.org/pdf/1811.11359.pdf>.
- Warde-Farley, David, Tom Van de Wiele, Tejas Kulkarni, Catalin Ionescu, Steven Hansen, and Volodymyr Mnih (2018). "Unsupervised Control Through Non-Parametric Discriminative Rewards." In: *CoRR* abs/[1811.11359](#).
- Watter, Manuel, Jost Tobias Springenberg, Joschka Boedecker, and Martin Riedmiller (2015). "Embed to Control: A Locally Linear Latent Dynamics Model for Control from Raw Images." In: *Advances in Neural Information Processing Systems (NeurIPS)*, pp. 2728–2736. arXiv: [1506.07365](#). URL: <https://arxiv.org/pdf/1506.07365.pdf%20http://arxiv.org/abs/1506.07365> (cit. on pp. [7](#), [8](#), [44](#)).
- Wawrzynski, Paweł (2009). "Real-time reinforcement learning by sequential Actor-Critics and experience replay." In: *Neural Networks* 22.10, pp. 1484–1497. DOI: [10.1016/j.neunet.2009.05.011](#). URL: <https://doi.org/10.1016/j.neunet.2009.05.011>.
- Williams, Grady, Nolan Wagener, Brian Goldfain, Paul Drews, James M Rehg, Byron Boots, and Evangelos A Theodorou (2017). "Information Theoretic MPC for Model-Based Reinforcement Learning." In: *IEEE International Conference on Robotics and*

- Automation (ICRA).* URL: <http://www.cc.gatech.edu/%7B~%7Dboots3/files/InformationTheoreticMPC.pdf>.
- Williams, Ronald J (1992). "Simple Statistical Gradient-Following Algorithms for Connectionist Reinforcement Learning." In: *Machine Learning*, pp. 229–256.
- Winograd, Terry (1972). *Understanding Natural Language*. Academic Press.
- Wu, Yifan, George Tucker, and Ofir Nachum (2019). "Behavior regularized offline reinforcement learning." In: *arXiv preprint arXiv:1911.11361*.
- Wu, Yifan, George Tucker, and Ofir Nachum (Nov. 2020). "Behavior Regularized Offline Reinforcement Learning." In: *International Conference on Learning Representations (ICLR)*. arXiv: [1911.11361](https://arxiv.org/abs/1911.11361). URL: <http://arxiv.org/abs/1911.11361>.
- Wu, Yuchen, Melissa Mozifian, and Florian Shkurti (2020). *Shaping Rewards for Reinforcement Learning with Imperfect Demonstrations using Generative Models*. arXiv: [2011.01298](https://arxiv.org/abs/2011.01298) [cs.R0].
- Xu, Danfei, Ajay Mandlekar, Roberto Martín-Martín, Yuke Zhu, Silvio Savarese, and Li Fei-Fei (2021). "Deep Affordance Foresight: Planning Through What Can Be Done in the Future." In: *International Conference on Robotics and Automation (ICRA)*. arXiv: [2011.08424v1](https://arxiv.org/abs/2011.08424v1). URL: <https://sites.google.com/stanford.edu/daf>.
- Yamanobe, Natsuki, Weiwei Wan, Ixchel G. Ramirez-Alpizar, Damien Petit, Tokuo Tsuji, Shuichi Akizuki, Manabu Hashimoto, Kazuyuki Nagata, and Kensuke Harada (Jan. 2018). "A Brief Review of Affordance in Robotic Manipulation Research." In: *Journal of the Robotics Society of Japan* 36, pp. 327–337. DOI: [10.7210/jrsj.36.327](https://doi.org/10.7210/jrsj.36.327).
- Yu, Tianhe, Aviral Kumar, Yevgen Chebotar, Karol Hausman, Sergey Levine, and Chelsea Finn (Dec. 2021). "Conservative Data Sharing for Multi-Task Offline Reinforcement Learning." In: *NeurIPS* 34.
- Zech, Philipp, Simon Haller, Safoura Rezapour Lakani, Barry Ridge, Emre Ugur, and Justus Piater (2017). "Computational models of affordance in robotics: a taxonomy and systematic classification." In: *Adaptive Behavior* 25.5, pp. 235–271. DOI: [10.1177/1059712317726357](https://doi.org/10.1177/1059712317726357). eprint: <https://doi.org/10.1177/1059712317726357>. URL: <https://doi.org/10.1177/1059712317726357>.
- Zhang, Chi, Sanmukh Rao Kuppannagari, and Viktor Prasanna (2021). {BRAC}+: *Going Deeper with Behavior Regularized Offline Reinforcement Learning*. URL: <https://openreview.net/forum?id=bMCfFepJXM>.
- Zhang, Marvin, Sharad Vikram, Laura Smith, Pieter Abbeel, Matthew J. Johnson, and Sergey Levine (Aug. 2019). "SOLAR: Deep Structured Representations for Model-Based Reinforcement Learning." In: *International Conference on Machine Learning (ICML)*. arXiv: [1808.09105](https://arxiv.org/abs/1808.09105). URL: <http://arxiv.org/abs/1808.09105>.

- Zhang, Richard, Phillip Isola, Alexei A Efros, Eli Shechtman, and Oliver Wang (2018). "The Unreasonable Effectiveness of Deep Features as a Perceptual Metric." In: *Conference on Computer Vision and Pattern Recognition (CVPR)* (cit. on p. 6).
- Zhang, Ruiyi, Bo Dai, Lihong Li, and Dale Schuurmans (Feb. 2020). "GenDICE: Generalized Offline Estimation of Stationary Values." In: *International Conference on Learning Representations (ICLR)*. arXiv: [2002.09072](https://arxiv.org/abs/2002.09072). URL: <http://arxiv.org/abs/2002.09072>.
- Zhang, Shangtong, Wendelin Boehmer, and Shimon Whiteson (2019). "Generalized Off-Policy Actor-Critic." In: *Advances in Neural Information Processing Systems 32*. Ed. by H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Álché-Buc, E. Fox, and R. Garnett. Curran Associates, Inc., pp. 2001–2011. URL: <http://papers.nips.cc/paper/8474-generalized-off-policy-actor-critic.pdf>.
- Zhang, Tianjun, Benjamin Eysenbach, Ruslan Salakhutdinov, Sergey Levine, and Joseph E. Gonzalez (2021). "C-Planning: An Automatic Curriculum for Learning Goal-Reaching Tasks." In: *ArXiv* abs/2110.12080.
- Zhao, Rui, Xudong Sun, and Volker Tresp (2019). "Maximum Entropy-Regularized Multi-Goal Reinforcement Learning." In: *International Conference on Machine Learning*, pp. 7553–7562.
- Zhao, Rui and Volker Tresp (2019). "Curiosity-Driven Experience Prioritization via Density Estimation." In: *CoRR* abs/1902.08039. arXiv: [1902.08039](https://arxiv.org/abs/1902.08039). URL: <http://arxiv.org/abs/1902.08039>.
- Zhou, Allan, Eric Jang, Daniel Kappler, Alexander Herzog, Mohi Khansari, Paul Wohlhart, Yunfei Bai, Mrinal Kalakrishnan, Sergey Levine, and Chelsea Finn (2019). "Watch, Try, Learn: Meta-Learning from Demonstrations and Reward." In: *CoRR* abs/1906.03352. arXiv: [1906.03352](https://arxiv.org/abs/1906.03352). URL: <http://arxiv.org/abs/1906.03352>.
- Zhu, Henry, Abhishek Gupta, Aravind Rajeswaran, Sergey Levine, and Vikash Kumar (Oct. 2019). "Dexterous Manipulation with Deep Reinforcement Learning: Efficient, General, and Low-Cost." In: *IEEE International Conference on Robotics and Automation (ICRA)*. Vol. 2019-May. Institute of Electrical and Electronics Engineers Inc., pp. 3651–3657. arXiv: [1810.06045](https://arxiv.org/abs/1810.06045). URL: <http://arxiv.org/abs/1810.06045>.
- Ziebart, Brian D, Andrew Maas, J Andrew Bagnell, and Anind K Dey (2008). "Maximum Entropy Inverse Reinforcement Learning." In: *AAAI Conference on Artificial Intelligence*, pp. 1433–1438. ISBN: 9781577353683 (ISBN). arXiv: [arXiv:1507.04888v2](https://arxiv.org/abs/1507.04888v2). URL: <https://www.aaai.org/Papers/AAAI/2008/AAAI08-227.pdf%20http://www.scopus.com/inward/record.url?eid=2-s2.0-57749097473%7B%5C&%7DpartnerID=40%7B%5C%7D5Chttp://www.aaai.org/Papers/AAAI/2008/AAAI08-227.pdf>.

Zucker, Matt, Nathan Ratliff, Anca D Dragan, Mihail Pivtoraiko, Matthew Klingensmith, Christopher M Dellin, J Andrew Bagnell, and Siddhartha S Srinivasa (2013). "CHOMP: Covariant Hamiltonian Optimization for Motion Planning." In: *IJRR*.